

Randomized Algorithms

Yannic Maus

Part I: Randomized algorithms:

Las Vegas algorithms (LV), Monte Carlo algorithms (MC)


Part II: Karger's min-cut algorithm

Part III: Small toolbox:

Probability boosting, Turn MC to LV

Linearity of expectation, Markov's inequality, with high probability

Part IV: Randomized Approximation algorithm for max-cut



Det. Algorithm = Function(Input)

Rand. Algorithm = Function (Input, Random Bits)

- **High level:** Your algorithm can flip coins
- **Example Quicksort:**
The algorithm flips a coin to decide which element to take as the pivot element.
Expected Runtime: $O(n \cdot \log n)$
Worst case runtime: $O(n^2)$

- The output of a randomized algorithm is a random variable
- The execution path of a randomized algorithm is a random variable

- The output of a randomized algorithm is a random variable
- The execution path of a randomized algorithm is a random variable

Think of input x as fixed:

1. Flip coins $r_1, r_2, \dots \in \{Heads, Tails\}$
2. Do some computation
3. Output $Alg(x, r_1, r_2, \dots)$

Possible Statements:

For all inputs x :

$$E[Running\ Time\ (Alg(x, r_1, r_2, \dots))] \leq 10|x| \quad (\text{expected running time})$$

For all inputs x :

$$\Pr(Alg(x, r_1, r_2, \dots) \text{ is correct}) \geq 0.3 \quad (\text{error probability})$$

Las Vegas (LV): Always correct, but may be slow

- output always correct
- running time is a random variable (one demands $E[\text{runtime}] < \infty$)

Monte Carlo (MC): Always fast, but may be incorrect

- output is a random variable, may be false
- runtime is bounded by something deterministically

memory aid: *MC = Mostly correct*


The question that we're asking: \forall inputs x :

- Fix runtime upper bound deterministically as asymptotic function $f(|x|)$
- Provide a lower bound for $\Pr(\text{Alg}(x, r_1, r_2, \dots) = \text{correct output for } x)$

Typical statement on a MC algorithm:


The algorithm has runtime $O(n^3)$ and its output is correct with probability 0.9.

Karger's min-cut algorithm



Definition: A **cut** of a graph $G = (V, E)$ is a partition of its vertices into two disjoint sets $U, W = V \setminus U \subseteq V$. $E(U, W)$ are the edges crossing the cut.

A **minimum cut (min-cut)** is a cut that **minimizes** the number of edges crossing the cut among all cuts.



*graphs are unweighted, undirected

Definition: A **cut** of a graph $G = (V, E)$ is a partition of its vertices into two disjoint sets $U, W = V \setminus U \subseteq V$. $E(U, W)$ are the edges crossing the cut.

A **minimum cut (min-cut)** is a cut that **minimizes** the number of edges crossing the cut among all cuts.

- There may be several minimum cuts.

Remark:

The **max-flow min-s-t-cut theorem** yields a deterministic (involved) algorithm to compute a min-s-t-cut. E.g., in $O(|E|^2|V|) = O(n^5)$, via the Edmonds-Karp Algorithm.

But this is an min-s-t-cut ... not a min-cut. What's the difference?

Karger's contraction algorithm

Input: Graph $G=(V,E)$


Output: Cut (U,W) of G

While $|V|>2$







 pick a random edge e in E

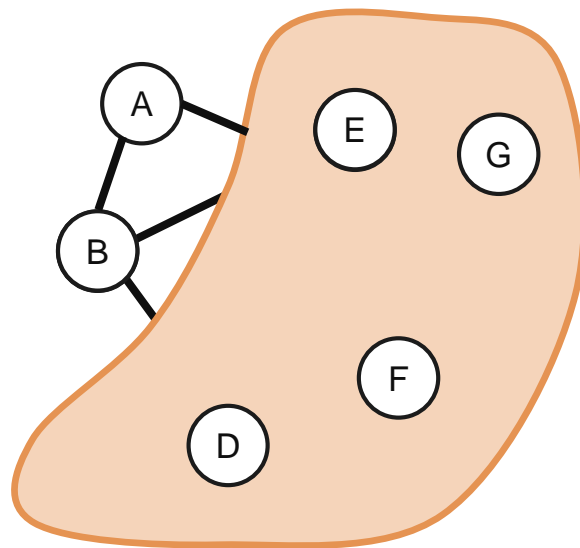

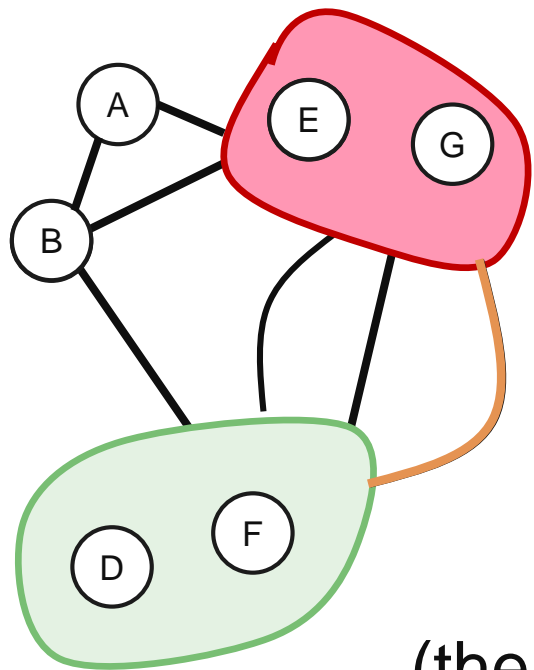
 contract e

 remove self-loops





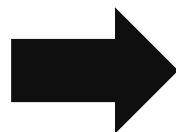

Output the cut induced by the two remaining vertices





(remove no self loops)

(the probability for remaining edges to become selected increases)



Final cut (U,W) of size 3

- After i steps, we have $n - i$ vertices remaining
- We repeat this for $n - 2$ steps, until we have exactly 2 vertices remaining
- The remaining 2 “super” vertices induce a cut

We want to show the following seemingly weak lemma:

Lemma: Karger contraction algorithm outputs a min-cut with probability at least $2/((n - 1)n)$.

Remark: This seems horrible, but indeed it is pretty good as we will see. It is much better than picking a random cut. There are exponentially ($2^{|V|} = 2^n$) many different cuts.

Intuitively Karger's algorithm is better than picking a random cut, because it is unlikely that we contract an edge of a minimum cut, simply because there are few such edges.

Lemma: Karger contraction algorithm outputs a min-cut with probability at least $2/((n-1)n)$.

Proof:

Consider an arbitrary min cut $(U, W = V \setminus U)$ with $C = E(U, W)$

- $e_1, e_2, e_3, \dots, e_{n-2}$: the edges contracted by Karger's algorithm
- E_i : event that e_i does not cross the cut C .


$$\begin{aligned}\Pr(\text{Karger returns cut } C) &= \Pr(E_1 \wedge E_2 \wedge E_3, \dots, \wedge E_{n-2}) \\ &= \Pr(E_1) \cdot \Pr(E_2 | E_1) \cdot \dots \Pr(E_{n-2} | E_1 \wedge \dots \wedge E_{n-3}) \\ &\dots (\text{we will show}) \dots \geq \frac{2}{n(n-1)}\end{aligned}$$

$$\Pr(\bar{E}_i \mid E_1 \wedge \cdots \wedge E_{i-1}) \leq \frac{\text{\#edges in cut } C}{\text{\#remaining edges after } i-1 \text{ contractions}} \leq \frac{2}{n-i+1}$$

The event that we contract an edge of C in the i -th step, given that we have not contracted any edge of C before

$$\Pr(E_i \mid E_1 \wedge \cdots \wedge E_{i-1}) = 1 - \Pr(\bar{E}_i \mid E_1 \wedge \cdots \wedge E_{i-1}) \geq (n-i-1)/(n-i+1)$$

$$\begin{aligned} \text{\#remaining edges} &\geq \text{\#remainingVertices} \cdot \frac{\text{minDegree}}{2} \\ &\geq (n - (i - 1)) \cdot \frac{\text{minDegree}}{2} \\ &\geq (n - i + 1) \cdot (\text{\# edges in cut } C)/2 \end{aligned}$$



Lemma: Karger's contraction algorithm outputs a min-cut with probability at least $2/(n-1)n$.

Proof:

$$\Pr(E_i \mid E_1 \wedge \cdots \wedge E_{i-1}) = 1 - \Pr(\bar{E}_i \mid E_1 \wedge \cdots \wedge E_{i-1}) \geq (n-i-1)/(n-i+1)$$

$$\begin{aligned} \Pr(\text{Karger returns cut } C) &= \Pr(E_1 \wedge E_2 \wedge E_3, \dots, \wedge E_{n-2}) \\ &= \Pr(E_1) \cdot \Pr(E_2 \mid E_1) \cdot \dots \Pr(E_{n-2} \mid E_1 \wedge \dots \wedge E_{n-3}) \\ &= \frac{n-2}{n} \cdot \frac{n-3}{n-1} \cdot \frac{n-4}{n-2} \cdot \dots \cdot \frac{3}{5} \cdot \frac{2}{4} \cdot \frac{1}{3} \geq \frac{2}{n(n-1)} \end{aligned}$$

end of proof

Theorem (Karger): $T = \frac{n(n-1)}{2} \cdot \log 1/\delta$ repetitions of **Karger's contraction algorithm** and returning the smallest cut you see during the process computes a min-cut with probability at least $1 - \delta$.

```
min-cut = ∞
Repeat for T times
    min-cut = min(min-cut, Karger-Contraction-Alg)
Return min-cut
```

One iteration
correct with prob.

$$p = \frac{2}{n(n-1)}$$

$$\Pr(\text{output is not a min-cut}) \leq (1 - p)^T \leq e^{-T \cdot p} = \delta.$$

$$1 - x \leq e^{-x}$$

Remark: *This algorithm only outputs the value of a min-cut.*

Of course we can also output a min-cut by remembering the best cut found.

There are many ways to actually implement Karger's algorithm with varying influence on the complexity.

One Option: Interpret Karger's algorithm as running Kruskal's MST algorithm with random edge weights.

- Recall that Kruskal with a union-find data structure maintains connected components of nodes that have been merged by a spanning tree. These components form the role of a super node in Karger's algorithm .

(the implementation is not the focus of this lecture)

A small Toolbox

For analyzing randomized algorithms

MC: Probability boosting (very important!)

Given: MC algorithm A, correct with probability $p > 0$

New MC algorithm B, correct with probability $\geq 1 - \delta > 0$

Algorithm B: Repeat algorithm A for $p^{-1} \log \left(\frac{1}{\delta} \right)$ iterations

Return “best solution”

Probability that none of the iterations is correct: $(1 - p)^i \leq e^{-p \cdot i} = \delta$

$$1 - x \leq e^{-x}, x \in \mathbb{R}$$

If you have an MC that is correct with probability 1%. Repeat it often enough and return the **best solution**, and you will have an MC algorithm that is correct with probability 99.9%.

Caveat: How to decide which solution is **best**?

In Karger's algorithm we saw an approach for probability boosting for maximization/minimization problems [return the largest/smallest solution].

If you can check whether an output is correct, one can transfer an MC algorithm into an LV algorithm:

Repeat MC algorithm until correct solution is found

This will always produces a correct solution (LV algorithm)

- Expected runtime depends on:
 - error probability of your MC (correct with probability p),
 - the runtime $f(n)$ of the MC algorithm, and
 - the runtime $h(n)$ of the checking procedure

If the correctness check is deterministic, **expected runtime** = $x \cdot (f(n) + h(n))$,
where x is the expected number of p -biased coin flips until you see heads ($x = 1/p$)
(geometric random variable)

Linearity of expectation: Let X_1, \dots, X_n be random variables and a_1, \dots, a_n real values. Then we have:

$$E[\sum a_i X_i] = \sum a_i E[X_i]$$

- Extremely powerful and important tool
- It does not matter whether the random variables X_i are dependent or not

(should be known from probability theory)

Markov inequality: If X is a nonnegative random variable and $a > 0$, then the probability that X is at least a is at most the expectation of X divided by a :

$$\Pr(X \geq a) \leq \frac{E[X]}{a}.$$

One prime application:

Consider an algorithm that should minimize some value X , and we have designed an algorithm that computes a small value for X , in expectation. Then, we obtain:

$$\Pr(X \geq 3 E[X]) \leq E[X]/(3E[X]) = 1/3$$

We obtain: $\Pr(X < 3 E[X]) = 1 - \Pr(X \geq 3E[X]) \geq 2/3$


*of course this works with other values than 3 as well

Definition (with high probability): An algorithm is correct **with high probability** if its output on an instance of size n is correct with probability $\geq 1 - \frac{1}{n}$.

(Typically, we want that algorithms that are correct w.h.p.)

In other words, the probability that the output is incorrect is at most $1/n$.
E.g., for an instance with 100 nodes we require that the input is false with probability at most 1%. On an input with 1000 nodes, we require that the input is false with probability at most 0.1%, etc.


Max-Cut



colors, but not a proper graph coloring

Definition: A **maximum cut (max-cut)** is a cut that **maximizes** the number of edges crossing the cut among all cuts.


- Max-cut is NP-complete (in contrast to min-cut), not proven in this lecture



Size of the cut?

Definition: A **maximum cut (max-cut)** is a cut that **maximizes** the number of edges crossing the cut among all cuts.

- Max-cut is NP-complete (in contrast to min-cut), not proven in this lecture



Size of the cut?

10 cut edges

Randomized Algorithm: Color each vertex randomly red/blue

How many cut-edges do we expect?

What is the probability for an edge to be a cut-edge?



$$\Pr(v = \text{blue} \wedge u = \text{red}) = \Pr(v = \text{blue}) \cdot \Pr(u = \text{red}) = 1/4$$



$$\Pr(v = \text{red} \wedge u = \text{blue}) = \Pr(v = \text{red}) \cdot \Pr(u = \text{blue}) = 1/4$$

$$\Pr(\text{edge } \{u, v\} \text{ is cut edge}) = 1/4 + 1/4 = 1/2$$

Randomized Algorithm: Color each vertex randomly red/blue

For each edge $e \in E$: random variable $X_e = 1$, iff e is cut edge, $X_e = 0$, otherwise

$$E[X_e] = 1 \cdot \Pr(X_e = 1) + 0 \cdot \Pr(X_e = 0) = 1 \cdot \frac{1}{2} = \frac{1}{2}$$

$$X = \sum_{e \in E} X_e \quad \text{Total number of cut edges}$$

$$E[X] = E[\sum X_e] = \sum E[X_e] = |E|/2$$

Lemma: Randomly assigning nodes to the partitions of a cut, in expectation produces $|E|/2$ cut edges.

How to produce a Monte Carlo algorithm?
(for which problem do we get a MC algorithm)

Lemma: Randomly assigning nodes to the partitions of a cut produces **at least $|E|/4$ cut edges**, with probability at least $1/3$.

Proof:

Let $Y = |E| - X$ be the number of monochromatic (non-cut edges).

$$E[Y] = |E| - E[X] = |E|/2.$$

$$\Pr\left(X \leq \frac{|E|}{4}\right) = \Pr\left(Y \geq \frac{3|E|}{4}\right) \leq \frac{E[Y]}{\frac{3|E|}{4}} = \frac{2}{3}. \quad \text{(Markov inequality)}$$

Theorem: For $\delta > 0$, there is a randomized MC algorithm that outputs a cut with at least $|E|/4$ cut edges with probability at least $1 - \delta$ in $O((|V| + |E|) \cdot \log_{\frac{3}{2}} 1/\delta)$ time.

Proof:

We repeat the previous algorithm $T = \log_{\frac{3}{2}}(1/\delta)$ times and output the largest cut that we see throughout. We obtain

$$\Pr\left(\text{output cut} < \frac{|E|}{4}\right) = \left(\frac{2}{3}\right)^T \leq \delta.$$

Runtime: The randomized flipping takes $O(|V|)$ steps. Checking the size of the cut in one iteration takes $O(|E|)$ steps.

Corollary: There is a randomized algorithm that w.h.p. outputs a cut with $|E|/4$ cut edges and has runtime $O((|V| + |E|) \cdot \log n)$.


Proof:

Use the previous theorem and set $\delta = 1/n$ to obtain that the error probability is at most $1/n$.

- Max-cut is NP-complete
- Our algorithm usually does not output an optimal solution
- Still, we get a constant approximation
- There is no PTAS for max-cut unless $P=NP$

Exercise: Show that the presented algorithm provides a 4-approximation.

Randomization is (a) great (tool)



- often simple algorithms
- often difficult analysis