

Approximation Algorithms

Yannic Maus



"I can't find an efficient algorithm, but neither can all these famous people."

What do you do? Quit your job & go home? Change problem?

figure by Stefan Szeider, TU Vienna

- **Option 1:** Code a “slow” approach, e.g., knapsack in $O(n \cdot \text{sumOfWeights})$,
Sometimes this is fast enough
- **Option 2:** Just code some approach without theoretical guarantees
This is pretty bad if we don't have any guarantee for our solution, isn't it?
- **Option 3:** Code simple approach, e.g., greedy, and prove theoretical guarantees

Approximation algorithms (this lecture)

- *Often run in polynomial time.*
- *They don't give you optimal solutions on all instances, but gives you solutions “close to optimal”.*
- *Provable guarantees on the quality of your solution*
- *Often simple greedy algorithm*


Approximation ratio $\rho(n)$ of an algorithm A : if for any n -sized input the algorithm A produces a solution with value C such that


- $\frac{C}{OPT} \leq \rho(n)$ for a **minimization problem**, and
- $\frac{C}{OPT} \geq \rho(n)$ for a **maximization problem**.

Minimization problem: $\rho(n) \geq 1$. **Maximization problem:** $\rho(n) \leq 1$.

There are different notions of approximation in the literature, e.g., being an additive term away from the optimal solution (above multiplicative factor), or OPT/C instead of C/OPT , or a mix of multiplicative and additive approximation.


Minimum Vertex Cover (MVC)






A vertex cover of an undirected graph G is a set S of vertices such that every edge in $E(G)$ is incident to at least one vertex in S . In other words, for every edge (u,v) in G , at least one of the vertices u or v is in the set S .


Goal: Compute a vertex cover of **minimum size**.



Not a valid cover.



Not a valid cover.
green edge not covered



Valid vertex cover,
but a really large one

Decision variant: Given graph G and k , is there a VC of G with size $\leq k$?

NP containment:


Polynomially verifiable whether a given subset $S \subseteq V$ is a valid vertex cover of size k .

Decision variant: Given graph G and k , is there a MVC of G with size $\leq k$?

NP-hardness:

Reduce 3-SAT to the decision variant of MVC (not part of this lecture)

Given 3-sat formular ϕ with m variables, l clauses, create graph G as follows:



An edge between “clause literal” each the corresponding variable literal

Exercise: 3-SAT formular ϕ is satisfiable iff G has a MVC of size $k = m + 2l$

How do you solve the problem?

Greedy!

“Vertex greedy” for MVC

Input: Graph $G = (V, E)$

Output: Vertex cover C

$C =$ empty set

$F = E(G)$

while F is not empty:


 choose largest degree vertex v in $G=(V-C, F)$

 add v to C


 remove all edges incident on v from F

return C


Example: Greedily adding vertices




Example: Greedily adding vertices




Example: Greedily adding vertices




Example: Greedily adding vertices




Example: Greedily adding vertices



Example: Greedily adding vertices



Example: Greedily adding vertices




vertex cover of size 8

Vertex Greedy does not give a constant approximation

Theorem: Greedily adding the largest degree vertex (in the graph induced by uncovered edges) is not a constant factor approximation for MVC.

This figure: $k = 3$

$k!$ vertices of degree k



For $x = k, k - 1, \dots, 1$: $k!/x$ vertices degree x

Algorithm: ApproximateVertexCover(G)

Input: Graph $G = (V, E)$

Output: Vertex cover C

C = empty set

$F = E(G)$

while F is not empty:


 choose any edge (u, v) from F

add u and v to C


 remove all edges incident on u and v from F

return C


Example: Greedily adding vertices




Example: Greedily adding vertices




Example: Greedily adding vertices



Example: Greedily adding vertices



Example: Greedily adding vertices



vertex cover of size 10 (worse than previous VC)

Theoretically, “greedy edge picking” has a better approximation guarantee (next slide) than “greedy vertex picking”

Theorem: Greedily adding both vertices of any still uncovered edge is a 2-approximation for the minimum vertex cover problem.

Proof:

F : the set of edges that our greedy algorithm picked. **F is a matching in G .** Why?
 $C = V(F)$ be the computed vertex cover.

- Clearly C is a vertex cover, as we continue adding vertices until all edges are covered.
- We have $|C| = 2|F|$ (no overlap, as F is a matching in G)

Let C_{OPT} be any optimal solution to MVC.

$|C_{OPT}| \geq |F|$, because C_{OPT} has to cover every edge, including all edges in F .

But no vertex of G can cover more than a single edge of F , as F is a matching.

$$\Rightarrow |C| = 2|F| \leq 2|C_{OPT}|.$$

- exact MVC is NP-complete
- “vertex greedy” does not give a constant approximation factor
- “edge greedy” gives 2-approximation
- that does not mean that “vertex greedy” is always worse than “edge greedy”

Hardness of approximation (not covered in this lecture):

- It is NP-hard to compute anything better than a $\sqrt{2} \approx 1.41$ –approximation
[Khot, Minzer, Shafra, 2017]
- It is conjectured that it is NP-hard to compute anything better than a $2-\epsilon$ approximation for any constant ϵ (unique games conjecture) [its details go well beyond the scope of this lecture!]

[Khot, Regev, 2008]

Set Cover

Set Cover: Definition & Example

Input:


- **Universe** $X = \{x_1, \dots, x_n\}$ of n elements
- **Collection of Sets** $S = \{S_1, \dots, S_k\}$, each $S_i \subseteq X$

Set Cover: a collection of sets (indices) $I \subseteq \{1, \dots, k\}$
s.t. all elements are covered, i.e.,

$$X \subseteq \bigcup_{i \in I} S_i$$

Goal: Select a minimum size set cover (minimize $|I|$).

(Minimum) set cover is NP-complete.



Optimal solution: S_1, S_3, S_5

How do you solve the problem?

Greedy!

```
GreedySetCover(Universe, Sets):  
    I = {}  
    while X is not empty:  
        MaxSet = argmax(Set in Sets, |Set ∩ X|)  
        I = I ∪ {MaxSet}  
        X = X - MaxSet  
    return I
```


Input:

- **Universe** $X = \{x_1, \dots, x_n\}$ of n elements
- **Collection of Sets** $S = \{S_1, \dots, S_k\}$, each $S_i \subseteq X$

Set Cover: a collection of sets (indices) $I \subseteq \{1, \dots, k\}$
s.t. all elements are covered, i.e.,

$$X \subseteq \bigcup_{i \in I} S_i$$

Goal: Select a minimum size set cover (minimize $|I|$).

(Minimum) set cover is NP-complete.

Proof (Correctness):

The greedy algorithm computes a valid cover, as we keep adding sets until all elements are covered.

Let C_{OPT} be an optimal cover, let $t = |C_{OPT}|$
Let X_k be the elements in iteration k . $X_0 = X$

Claim: For all $0 \leq k$, the set X_k can be covered with t sets.

Proof: The original set X can be covered with t sets, so the same is true for $X_k \subseteq X$

In step k , there exists a set that covers at least $|X_k|/t$ elements (pigeonhole principle)

\Rightarrow in step k the greedy algorithm is going to pick a set of size at least $|X_k|/t$

For all k , we have $|X_{k+1}| \leq \left(1 - \frac{1}{t}\right) |X_k|$

By induction for all $k \geq 0$: $|X_k| \leq \left(1 - \frac{1}{t}\right)^k |X_0| = \left(1 - \frac{1}{t}\right)^k \cdot |X|$

Let C_{OPT} be an optimal cover, let $t = |C_{OPT}|$
Let X_k be the elements in iteration k . $X_0 = X$

By induction for all $k \geq 0$: $|X_k| \leq \left(1 - \frac{1}{t}\right)^k |X_0| = \left(1 - \frac{1}{t}\right)^k \cdot |X|$

When do we stop? How many sets do we choose?

We stop when $X_k = \emptyset$ ($|X_k| < 1$), chosen at most k sets

For $k^* = t \cdot (\lceil \log |X| \rceil + 1)$, we obtain

$$|X_{k^*}| \leq \left(1 - \frac{1}{t}\right)^{k^*} \cdot |X| \leq e^{-\frac{k^*}{t}} \cdot |X| = e^{-\frac{k^*}{t} + \log |X|} \leq e^{-1} < 1$$

At most $k^* = t \cdot (\lceil \log |X| \rceil + 1)$ sets, so we have a $(\lceil \log |X| \rceil + 1)$ -approximation.

Theorem: Greedily adding the set that covers most uncovered elements is a $(\lceil \log |X| \rceil + 1)$ -approximation for the (minimum) set cover problem

- This is not a constant approximation factor! The more elements, the worse the approximation!
- Still, the problem and algorithm appear are used quite often

Hardness of approximation (not covered in this lecture):

- It is known that computing a constant-factor approximation is NP-hard
→ (difficult) research area hardness of approximation
[Raz, Safra '97]
- Computing a better than $(1-o(1))$ $\ln n$ -approximation is NP-hard
[Alon, Moshkovitz, Safra '06], [Dinur, Steurer '13]

Partition

Input:

n positive integers s_1, \dots, s_n

Goal:

Partition the set of integers (integer indices) into two sets $A, B \subseteq \{1, \dots, n\}$ to minimize

$$\max \left\{ \sum_{i \in A} s_i, \sum_{i \in B} s_i \right\}$$

- NP-complete
- Dynamic programming: $O(n \cdot \sum s_i)$ (does not contradict NP-hardness)
- Bruteforce by trying all combinations: $O(2^n)$

2-Approximation: Any distribution is a 2-approximation. Not very helpful



Polynomial time approximation scheme (PTAS): For each $\epsilon > 0$:

- Computes a $(1 \pm \epsilon)$ -approximation
 - Runtime is polynomial in input for fixed ϵ .
-
- The runtime can be different for different ϵ
 - So $O(n^{1/\epsilon})$ is fine, e.g. for $\epsilon = 0.01$, this is $O(n^{100})$.
 - Even $n^{\exp(\frac{1}{\epsilon})}$ is also fine, but maybe not useful in practice

Fix: $m = \left\lceil \frac{1}{\epsilon} \right\rceil - 1$

Order from largest to smallest $s_1 \geq s_2 \geq \dots \geq s_n$.

Compute an **optimal solution** (A,B) for the first m elements.

Greedy for the rest:

For $i=m+1, \dots, n$

 If $\text{weight}(A) \leq \text{weight}(B)$:

 add i to A,

 else

 add i to B.

Runtime: $O(n \cdot \log n + 2^{\frac{1}{\epsilon}+1} + n)$

Notation:

A, B : Sets at the end of the algorithm

A_k, B_k : Sets after the k -th step

Proof: Wlog assume at the end we have $weight(A) \geq weight(B)$

Let s_k be the last element added to A .

Look at the snapshot after adding k :

We only need to prove $w(A) = w(A_k) \leq (1 + \epsilon)OPT$

but we don't know OPT , how can we compare to it?

As we do not know OPT , we show that our solution is even a good approximation of L

$$L = \frac{1}{2} \sum s_i \leq OPT$$

lower bound for OPT (for a maximization problem we would need an upper bound)

Case 1 (k was added in the first phase) $\Rightarrow k \leq m$

After the first phase, A was optimal for the smaller problem of adding the first m elements. Later, we never add anything to A. We obtain:

$$w(A) = w(A_k) \leq OPT(s_1, \dots, s_m) \leq OPT(s_1, \dots, s_n)$$

→ Approximation ratio in this case is 1

Case 2 (k was added in the second phase) $\Rightarrow k > m$

Claim: $s_k \leq 2L/(m+1)$

Proof: As $s_1 \geq s_2 \geq \dots \geq s_k$ we get:

$$2L \geq \sum_{i=1}^k s_i \geq \sum_{i=1}^k s_k \geq (m+1) \cdot s_k,$$

claim follows by dividing by $m+1$.

$w(A) \leq w(B_{k-1}) + s_k$ (we added k to A because of $w(A_{k-1}) \leq w(B_{k-1})$, and never changed A afterwards)

$$w(A) \leq w(B_{k-1}) + s_k \leq w(B) + s_k = 2L - w(A) + s_k$$

$$w(A) \leq L + 0.5 \cdot s_k \leq (1 + \epsilon)L \leq (1 + \epsilon)OPT$$

Theorem: For any $\epsilon > 0$ there exists an algorithm that computes a $(1 + \epsilon)$ -approximation of the partition problem in time $O(n \cdot \log n + 2^{\frac{1}{\epsilon}+1} + n)$.

- Often simple greedy algorithms provide good approximations, but not always
- Sometimes it is hard to compare the algorithms performance with OPT, instead compare with a “lower bound” on OPT (as the L value in the partition proof, or $|OPT| \geq |Matching|$ in MVC)
- PTAS: can approximate arbitrarily well, but one pays for it in terms of runtime
- Not discussed: often approximation also makes sense for problems that are not NP-complete

In practice: *If you use a “heuristic” or “greedy” approach, you should ask yourself whether you provide a theoretical guarantee for it. Unfortunately, often this is not possible (e.g., in many machine learning algorithms)*