

A*-Algorithm


Yannic Maus



A*-algorithm was introduced in 1968:

Peter E. Hart, Nils, J. Nilsson, Bertram Raphael –


*A **Formal** Basis for the **Heuristic** Determination of **Minimum Cost Paths***



A*-algorithm was introduced in 1968:

Peter E. Hart, Nils, J. Nilsson, Bertram Raphael –

*A **Formal** Basis for the **Heuristic** Determination of **Minimum Cost Paths***




- Why do we need another shortest path algorithm?

There are many shortest path algorithms: BFS, Dijkstra, Floyd Warshall, ...

- The A*-Algorithm and its heuristic (proof of correctness)
- Advantages/Disadvantages


A*-Algorithm, yet another shortest path algorithm?

A* = A star is a **shortest path algorithm**




A*-Algorithm, yet another shortest path algorithm?

A* = A star is a **shortest path algorithm**



A*-Algorithm, yet another shortest path algorithm?

A* = A star is a **shortest path algorithm**




There are many shortest path algorithms

- **BFS:** unweighted edges
- **Dijkstra:** weighted edges
- **Floyd Warshall:** weighted edges
- **Bellman-Ford:** weighted edges, even negative weights

A*-Algorithm, yet another shortest path algorithm?

A* = A star is a **shortest path algorithm**




There are many shortest path algorithms

- **BFS:** unweighted edges
 - **Dijkstra:** weighted edges
 - **Floyd Warshall:** weighted edges
 - **Bellman-Ford:** weighted edges, even negative weights
- } Recap (A* is more involved but similar)

Breadth First Search (unweighted graph):


Look at paths of length 1,
then length 2,
then length 3,
etc.,
until a path is found:



guaranteed shortest hop path.

Breadth First Search (unweighted graph):


Look at paths of length 1,
then length 2,
then length 3,
etc.,
until a path is found:



guaranteed shortest hop path.


Breadth First Search (unweighted graph):

Look at paths of length 1,
then length 2,
then length 3,
etc.,
until a path is found:




guaranteed shortest hop path.

guaranteed shortest hop path.



Breadth First Search (unweighted graph):


Look at paths of length 1,
then length 2,
then length 3,
etc.,
until a path is found:



guaranteed shortest hop path.


Breadth First Search (unweighted graph):

Look at paths of length 1,
then length 2,
then length 3,
etc.,
until a path is found:




guaranteed shortest hop path.

- If we use BFS to find the path (disregarding weights), we would use **queue** to enqueue each path



BFS

Expand v and queue this path



BFS from s to t :

create a **queue** of paths (a vector), q

$q.enqueue(s \text{ (path)})$

while q is not empty and t is not yet visited:

$path = q.dequeue()$

$v =$ last element in path

 mark v as visited

 if $v == t$ then stop

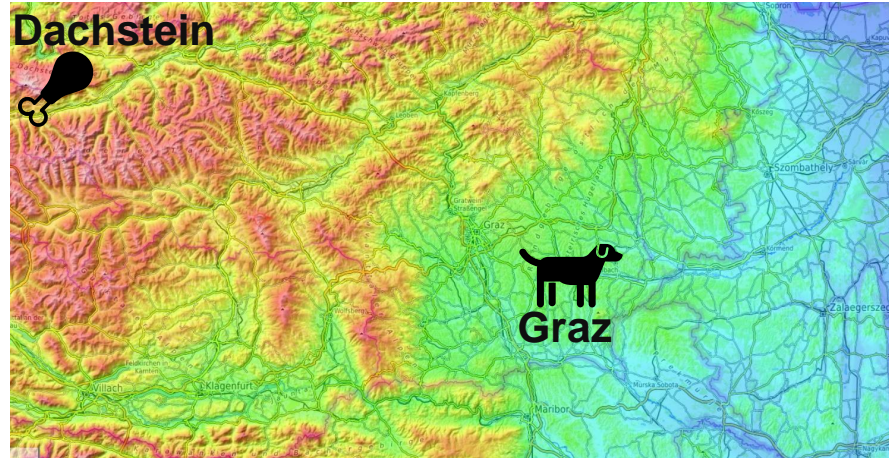
 for each unvisited neighbor of v :

 make new path with v 's neighbor as last element

$q.enqueue(\text{new path})$

QUEUE = FIFO

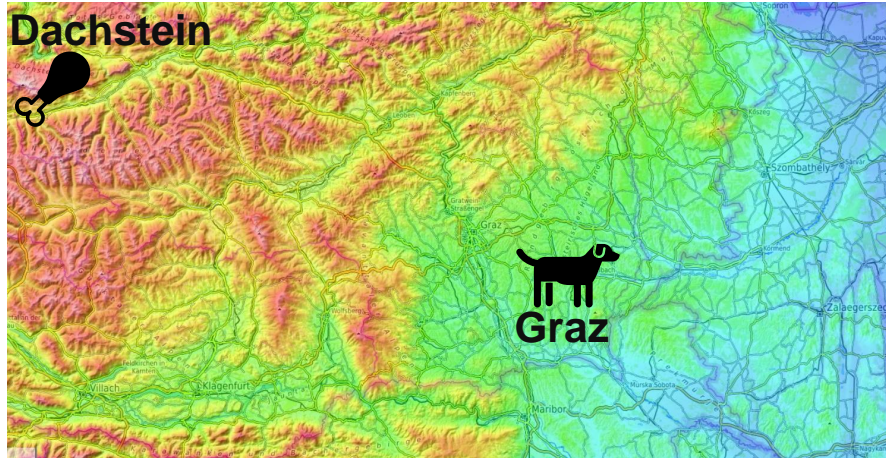
BFS is often not enough ...




(we want weights)



BFS is often not enough ...




(we want weights)




BFS does not work with weights

- If we use BFS to find the path (disregarding weights), we would use **queue** to enqueue each path



BFS

Expand v and queue this path



- [illegible]

- ## Cost of path 22!

Dijkstra from s to t :

create a **priority queue** of paths (a vector), q

$q.enqueue(s \text{ (path)})$

while q is not empty and t is not yet visited:

$path = q.dequeue()$ **//order determined by costs of path**

v = last element in path

 mark v as visited

 if $v == t$ then stop

 for each unvisited neighbor of v :

 make new path with v 's neighbor as last element **//costs!**

$q.enqueue(\text{new path})$

BFS from s to t :

create a ~~priority~~ queue of paths (a vector), q

$q.enqueue(s \text{ (path)})$

while q is not empty and t is not yet visited:

$path = q.dequeue()$ ~~//order determined by costs of path~~

v = last element in path

 mark v as visited

 if $v == t$ then stop

 for each unvisited neighbor of v :

 make new path with v 's neighbor as last element ~~//costs!~~

$q.enqueue(\text{new path})$


Dijkstra's algorithm is what we call a "**greedy**" algorithm.

Take what's **best at the given time**, and it finds an **optimal solution!**
(for many problems greedy algorithms don't give optimal solutions)

Edsgar Dijkstra was very influential

Countless contributions to computer science

1930-2002





Yannic Maus




Dijkstra Award (in Distributed Computing)

- algorithms in large networks
- massively parallel computing (MapReduce by google, ..)
- blockchain, ...
- ...




Dijkstra's is great

but we can do better?




Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.




Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.




Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.




Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.



Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.




But, we know something about how to get to Dachstein, we need to go Northwest from Graz!

*This is more information! Let's not only prioritize by weights, but also give some priority to the **direction** we want to go.*

Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.



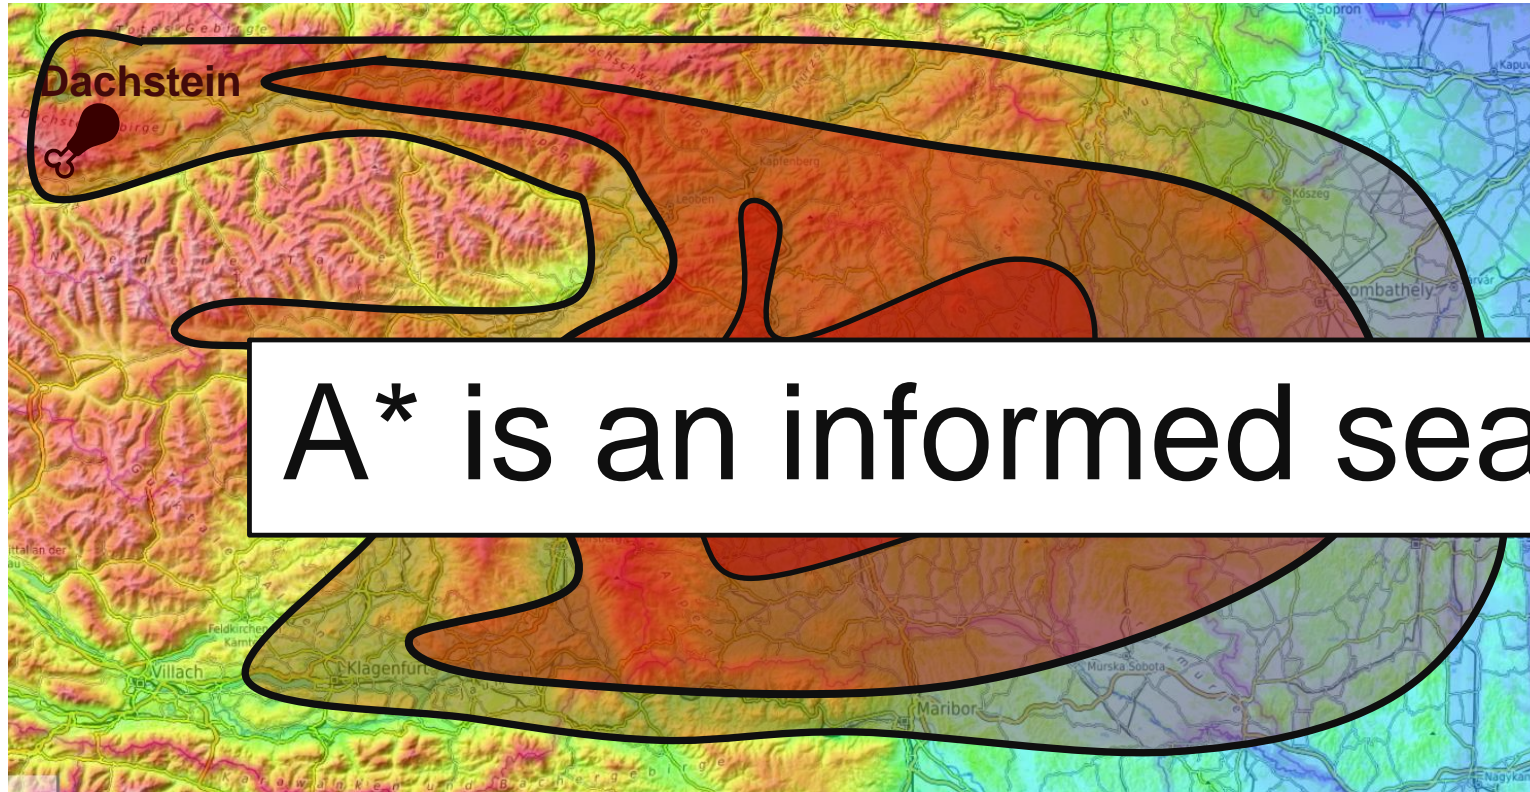
But, we know something about how to get to Dachstein, we need to go Northwest from Graz!

*This is more information! Let's not only prioritize by weights, but also give some priority to the **direction** we want to go.*

E.g., we will add more information based on a heuristic, (which could be direction in the case of a street map.)

Why do we search in all directions?

If we want to travel from Graz to the Dachstein, Dijkstra's algorithm will look at path distances around Graz.




But, we know something about how to get to Dachstein, we need to go Northwest from Graz!

A* is an informed search algorithm

It's not just about distance, it also gives some priority to the direction we want to go.

E.g., we will add more information based on a heuristic, (which could be direction in the case of a street map.)

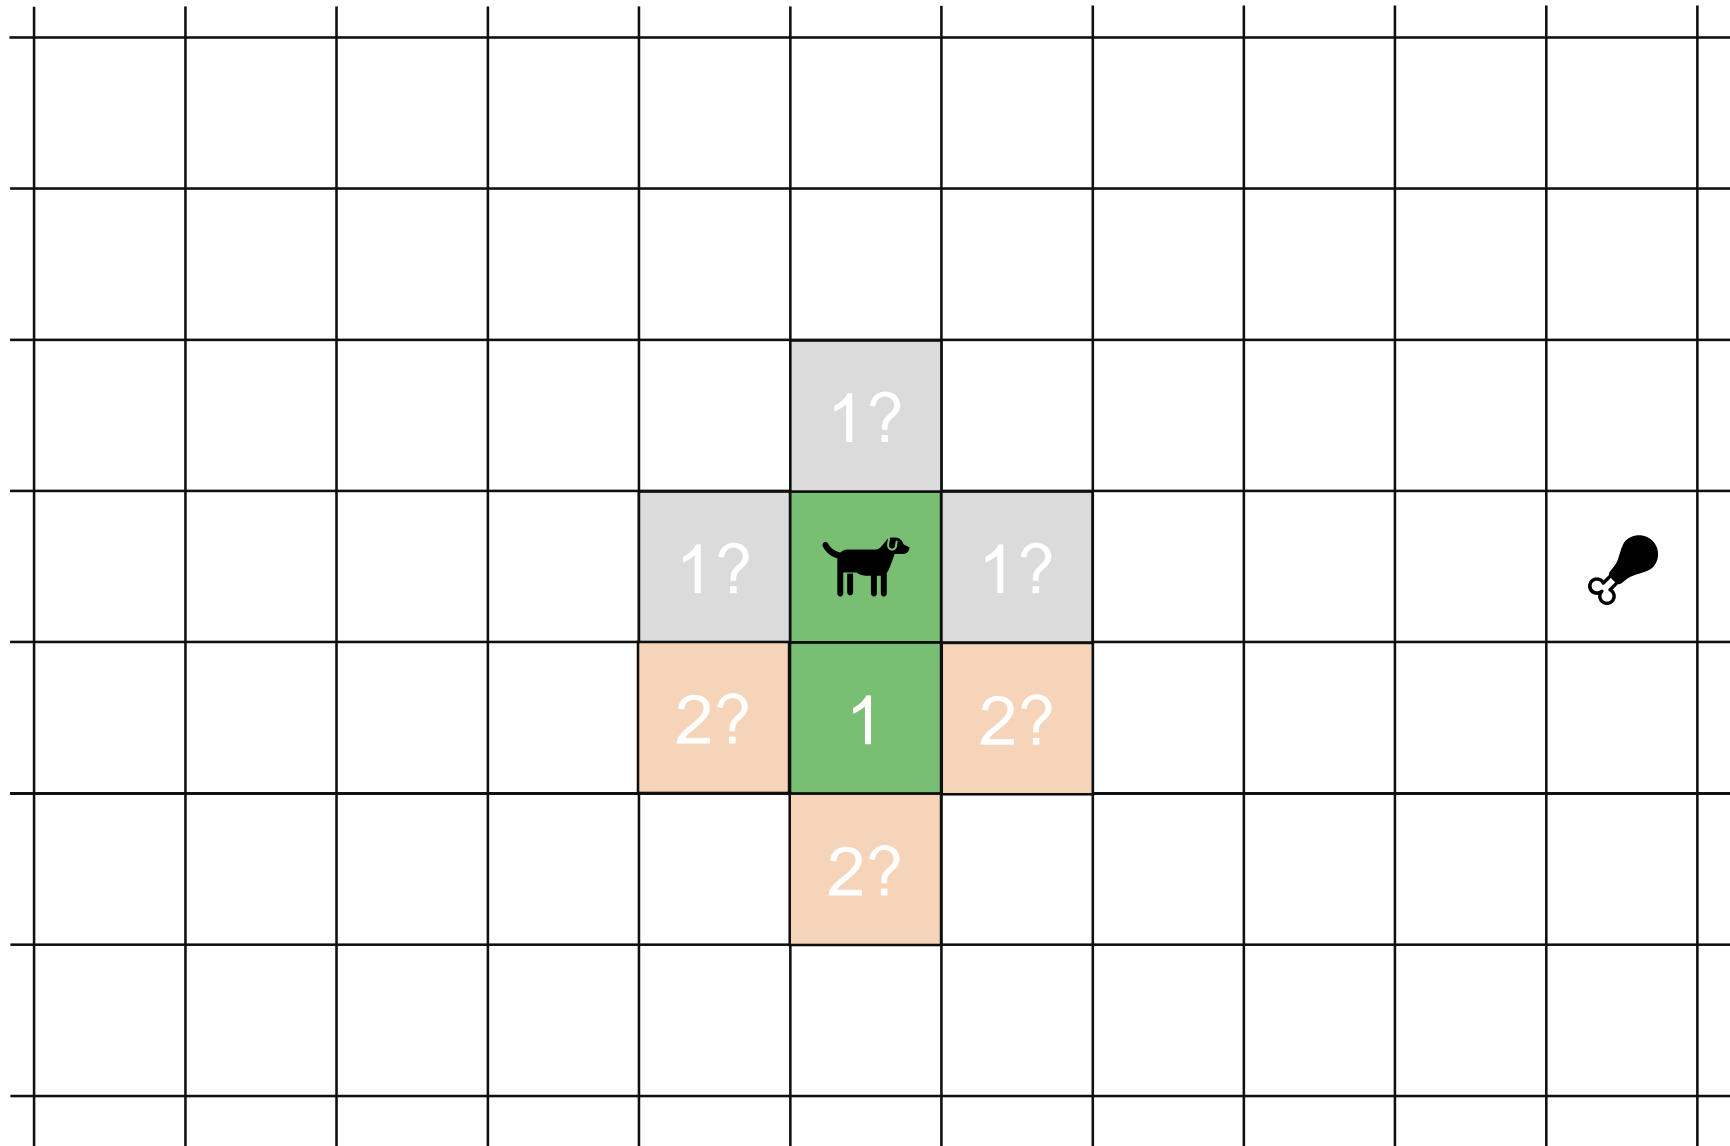
Dijkstra where each edge has cost 1



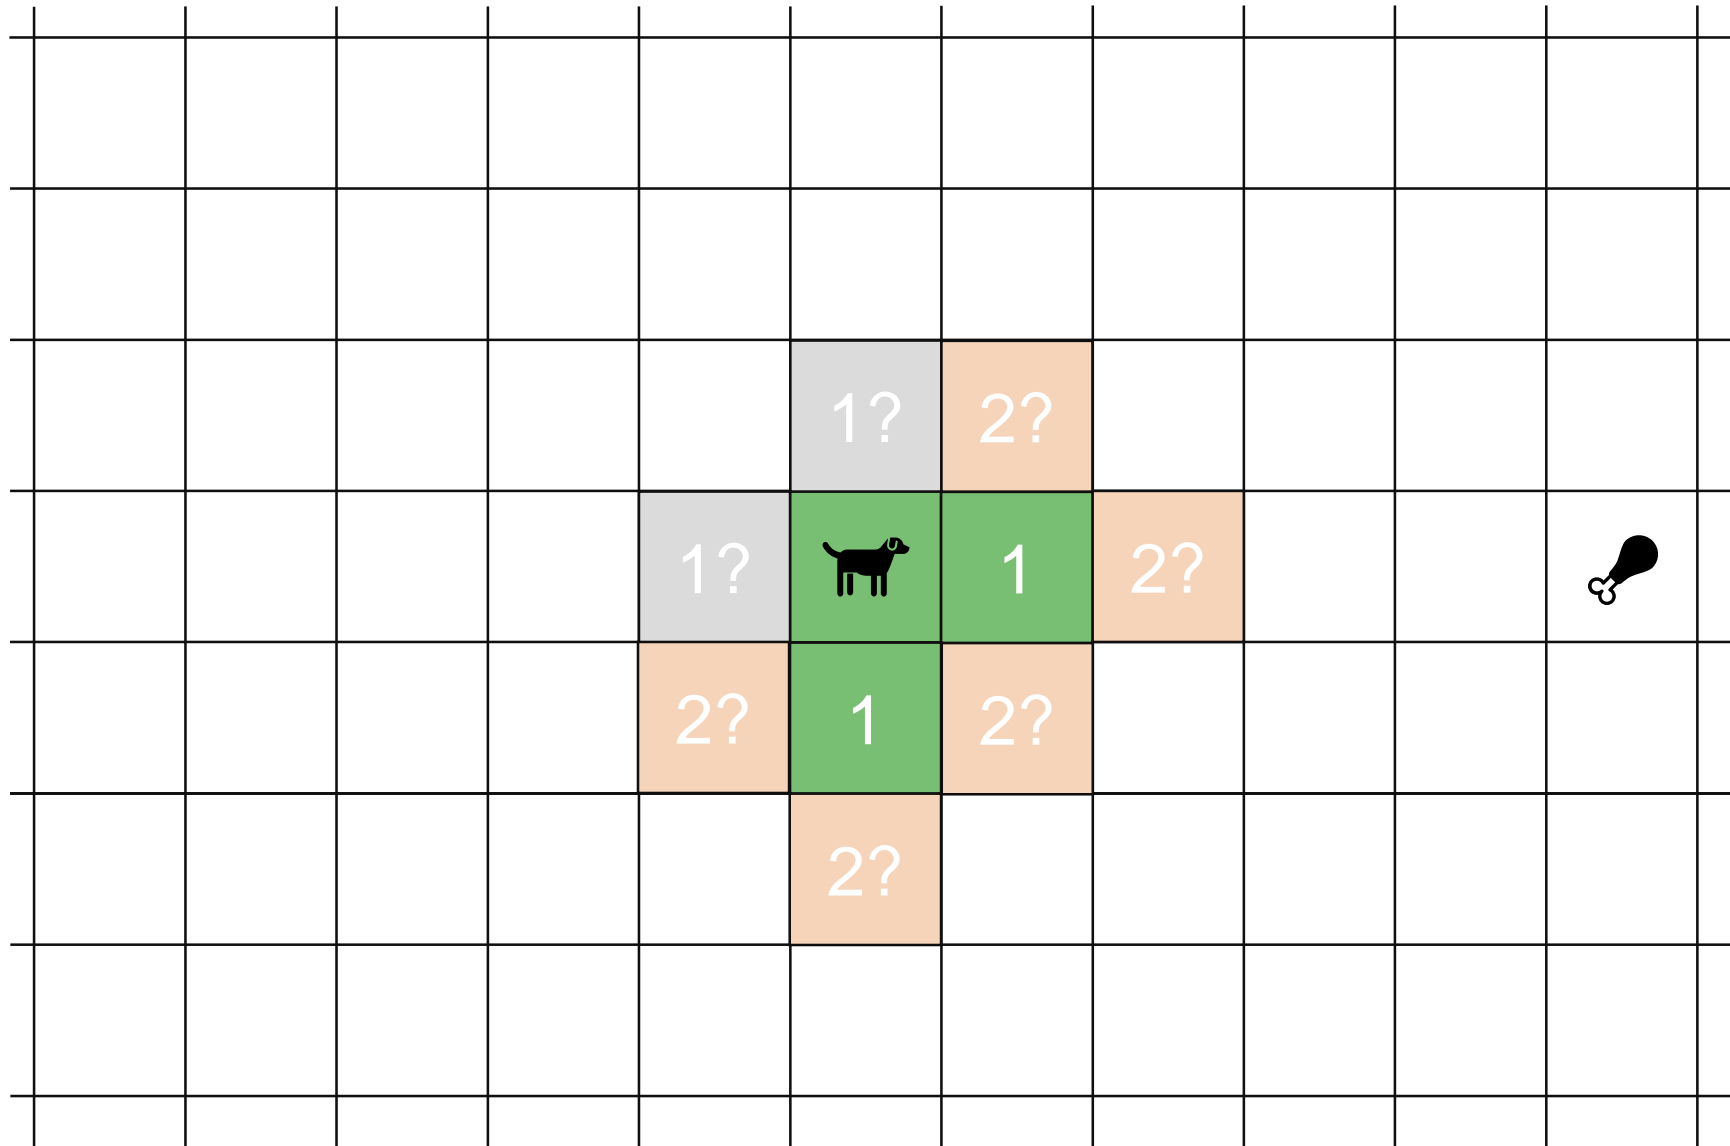
Greedyly continue
at smallest

$\text{distance}(s, u)$

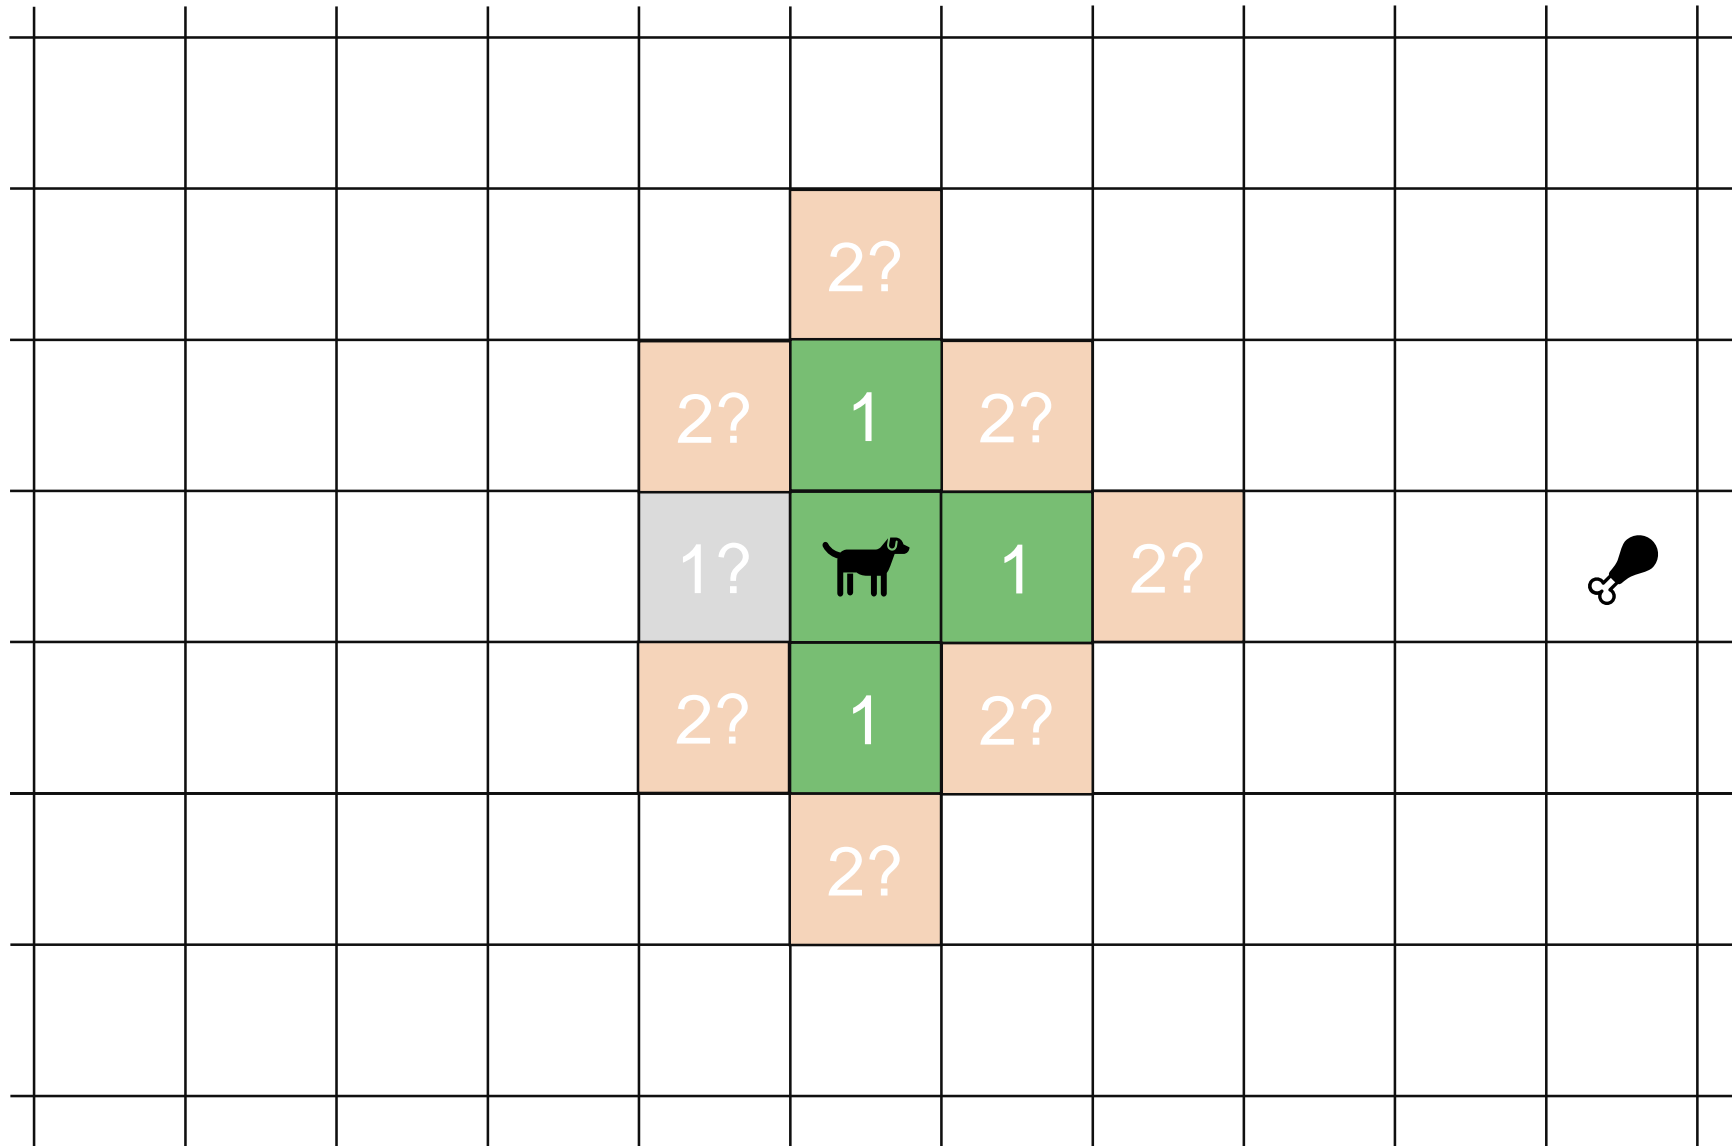
Dijkstra where each edge has cost 1



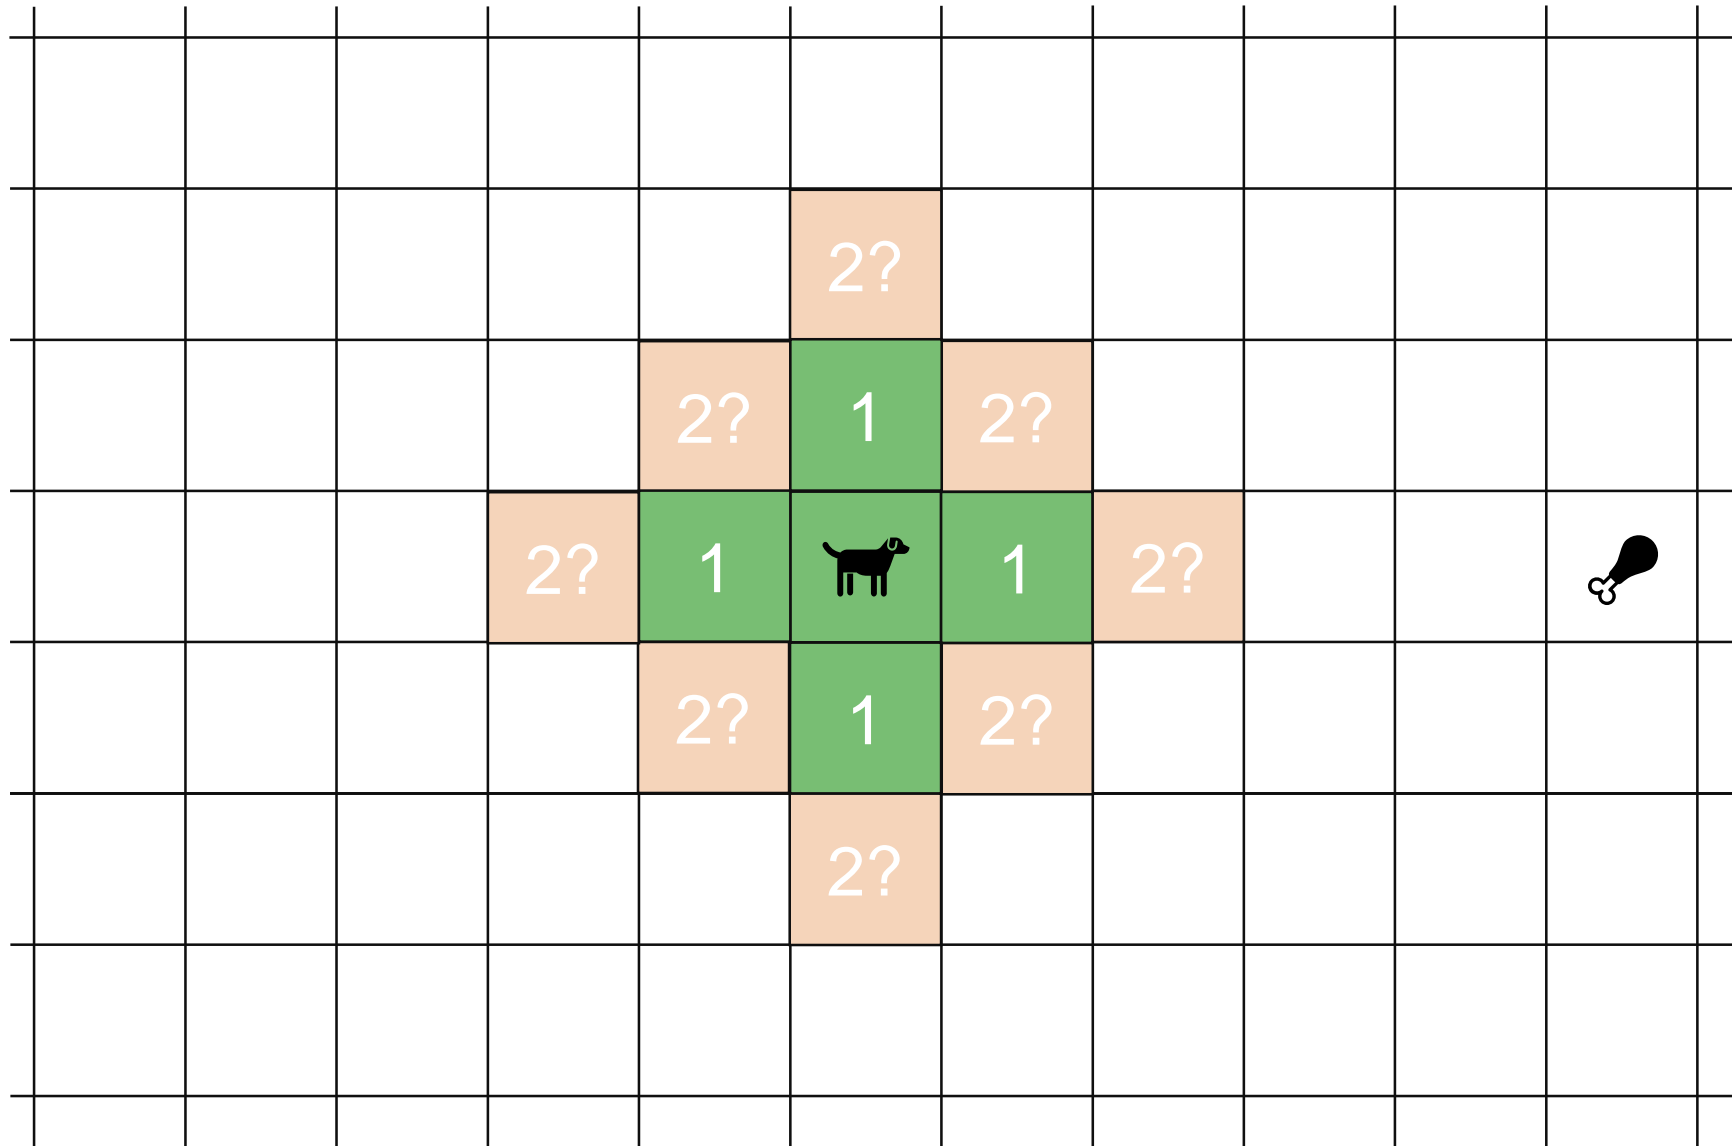
Dijkstra where each edge has cost 1



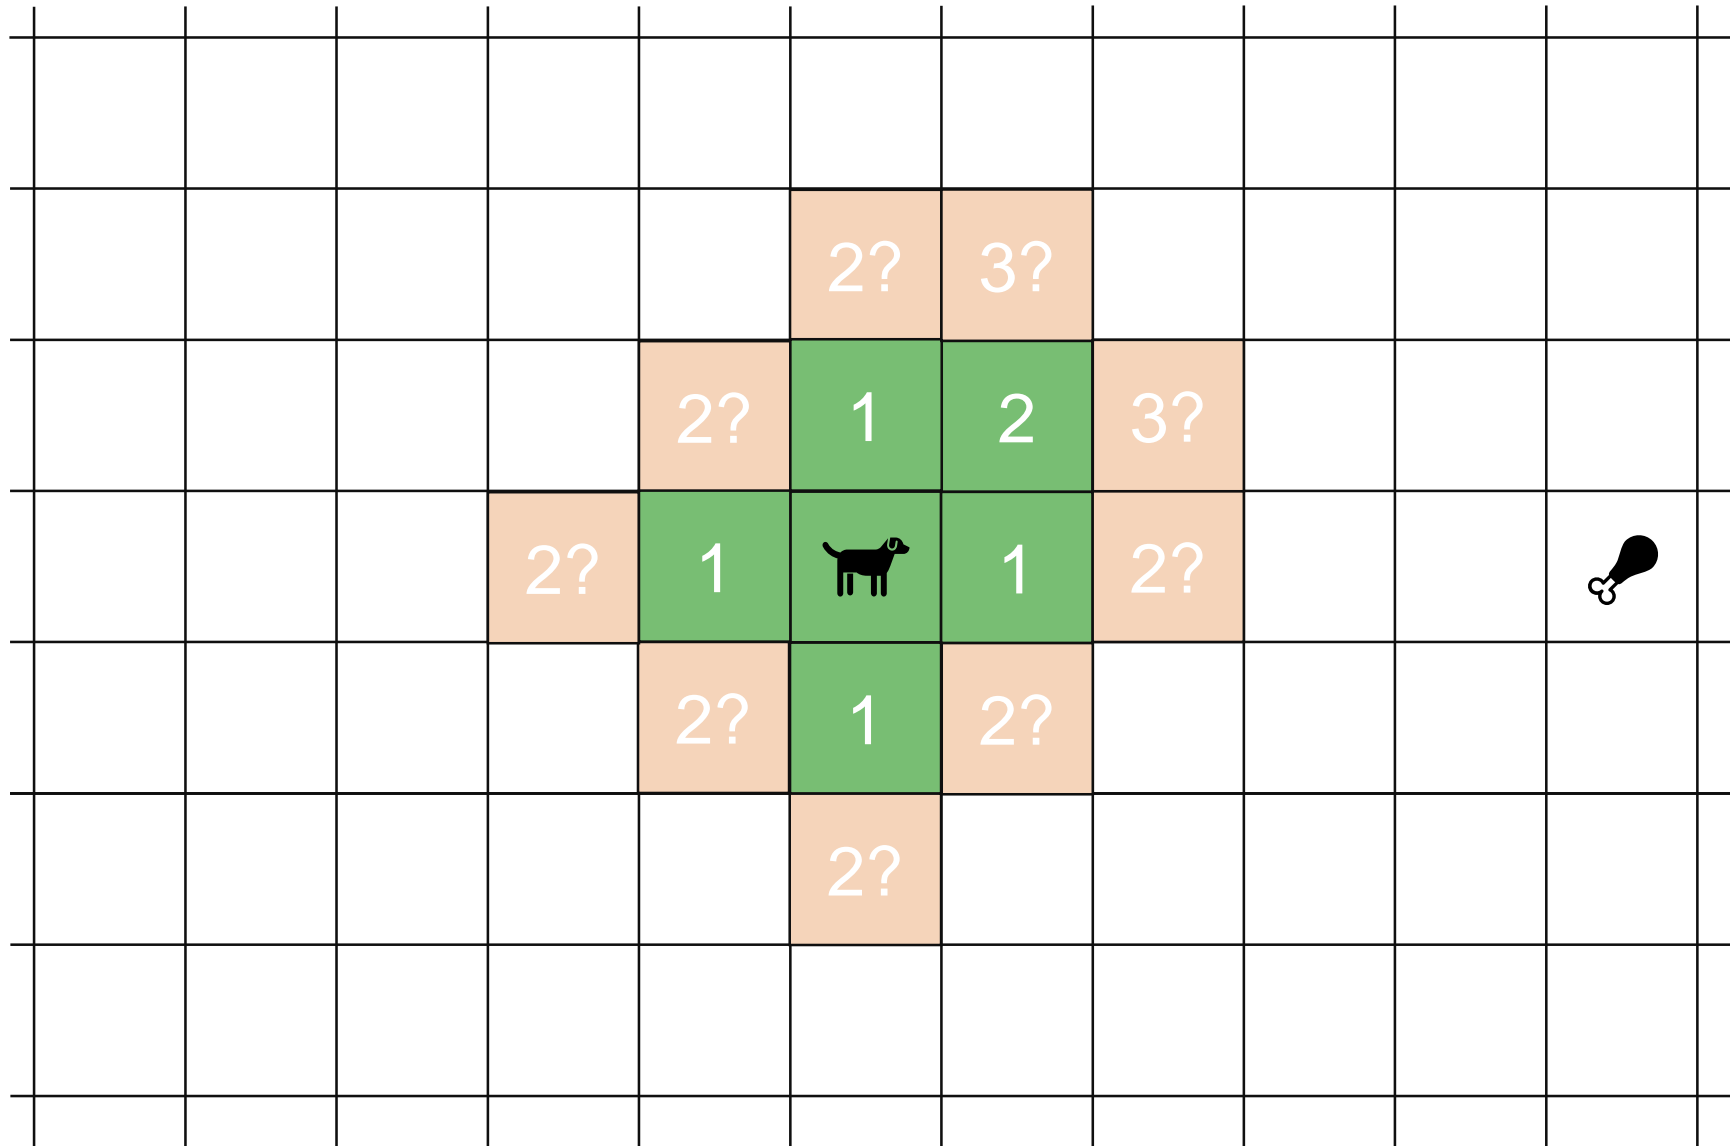
Dijkstra where each edge has cost 1



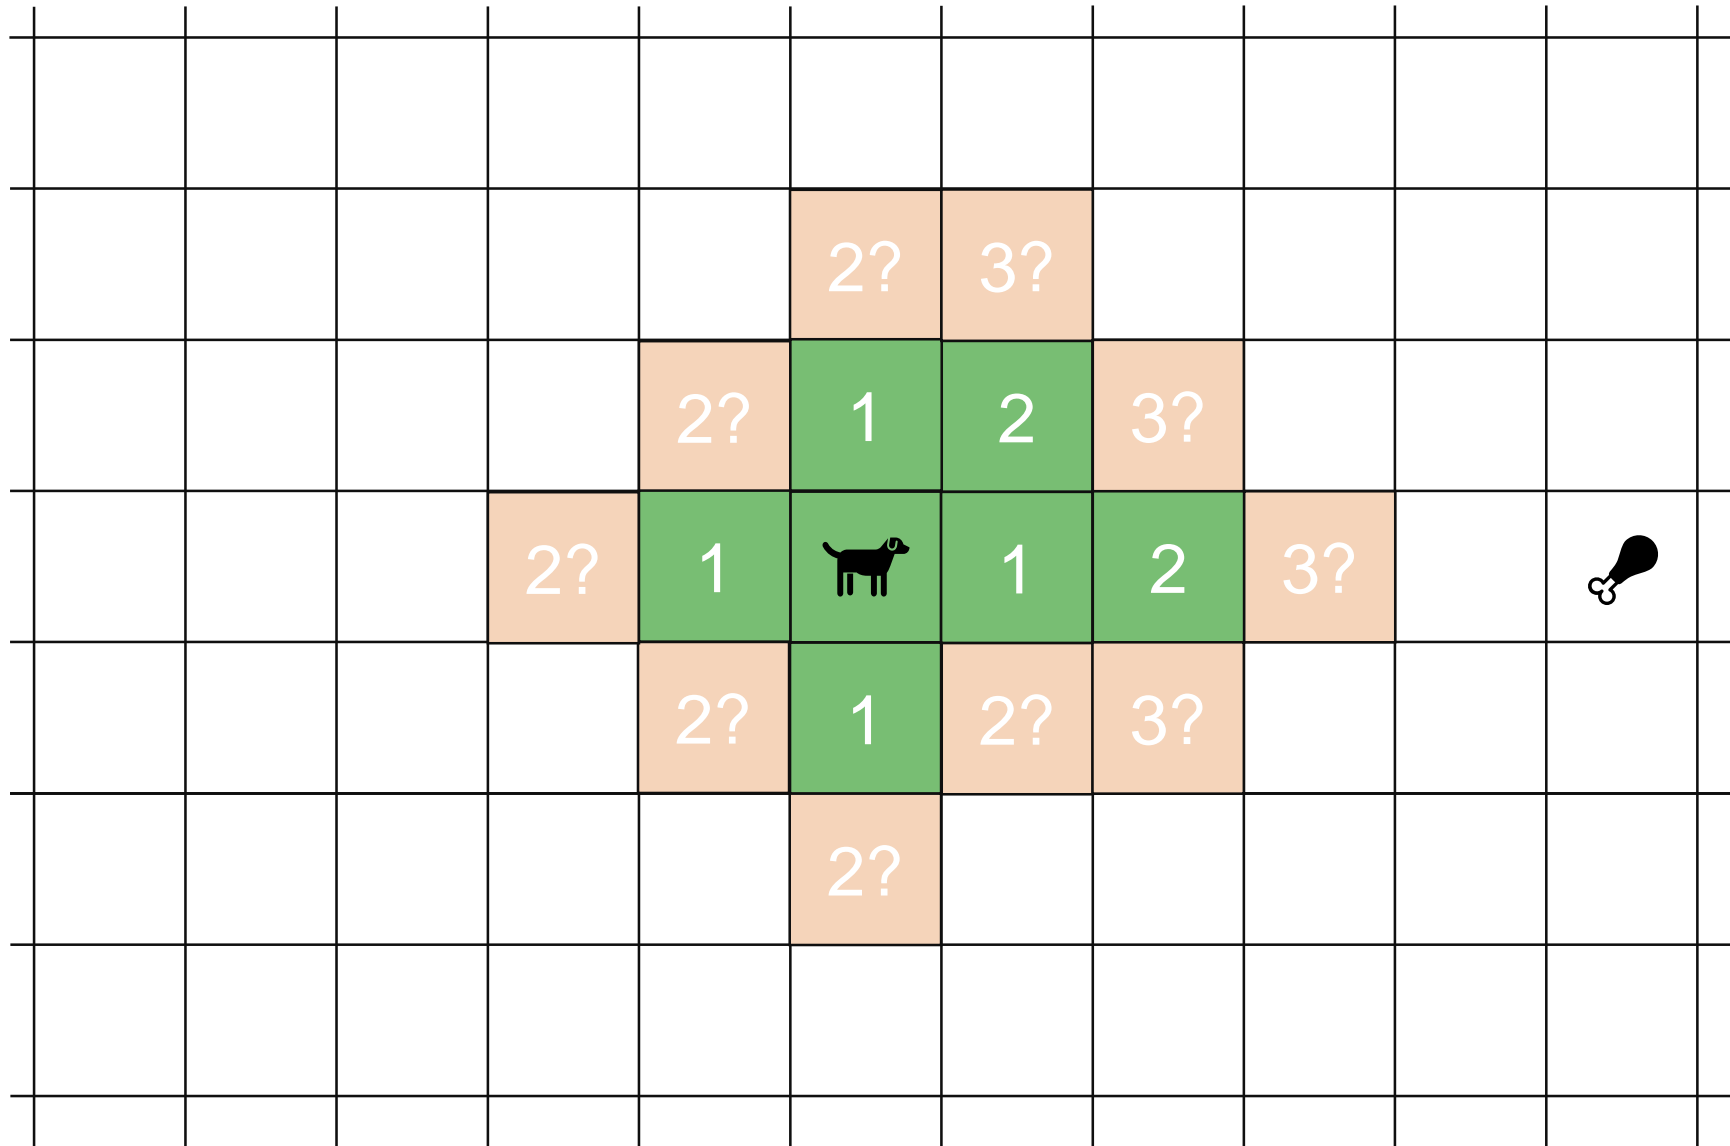
Dijkstra where each edge has cost 1



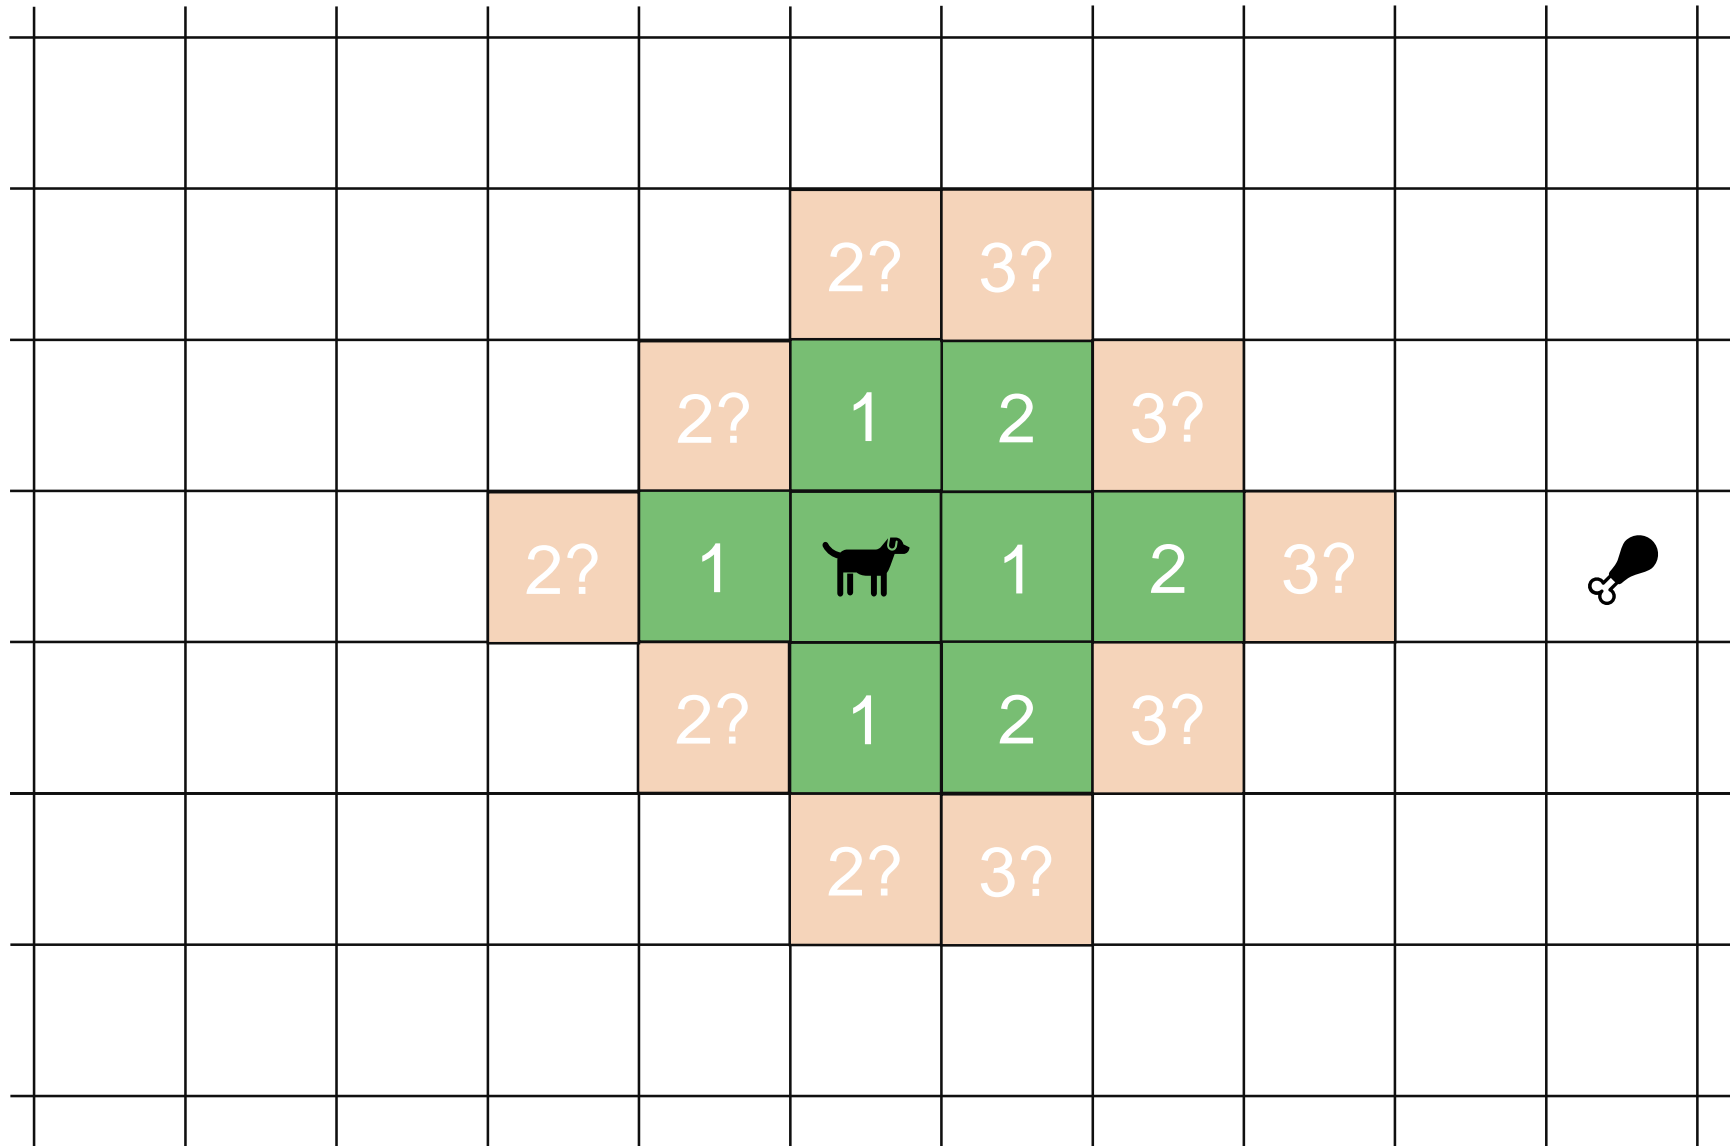
Dijkstra where each edge has cost 1



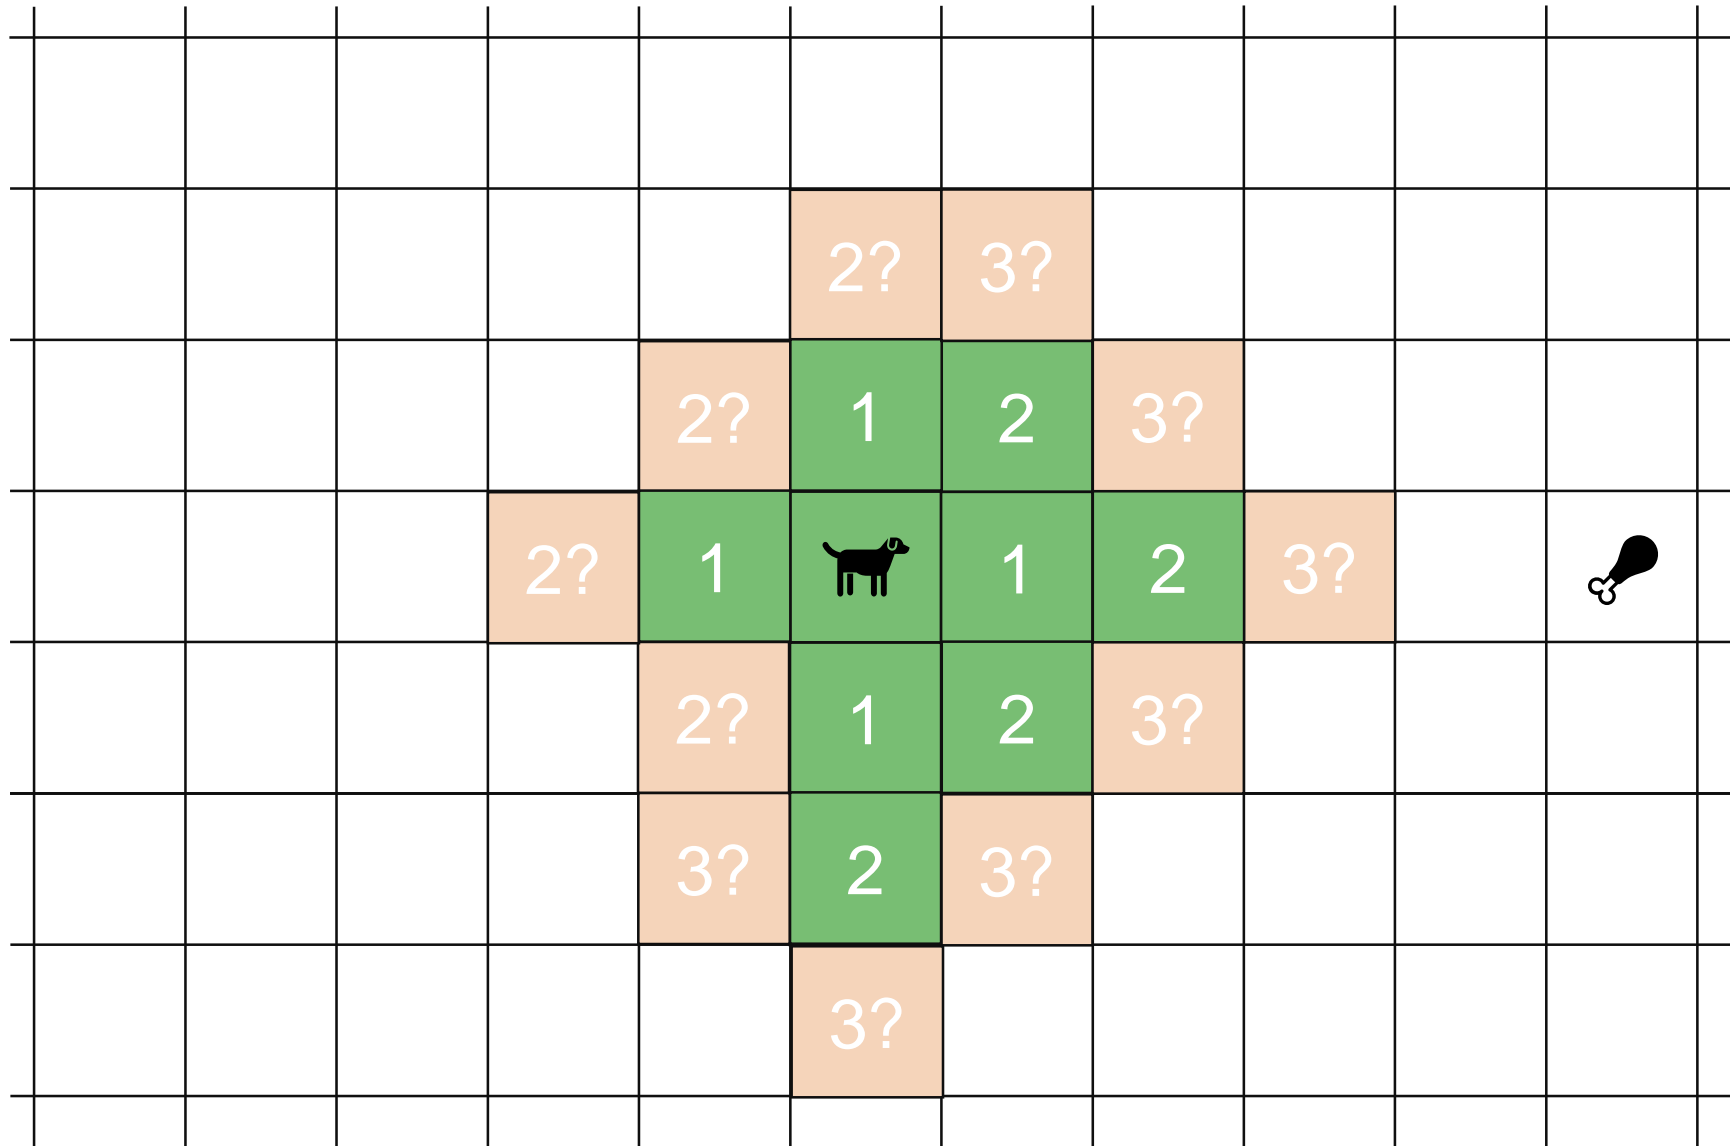
Dijkstra where each edge has cost 1



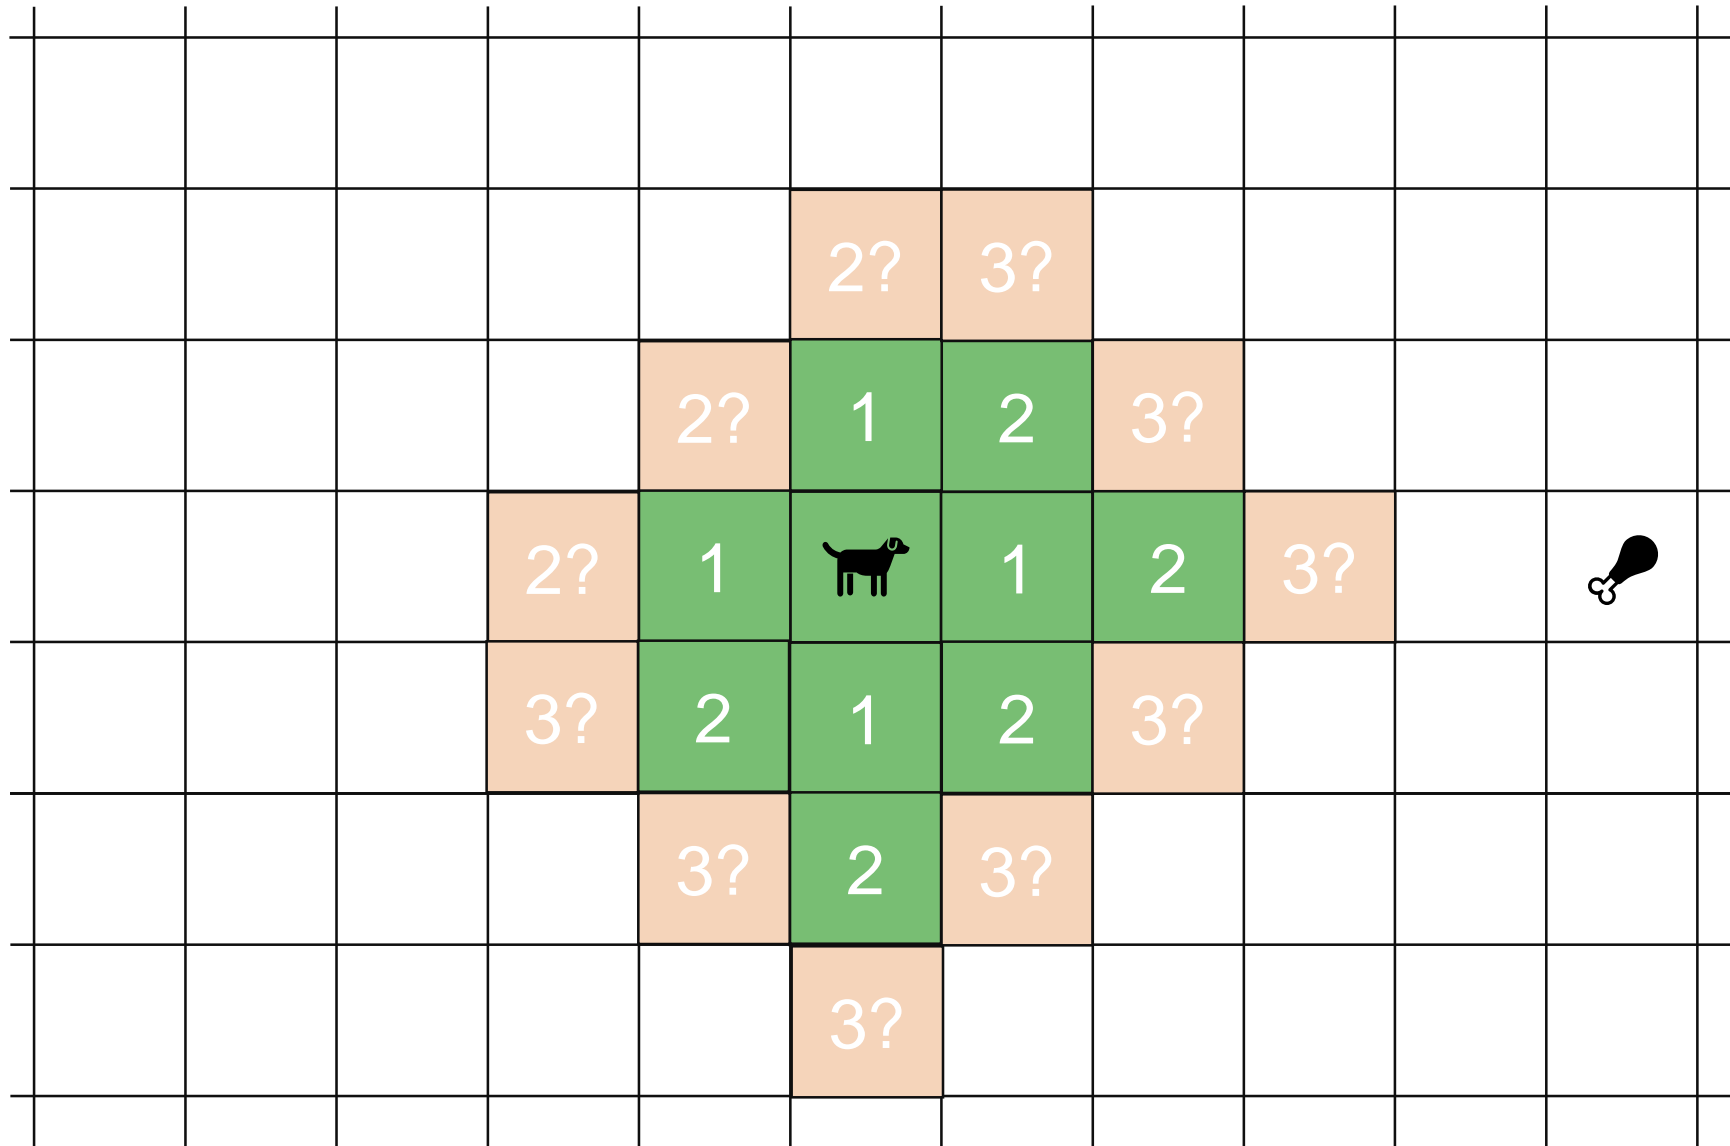
Dijkstra where each edge has cost 1



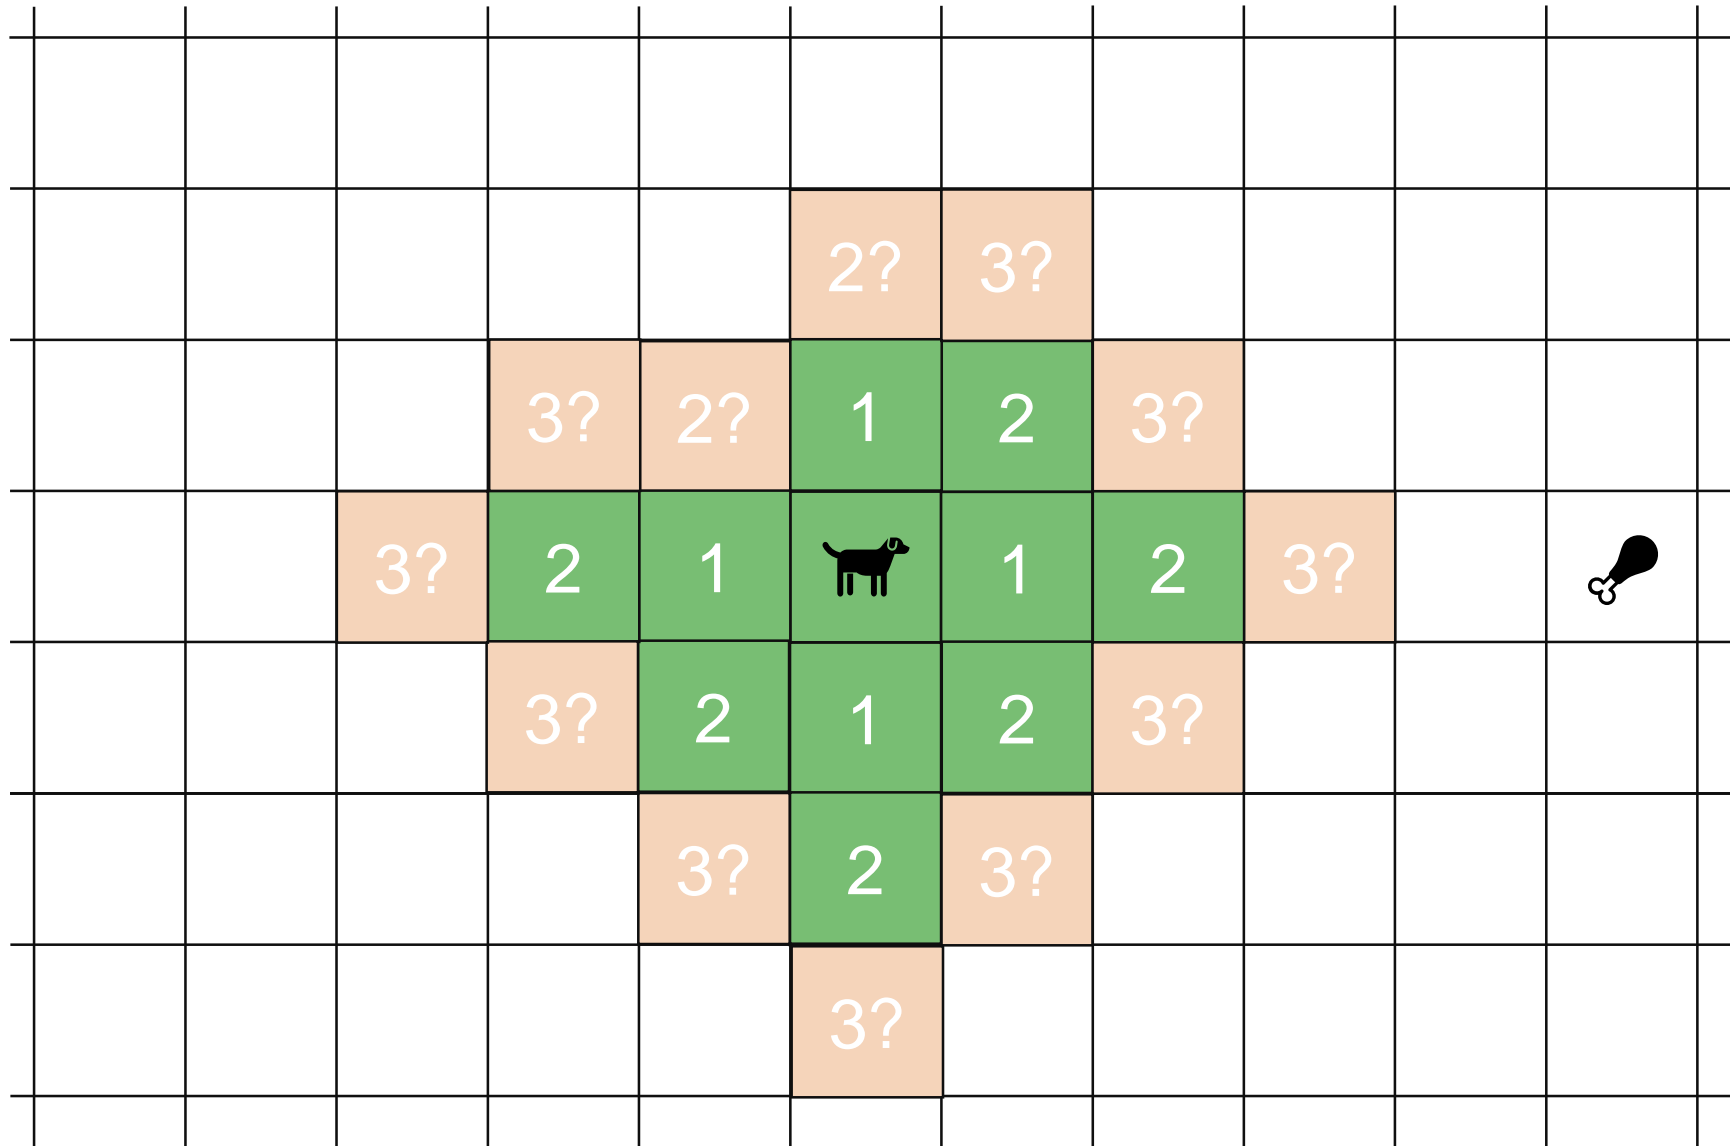
Dijkstra where each edge has cost 1



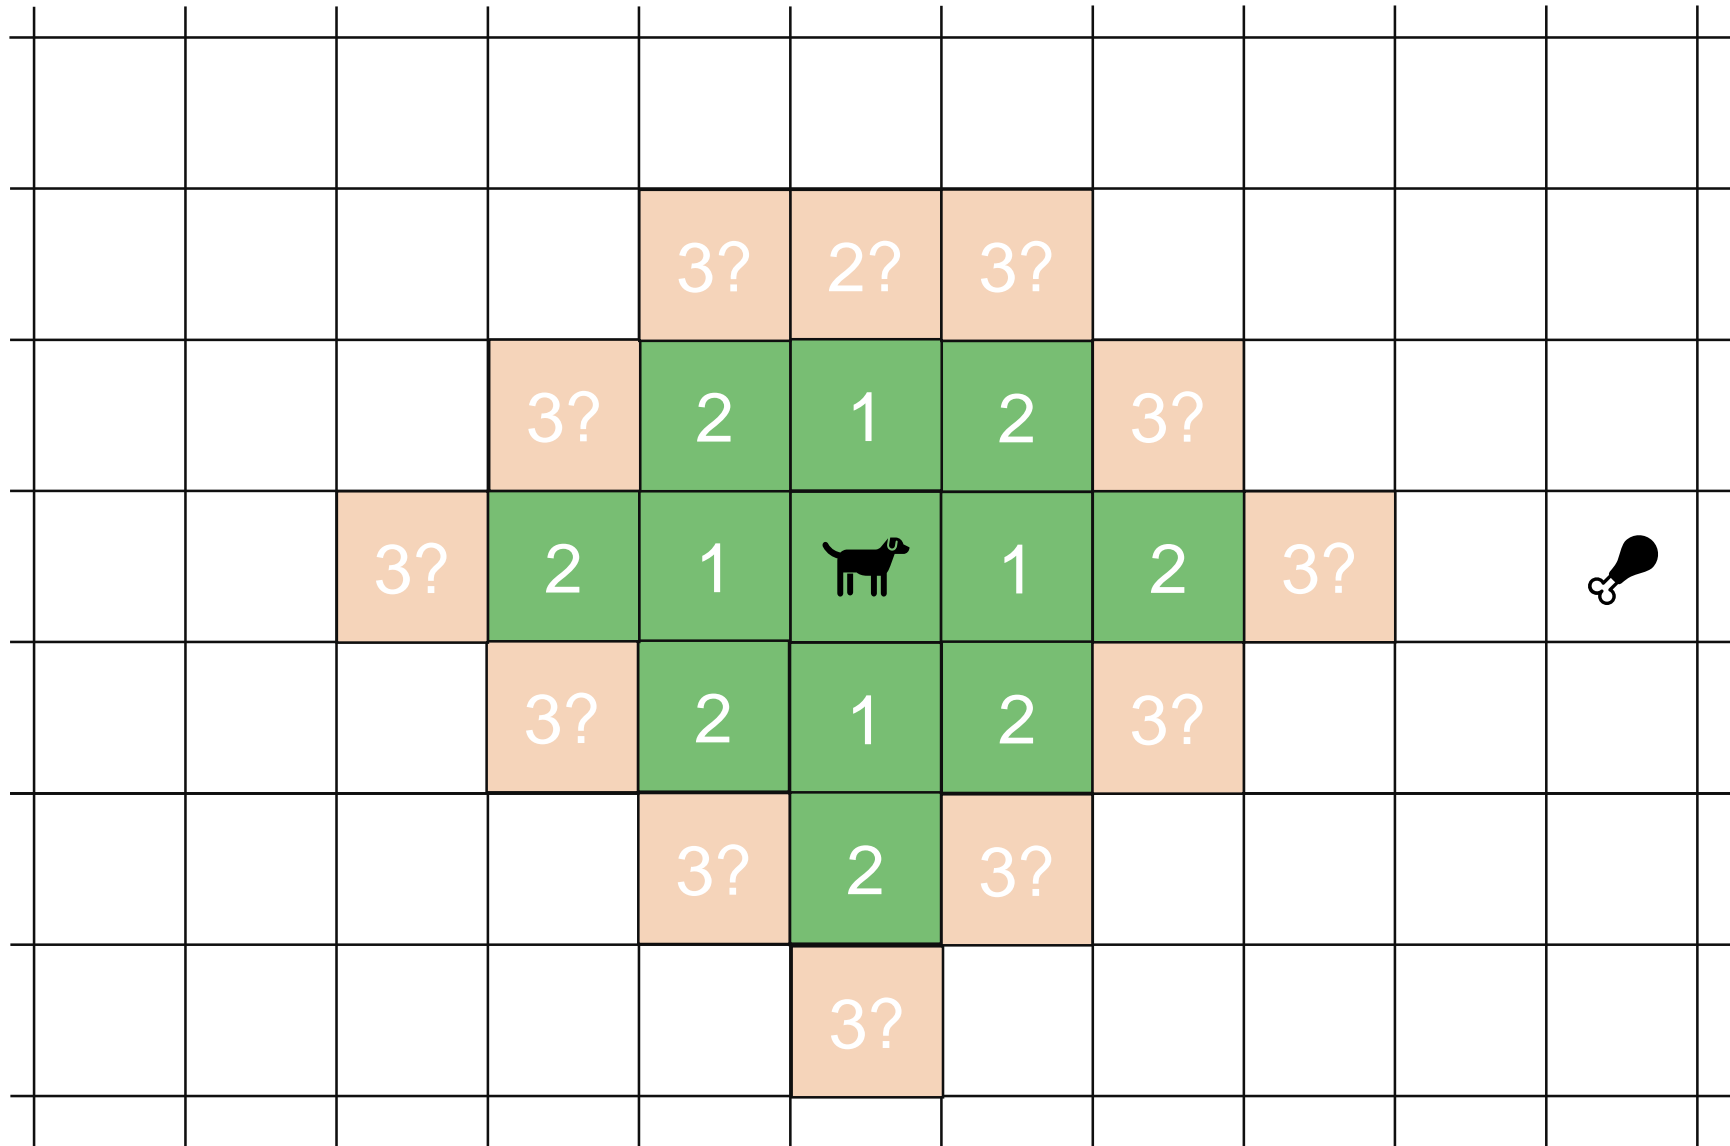
Dijkstra where each edge has cost 1



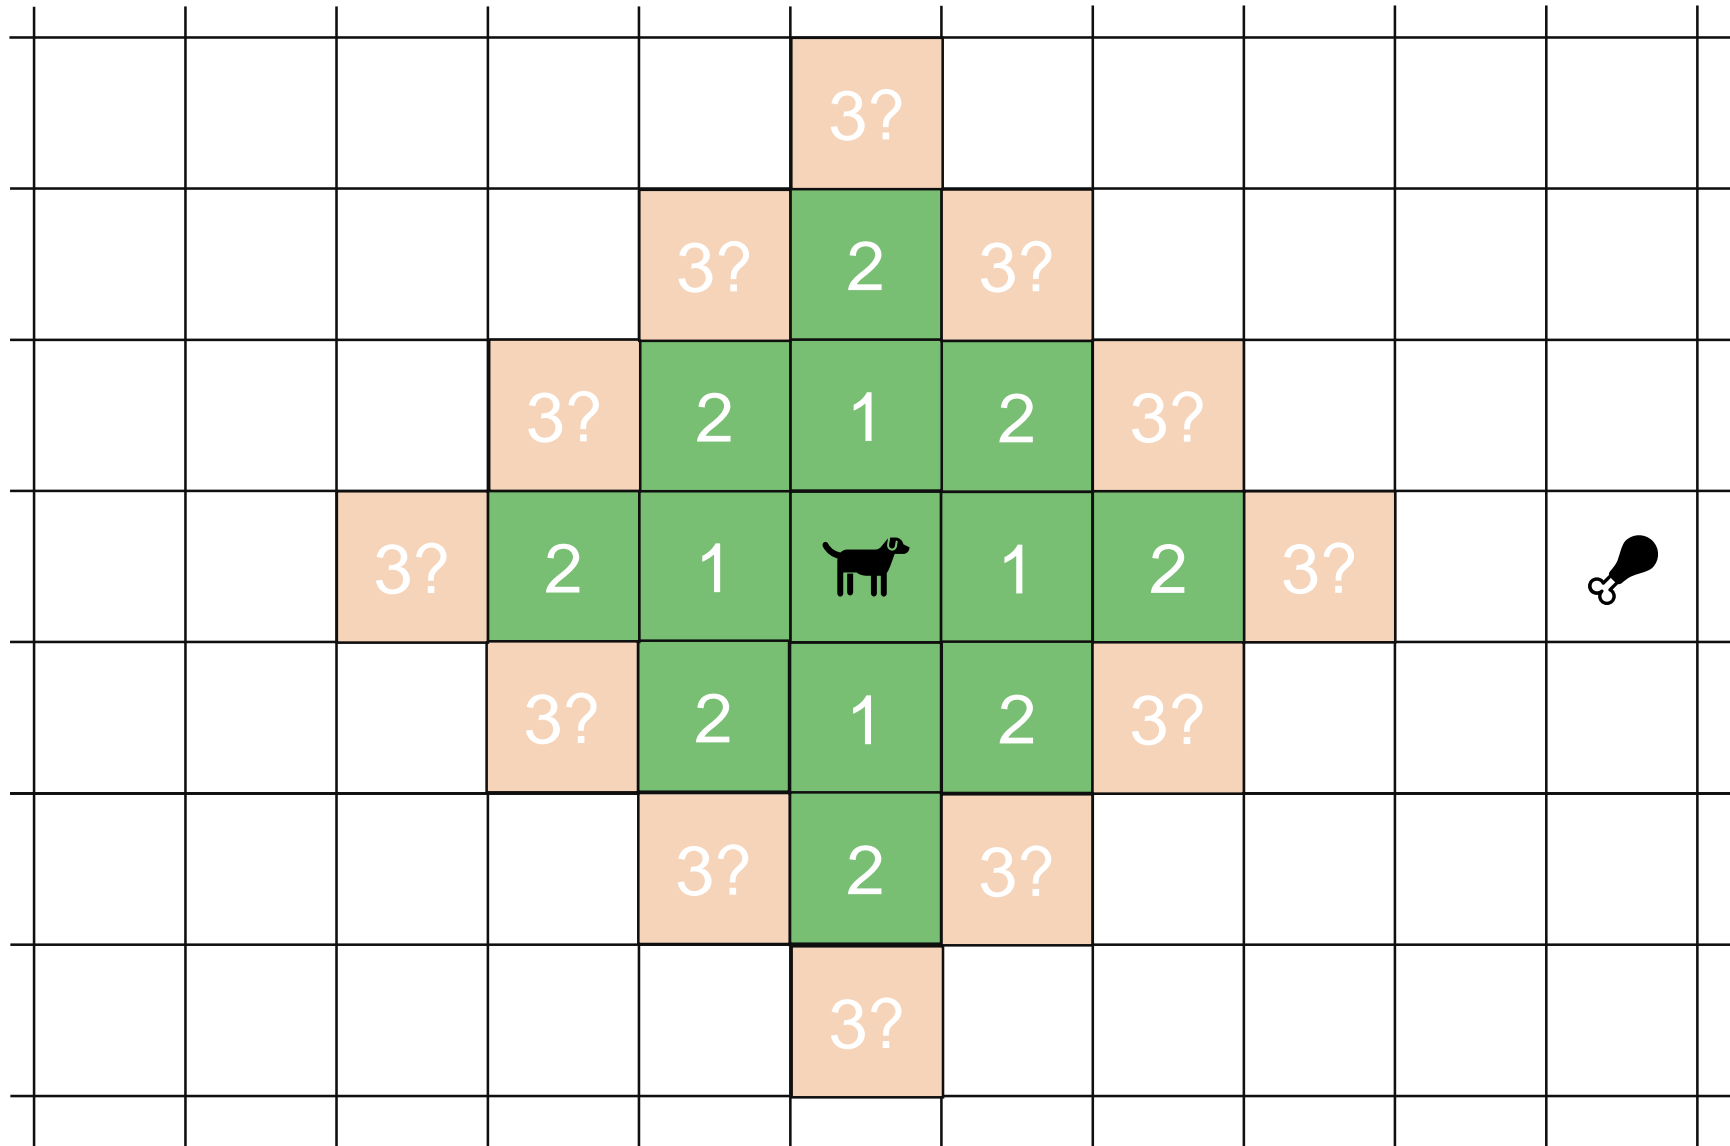
Dijkstra where each edge has cost 1



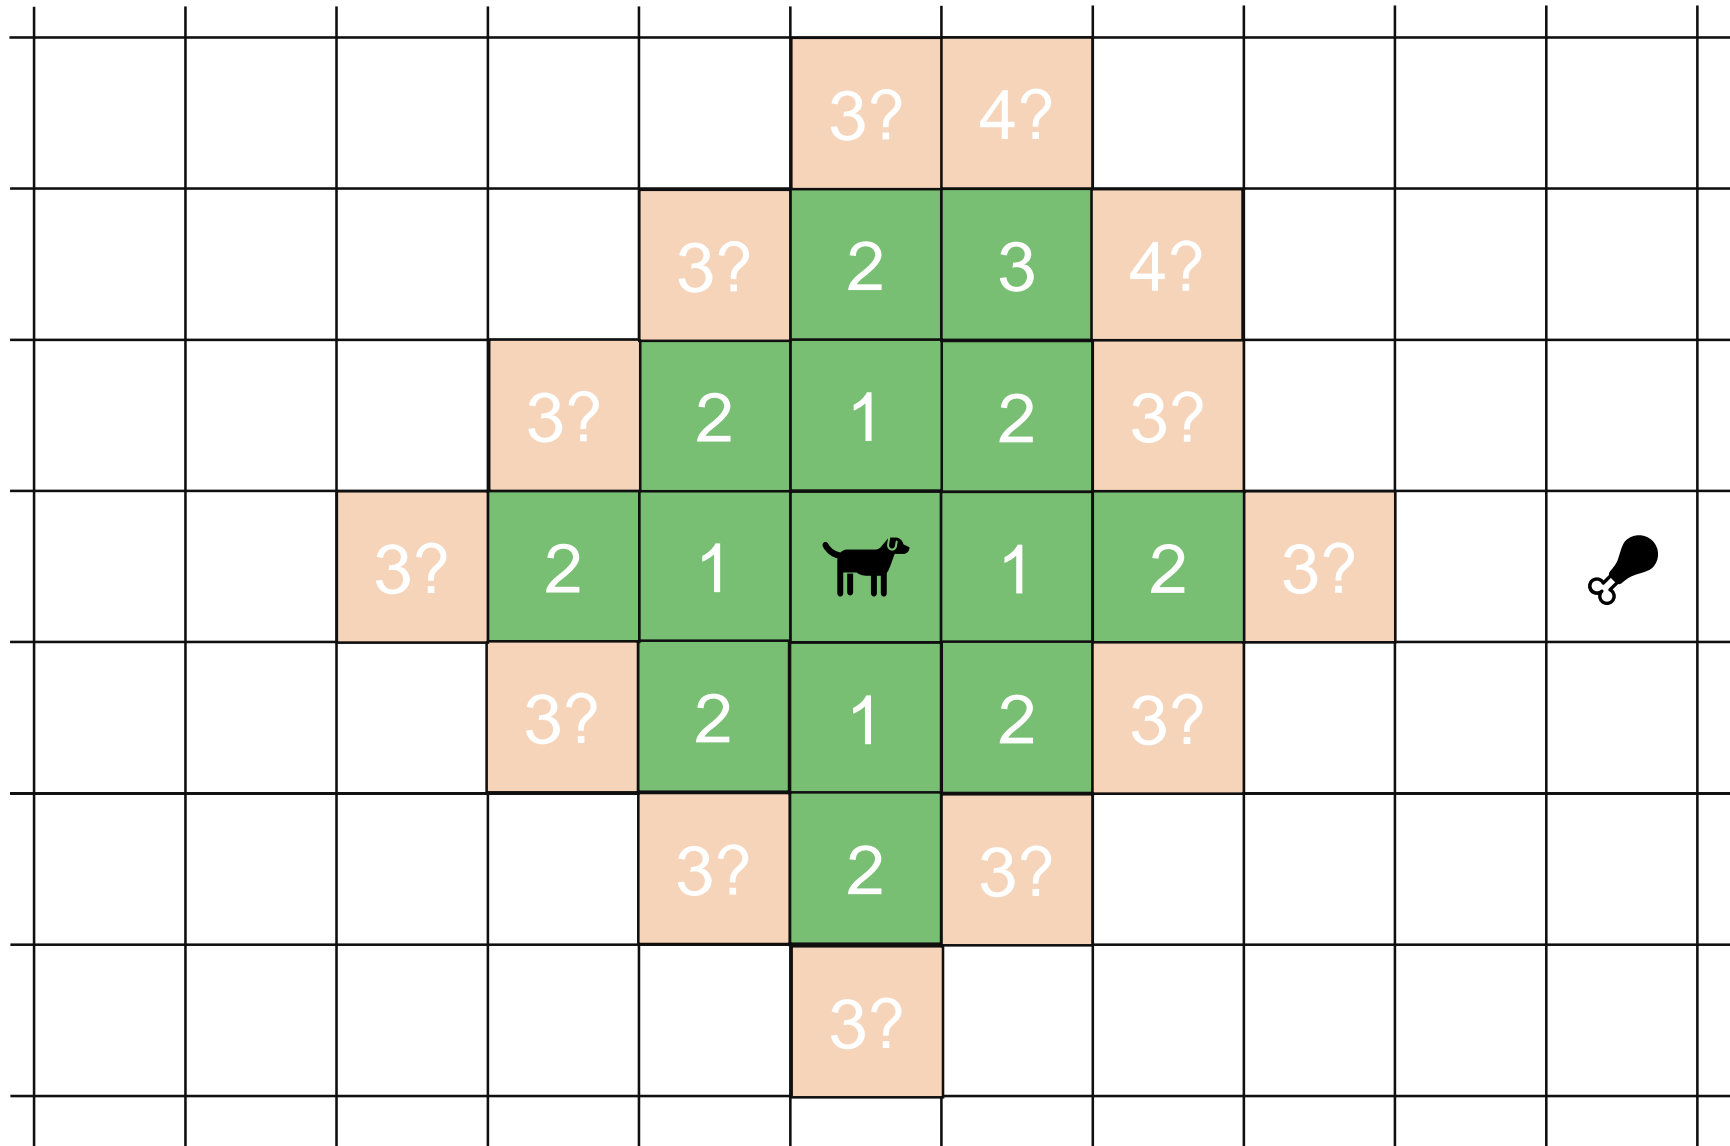
Dijkstra where each edge has cost 1



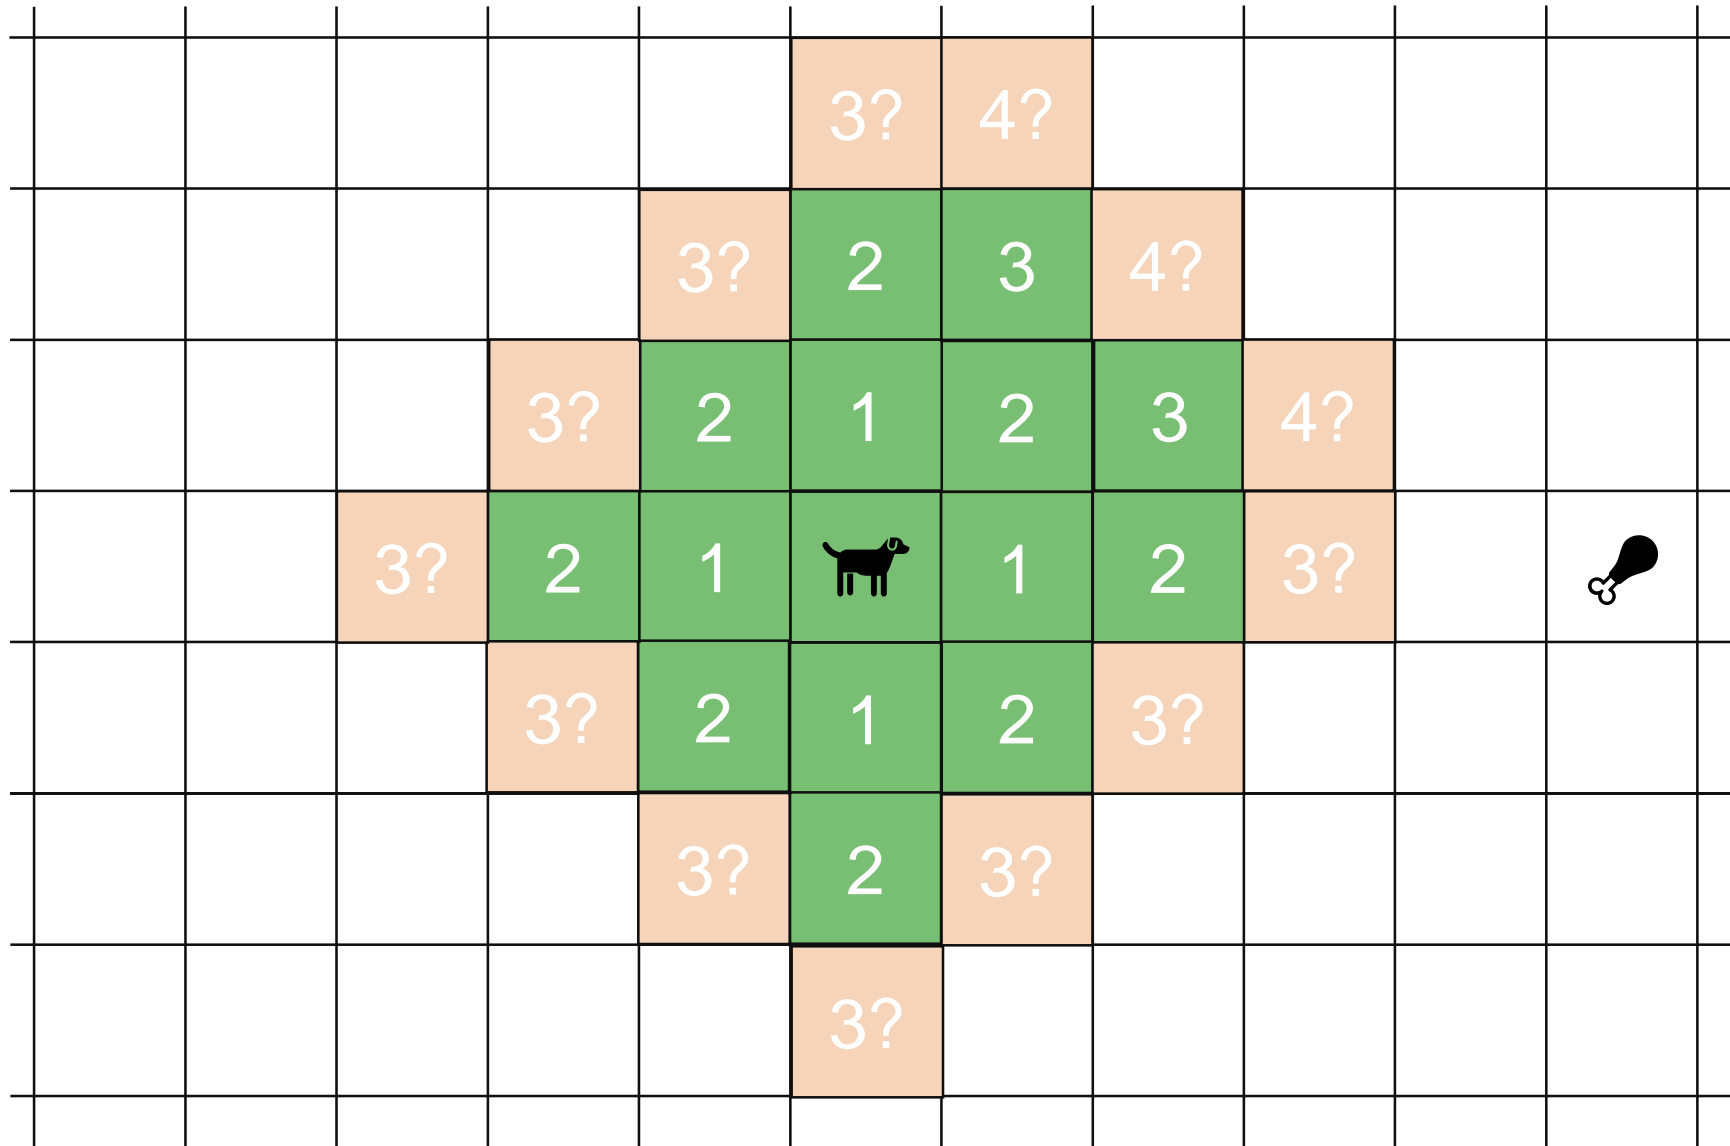
Dijkstra where each edge has cost 1



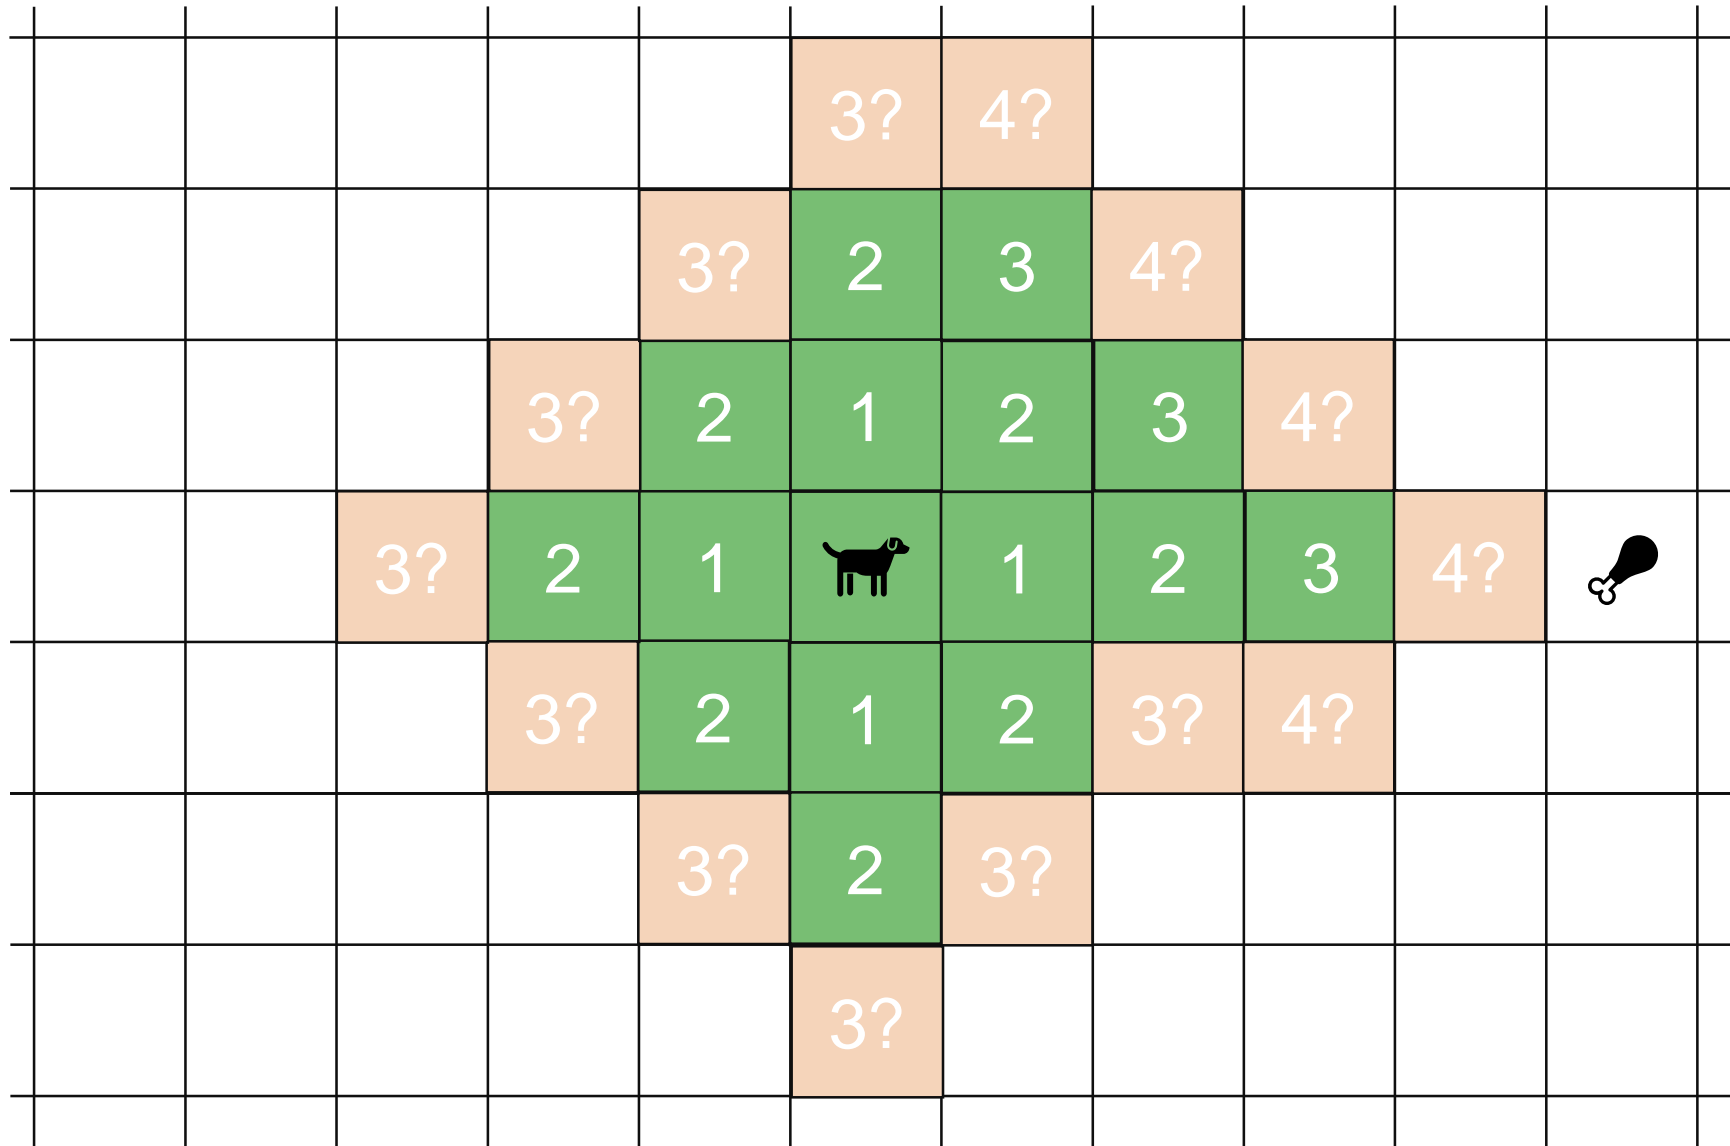
Dijkstra where each edge has cost 1



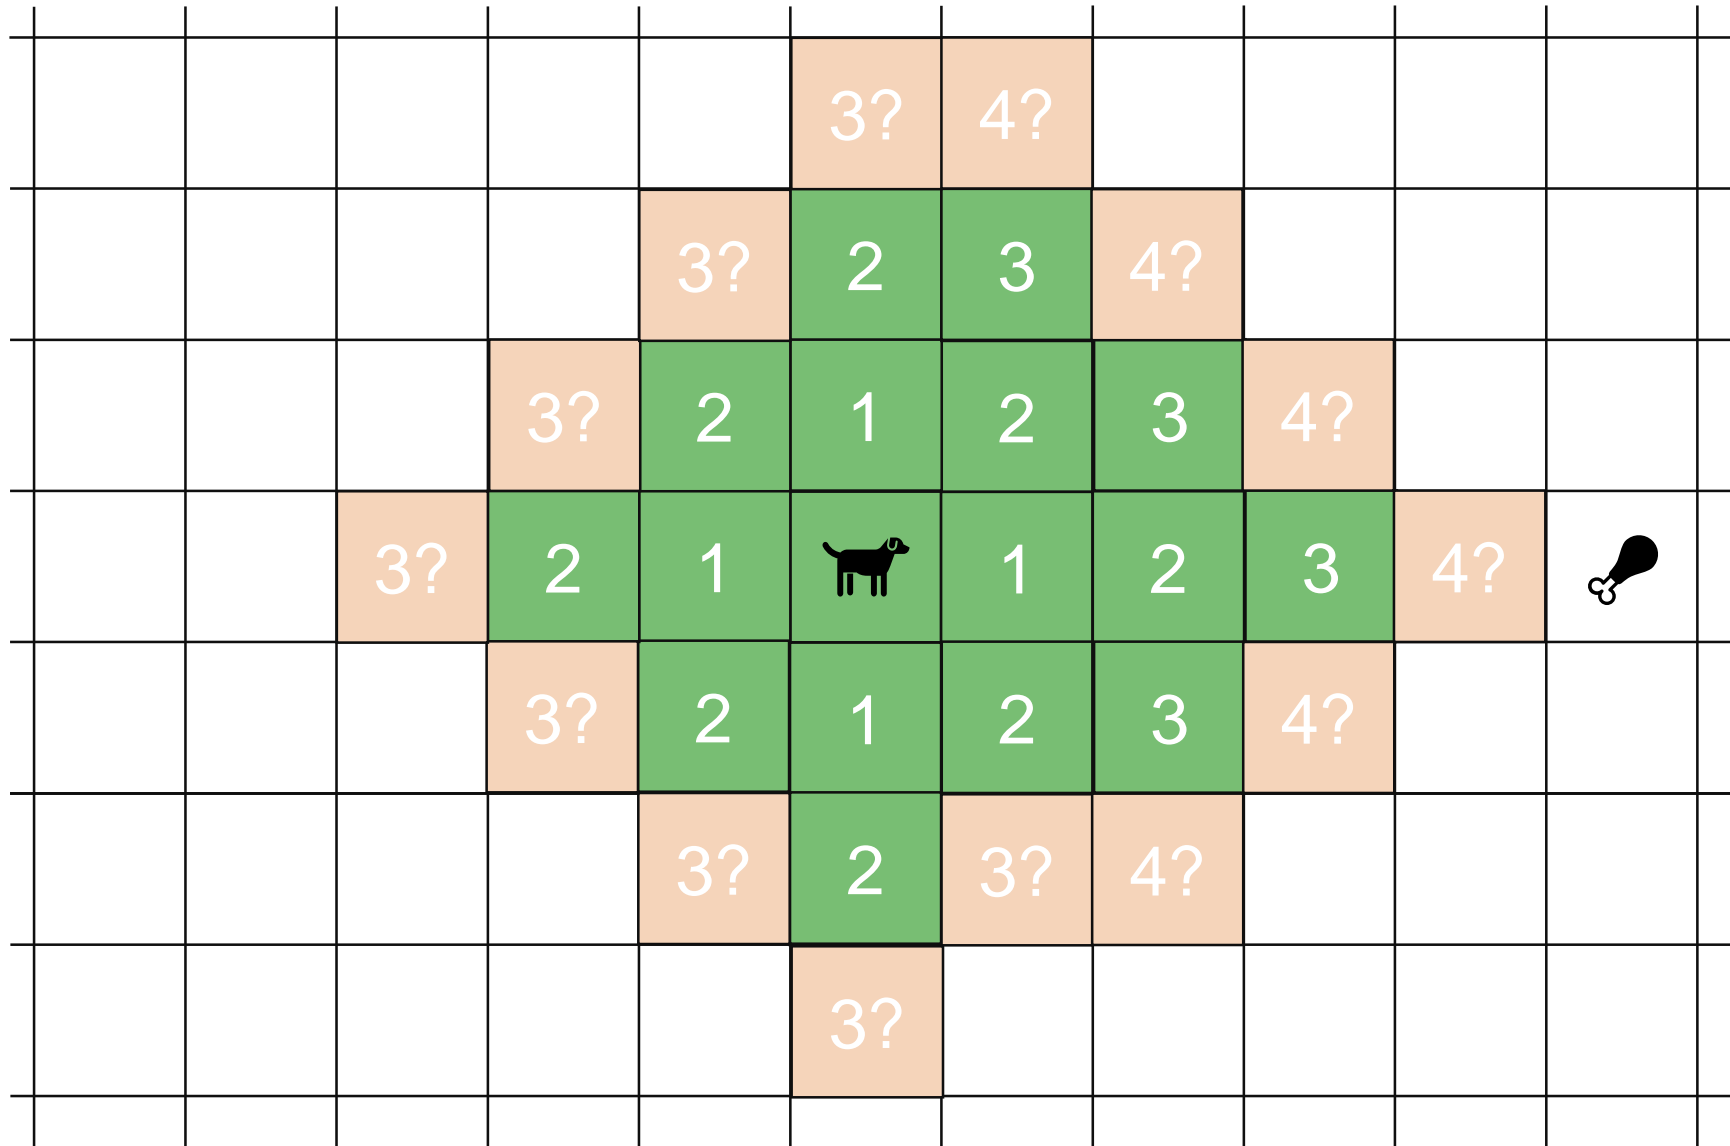
Dijkstra where each edge has cost 1



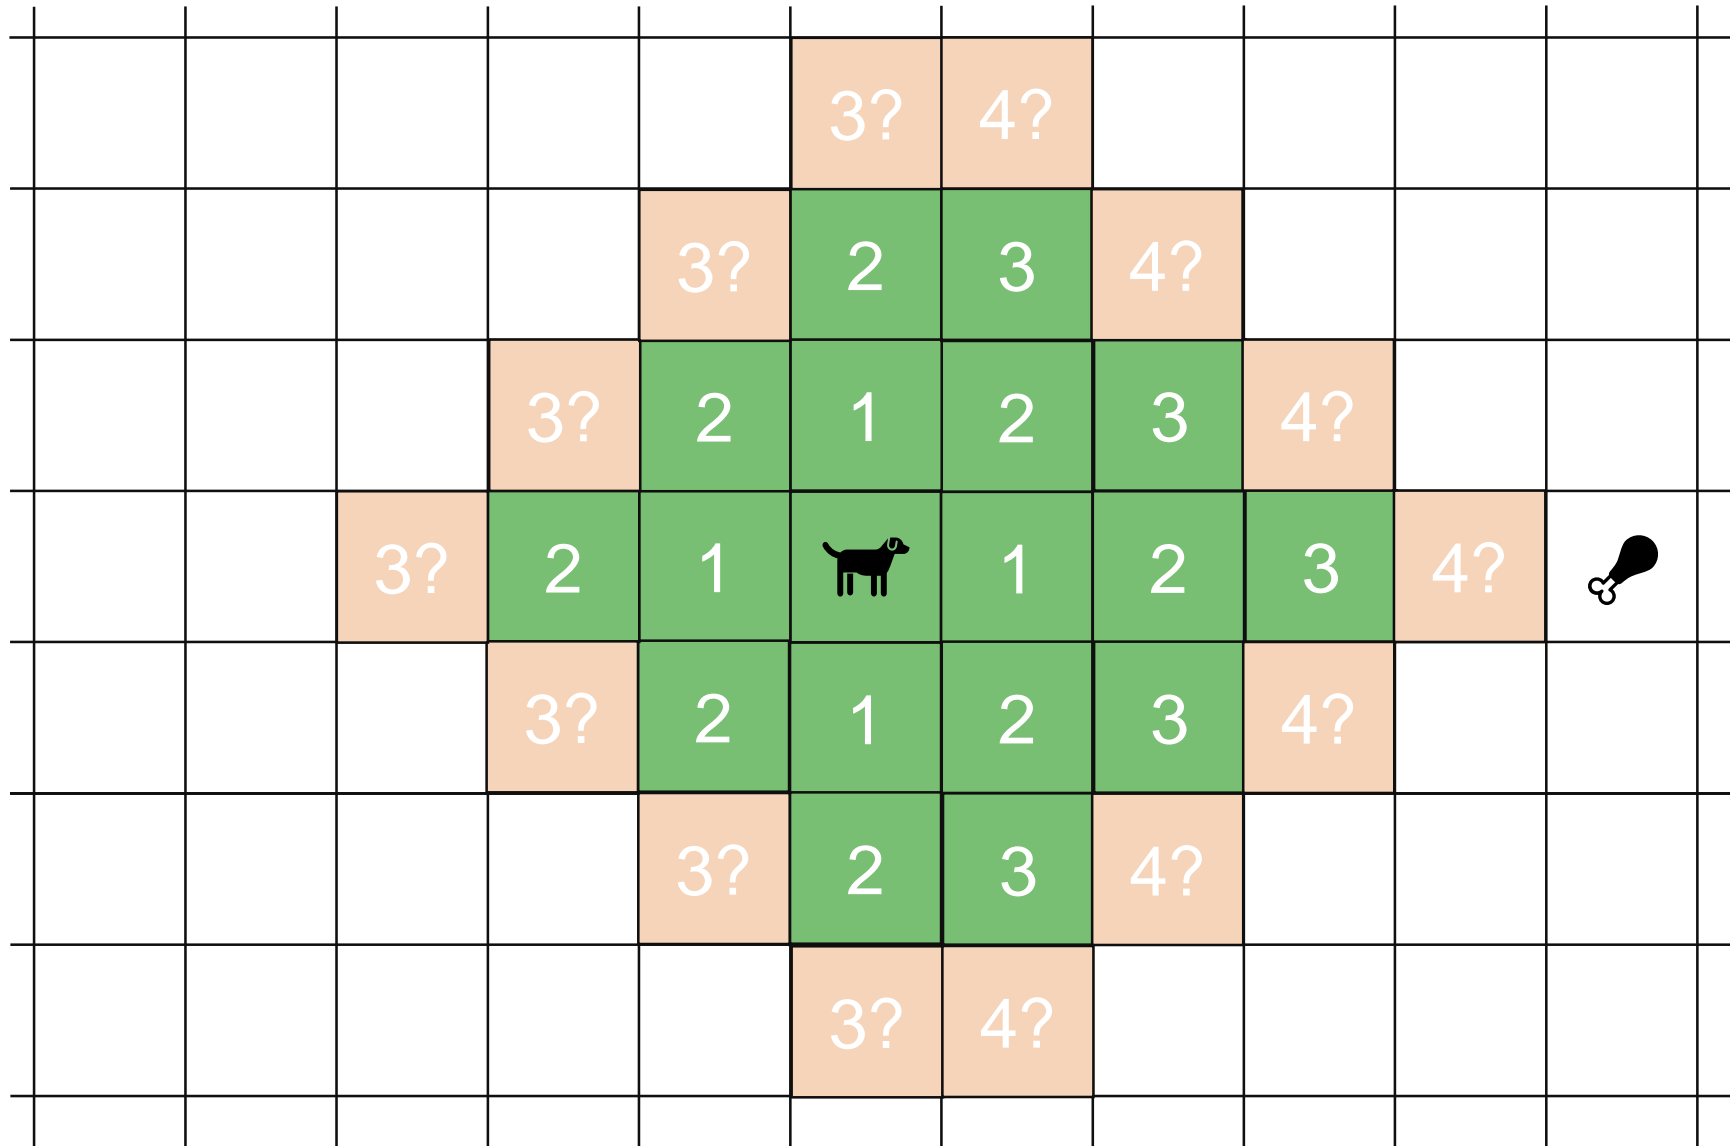
Dijkstra where each edge has cost 1



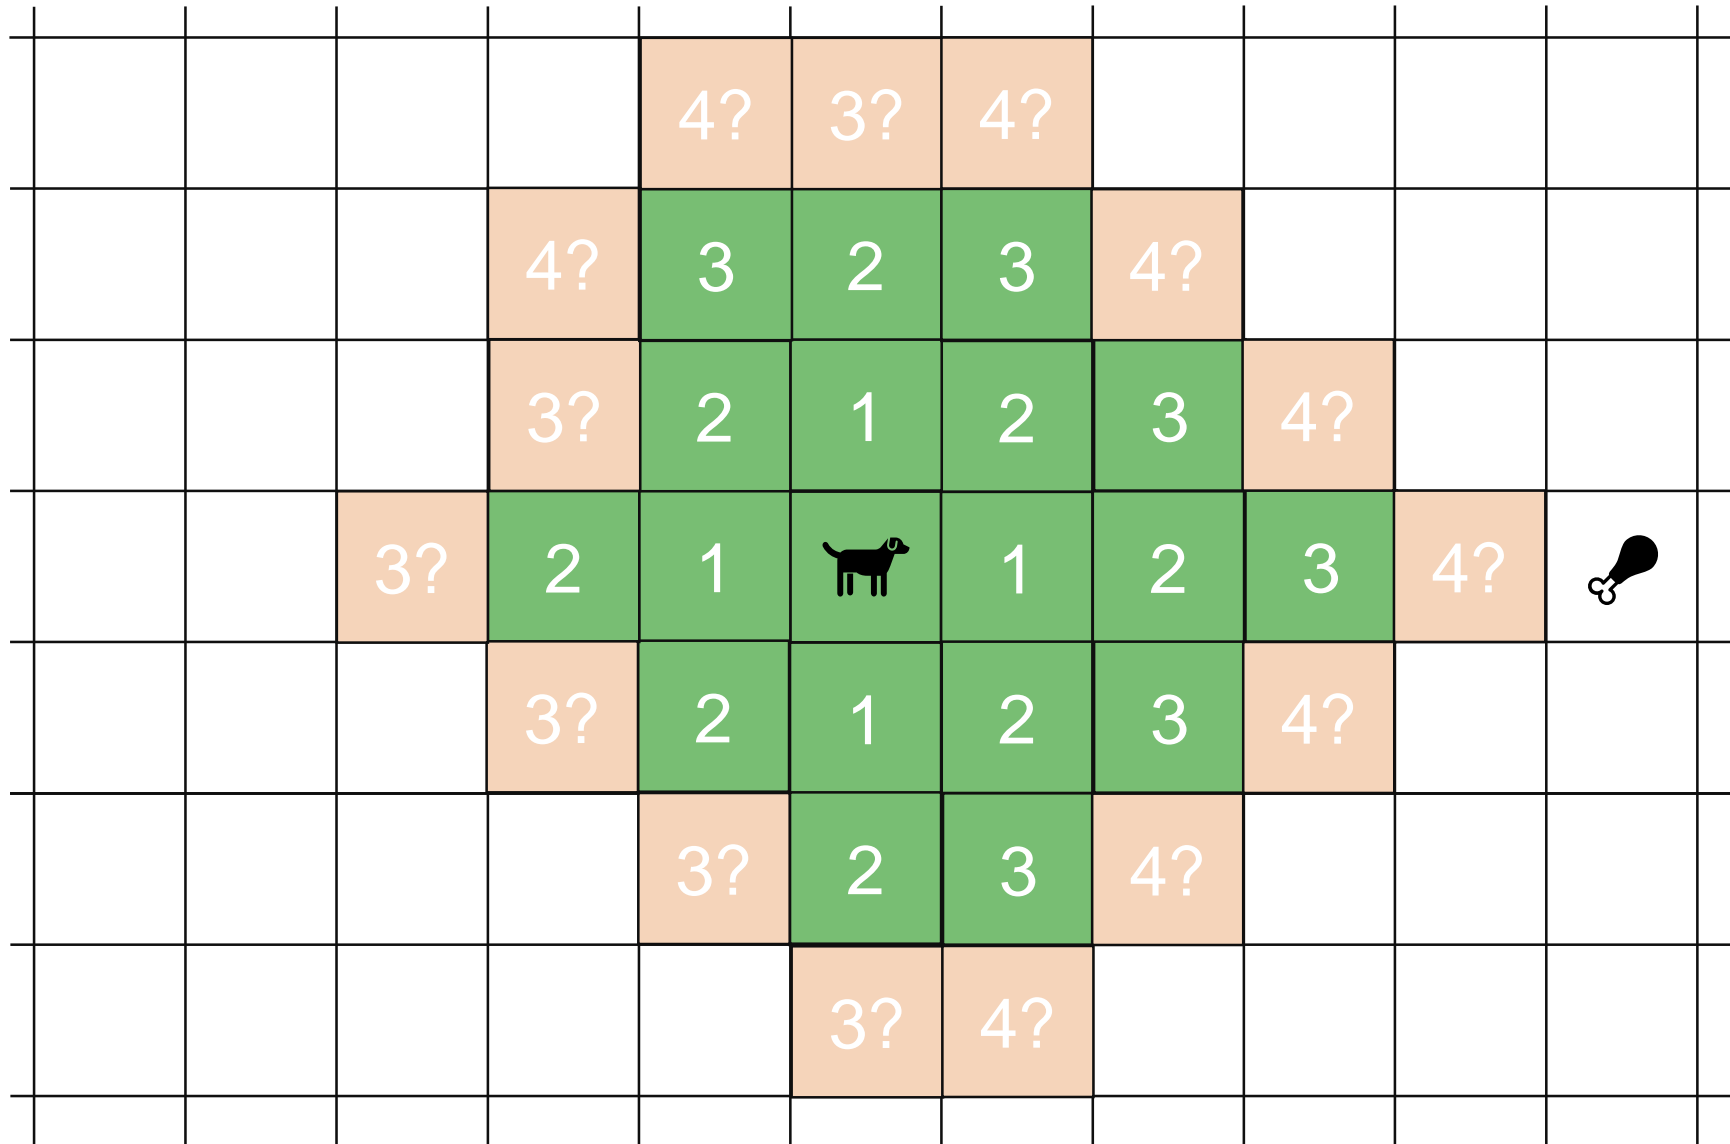
Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1



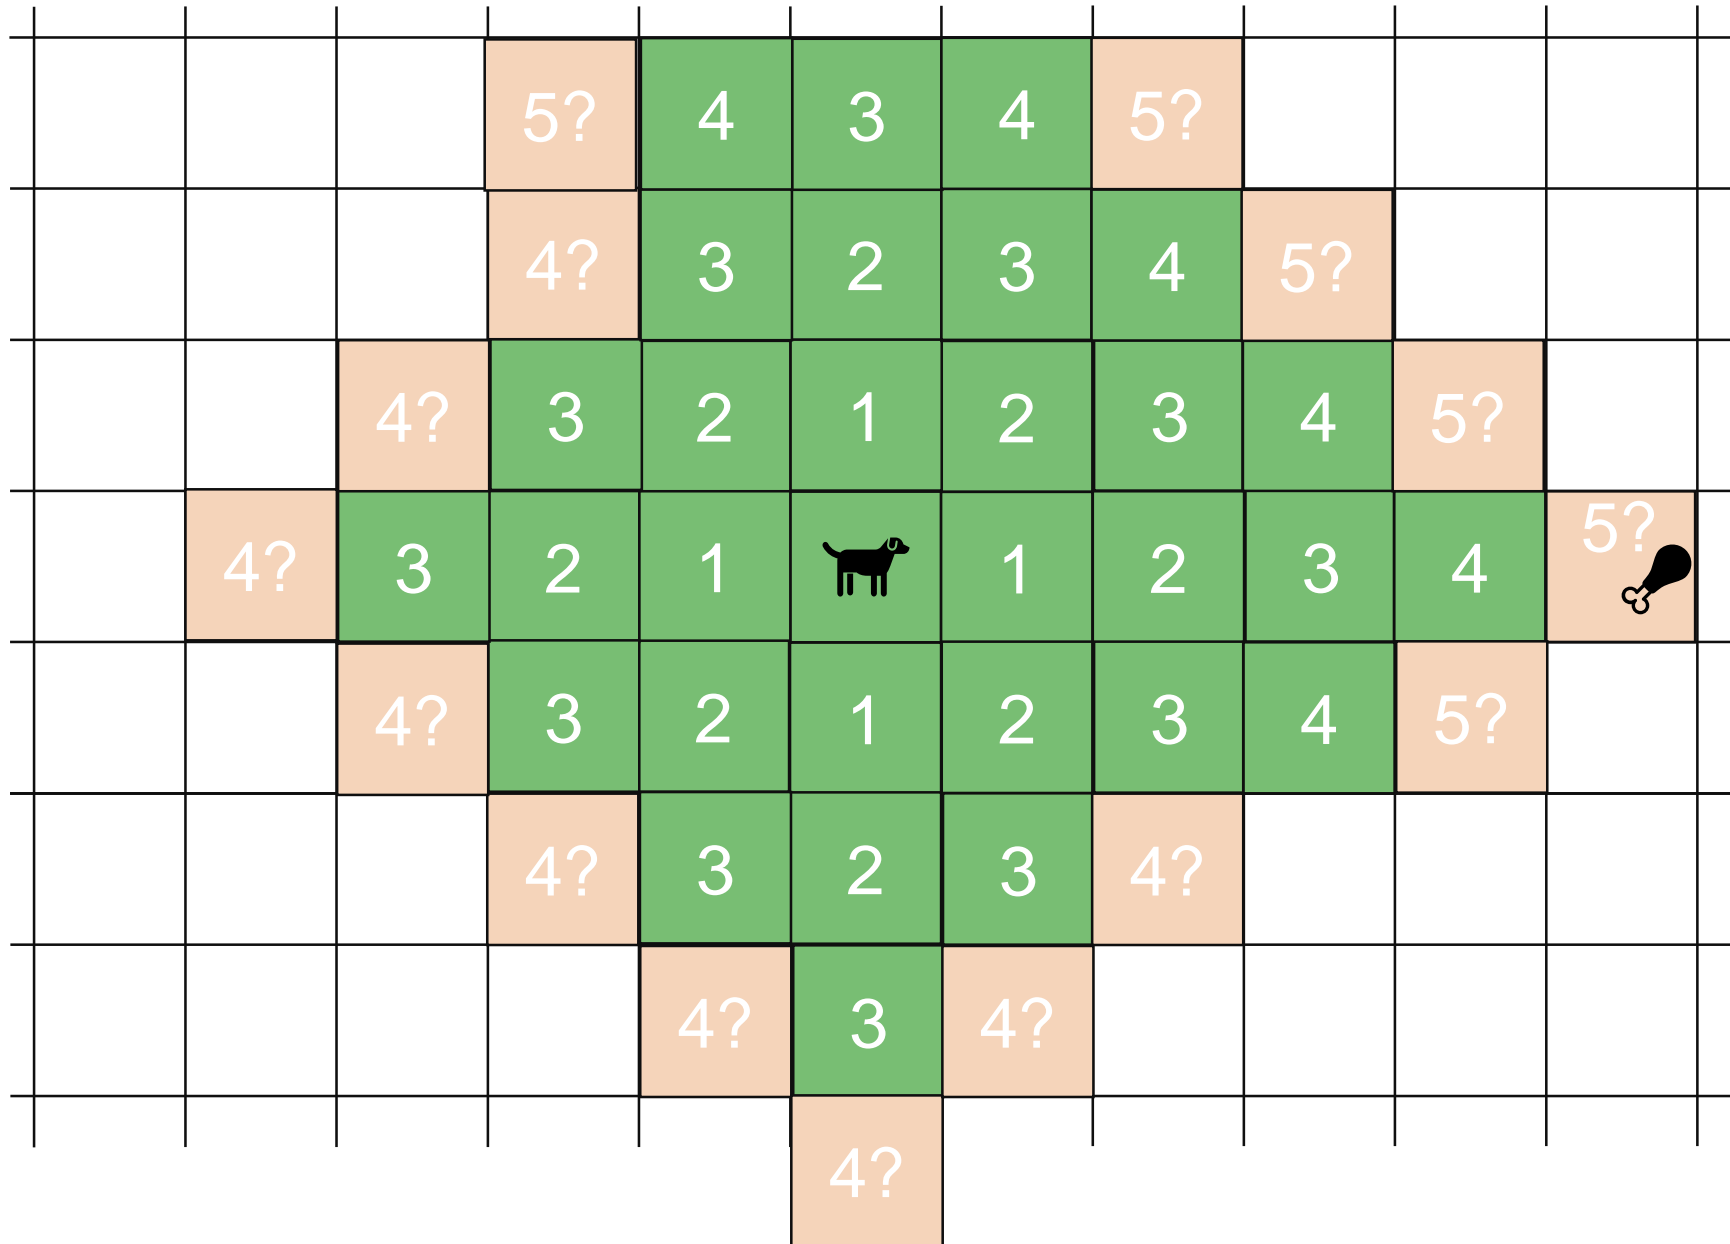
Dijkstra where each edge has cost 1



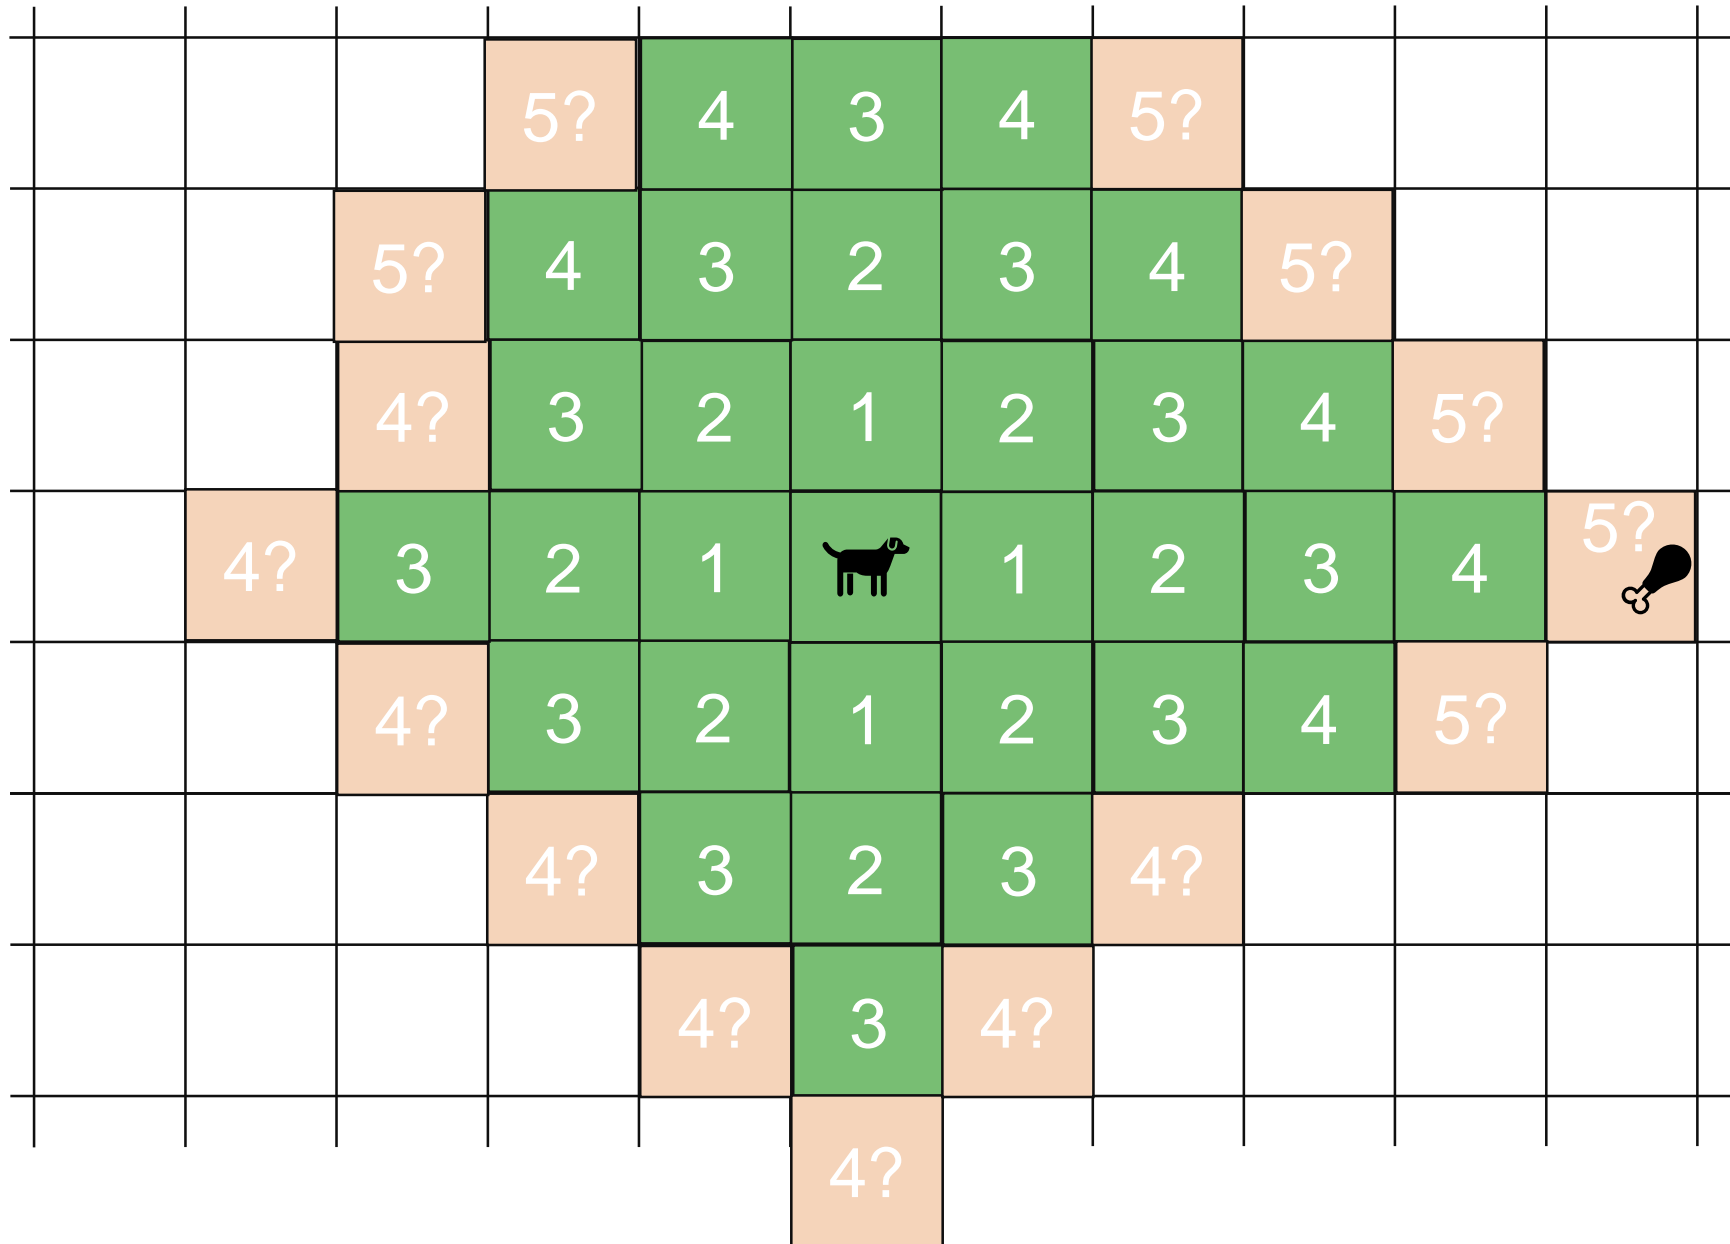
Dijkstra where each edge has cost 1



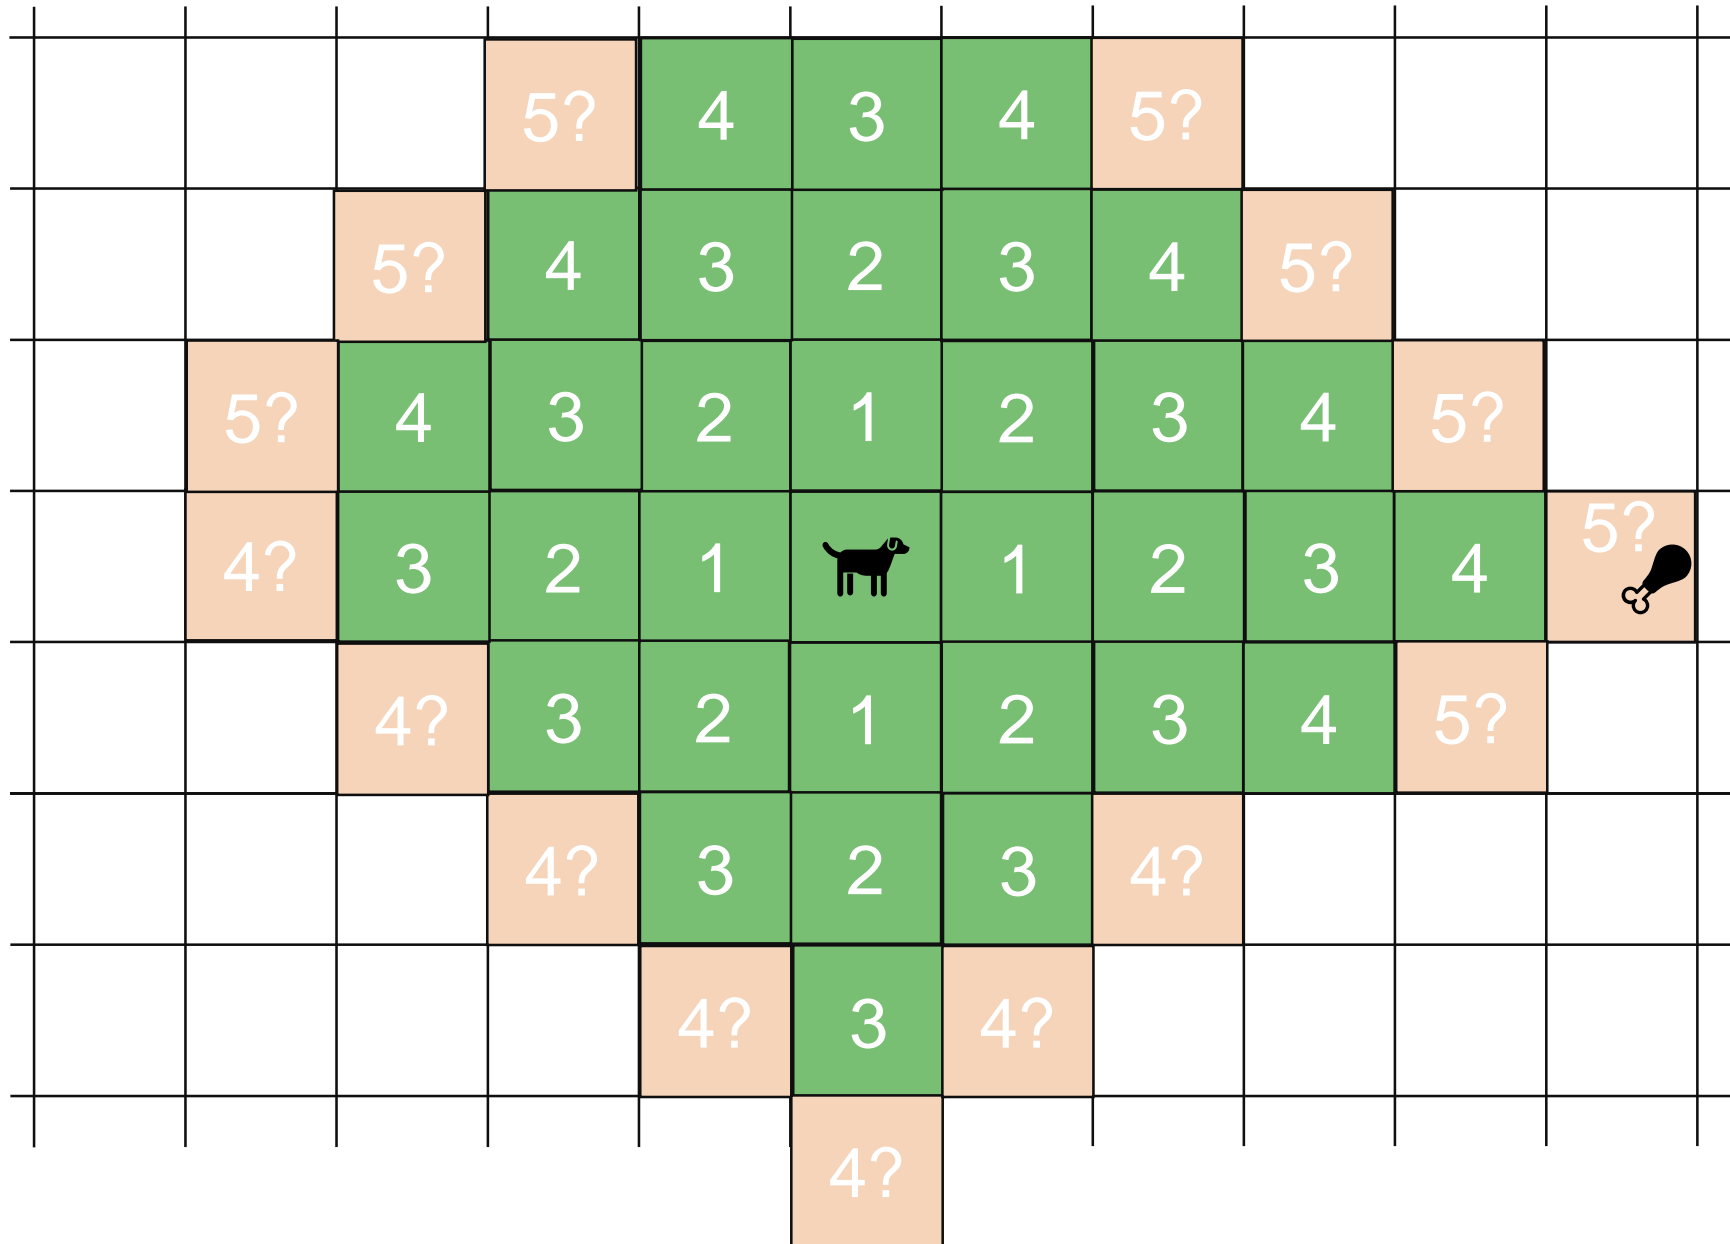
Dijkstra where each edge has cost 1



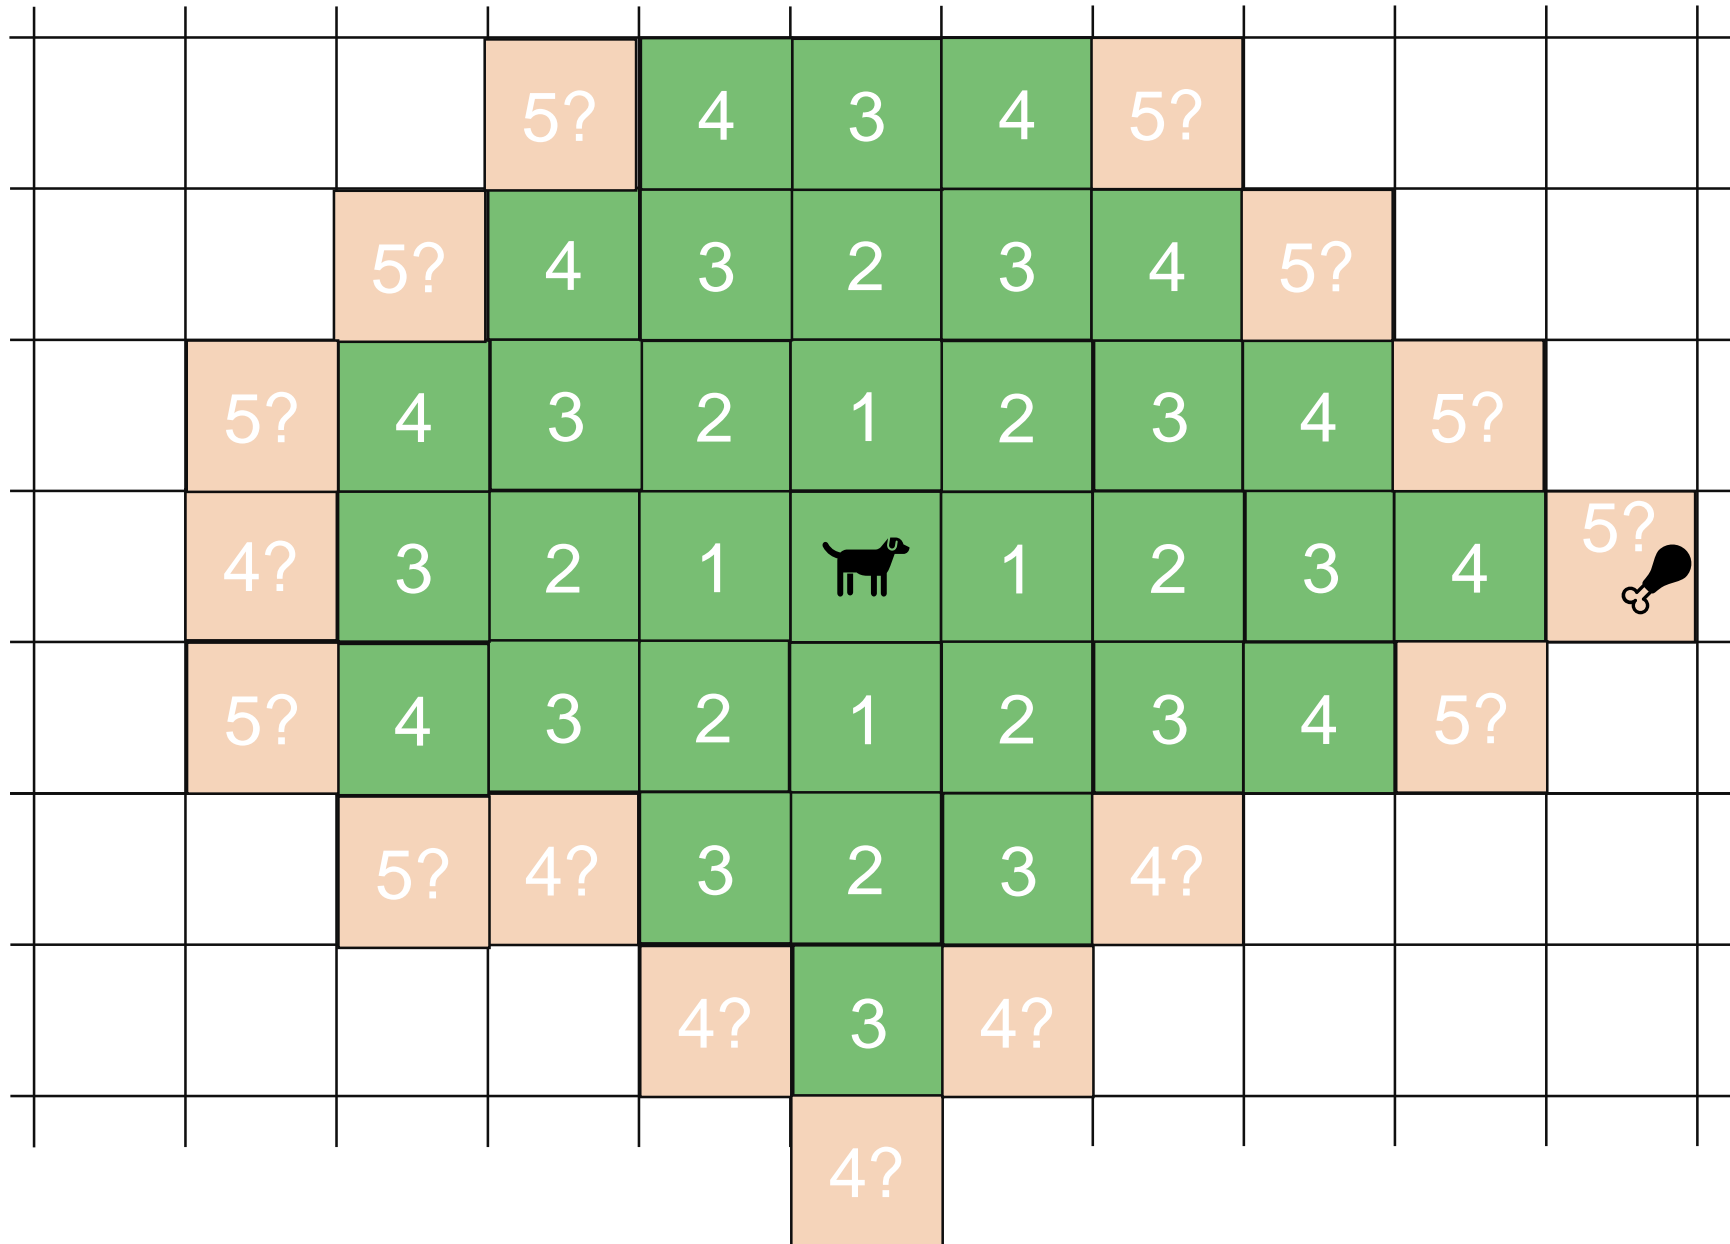
Dijkstra where each edge has cost 1



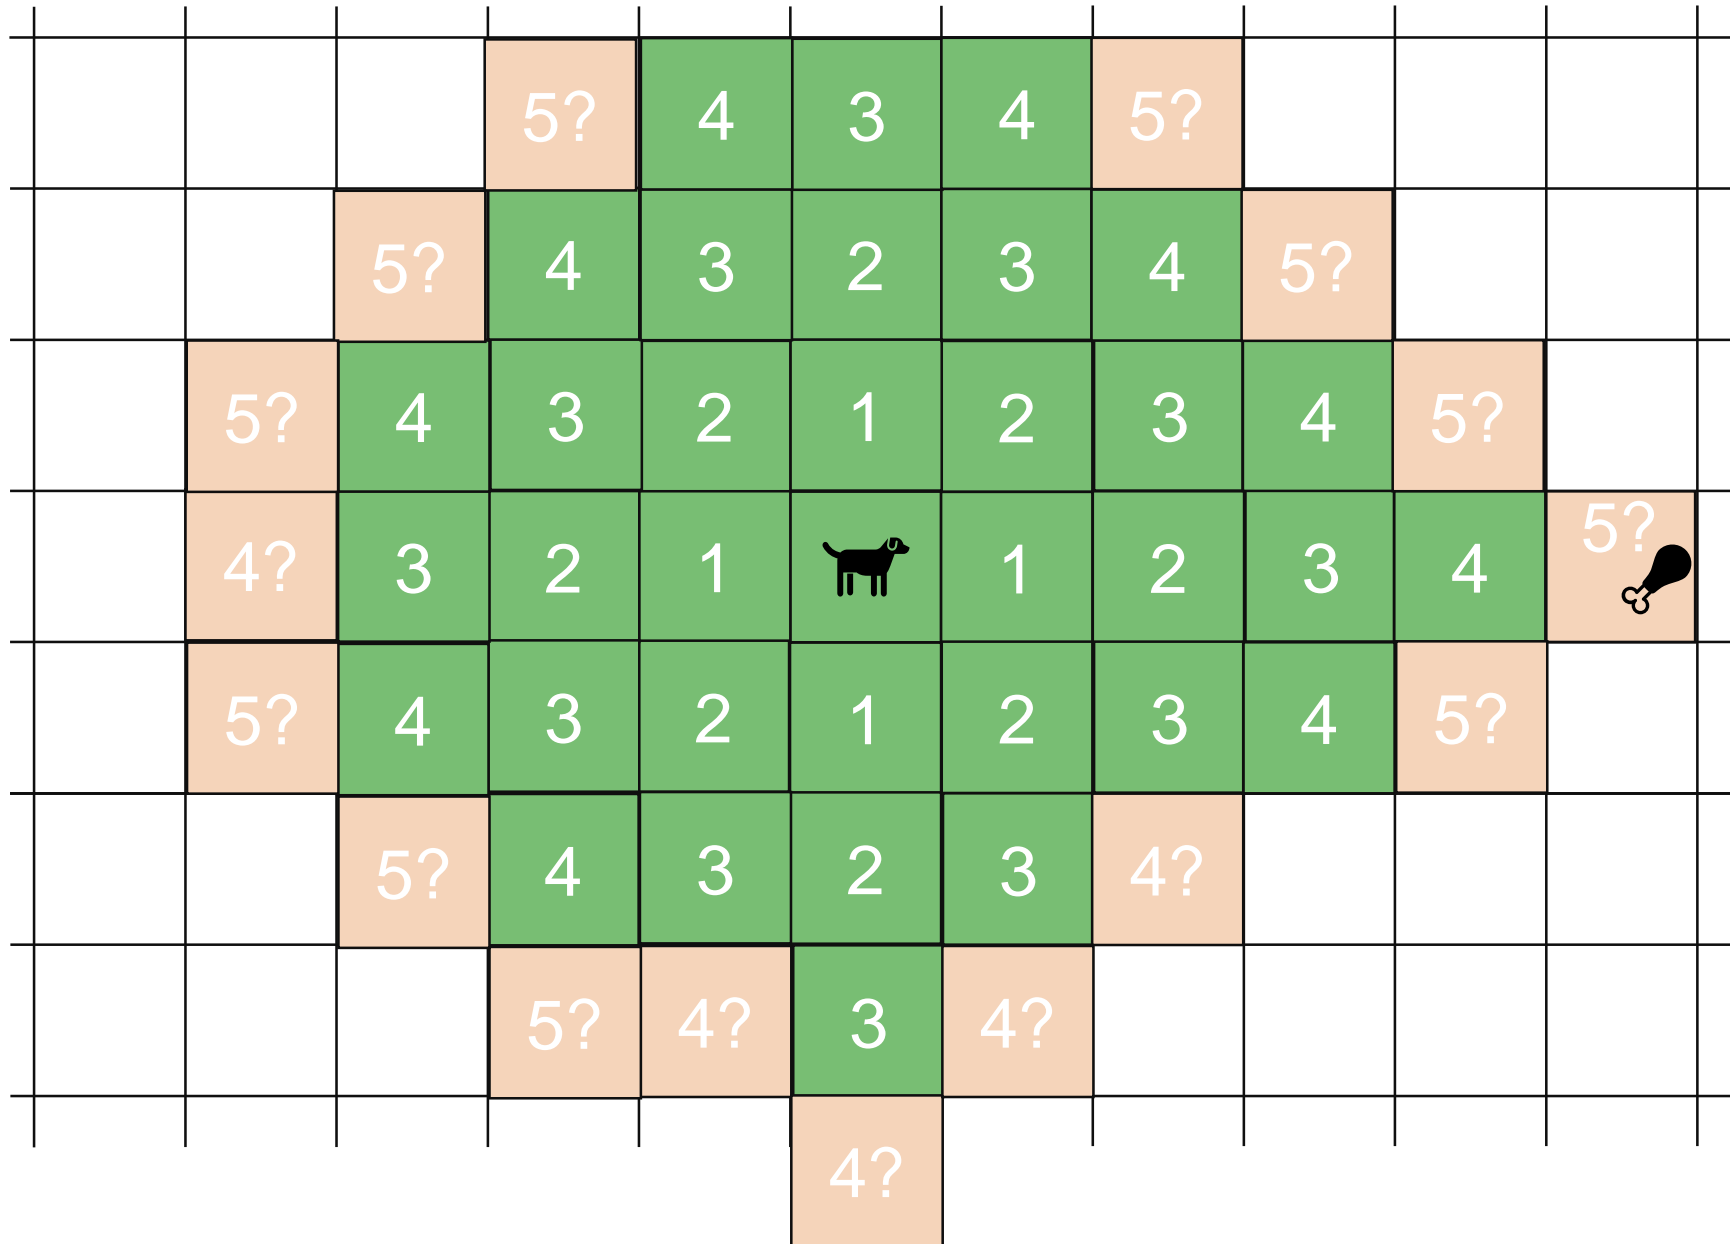
Dijkstra where each edge has cost 1



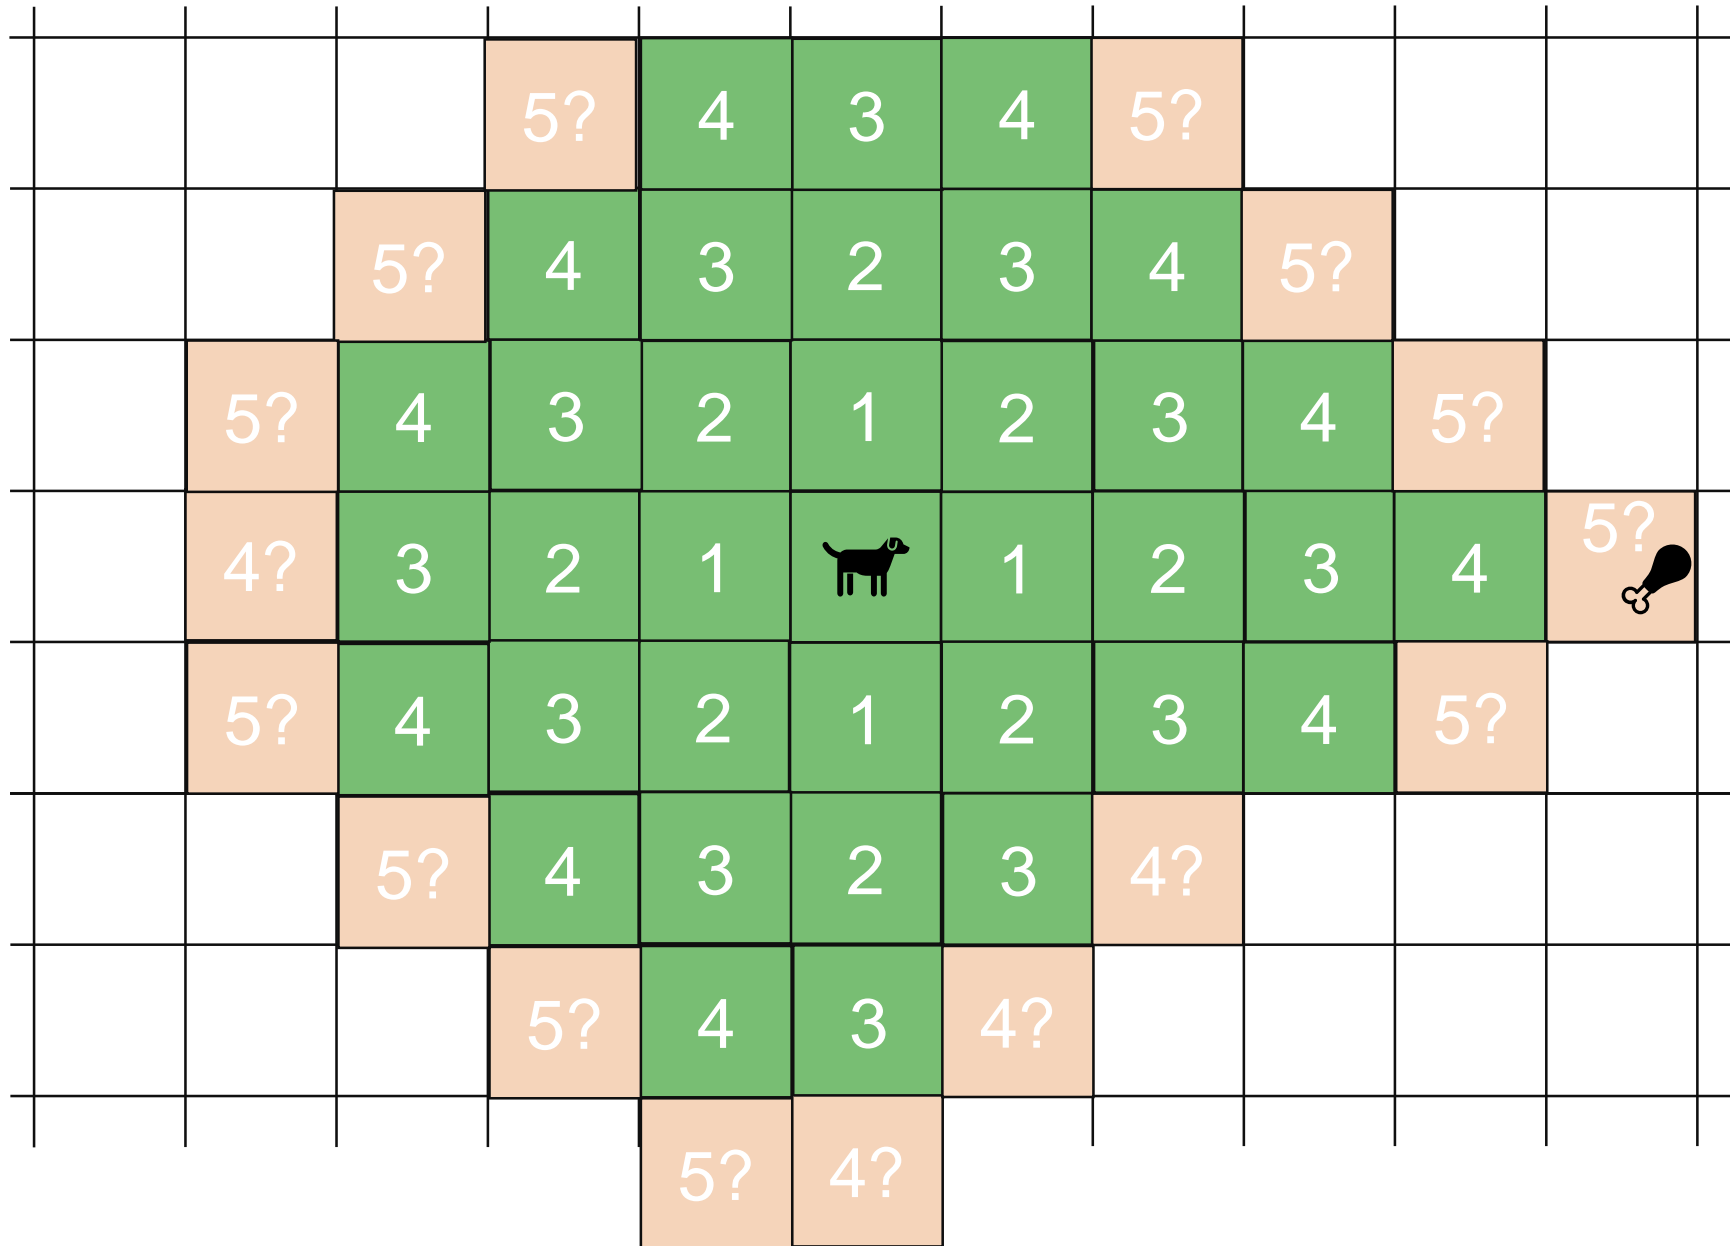
Dijkstra where each edge has cost 1



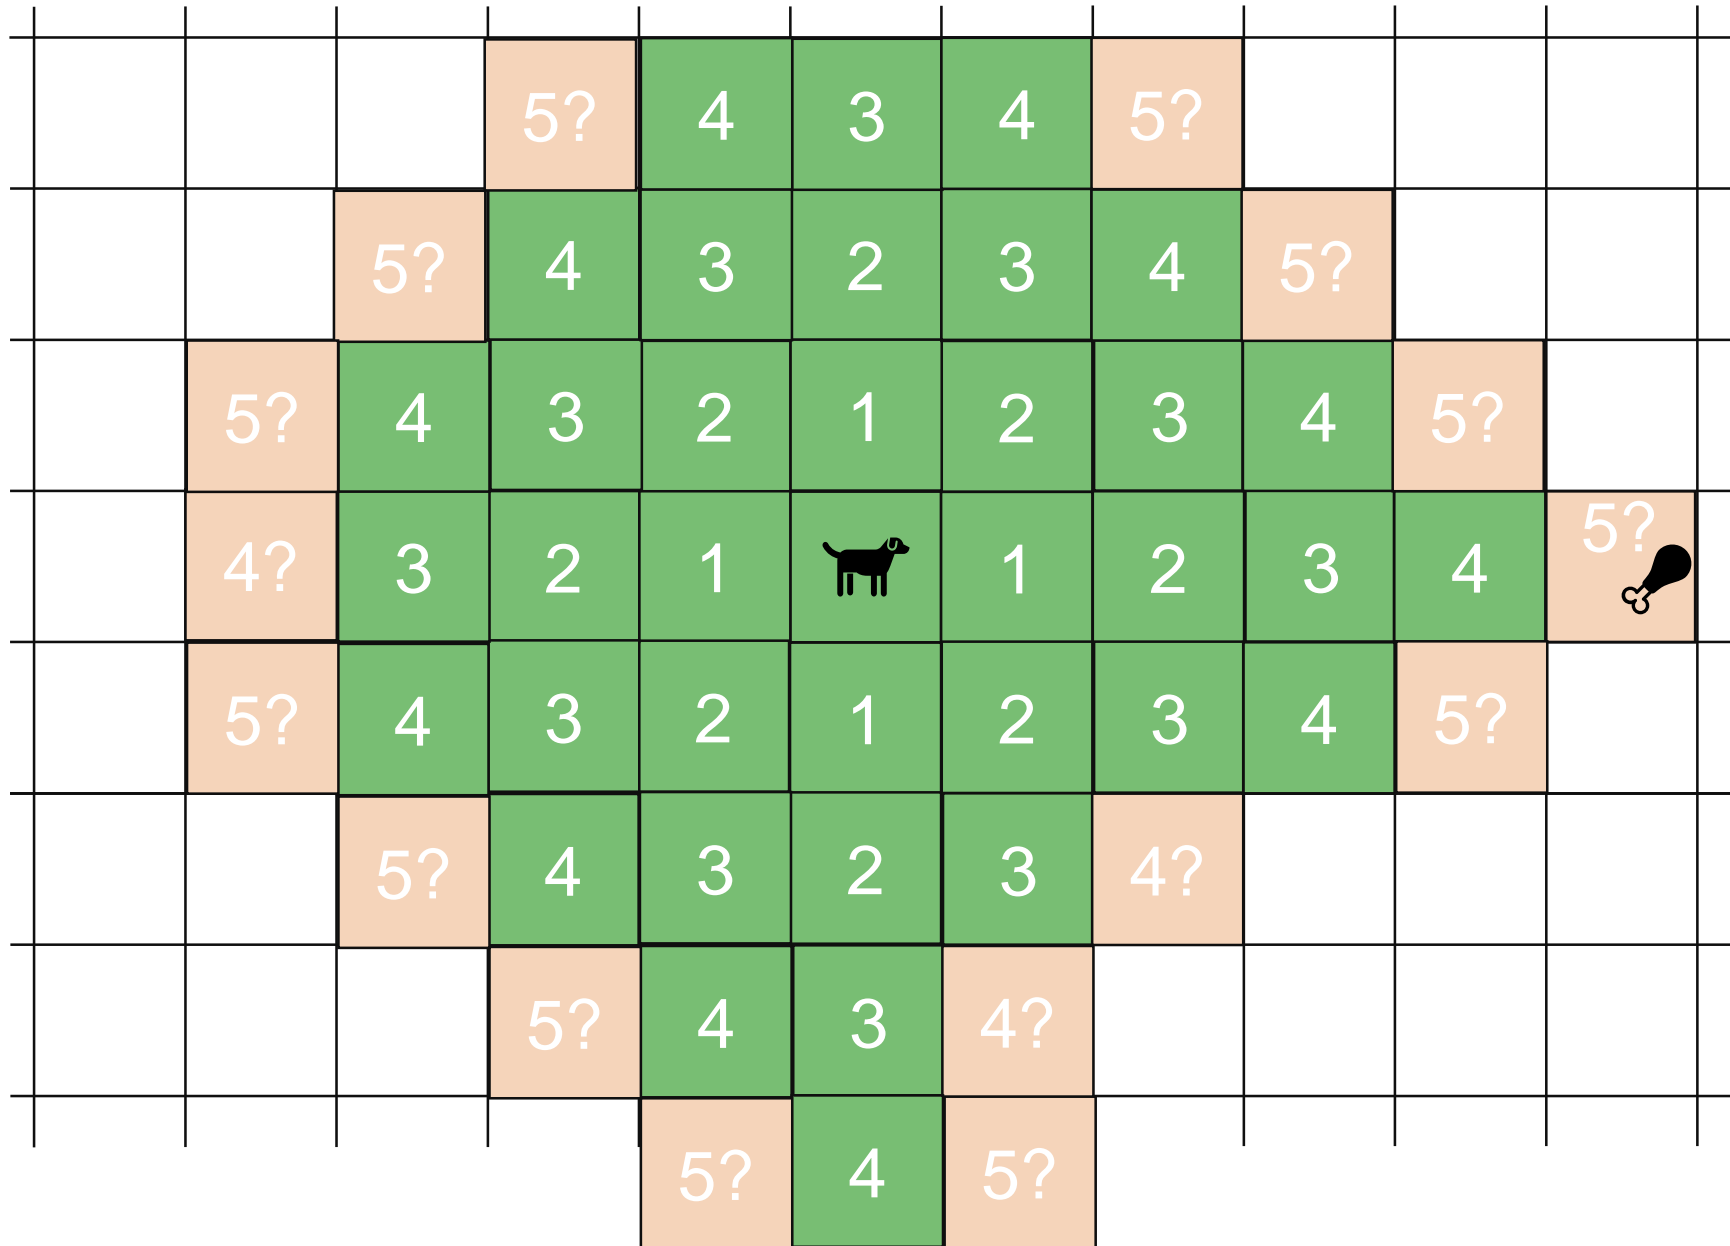
Dijkstra where each edge has cost 1



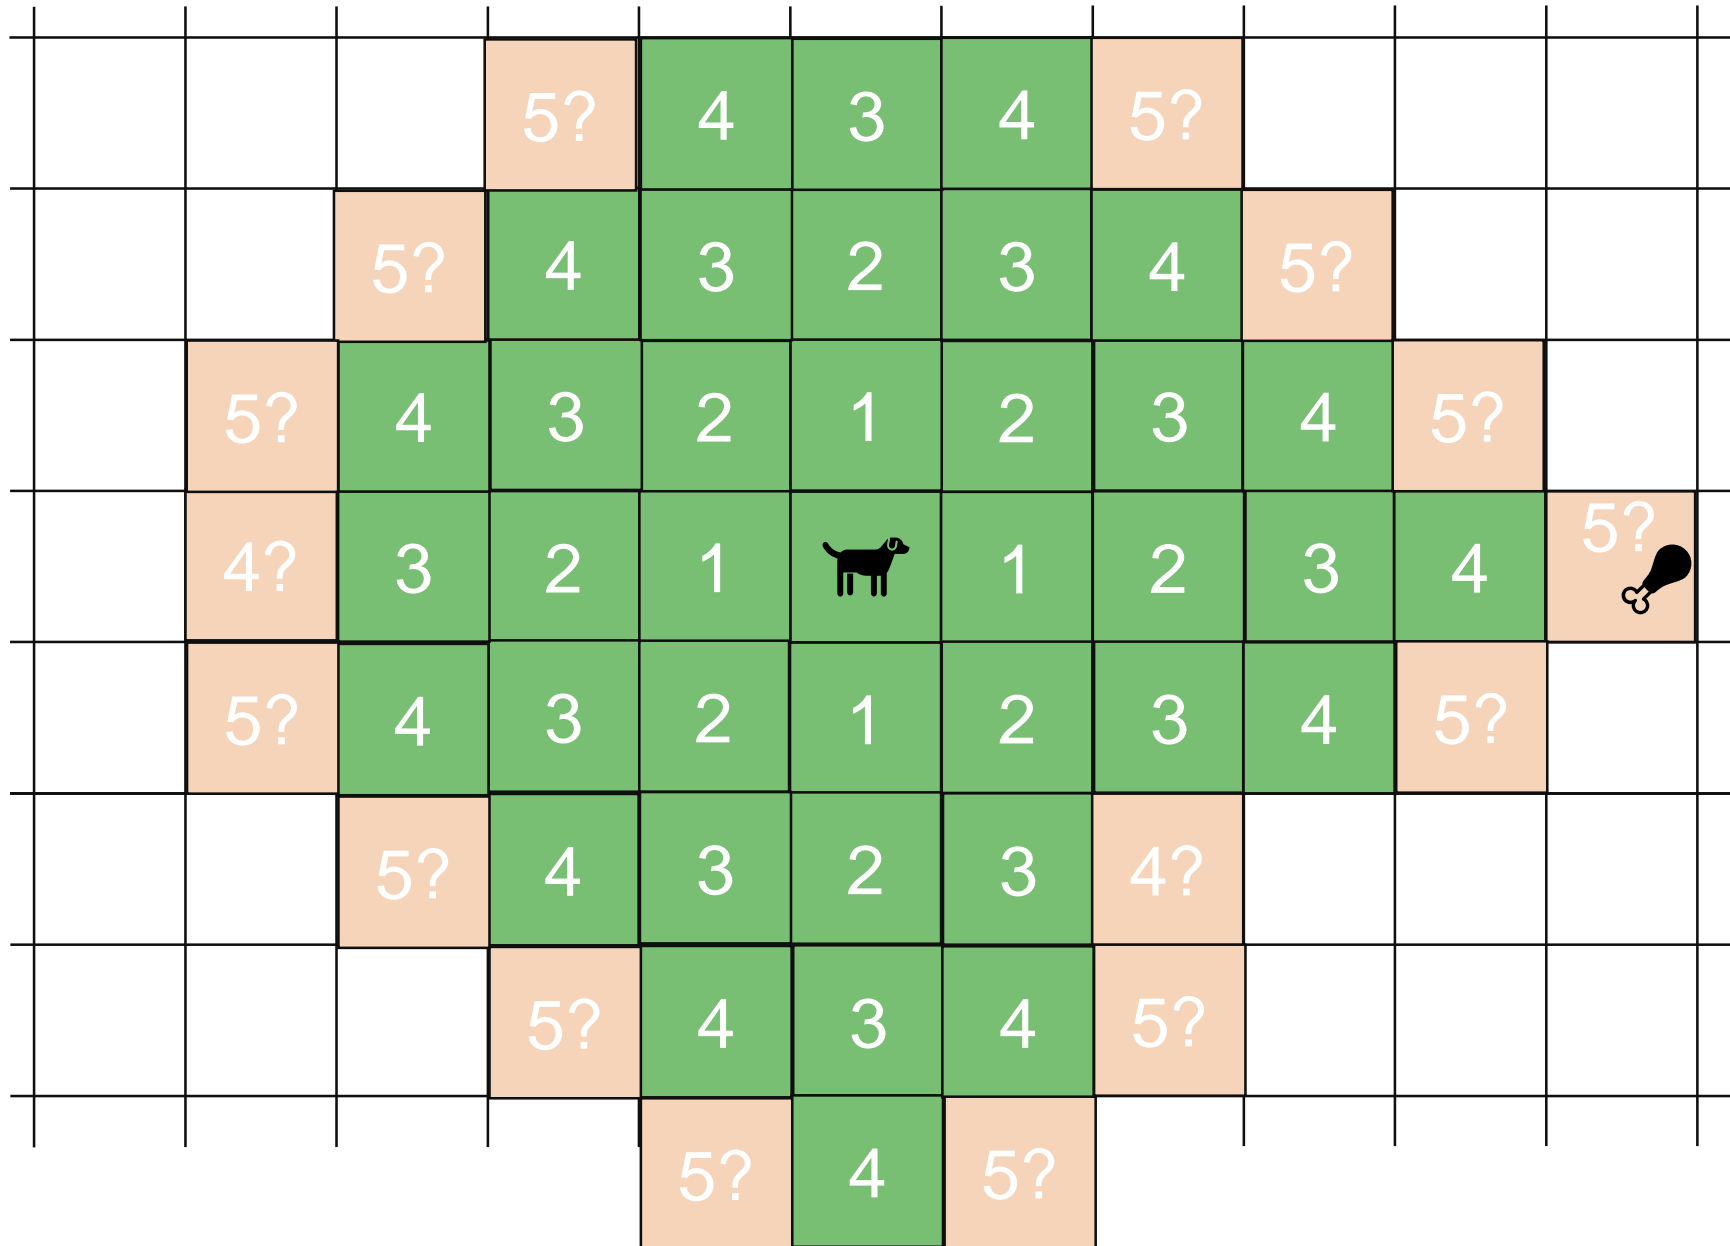
Dijkstra where each edge has cost 1



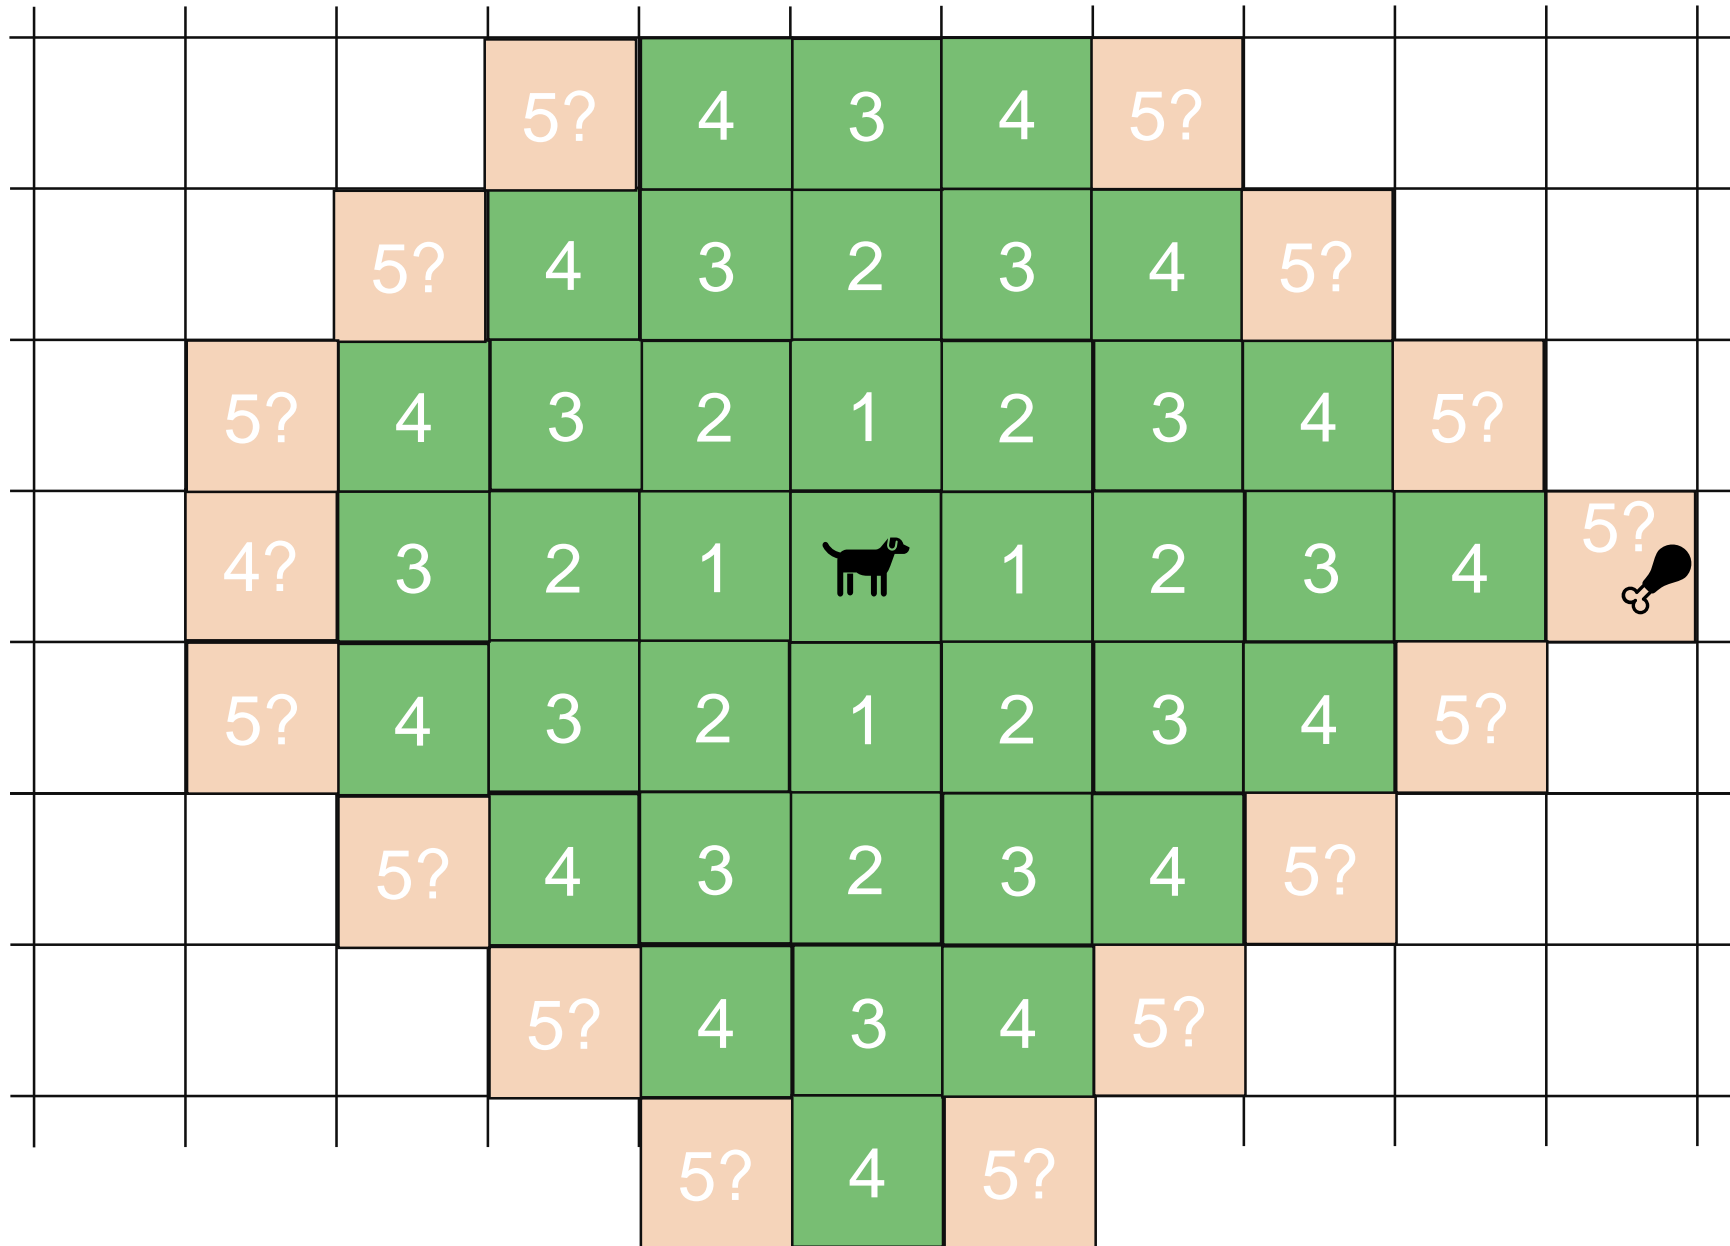
Dijkstra where each edge has cost 1



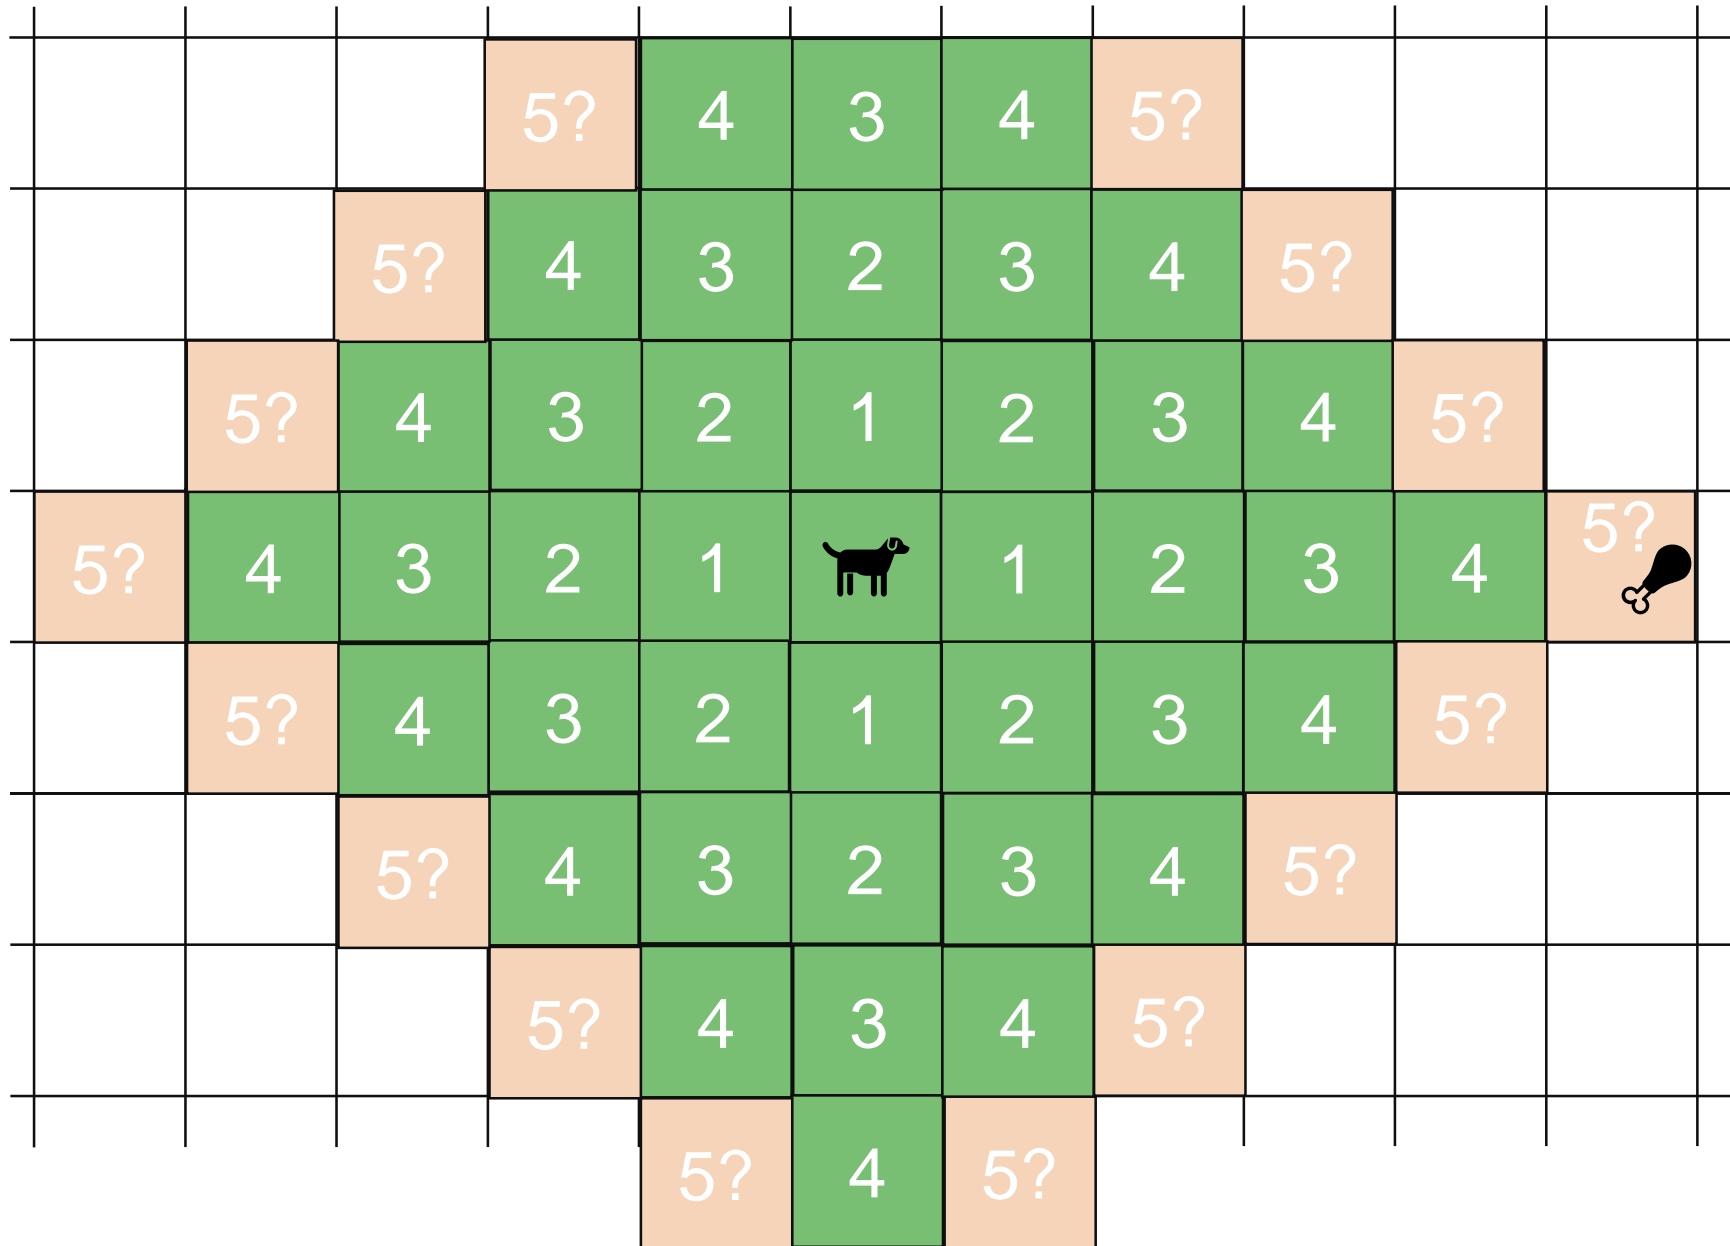
Dijkstra where each edge has cost 1




Dijkstra where each edge has cost 1



Dijkstra where each edge has cost 1



Dijkstra where each edge has cost 1




Why do we even look into this direction?


We haven't taken into account the cost from that node to the target

Very far

...very far ...




$$\text{priority}(u) = \boxed{\text{distance}(s, u)} + \boxed{\text{futureCost}(u, t)}^*$$



priority of the path that ends in u


*needs a revision, often we do not know $\boxed{\text{futureCost}(u, t)}$

A^* : Future cost = row + column distance



Can be computed if we know row and column of the target t and of v

A*: Future cost = row + column distance




Greedy continue
at smallest

distance(s, u)

+

futureCost(u, t)

A*: Future cost = row + column distance




Greedy continue
at smallest

distance(s, u)

+

futureCost(u, t)

A*: Future cost = row + column distance




Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance




Greedy continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance




Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance




Greedy continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance




Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance




Greedy continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



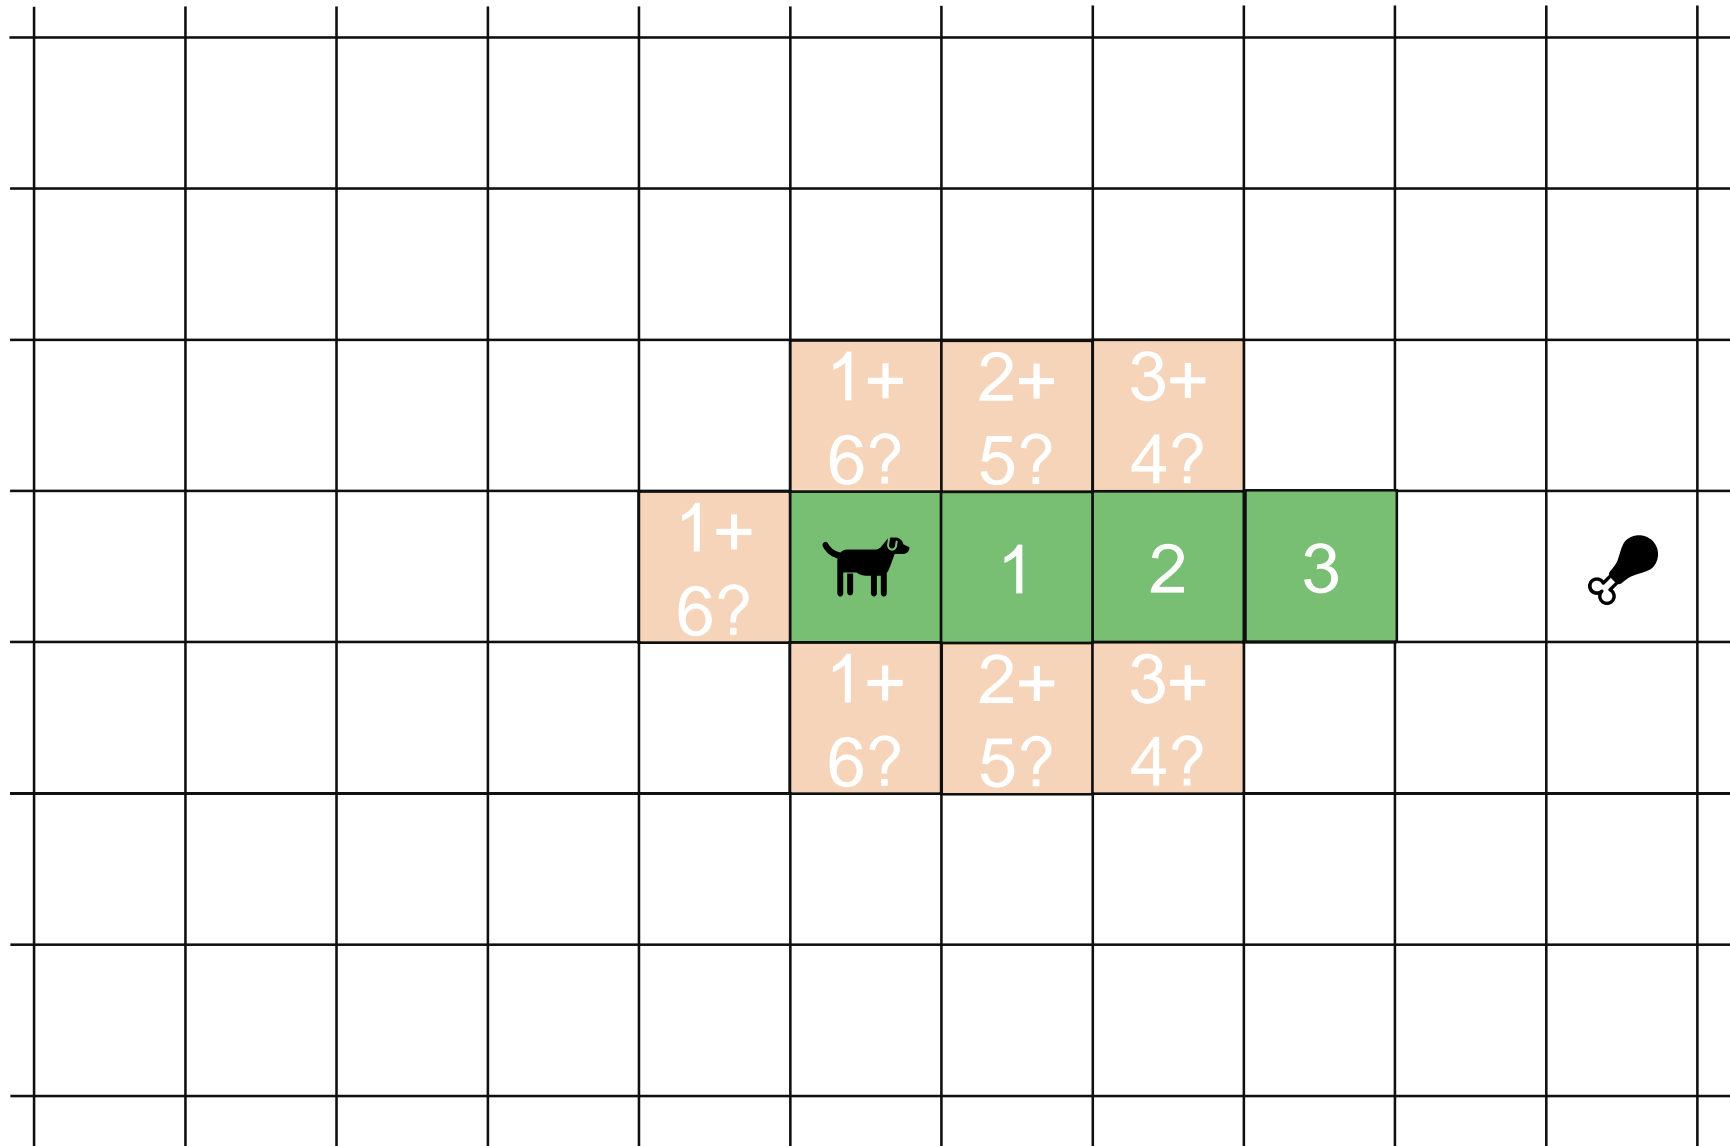
Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



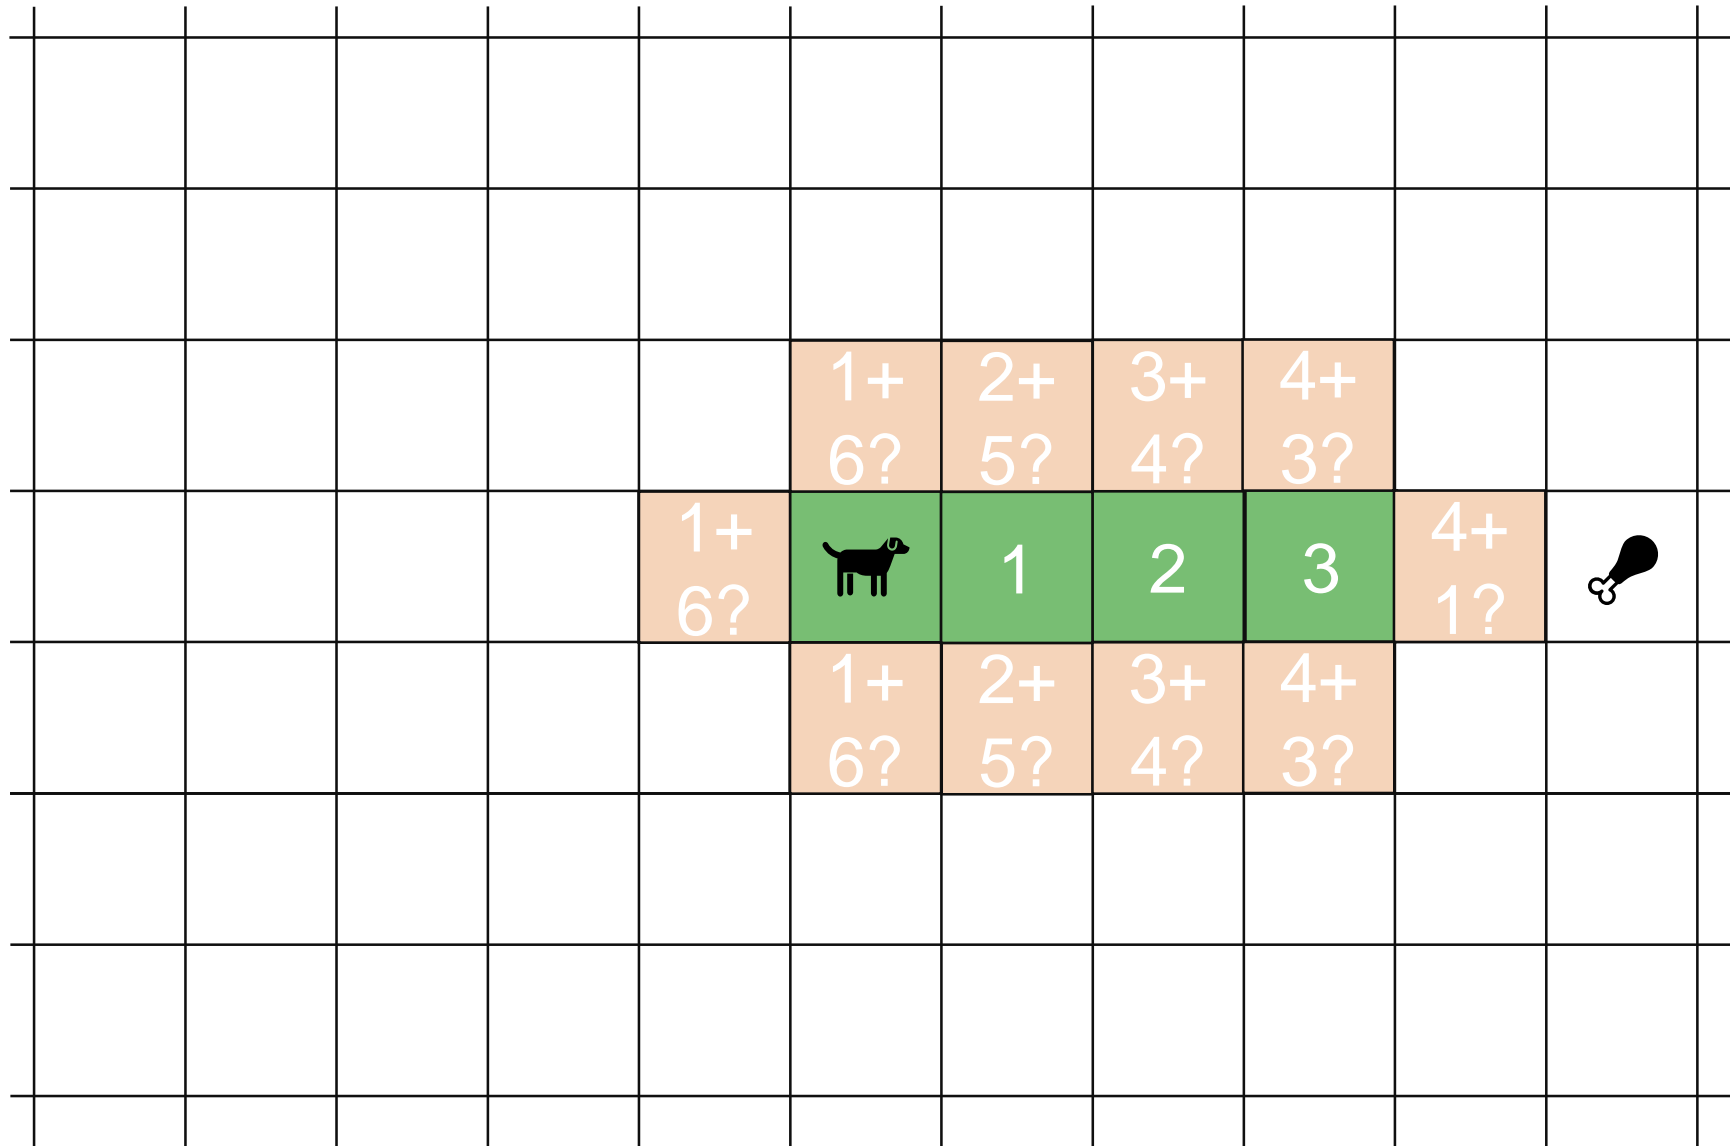
Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



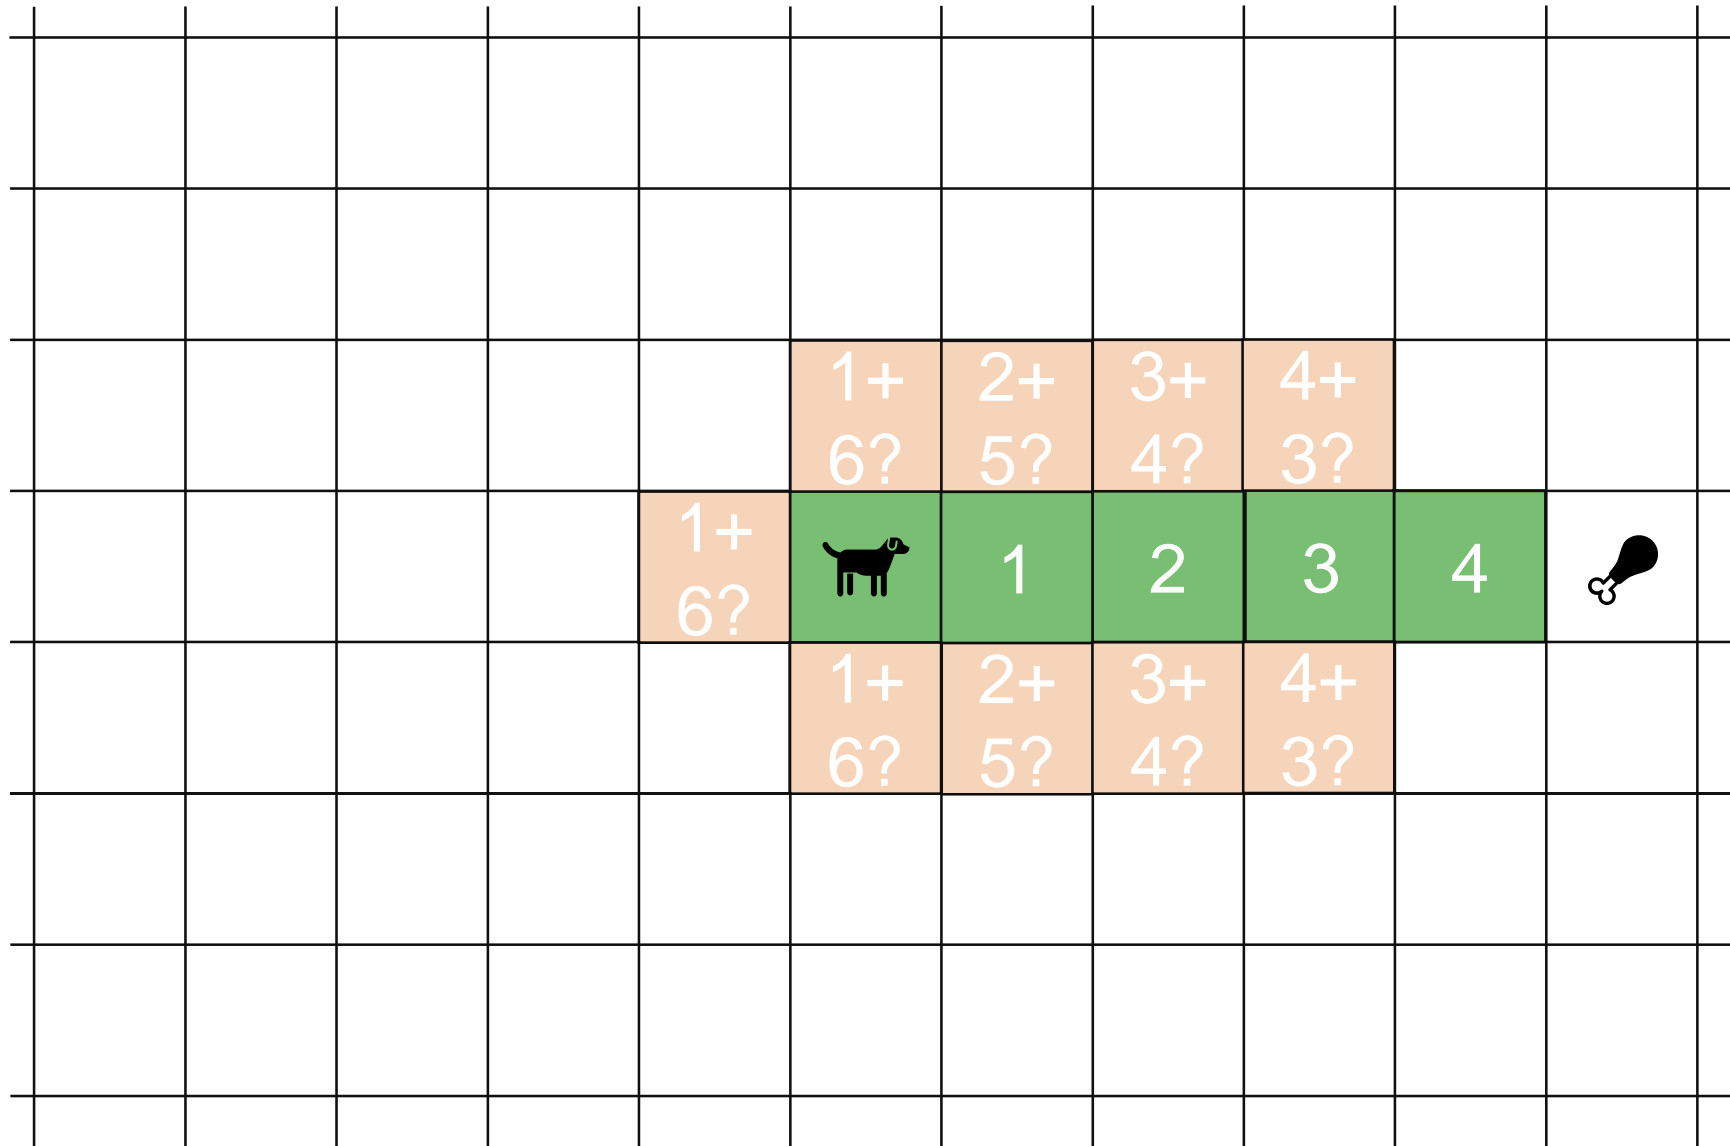
Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



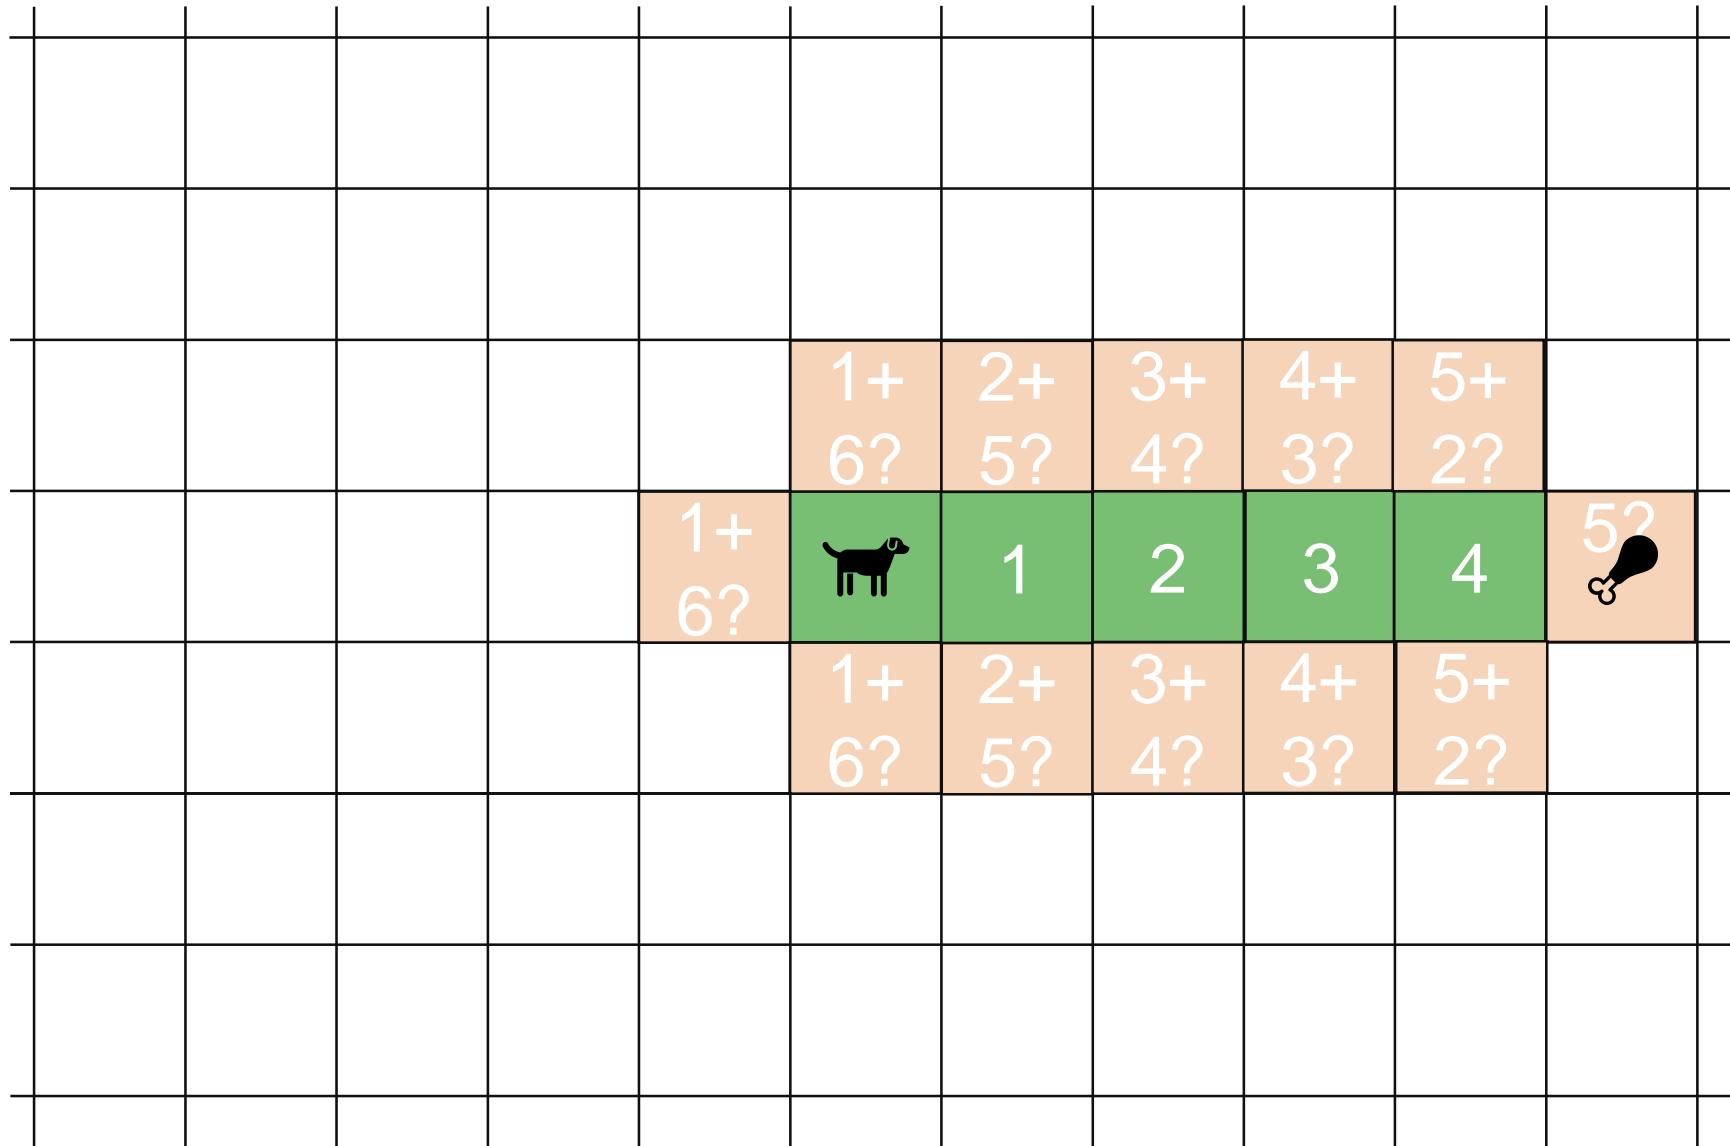
Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



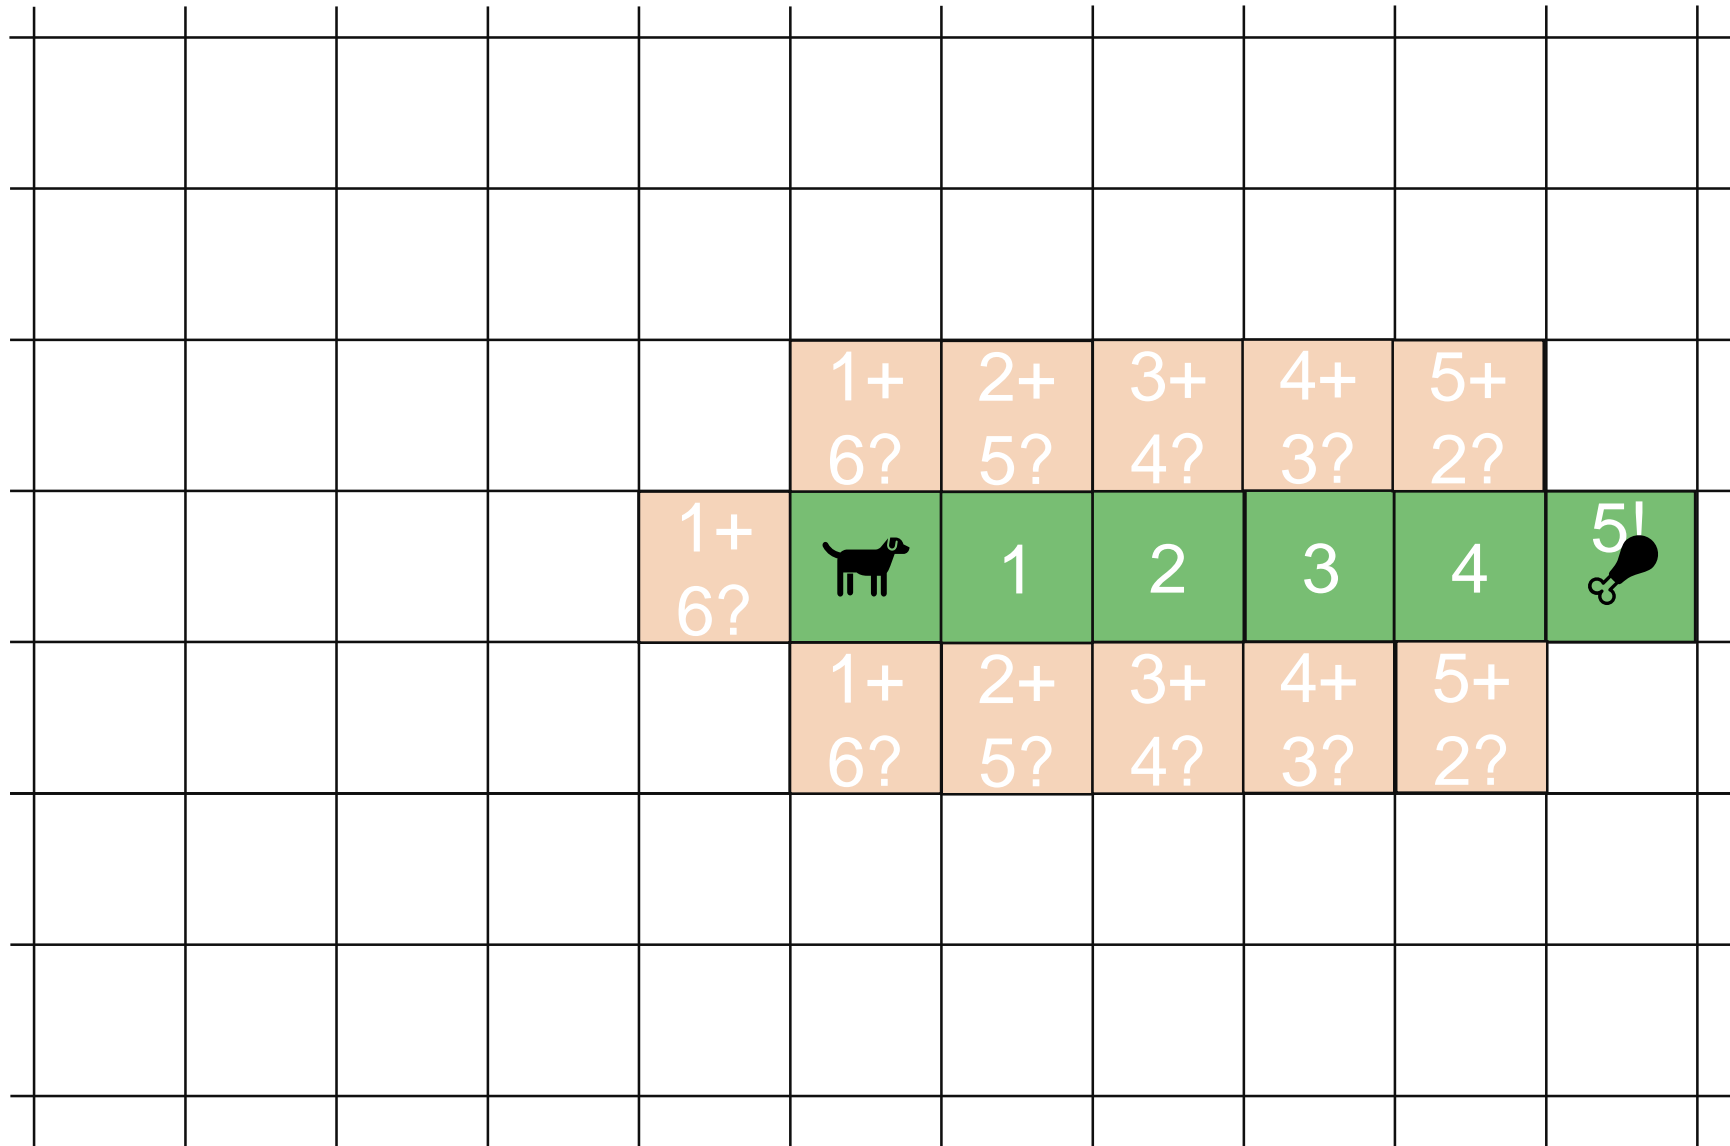
Greedyly continue
at smallest

distance(s, u)

+

futureCost(u,t)

A*: Future cost = row + column distance



Greedyly continue
at smallest

distance(s, u)


+

futureCost(u,t)

A bit easy

A*: Use row + column distance

row+column distance \neq
the real future cost



Greedy continue
at smallest

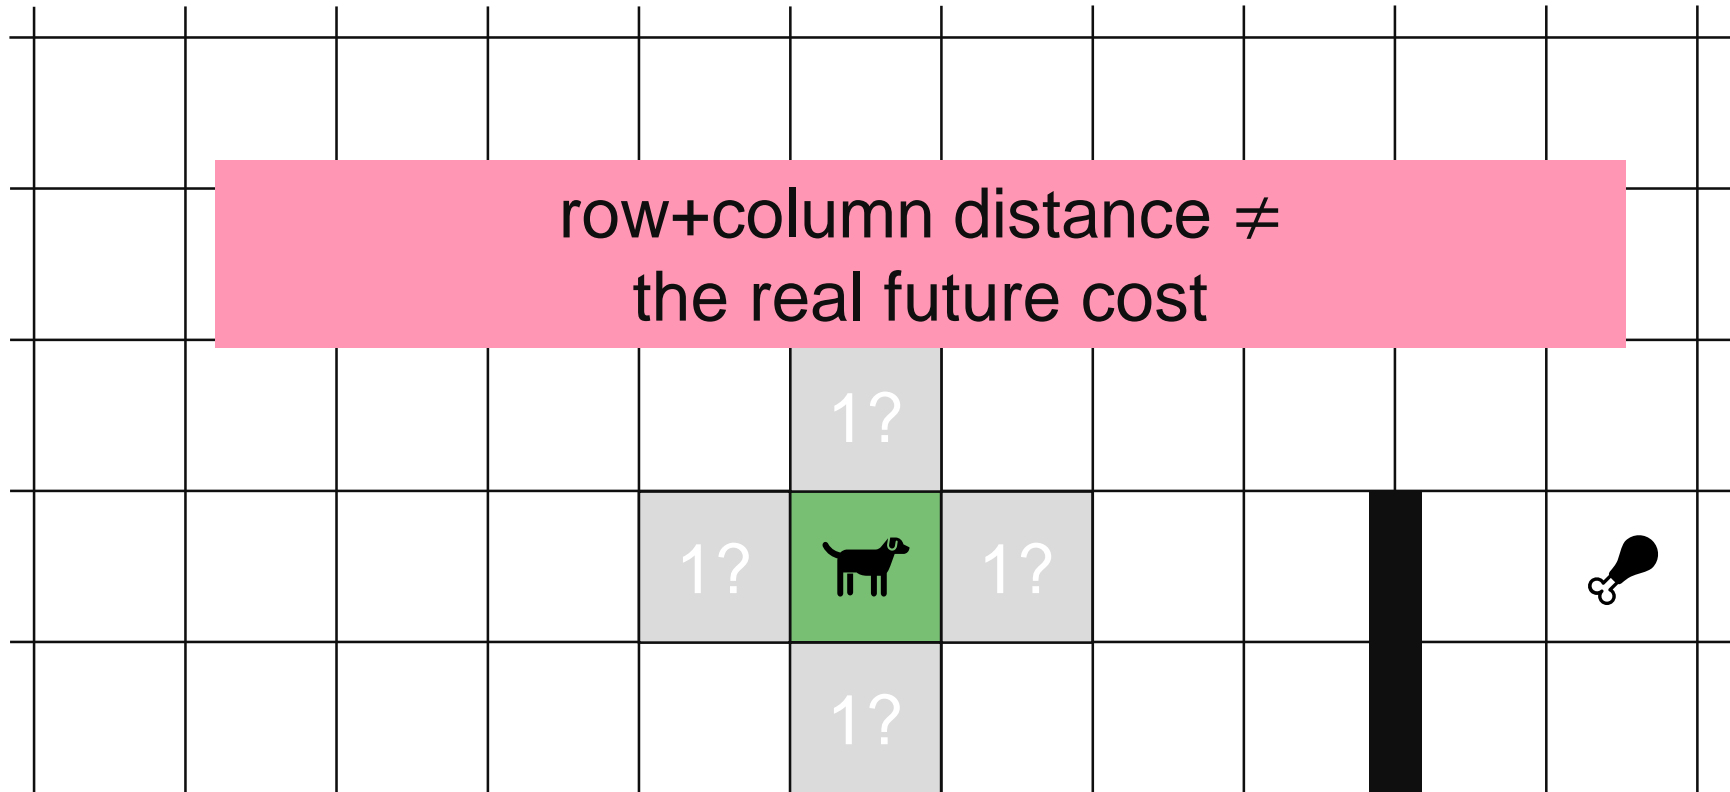
distance(s, u)

+


~~futureCost(u, t)~~

wanna be
futureCost(u, t)

A*: Use row + column distance



This is the A* Algorithm!
(Dijkstra with different path priorities)



Greedy continue
at smallest



distance(s, u)

+

~~futureCost(u, t)~~

wanna be
futureCost(u, t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




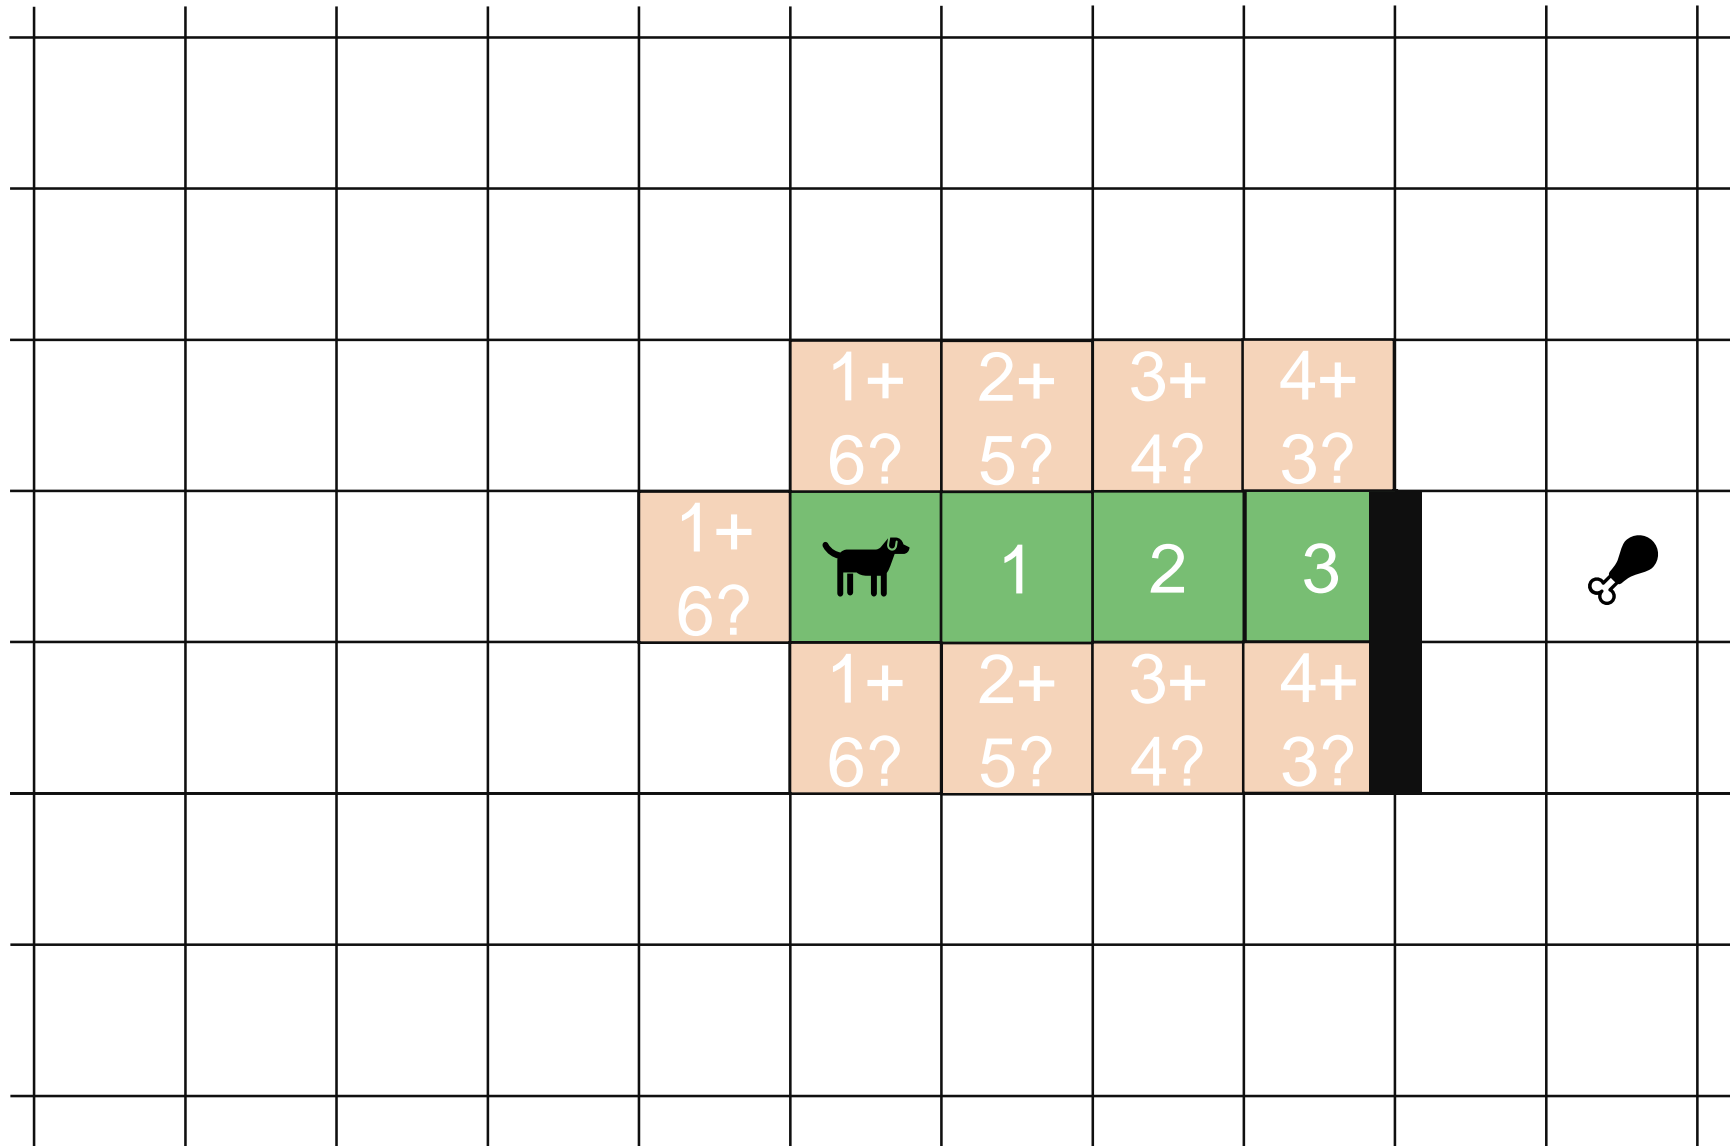
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




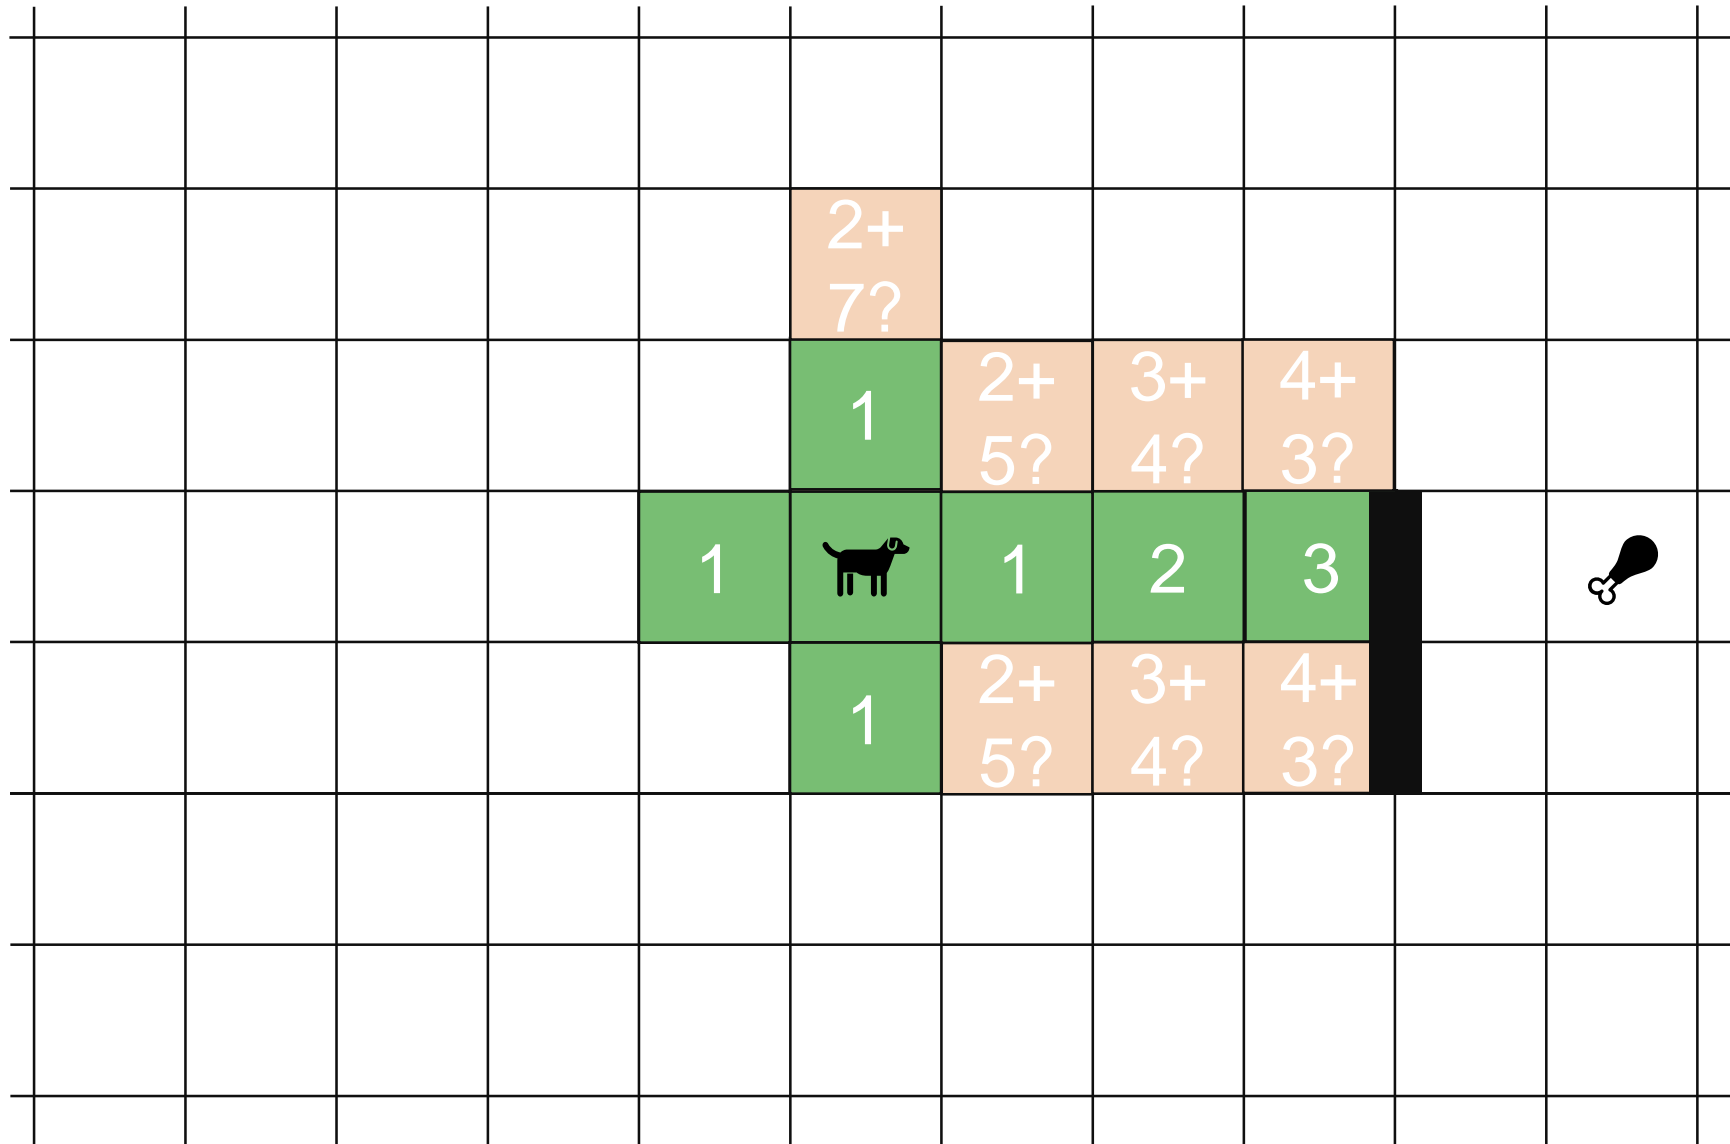
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




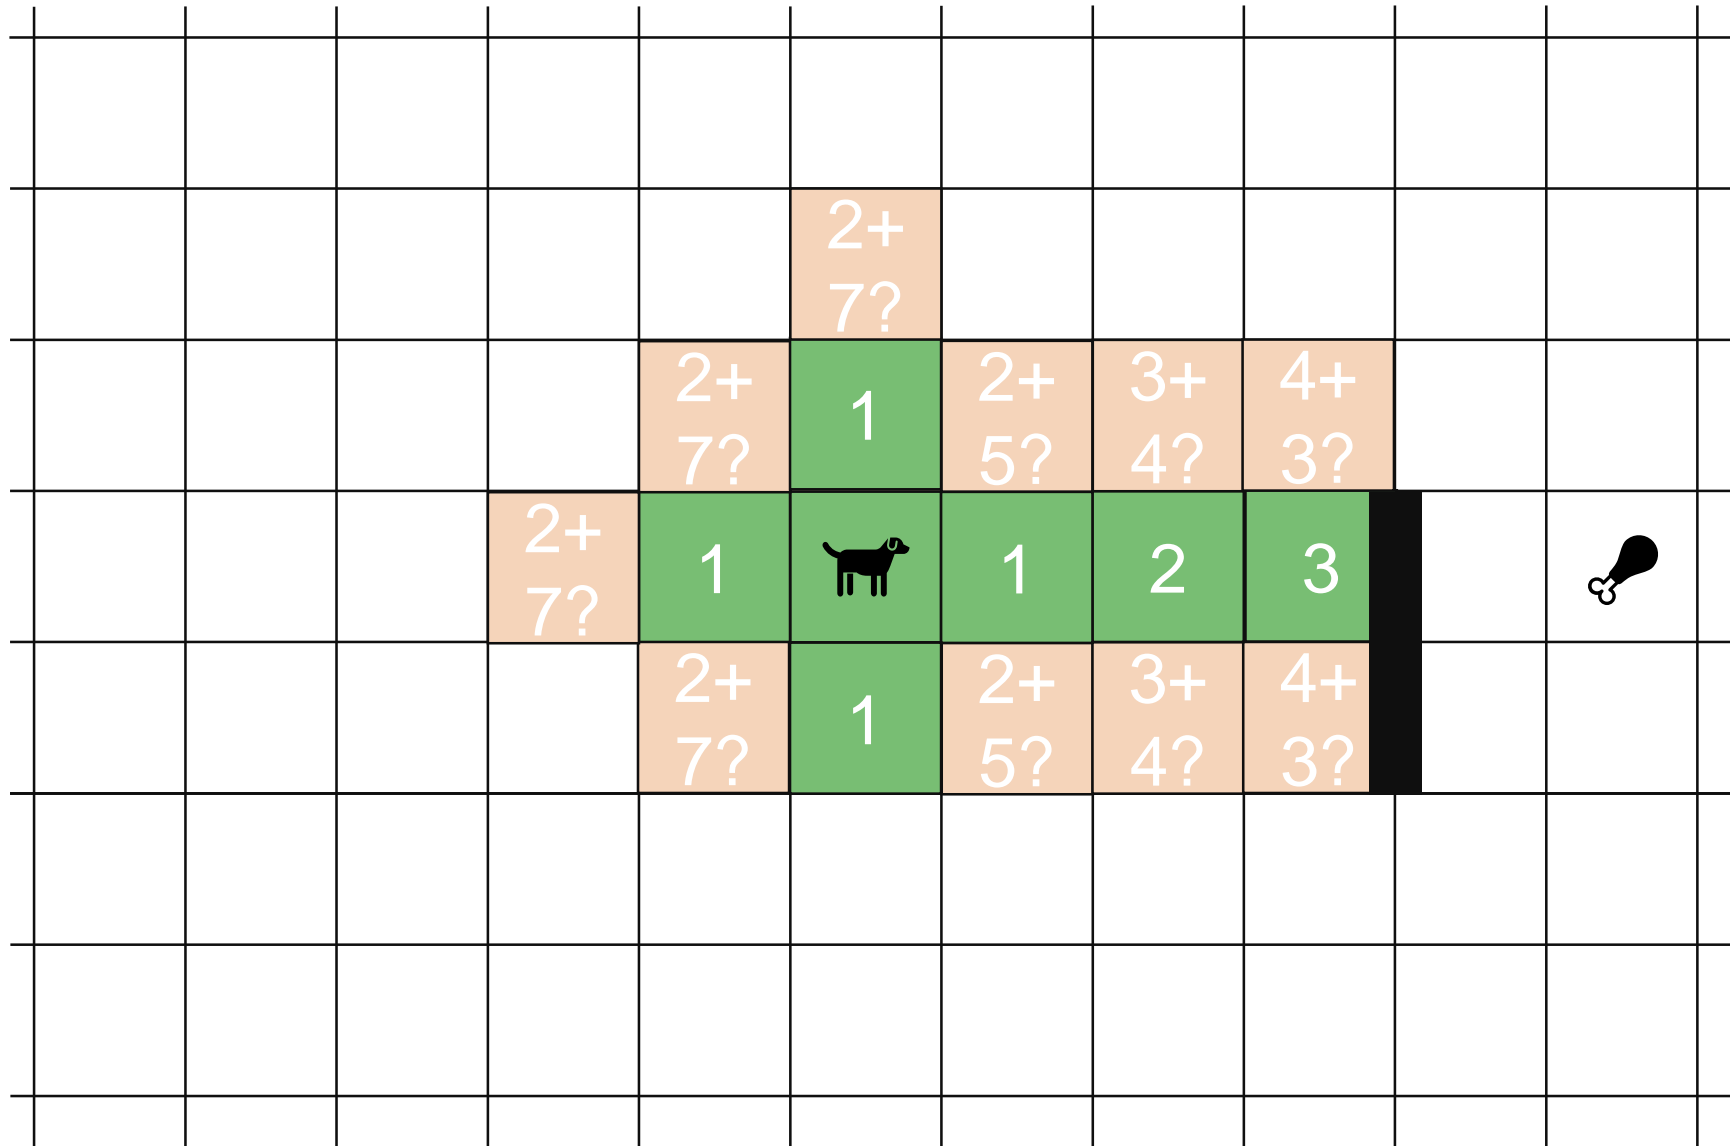
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




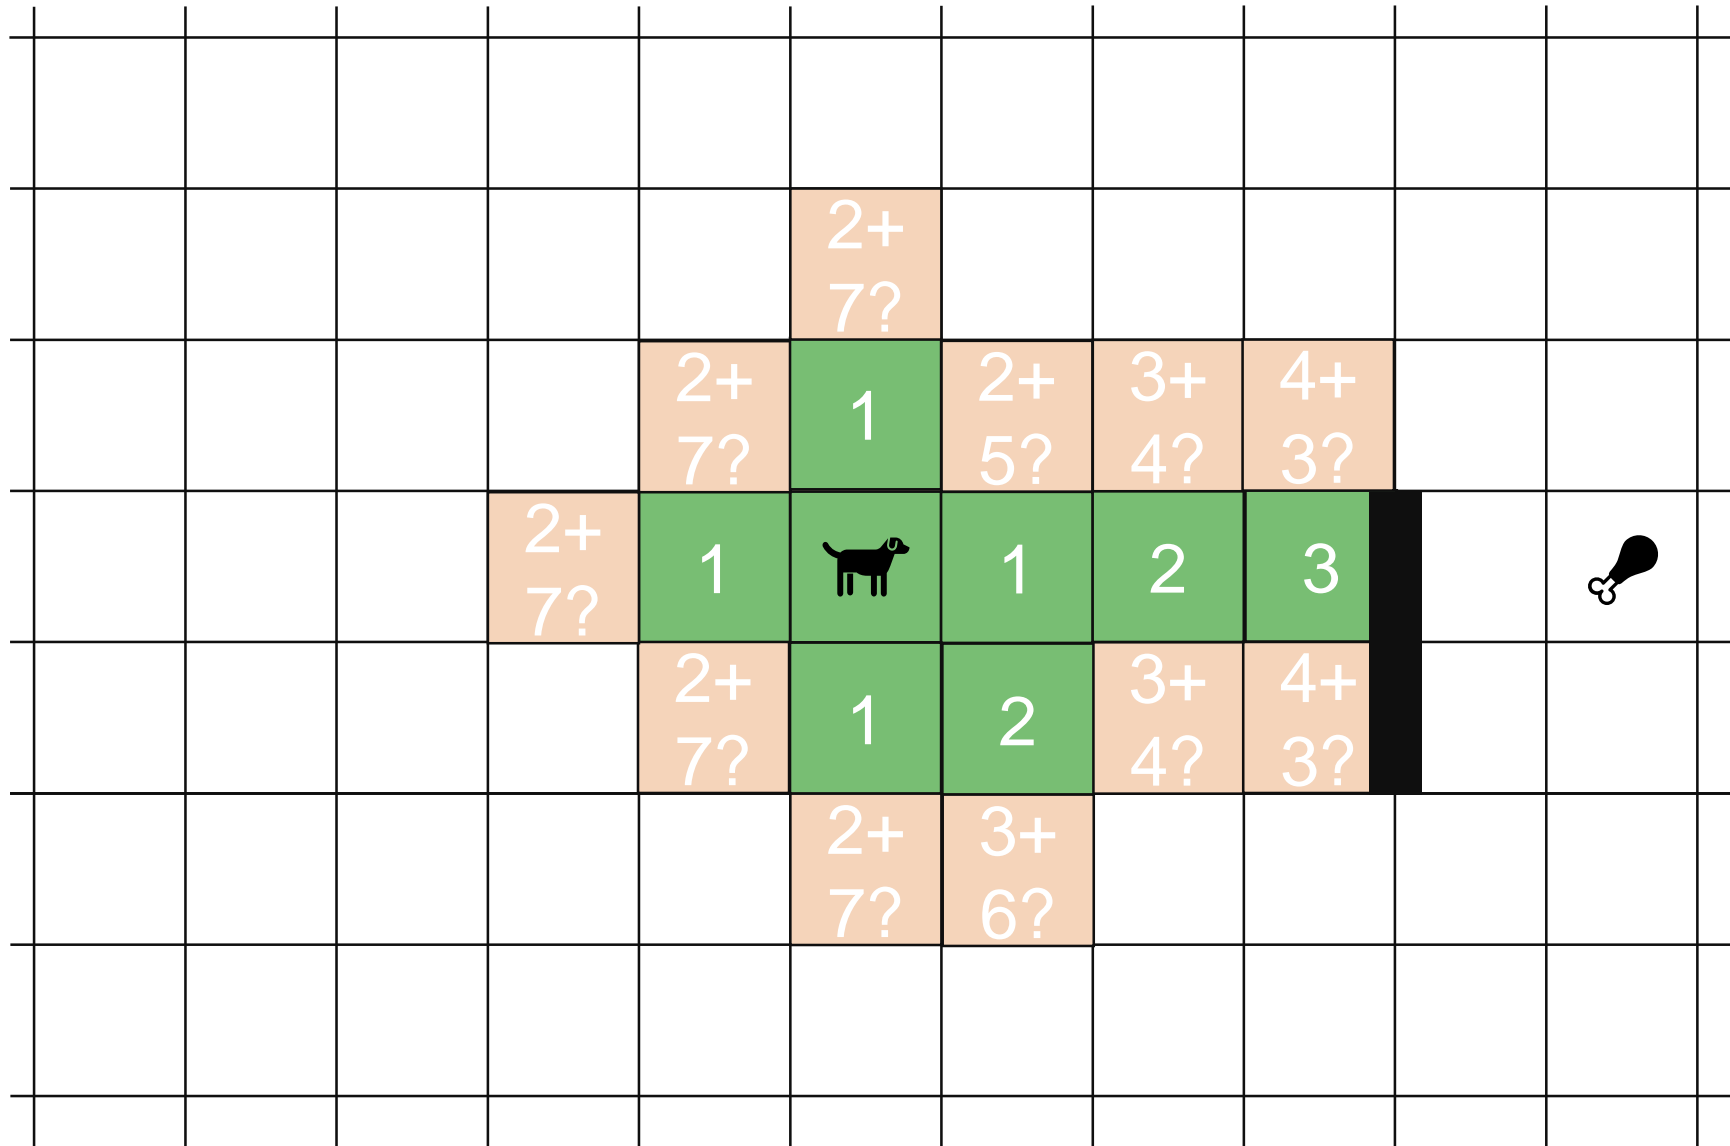
Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




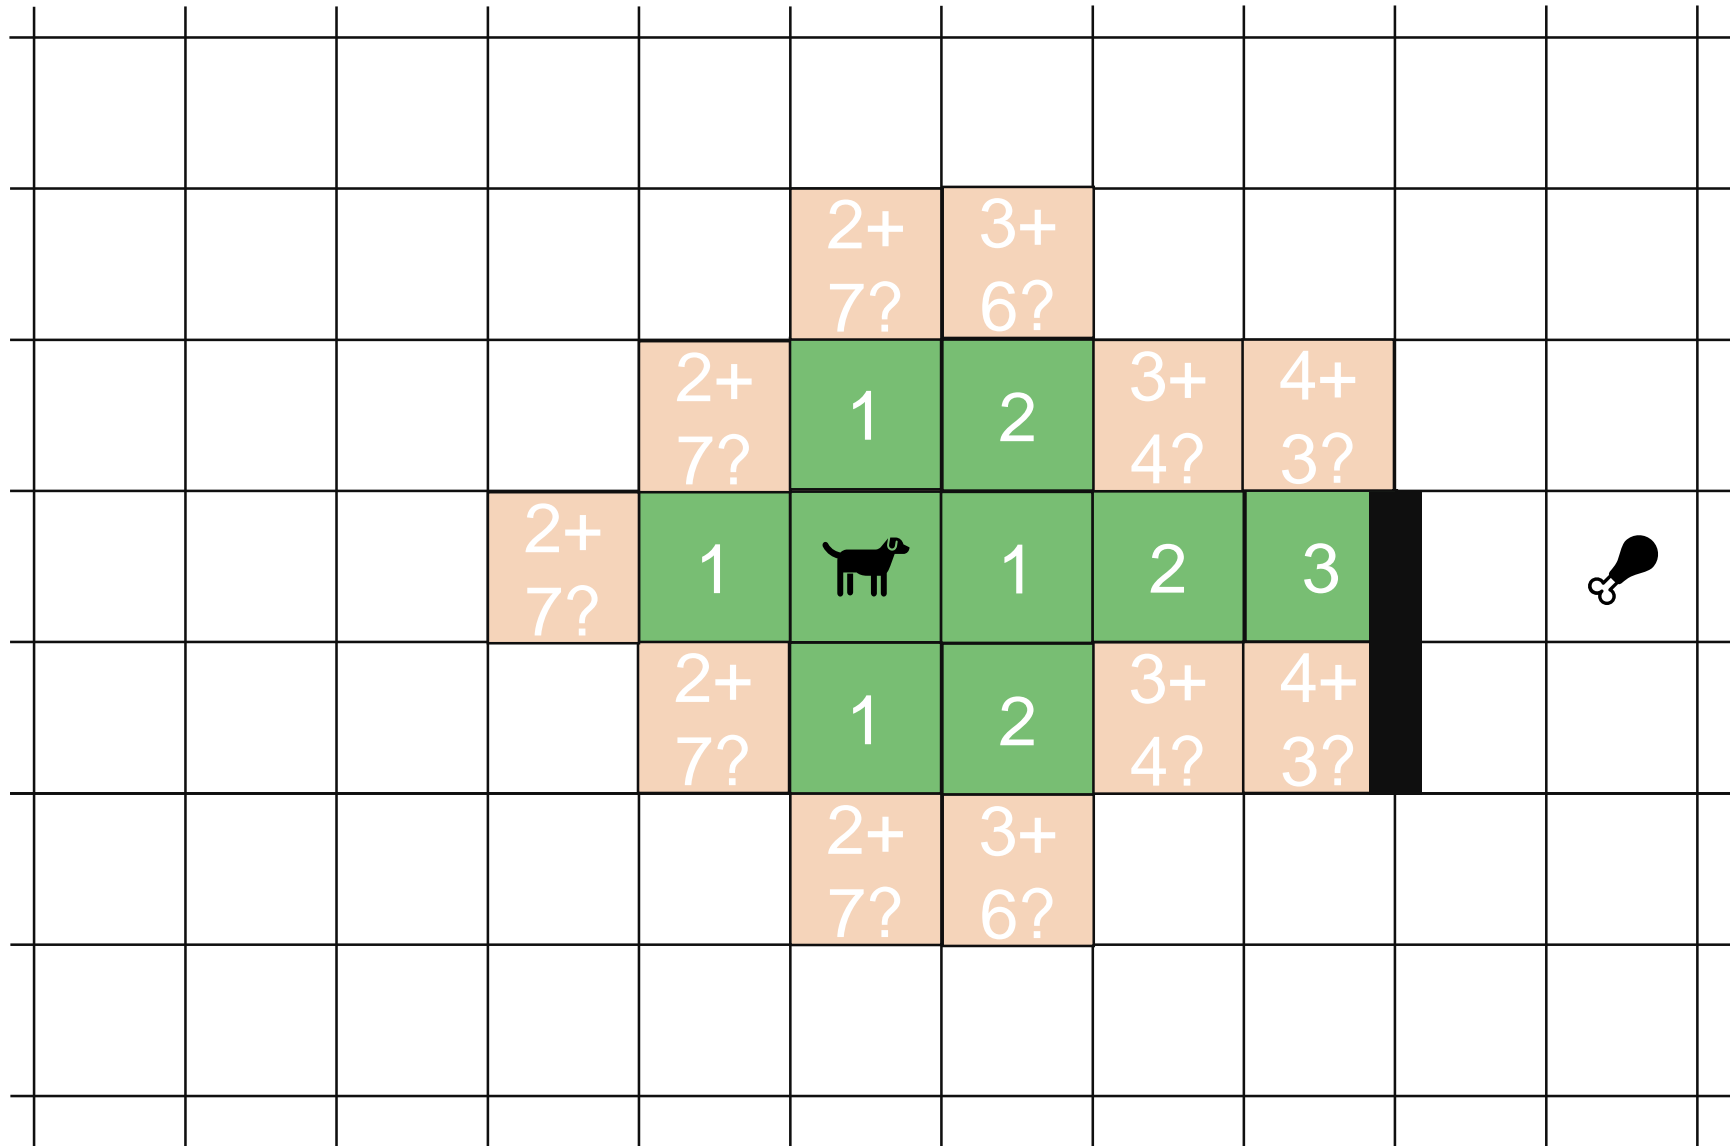
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




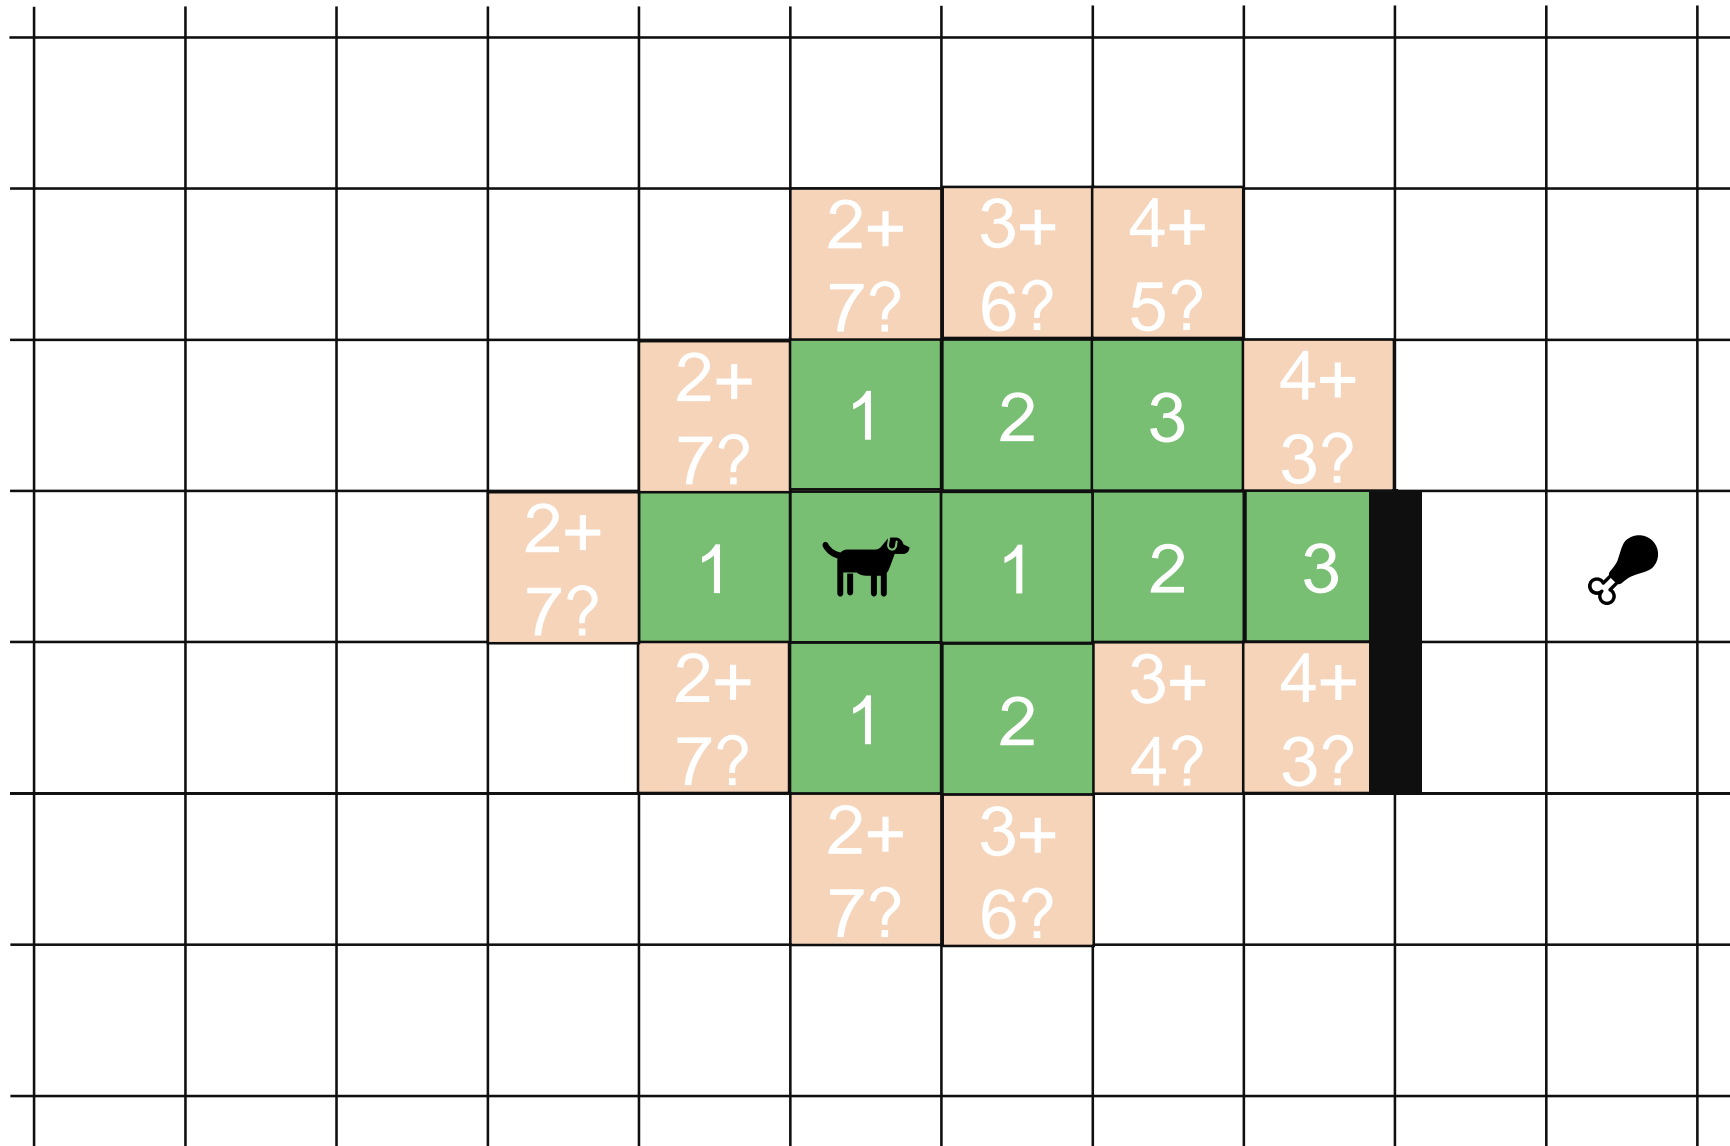
Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




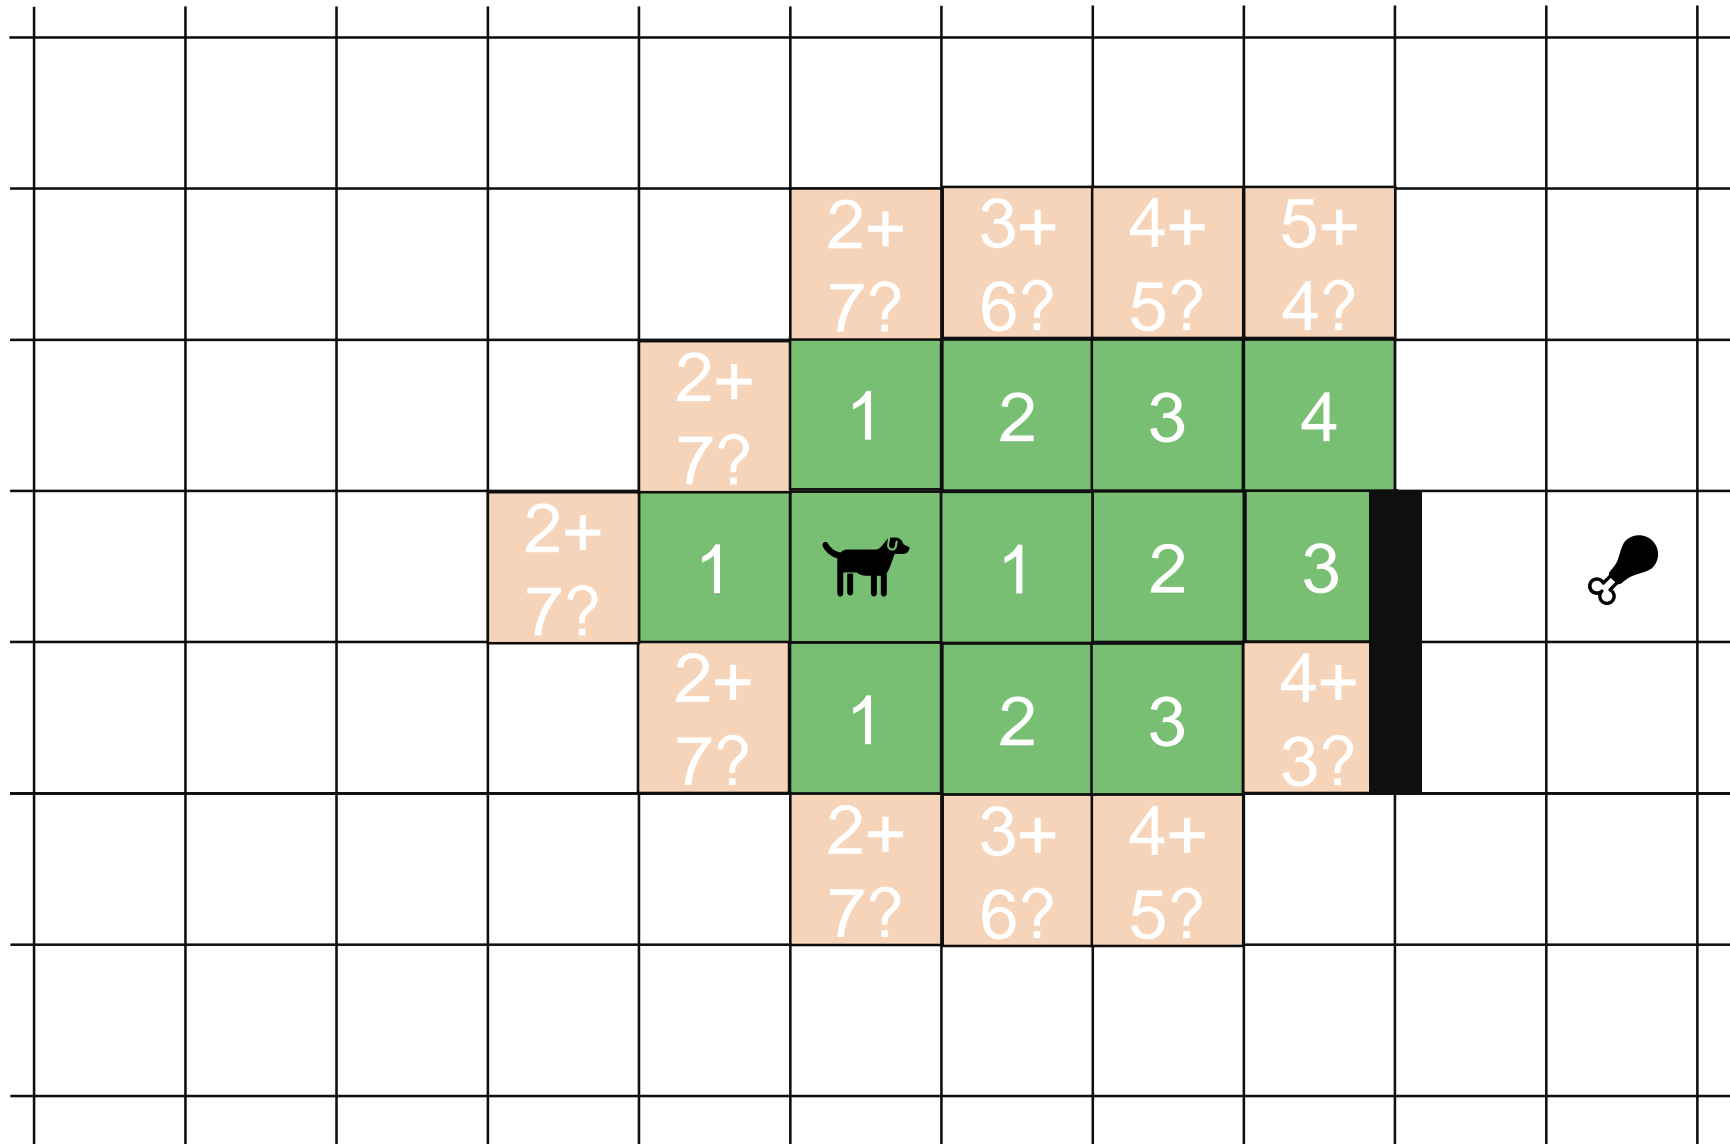
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




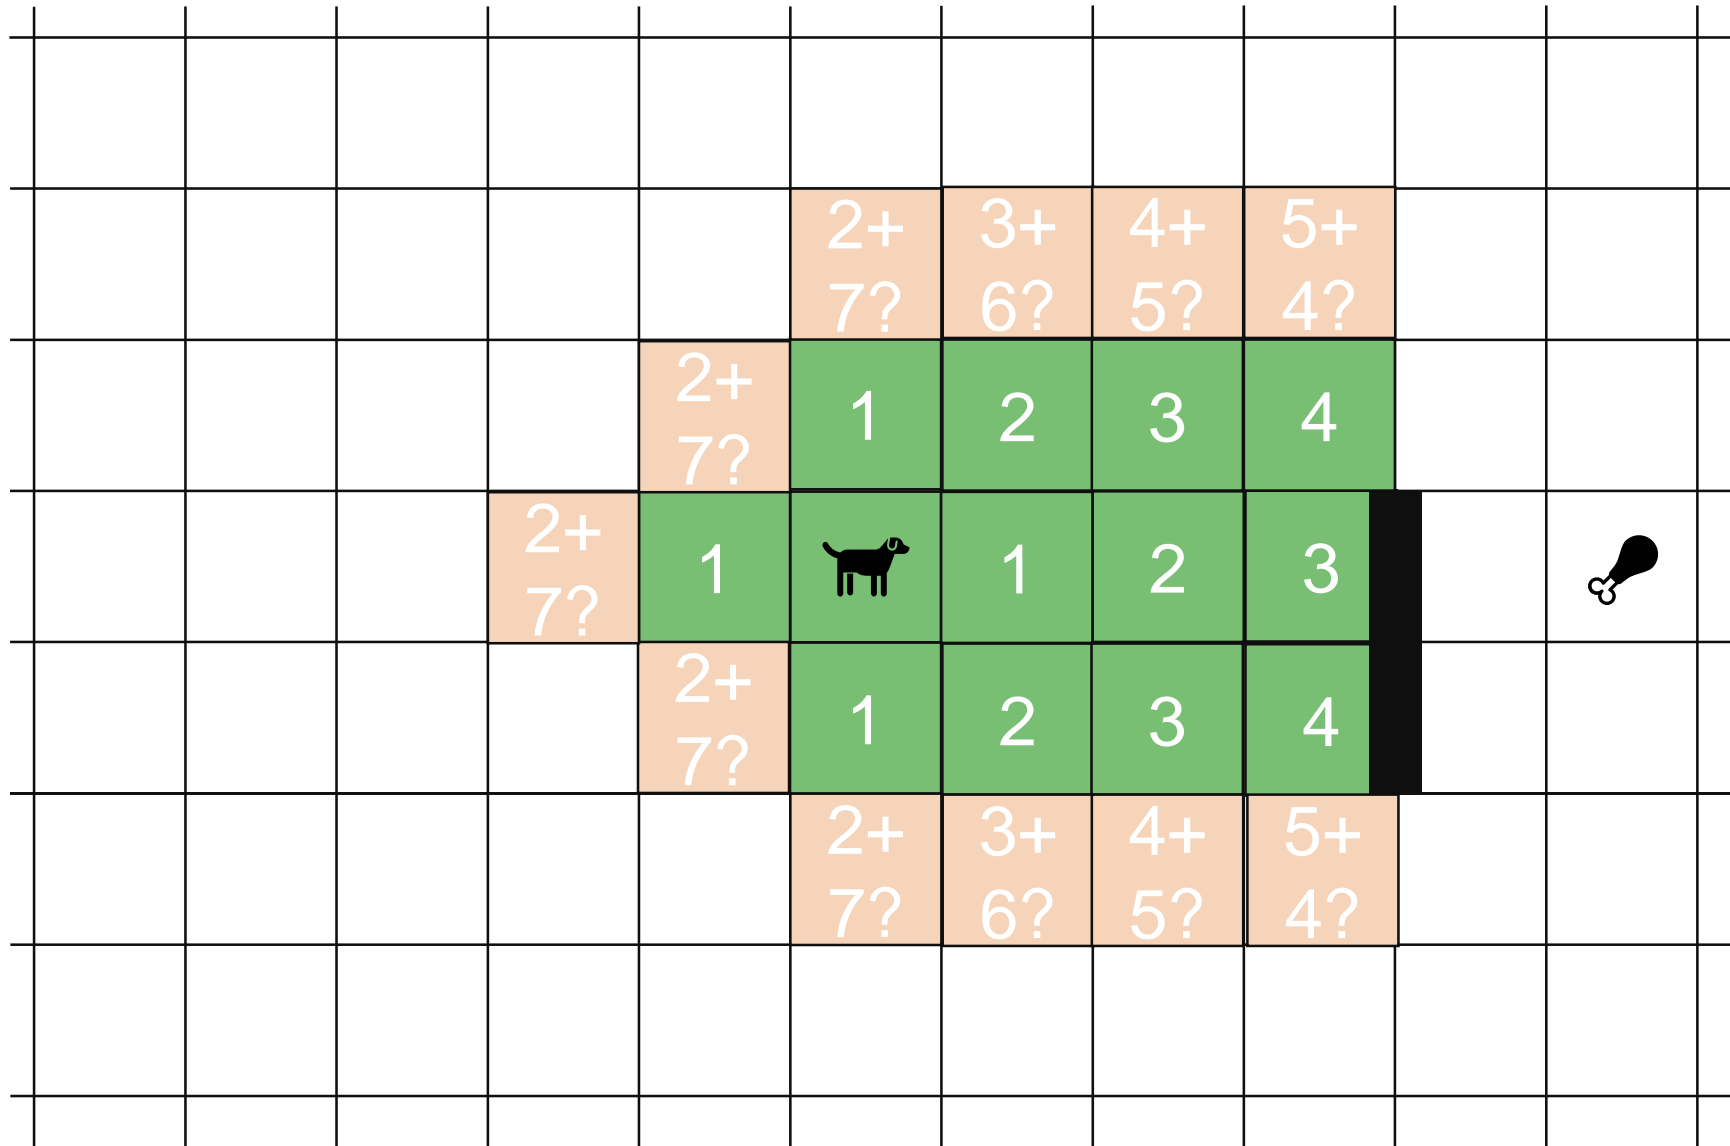
Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




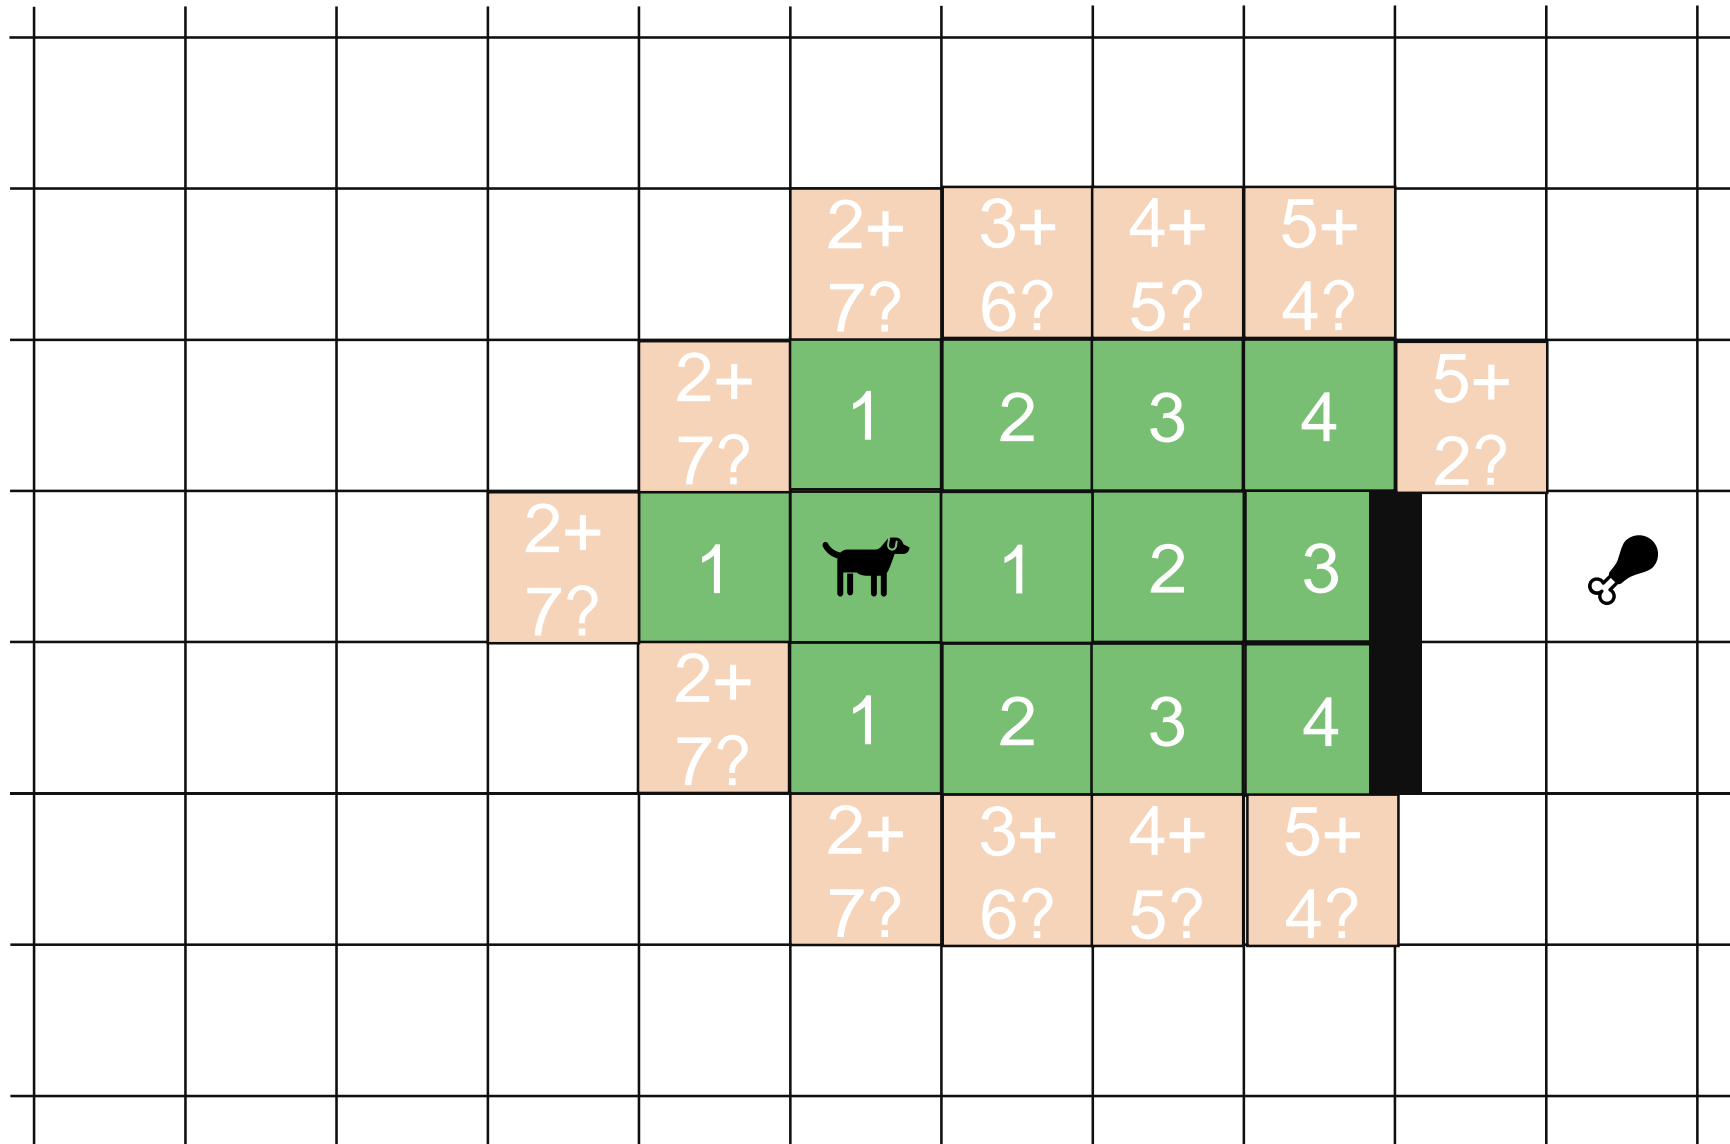
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




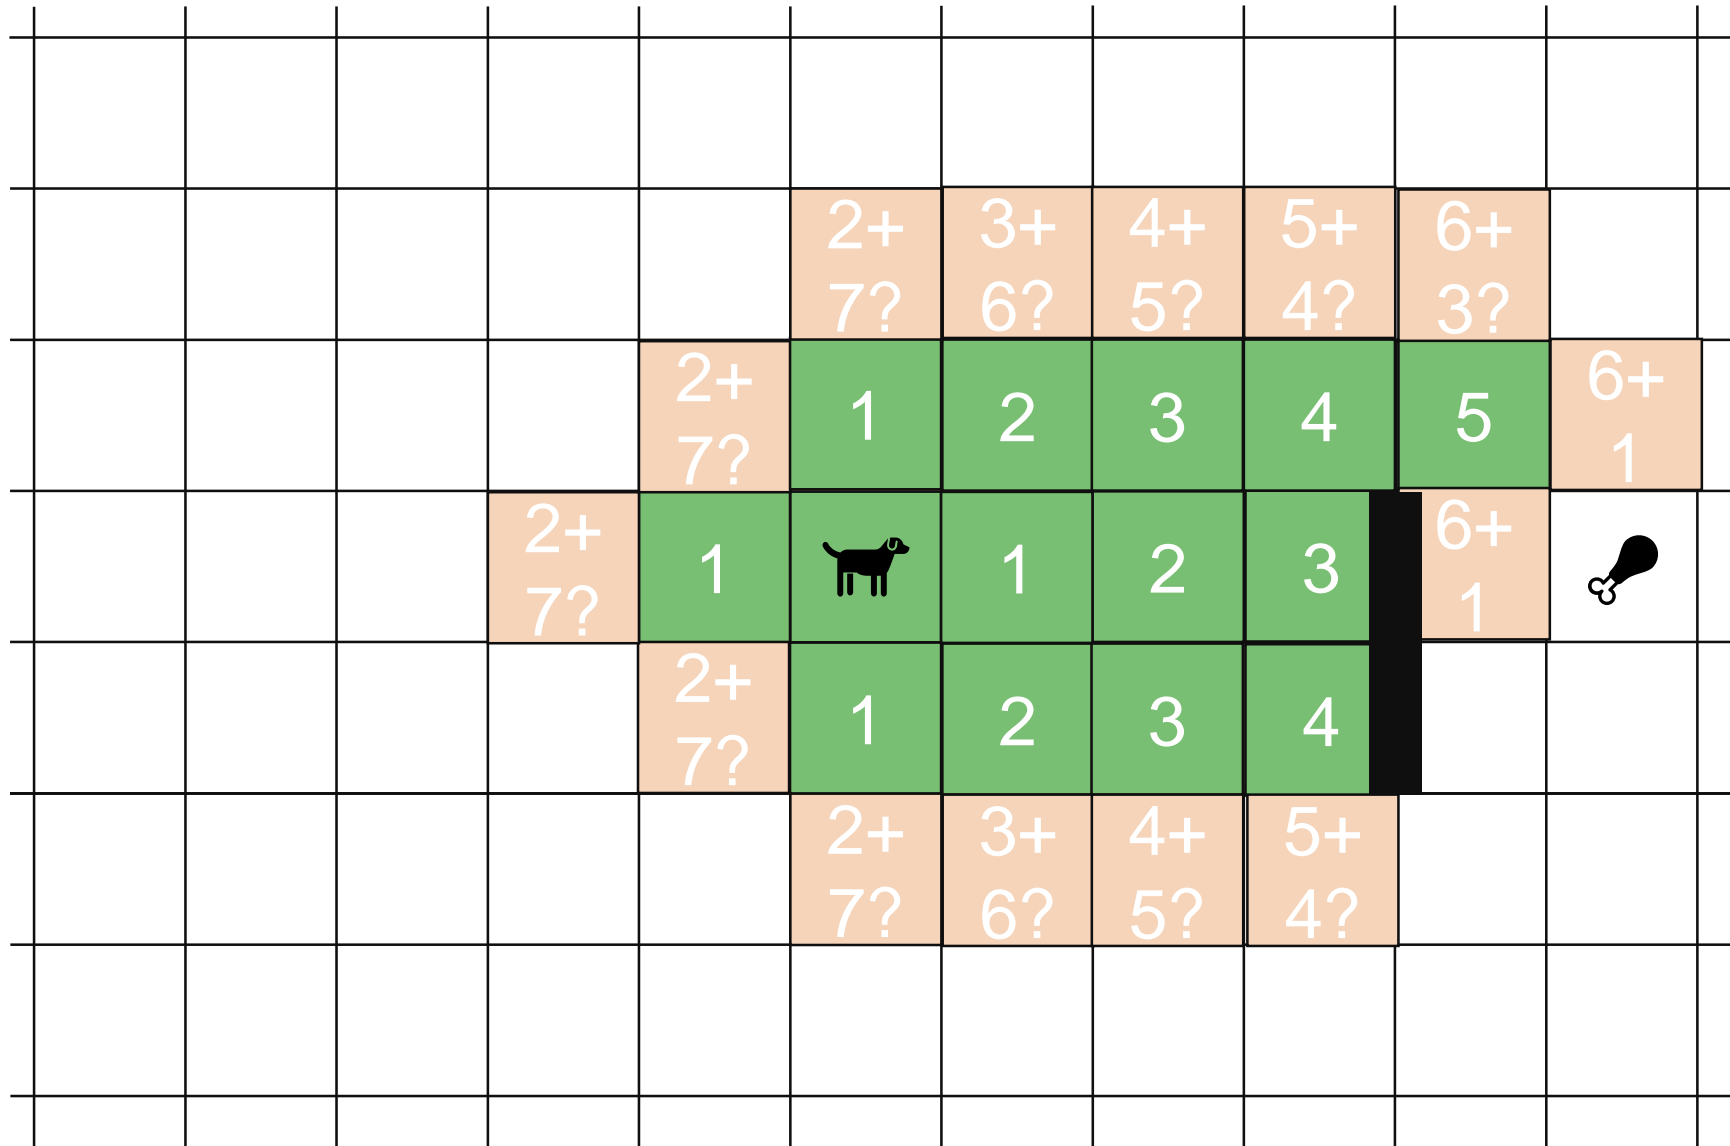
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




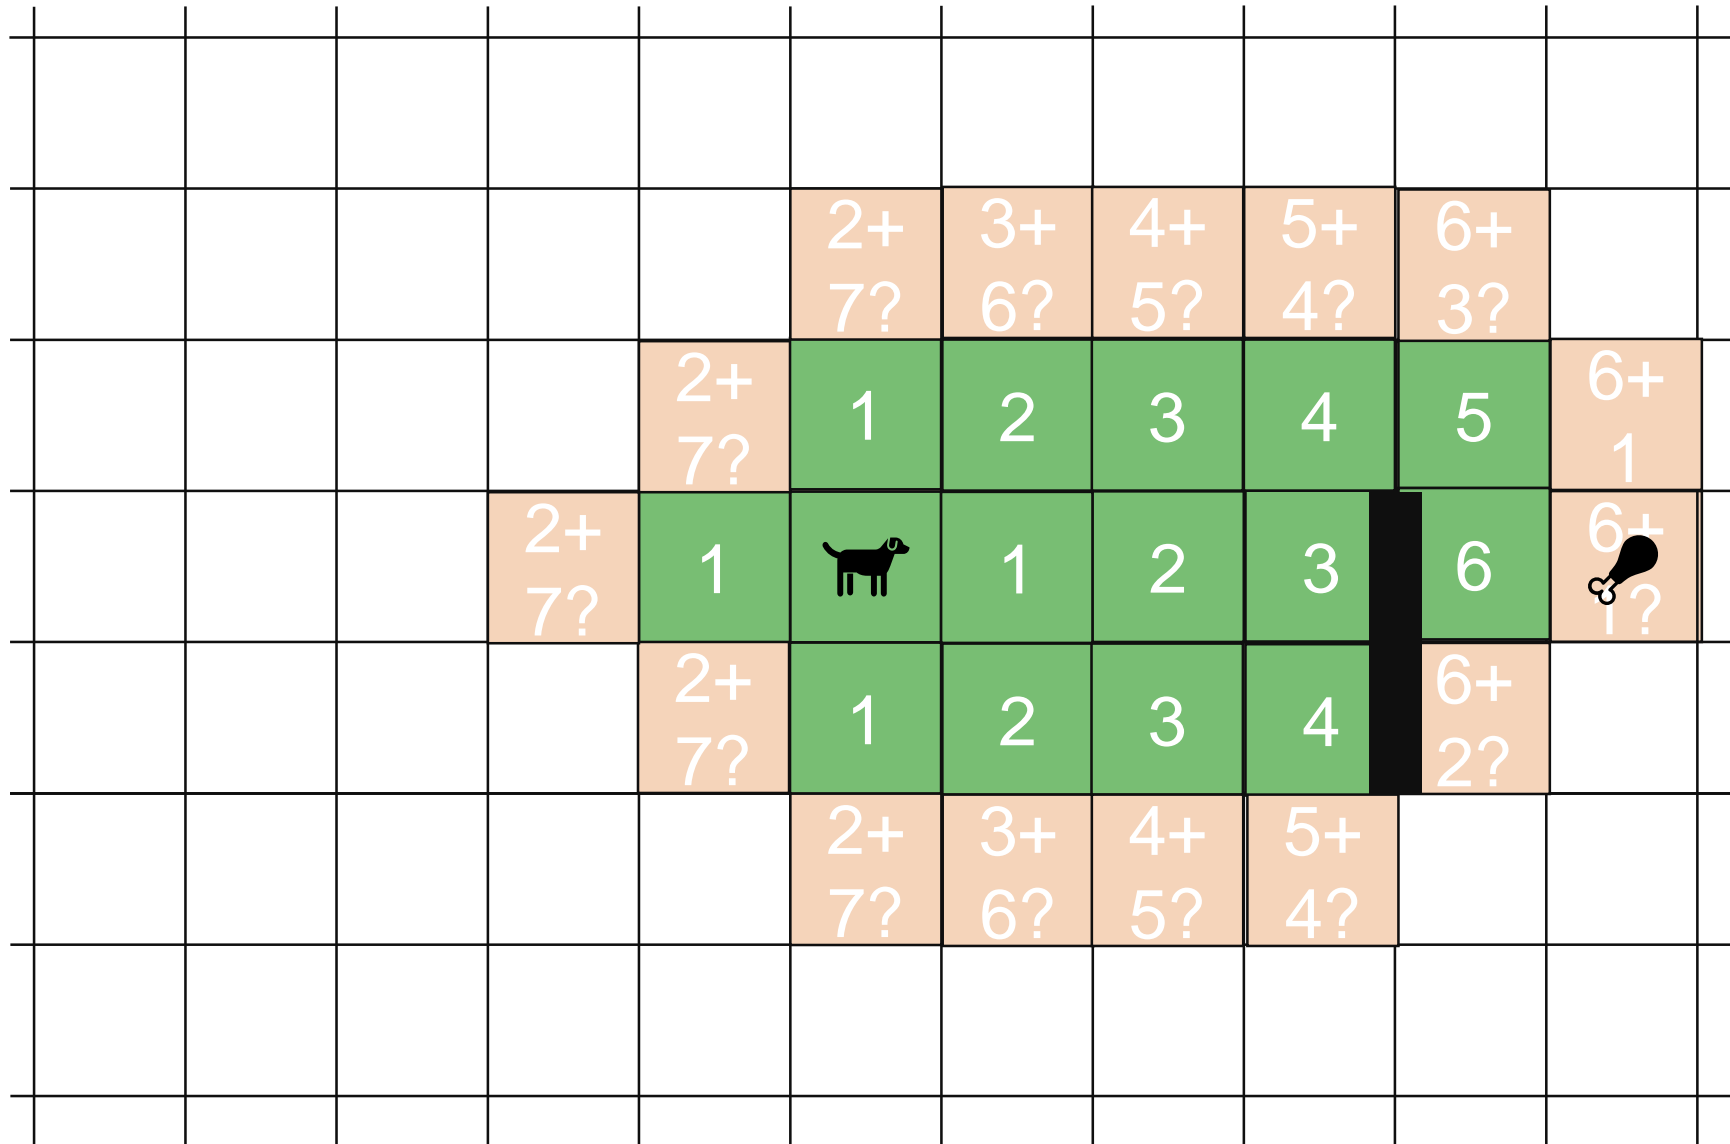
Greedyly continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




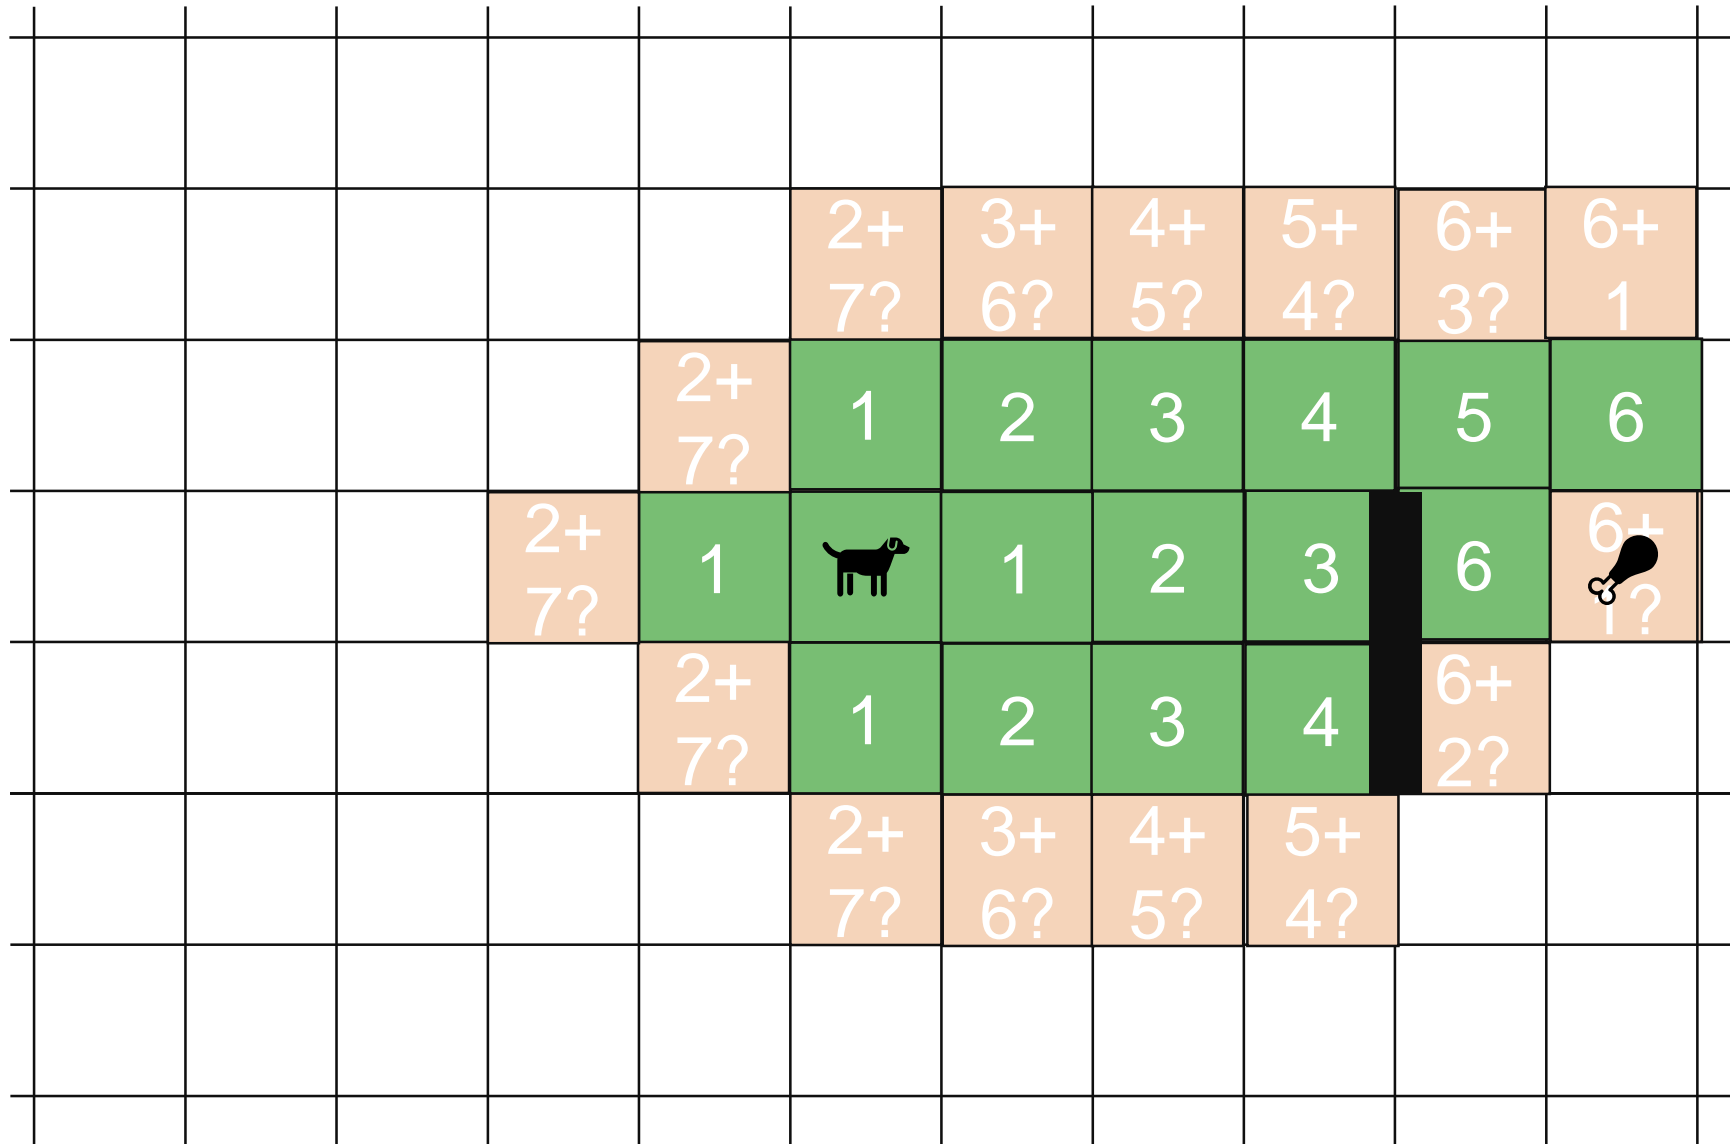
Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance




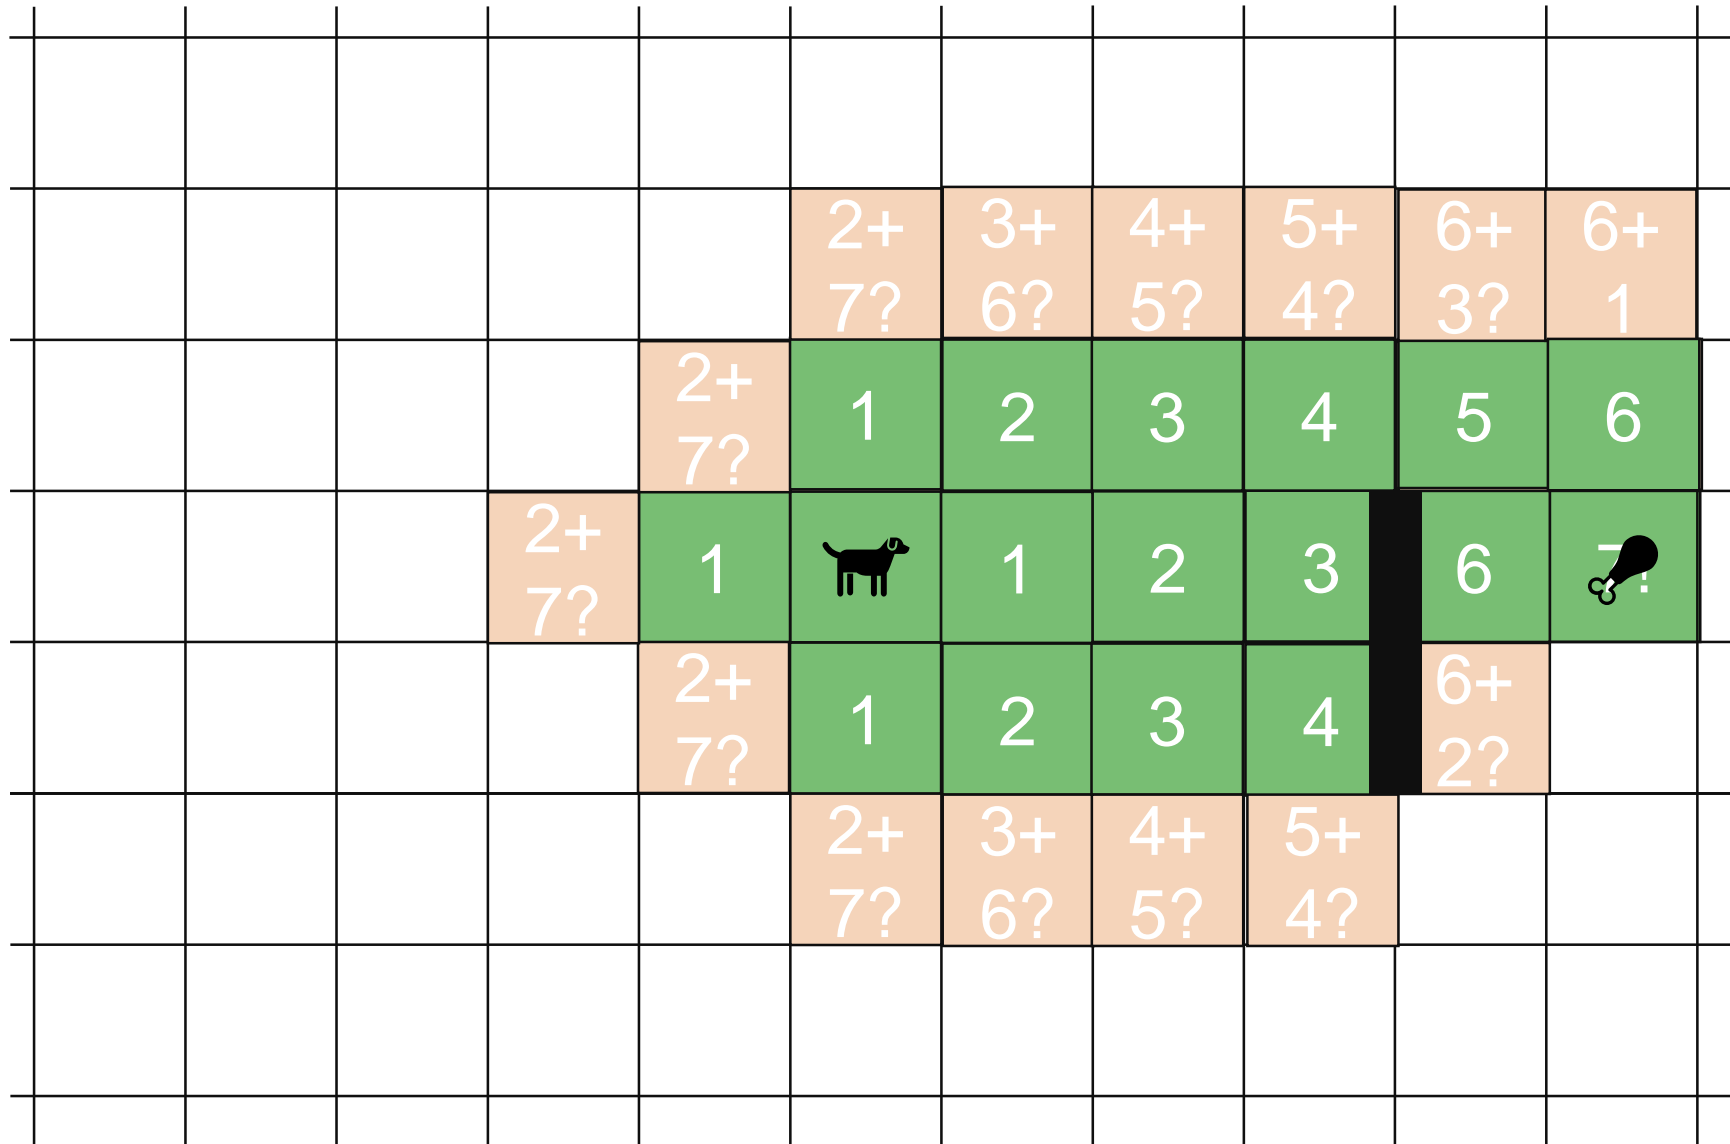
Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

A*: Use row + column distance





Greedy continue
at smallest

distance(s, u)

+

wanna be
futureCost(u,t)

What Dijkstra would have done




Greedy continue
at smallest

distance(s, u)

+


~~wait a bit
future cost (u, t)~~

$$\text{priority}(u) = \boxed{\text{distance}(s, u)} + \boxed{\text{wanna be futureCost}(u, t)}^*$$



priority of the path that ends in u

$$\text{priority}(u) = \boxed{\text{distance}(s, u)} + \boxed{\text{heuristic}(u, t)}$$




priority of the path that ends in u

depends on the problem

A heuristic is like a good educated guess.

$$\text{priority}(u) = \boxed{\text{distance}(s, u)} + \boxed{\text{heuristic}(u, t)}$$



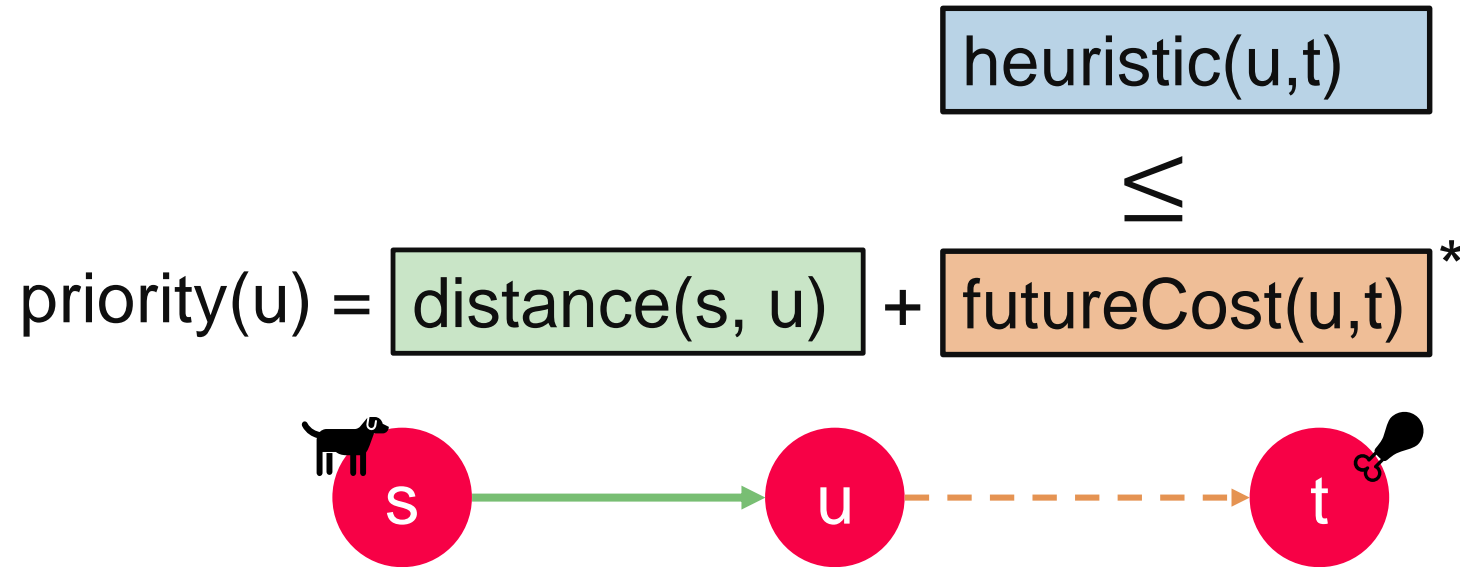
priority of the path that ends in u

depends on the problem


A heuristic is like a good educated guess.

! A* only works if your heuristic is good!
! good, if it underestimates the future cost !

Definition: A heuristic is **admissible** if it underestimates the future cost, that is,
 $h(u) = h(u, t) \leq d(u, t)$ holds.




Definition: A heuristic is **admissible** if it underestimates the future cost, that is, $h(u) = h(u, t) \leq d(u, t)$ holds.

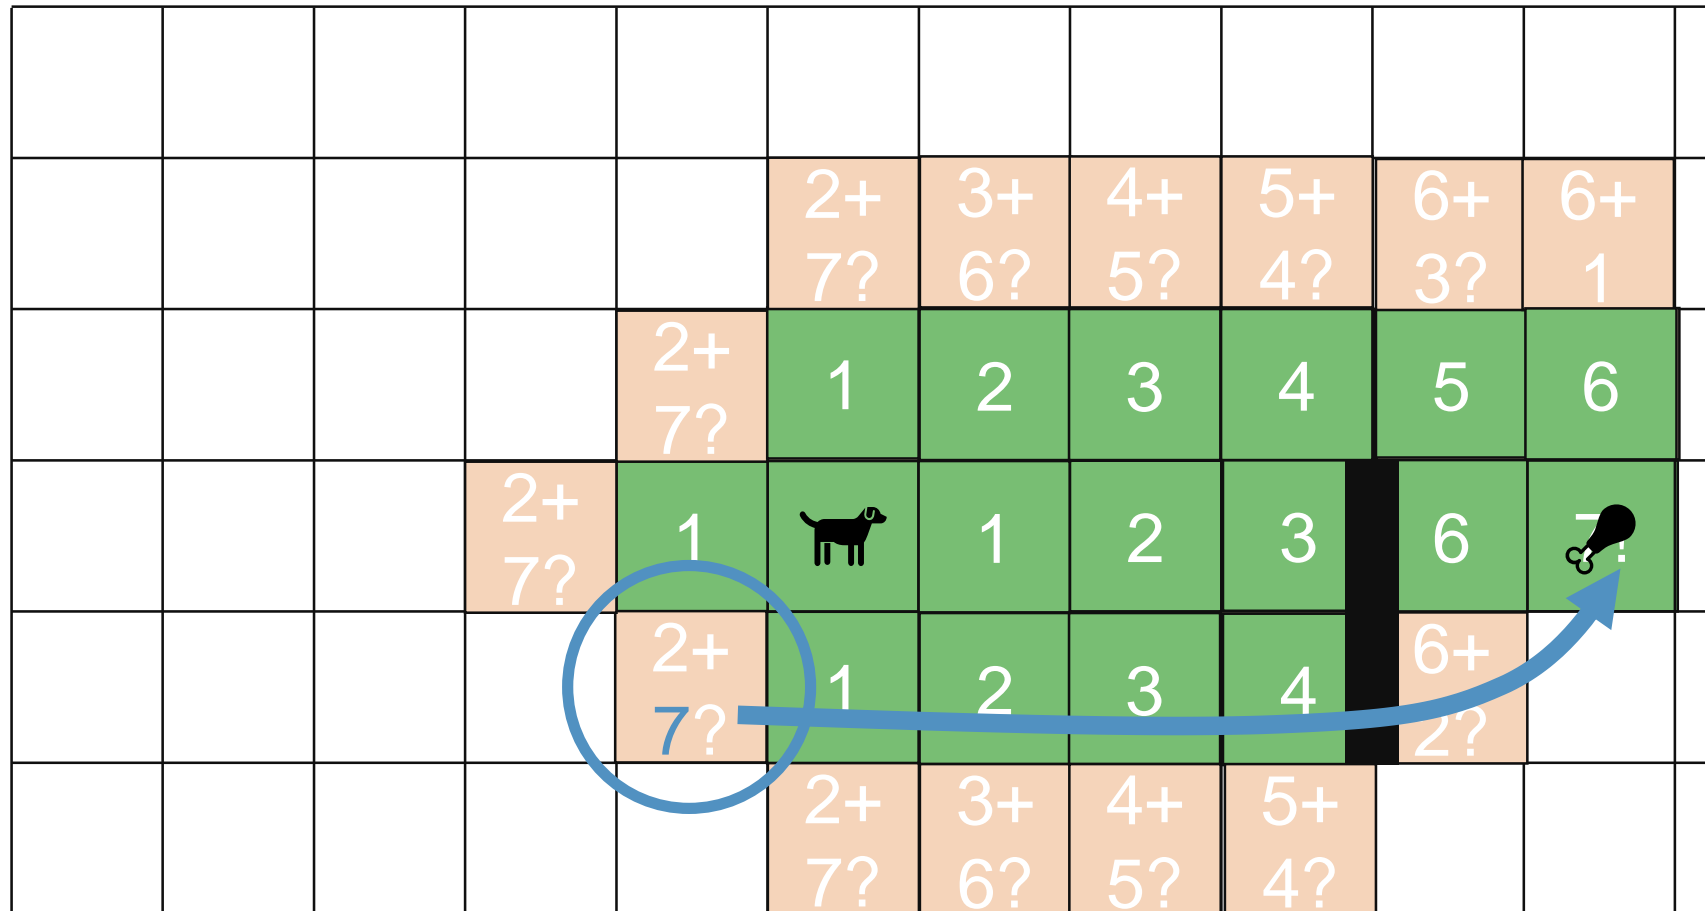
$$\text{priority}(u) = \boxed{\text{distance}(s, u)} + \boxed{\text{futureCost}(u, t)}^* \leq \boxed{\text{heuristic}(u, t)}$$



Example for an admissible heuristic:

As the crow flies




“Ignore path, as even if everything goes perfect (from here onwards), **this path** is still terrible...”



Input: $G(V, E, W)$ start point:  s , end point:  t

Initialize $S = \{s\}$, $g(s) = 0$, $g(V \setminus N(s)) = \infty$, $g(v) = w(s, v)$, $v \in N(s)$



S : expanded/closed vertices
 $V \setminus S$: the open vertices

s

$g(v) +$
 $h(u)?$

$g(v)$ is the length of the best
known (!) path from s to v

Input: $G(V, E, W)$ start point:  s , end point:  t

Initialize $S = \{s\}$, $g(s) = 0$, $g(V \setminus N(s)) = \infty$, $g(v) = w(s, v)$, $v \in N(s)$

While $t \notin S$ **do**

$u = \operatorname{argmin}_{u \in V \setminus S} \{g(u) + h(u, t)\}$

For v s.t. $\{u, v\} \in E$ **do**

$\text{temp} = \min\{g(v), g(u) + w(u, v)\}$

If $\text{temp} < g(v)$ then:

$g(v) = \text{temp}$

$S = S \setminus \{v\}$

//does nothing if $v \notin S$

$S = S \cup \{u\}$

$g(v) +$
 $h(u)?$

$g(v)$ is the length of the best
known (!) path from s to v

$h(u, t)$ is a heuristic guess for
the path from u to t .

S might decrease!

Input: $G(V, E, W)$ start point:  s , end point:  t

Initialize $S = \{s\}$, $g(s) = 0$, $g(V \setminus N(s)) = \infty$, $g(v) = w(s, v)$, $v \in N(s)$

While $t \notin S$ do

→ $u = \operatorname{argmin}_{u \in V \setminus S} \{g(u) + h(u, t)\}$

For v s.t. $\{u, v\} \in E$ do

$\text{temp} = \min\{g(v), g(u) + w(u, v)\}$

If $\text{temp} < g(v)$ then:

$g(v) = \text{temp}$

→ $S = S \setminus \{v\}$

//does nothing if $v \notin S$

$S = S \cup \{u\}$

S might decrease!

$g(v) + h(u)?$

$g(v)$ is the length of the best known (!) path from s to v

$h(u, t)$ is a heuristic guess for the path from u to t .

Analysis of A^*


Does it always terminate?

Does it always find an optimal path?


Runtime? Space consumption?




path finding with a heuristic




Downside: A node might be expanded multiple times




Downside: A node might be expanded multiple times



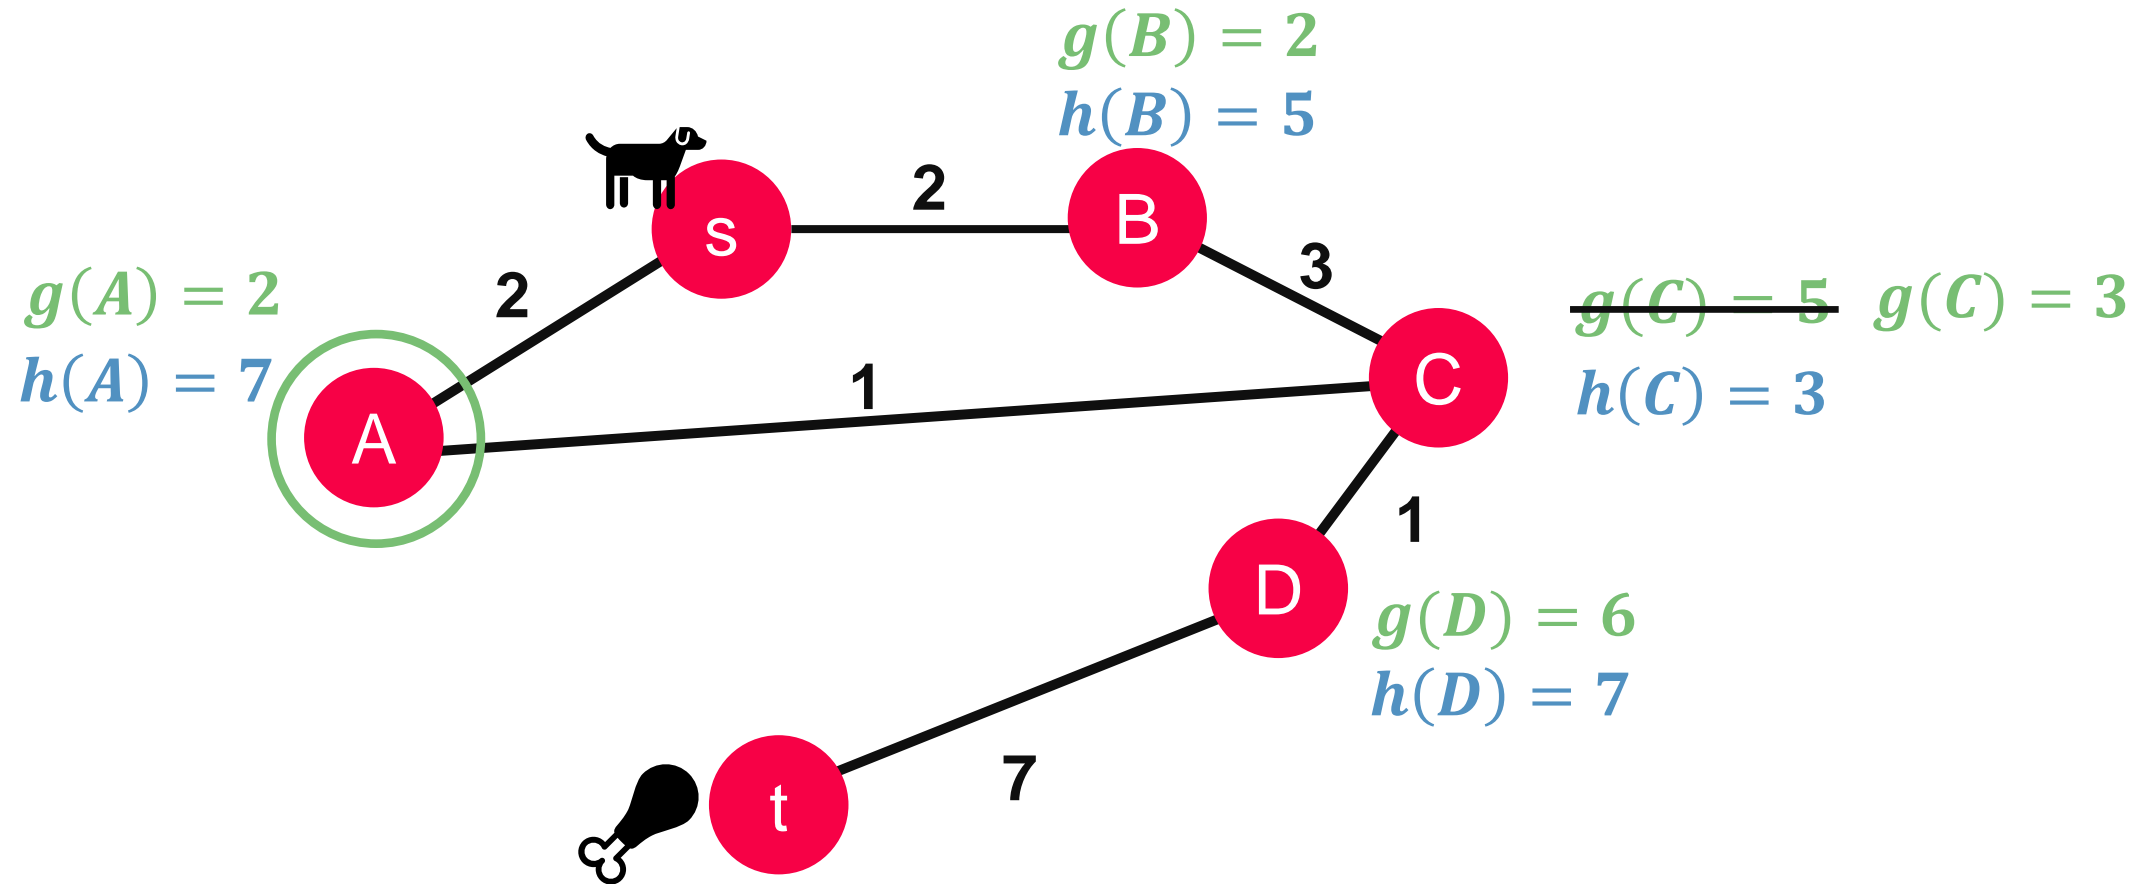
Downside: A node might be expanded multiple times




Downside: A node might be expanded multiple times




Downside: A node might be expanded multiple times



Downside: A node might be expanded multiple times



Downside: A node might be expanded multiple times




Order of expansion: S, B, C, A, C, ...

Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,


Example of Dijkstra:



Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,


Example of Dijkstra:



Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,


Example of Dijkstra:



Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,

Example of Dijkstra:




(in Dijkstra a node is only expanded once)

Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,

Example of Dijkstra:




(in Dijkstra a node is only expanded once)

- in Dijkstra v will not be expanded before $g(v)$ is finalized,

Why can a node be expanded more than once in A*?

In A*, similar to Dijkstra the $g(v)$ value can change,

Example of Dijkstra:



(in Dijkstra a node is only expanded once)

- in Dijkstra v will not be expanded before $g(v)$ is finalized,
- in A* a node might be expanded **before** its $g(v)$ value is **finalized**!

Proof:

- In each iteration of A* a new *acyclic path* is generated because:
 - Node added the first time: new path
 - node is “promoted” (If $temp < g(v)$): path is new because it is shorter



Proof:

- In each iteration of A* a new *acyclic path* is generated because:
 - Node added the first time: new path
 - node is “promoted” (If $temp < g(v)$): path is new because it is shorter
- $\#steps \leq \#acyclicPaths$



Proof:

- In each iteration of A* a new *acyclic path* is generated because:
 - Node added the first time: new path
 - node is “promoted” (If $temp < g(v)$): path is new because it is shorter
- $\#steps \leq \#acyclicPaths$
- $\#acyclicPaths$ starting from s is finite




Is A* Guaranteed to Terminate? **Yes!**

Proof:

- In each iteration of A* a new *acyclic path* is generated because:
 - Node added the first time: new path
 - node is “promoted” (If $temp < g(v)$): path is new because it is shorter
- $\#steps \leq \#acyclicPaths$
- $\#acyclicPaths$ starting from s is finite

This also means we find t



The A*-Algorithm with an **admissible** heuristic is:


- Complete: if a path to t exists, it finds one ✓
- Optimal: if a path to t exists, it always finds an optimal path



Wikipedia: When A^* terminates its search, it has found a path from start to goal whose actual cost is lower than the **estimated cost** of any path from start to goal through any open node (nodes in $V \setminus S$). When the *heuristic is admissible*, those **estimates are optimistic**, so A^* can safely ignore those nodes because they cannot possibly lead to a cheaper solution than the one it already has.



Wikipedia: When A^* terminates its search, it has found a path from start to goal whose actual cost is lower than the **estimated cost** of any path from start to goal through any open node (nodes in $V \setminus S$). When the *heuristic is admissible*, those **estimates are optimistic**, so A^* can safely ignore those nodes because they cannot possibly lead to a cheaper solution than the one it already has.




Don't expand u because large

$$\boxed{g(u)} + \boxed{\text{heuristic}(u,t)}$$

h is admissible, so better than the shortest path from u to t



Wikipedia: When A^* terminates its search, it has found a path from start to goal whose actual cost is lower than the **estimated cost** of any path from start to goal through any open node (nodes in $V \setminus S$). When the *heuristic is admissible*, those **estimates are optimistic**, so A^* can safely ignore those nodes because they cannot possibly lead to a cheaper solution than the one it already has.




Don't expand u because large

$$g(u) + \text{heuristic}(u,t)$$



Wikipedia: When A^* terminates its search, it has found a path from start to goal whose actual cost is lower than the **estimated cost** of any path from start to goal through any open node (nodes in $V \setminus S$). When the *heuristic is admissible*, those **estimates are optimistic**, so A^* can safely ignore those nodes because they cannot possibly lead to a cheaper solution than the one it already has.



Don't expand u because large


$$\boxed{g(u)} + \boxed{\text{heuristic}(u,t)} \neq$$

$$\boxed{\text{distance}(s, u)} + \boxed{\text{heuristic}(u,t)}$$

Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an ***open*** node u on P with $g(u) = \text{distance}(s, u)$


Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an ***open*** node u on P with $g(u) = \text{distance}(s, u)$

Proof: Let $P = (s = v_0, v_1, v_2, \dots, v = v_k)$.



Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an ***open*** node u on P with $g(u) = \text{distance}(s, u)$


Proof: Let $P = (s = v_0, v_1, v_2, \dots, v = v_k)$.



$$C = \{v_i \in P \mid v_i \text{ closed}, g(v_i) = d(s, v_i)\} \neq \emptyset,$$

Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an ***open*** node u on P with $g(u) = \text{distance}(s, u)$

Proof: Let $P = (s = v_0, v_1, v_2, \dots, v = v_k)$.




$$C = \{v_i \in P \mid v_i \text{ closed}, g(v_i) = d(s, v_i)\} \neq \emptyset,$$

- Let v^* , the vertex in C with highest index. $v^* \neq v$.
- Let u be the successor of v^* in P (possibly $u = v$).

Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an **open** node u on P with $g(u) = \text{distance}(s, u)$

Proof: Let $P = (s = v_0, v_1, v_2, \dots, v = v_k)$.



$$C = \{v_i \in P \mid v_i \text{ closed}, g(v_i) = d(s, v_i)\} \neq \emptyset,$$

- Let v^* , the vertex in C with highest index. $v^* \neq v$.
- Let u be the successor of v^* in P (possibly $u = v$).

P is optimal path

$$g(u) \leq g(v^*) + w(v^*, u) = d(s, v^*) + w(v^*, u) = d(s, u) \leq g(u)$$


v^* expanded

definition of v^*

always

Lemma 1: Always, for every *open* node v and every optimal path P from s to v , there exists an **open** node u on P with $g(u) = \text{distance}(s, u)$

Proof: Let $P = (s = v_0, v_1, v_2, \dots, v = v_k)$.



$$C = \{v_i \in P \mid v_i \text{ closed}, g(v_i) = d(s, v_i)\} \neq \emptyset,$$

- Let v^* , the vertex in C with highest index. $v^* \neq v$.
- Let u be the successor of v^* in P (possibly $u = v$).

P is optimal path

$$g(u) \leq g(v^*) + w(v^*, u) = d(s, v^*) + w(v^*, u) = d(s, u) \leq g(u)$$

v^* expanded

definition of v^*

always

u has to be open by definition of C

Corollary: Suppose h is admissible and A^* has not terminated. Then, for any optimal path P from s to t , there is an open node u with

$$g(u) + \text{heuristic}(u,t) \leq \text{distance}(s,t)$$

Corollary: Suppose h is admissible and A^* has not terminated. Then, for any optimal path P from s to t , there is an open node u with

$$g(u) + \text{heuristic}(u,t) \leq \text{distance}(s,t)$$

Proof: By Lemma 1, we have open node $u \in P$ with $g(u) = \text{distance}(s, u)$

Corollary: Suppose h is admissible and A^* has not terminated. Then, for any optimal path P from s to t , there is an open node u with

$$g(u) + \text{heuristic}(u,t) \leq \text{distance}(s,t)$$

Proof: By Lemma 1, we have open node $u \in P$ with $g(u) = \text{distance}(s, u)$

$$g(u) + \text{heuristic}(u,t) = \text{distance}(s, u) + \text{heuristic}(u,t)$$

Corollary: Suppose h is admissible and A^* has not terminated. Then, for any optimal path P from s to t , there is an open node u with

$$g(u) + \text{heuristic}(u,t) \leq \text{distance}(s,t)$$

Proof: By Lemma 1, we have open node $u \in P$ with $g(u) = \text{distance}(s, u)$

$$g(u) + \text{heuristic}(u,t) = \text{distance}(s, u) + \text{heuristic}(u,t)$$

$$\begin{array}{c} h \text{ admissible} \end{array} \rightarrow \leq \text{distance}(s, u) + \text{futureCost}(u,t)$$

Corollary: Suppose h is admissible and A^* has not terminated. Then, for any optimal path P from s to t , there is an open node u with

$$g(u) + \text{heuristic}(u,t) \leq \text{distance}(s,t)$$

Proof: By Lemma 1, we have open node $u \in P$ with $g(u) = \text{distance}(s, u)$

$$g(u) + \text{heuristic}(u,t) = \text{distance}(s, u) + \text{heuristic}(u,t)$$

h admissible

$$\leq \text{distance}(s, u) + \text{futureCost}(u,t)$$

u on optimal path P

$$= \text{distance}(s,t)$$

The A*-Algorithm with an **admissible** heuristic is:

- Complete: if a path to t exists, it finds one ✓
- **Optimal: if a path to t exists, it always finds an optimal path**

The A*-Algorithm with an **admissible** heuristic is:

- Complete: if a path to t exists, it finds one ✓
- **Optimal: if a path to t exists, it always finds an optimal path**

Proof (Optimality):

Suppose A* terminates at t with a suboptimal path, i.e., in the last step we expanded t with

$$\boxed{g(t)} + \overbrace{\boxed{\text{heuristic}(t,t)}}^0 > \boxed{\text{distance}(s, t)}$$

Theorem 1

The A*-Algorithm with an **admissible** heuristic is:

- Complete: if a path to t exists, it finds one ✓
- **Optimal: if a path to t exists, it always finds an optimal path**

Proof (Optimality):

Suppose A* terminates at t with a suboptimal path, i.e., in the last step we expanded t with

$$\boxed{g(t)} + \overbrace{\boxed{\text{heuristic}(t,t)}}^0 > \boxed{\text{distance}(s, t)}$$

But, by the corollary, there existed just before “expanding t ”, there is an open node u on an optimal path with

$$\boxed{g(u)} + \boxed{\text{heuristic}(u,t)} \leq \boxed{\text{distance}(s, t)}$$



Contradiction!

Informal Remark: A^* with an admissible heuristic is also *Optimally efficient*. There is no other algorithm that finds the solution faster with the same heuristic (A^* expands a minimal number of vertices)

The remark is not about the number of expansions!

We have seen: A node might be expanded several times with an admissible heuristic

Def: A heuristic is **consistent** if for every edge $\{u, v\} \in E$ we have

$$h(u) \leq w(u, v) + h(v)$$

With an **admissible** and **consistent** heuristic, A* is guaranteed to find an optimal path without expanding any node more than once.


We have seen: A node might be expanded several times with an admissible heuristic

Def: A heuristic is **consistent** if for every edge $\{u, v\} \in E$ we have

$$h(u) \leq w(u, v) + h(v)$$

With an **admissible** and **consistent** heuristic, A* is guaranteed to find an optimal path without expanding any node more than once.

“as the crow flies”-heuristic is consistent



We have seen: A node might be expanded several times with an admissible heuristic


Def: A heuristic is **consistent** if for every edge $\{u, v\} \in E$ we have

$$h(u) \leq w(u, v) + h(v)$$

With an **admissible** and **consistent** heuristic, A* is guaranteed to find an optimal path without expanding any node more than once.

Heuristic in previous example was not consistent

$$h(A) = 7 \not\leq w(A, C) + h(C) = 4$$



$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

Dijkstra

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

**Only possible heuristic,
might be faster than Dijkstra**

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

Only finds best paths, super fast!
Requires perfect knowledge! 

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

**Also fast (not admissible),
but might not find the paths!**



Remarks:

- path construction similar to Dijkstra (we will not detail on this):
 - always remember a pointer to the best (known) predecessor
 - \rightarrow sufficient to reconstruct the path to the start node s .
 - A^* works in directed graphs

(no comment)
(Wikipedia)



- It depends on...
 - the heuristic
 - Implementation of `argmin` (e.g. binary heap)
 - Implementation of set S (e.g., as array)

worst case



- It depends on...
 - the heuristic
 - Implementation of argmin (e.g. binary heap)
 - Implementation of set S (e.g., as array)

Bottom line:

- worst case runtime is not interesting for A^*
- Often much better than Dijkstra
- In applications **space** is the bottleneck ($g(v) + h(v)$ is stored for each visited v)

worst case



- It depends on...
 - the heuristic
 - Implementation of argmin (e.g. binary heap)
 - Implementation of set S (e.g., as array)

Bottom line:

- worst case runtime is not interesting for A^*
- Often much better than Dijkstra
- In applications **space** is the bottleneck ($g(v) + h(v)$ is stored for each visited v)

15-Puzzle: Search space has a node for each configuration:
 $16! = 20.922.789.888.000$ vertices!

worst case



- It depends on...
 - the heuristic
 - Implementation of argmin (e.g. binary heap)
 - Implementation of set S (e.g., as array)

worst case




Bottom line:

- worst case runtime is not interesting for A^*
- Often much better than Dijkstra
- In applications **space** is the bottleneck ($g(v) + h(v)$ is stored for each visited v)

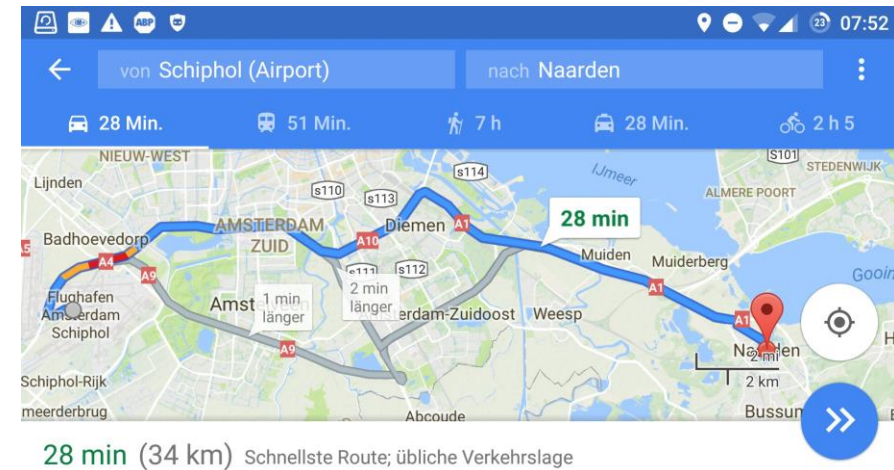
15-Puzzle: Search space has a node for each configuration:
 $16! = 20.922.789.888.000$ vertices!

→ Memory bounded heuristic search: Iterative deepening A^* , ...




Why doesn't Google Maps Pre-Compute Paths?

- How many nodes are in the Google Maps graph?
Answer: About $N = 75$ million
- How many sets of paths would they need to generate?
Answer: (roughly) N^2
- How long would that take?
Answer: 6×10^{15} seconds,
or... 190 million years
(numbers might be outdated)



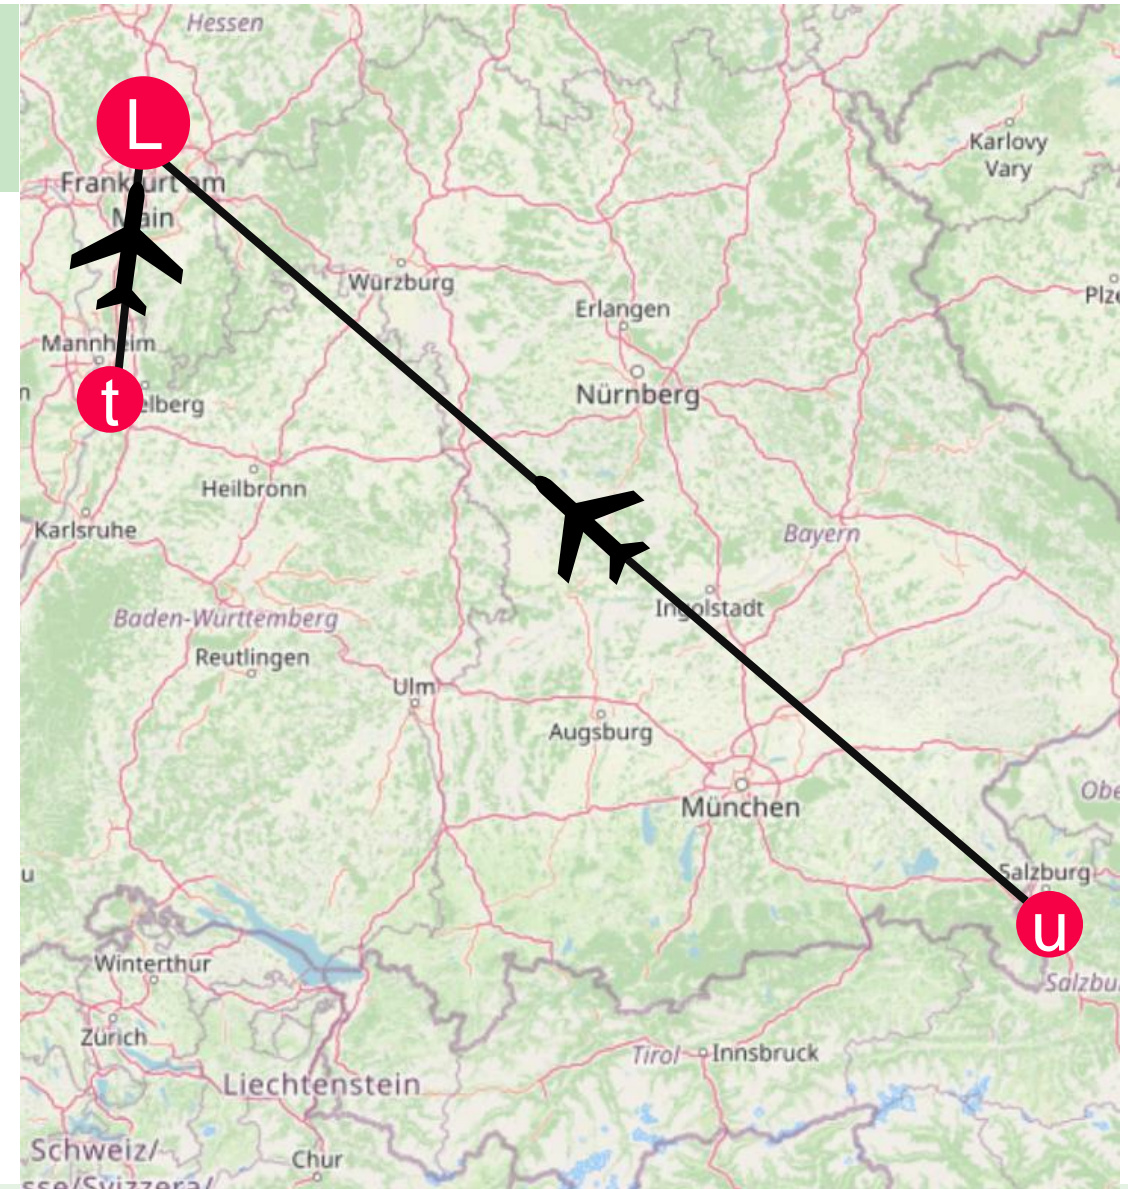
What Heuristics Could Google Maps Use?

- as the crow flies (straight-line distance)
 - calculate the straight-line distance from u to t , and divide by the speed on the fastest highway




Landmark heuristic with landmark L :

$$h(u, t) = |d(u, L) - d(t, L)| \leq d(u, t)$$




Landmark heuristic with landmark L :

$$h(u, t) = |d(u, L) - d(t, L)| \leq d(u, t)$$



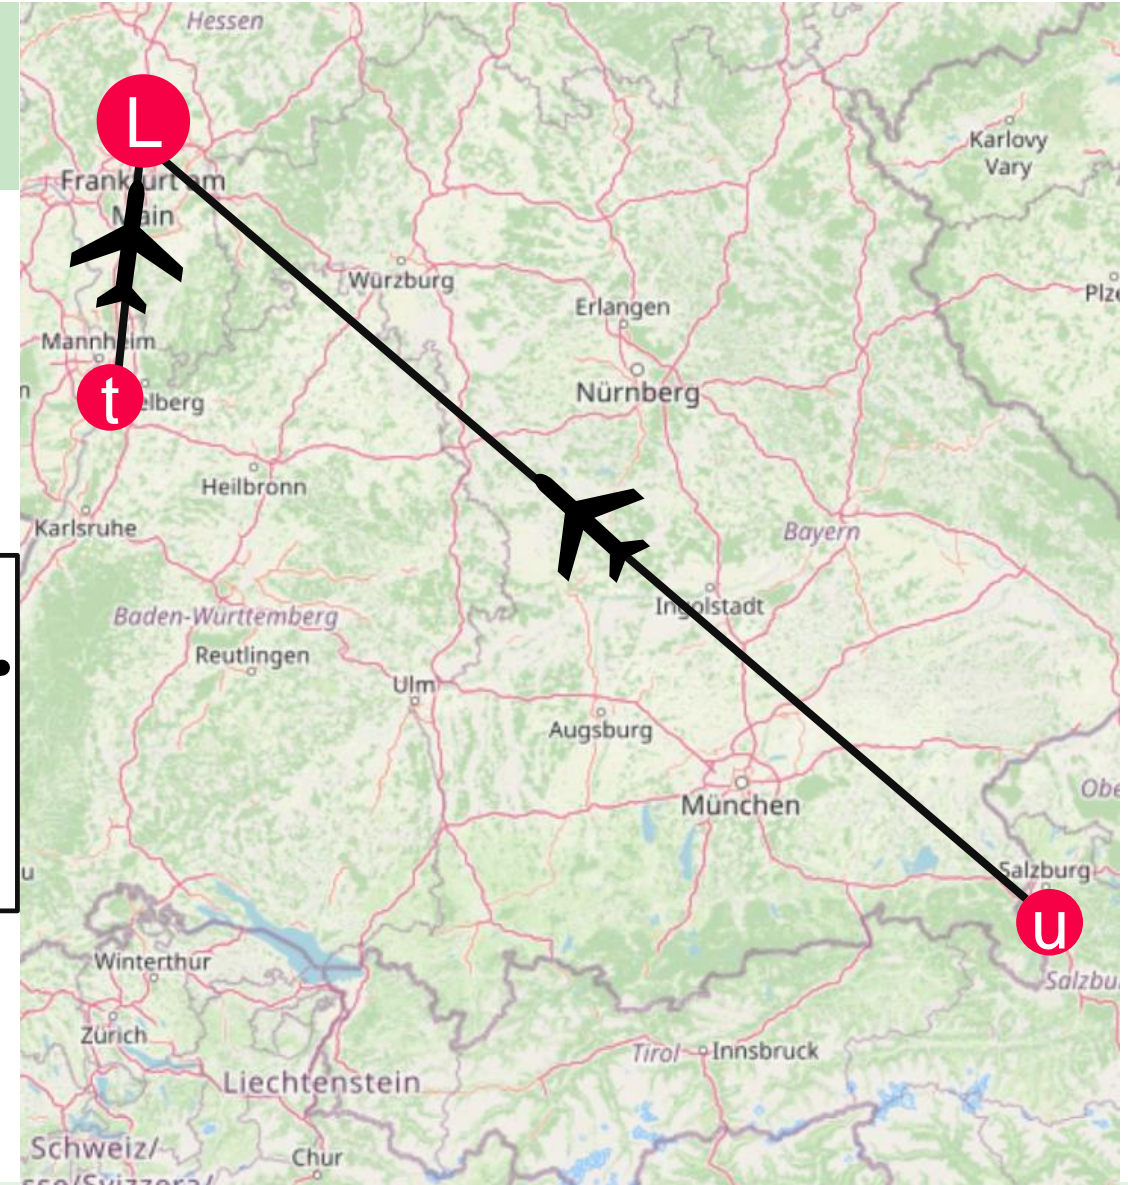
Landmark heuristic with landmark L :

$$h(u, t) = |d(u, L) - d(t, L)| \leq d(u, t)$$




What does Google probably do?

- all of these heuristics and more?
- One can use multiple heuristics and choose the max



Landmark heuristic with landmark L :

$$h(u, t) = |d(u, L) - d(t, L)| \leq d(u, t)$$

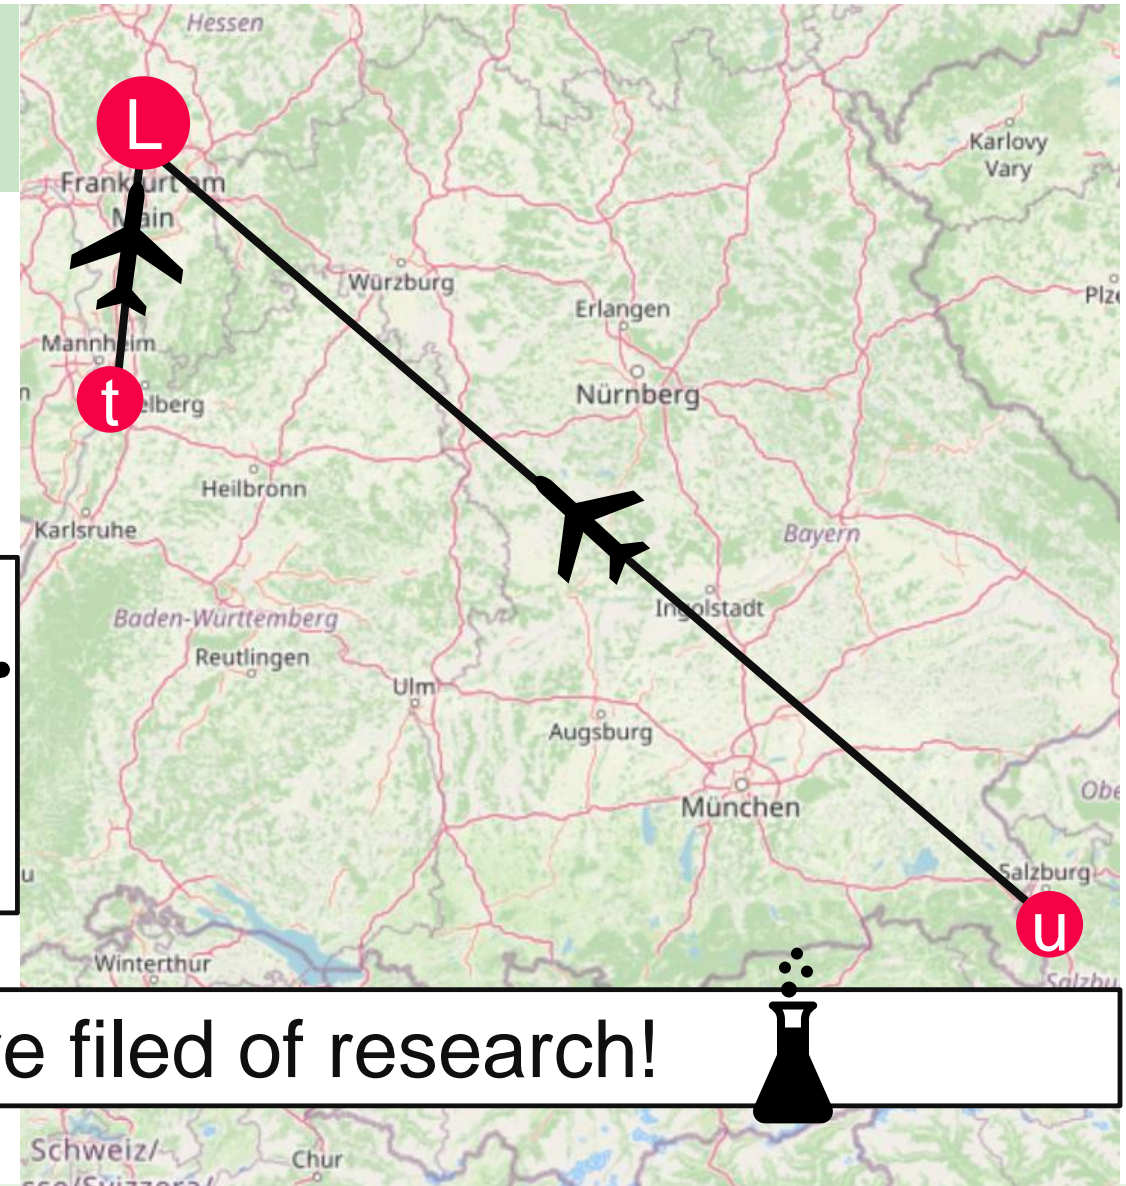


What does Google probably do?

- all of these heuristics and more?
- One can use multiple heuristics and choose the max



Finding good heuristics is an active field of research!



One source and One Destination

- Use A* Search Algorithm (For Unweighted as well as Weighted Graphs)

The Heuristic matters!

One source and One Destination

- Use A* Search Algorithm (For Unweighted as well as Weighted Graphs)

The Heuristic matters!

One Source, All Destination

- BFS (For Unweighted Graphs)
- Use Dijkstra (For Weighted Graphs without negative weights)
- Use Bellman Ford (For Weighted Graphs with negative weights)

One source and One Destination

- Use A* Search Algorithm (For Unweighted as well as Weighted Graphs)

The Heuristic matters!

One Source, All Destination

- BFS (For Unweighted Graphs)
- Use Dijkstra (For Weighted Graphs without negative weights)
- Use Bellman Ford (For Weighted Graphs with negative weights)

Between every pair of nodes

- Floyd-Warshall
- Johnson's Algorithm (not discussed)

Link to animated version of A*: <https://algorithms.discrete.ma.tum.de/>

Thank you