Information Security

System Security 2 - Side Channels and Microarchitectural Attacks

November 17, 2023

Side-channel Attacks

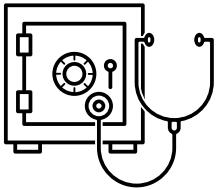

- Safe software infrastructure does not mean safe execution
- Information leaks because of the underlying hardware
- Exploit unintentional information leakage by side-effects

Power consumption

Execution time

CPU caches

Side channels in Software

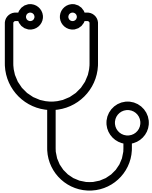

- Side channels also exist in software
- Can be used for attacks
- Usually timing differences

Example: PIN Comparison

- Trivial approach: Compare each digit until a difference

```
int check_pin(char* input) {
    const char* correct = "1234";
    for(int i = 0; i < 4; i++) {
        if(correct[i] != input[i]) {
            // digit differs, abort
            return ERROR;
        }
    }
    // PIN is correct
    return OK;
}
```

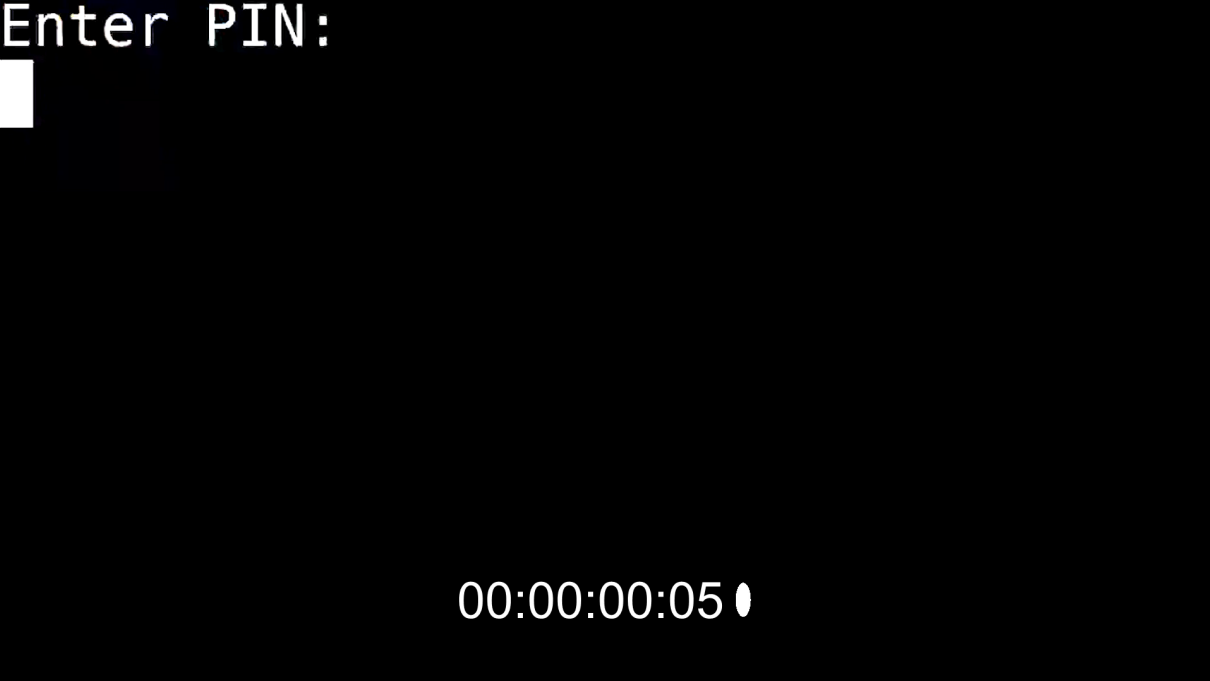
Enter PIN:
00:00:00:05

Example: PIN Comparison

- Measuring the execution times for different PINs

PIN	Time
0000	
1000	
2000	
3000	
...	...

- If digit is correct, next digit is checked → longer execution time
- 10 tries (maximum) to get a digit

Example: PIN Comparison

- Measuring the execution times for different PINs

PIN	Time
1000	
1100	
1200	
1300	
...	...

- Repeat for every digit
- Longest execution time reveals correct digit

Example: PIN Comparison

- Maximum 10 measurements per digit
- 4-digit PIN: 40 tries
- Brute force: 10 000 tries
- Simple side channel reduces tries by factor 250

Example: PIN Comparison

- Many functions can be implemented with constant runtime

```
int check_pin(char* input) {
    const char* correct = "1234";
    int same = 0;
    for(int i = 0; i < 4; i++) {
        same |= correct[i] - input[i];
    }
    return (same == 0);
}
```

- Sometimes, there is still a side channel in the hardware

Architecture and Microarchitecture

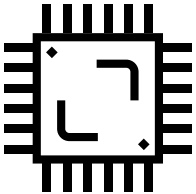

- Instruction Set Architecture (ISA) is an abstract model of a computer (x86, ARMv8, SPARC, . . .)
- Serves as the interface between hardware and software
- Microarchitecture is an actual implementation of the ISA

●●●

Microarchitectural Elements

- Modern CPUs contain multiple microarchitectural elements

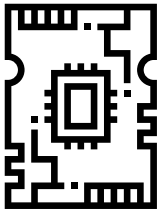

Caches and buffer

Predictor

- Transparent for the programmer
- Optimize program execution
- Timing differences $\rightarrow$ side-channel leakage

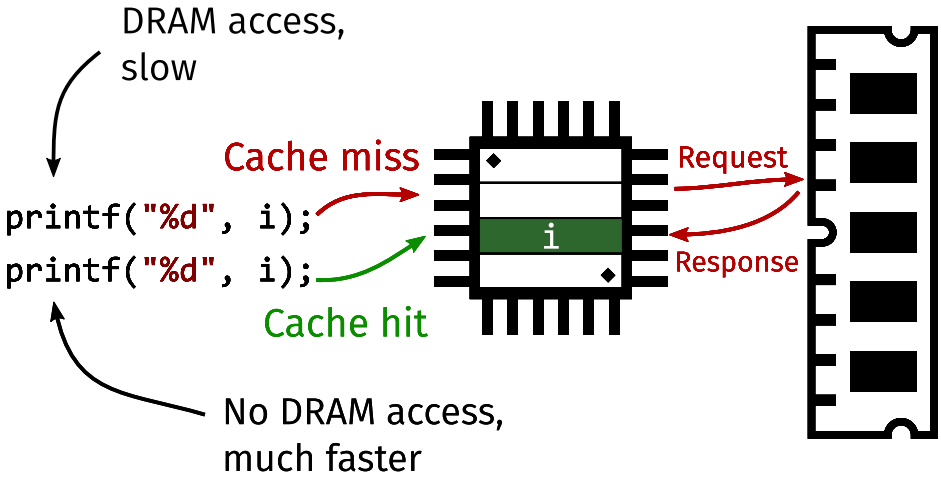
DRAM access,
slow
Cache miss
printf("%d", i);
printf("%d", i);
Cache hit
No DRAM access,
much faster
i
Request
Response

Memory Access Latency

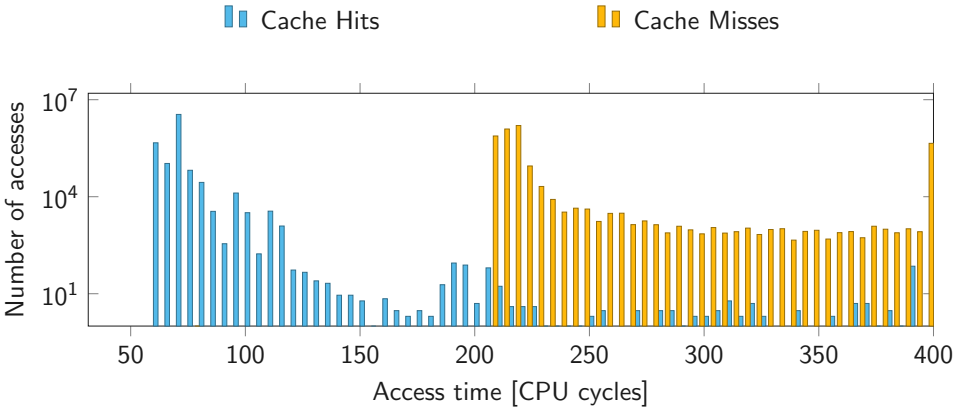

Cache hierarchy

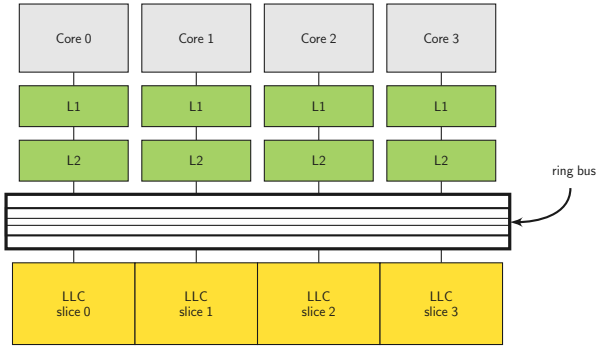

- L1 and L2 are private
- Last-level cache is
  - divided into slices
  - shared across cores
  - inclusive

Set-associative Last-level Cache

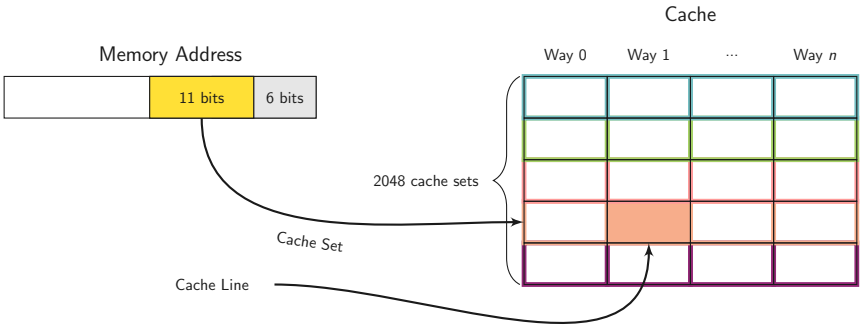

- Location in cache depends on the physical address of data
- Bits 6 to 16 determine the cache set
- A cache set has multiple ways to store the data
- A way inside a cache set is a cache line, determined by the cache replacement policy

Flush+Reload

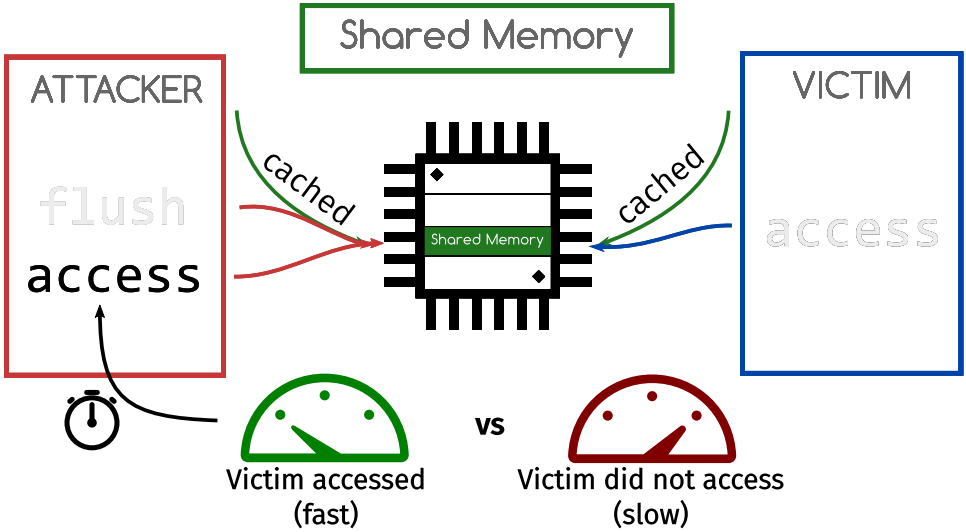

Flush+Reload

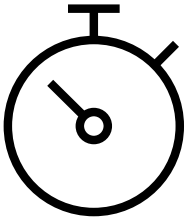

```
struct shared_data[256];
[...]
return shared_data[84];
[...]
```

- Flush+Reload over memory locations

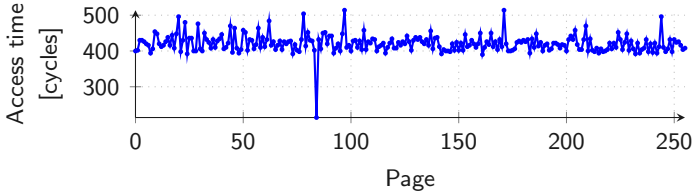

- Accessed index results in faster access time

Flush+Reload on Keystrokes

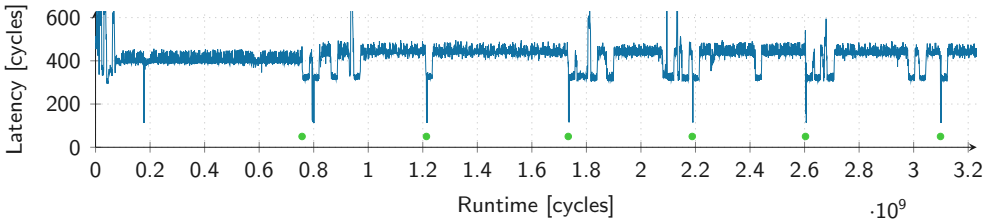

- Key presses trigger code execution in shared library (e.g., `libgdk`)
- Flush+Reload does not reveal actual key, only time difference between keys
- $\rightarrow$ Recover text with machine learning

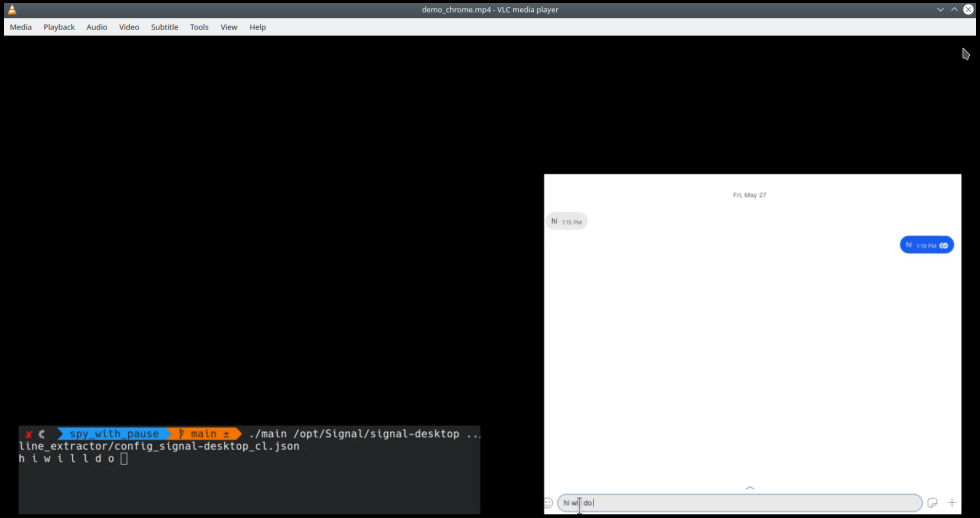
demo_chrome.mp4 - VLC media player
Media Playback Audio Video Subtitle Tools View Help
spy_with_pause main ± ./main /opt/Signal/signal-desktop ...
line_extractor/config_signal-desktop_cl.json
h i w i l l d o
Fri, May 27
hi 1:15 PM
hi 1:15 PM
hi will do

A Double Fetch

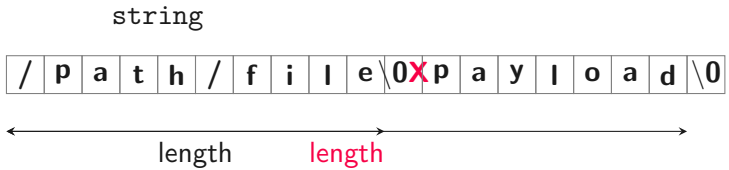

Thread 1

```
strcpy(string, "/path/file\0payload");
open(string, O_CREAT);

// <switch to kernel>

int len = strlen(string);
char* local = malloc(len + 1);

strcpy(local, string);
// <memory corruption>
```

Thread 2

schedule

schedule

```
string[10] = 'X';
```

Flush+Reload

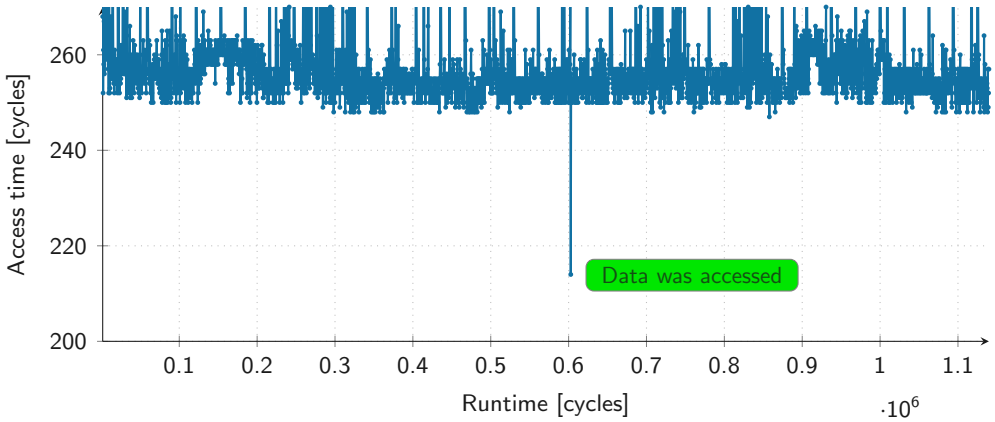

Double-fetch Detection

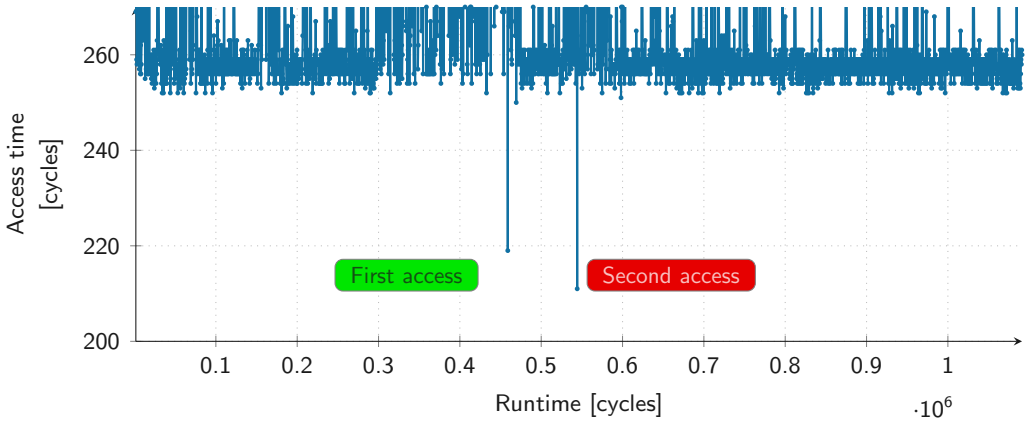

Cache-based Trigger

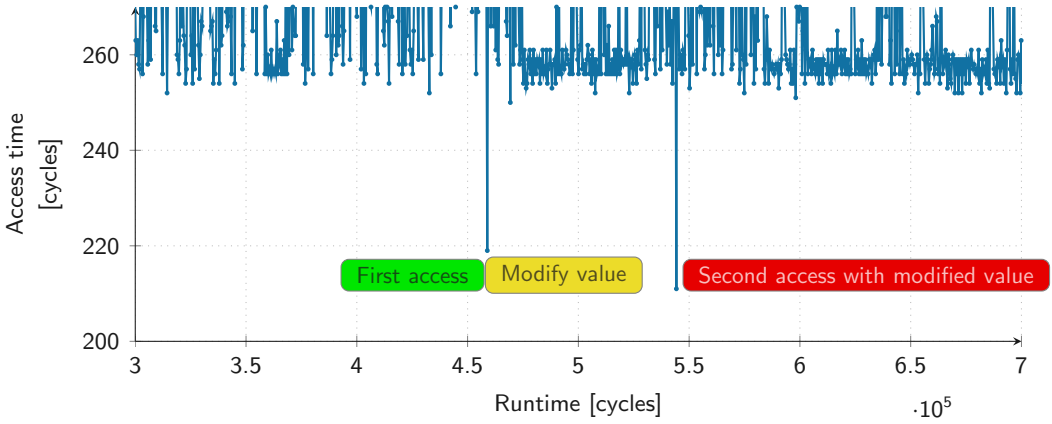

Prime+Probe

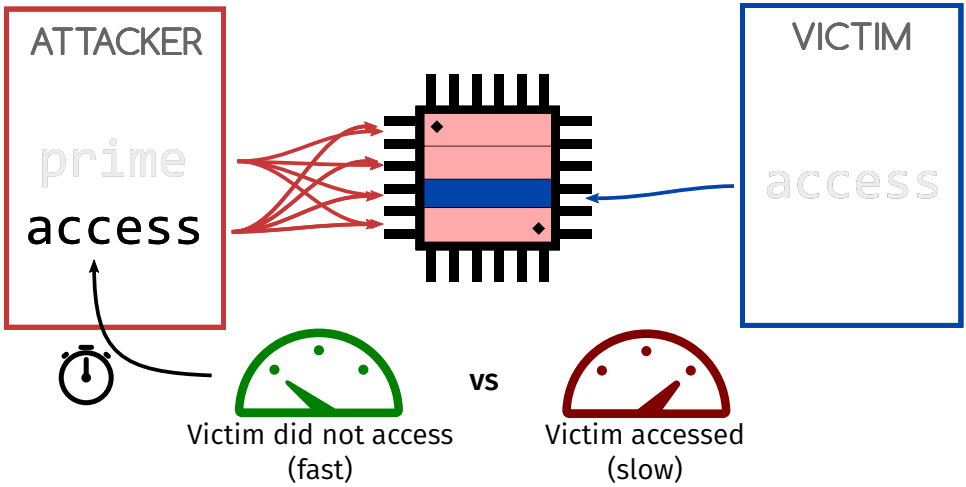
ATTACKER
prime
access
VICTIM
access
vs
Victim did not access
(fast)
Victim accessed
(slow)

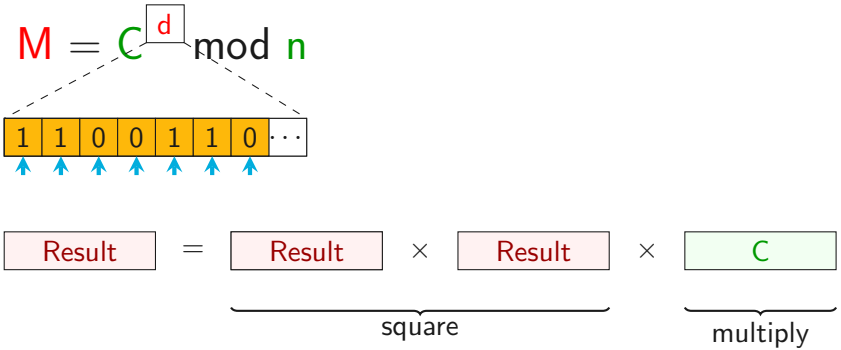
M = C^d mod n
1 1 0 0 1 1 0 ···
Result = Result × Result × C
square
multiply

Measured Trace

Raw Prime+Probe trace...

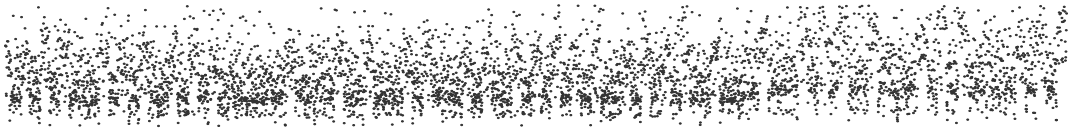

Measured Trace

...processed with a simple moving average...

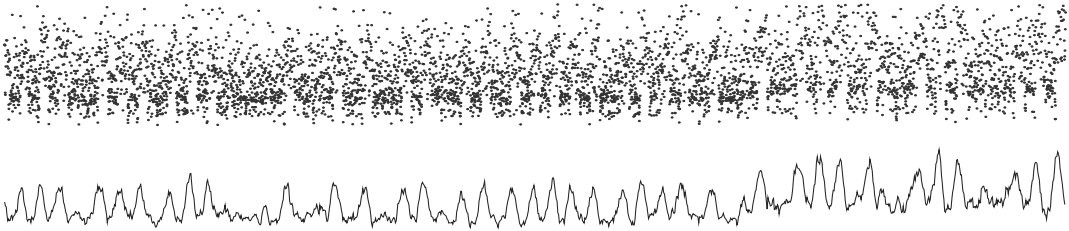

Measured Trace

...allows to clearly see the bits of the exponent

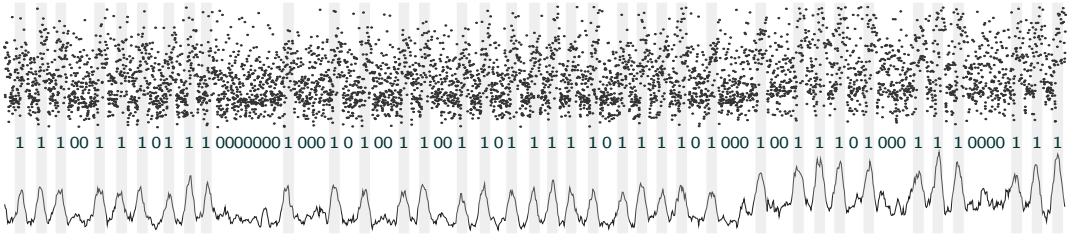

Intel claiming it is out of scope

Side-channel Researcher

Covert channel

What is a covert channel?

- Two programs would like to communicate but are not allowed to do so
  - either because there is no communication channel...
  - ...or the channels are monitored and programs are stopped on communication attempts
- Use side channels and stay stealthy

Covert channel

Sending Data

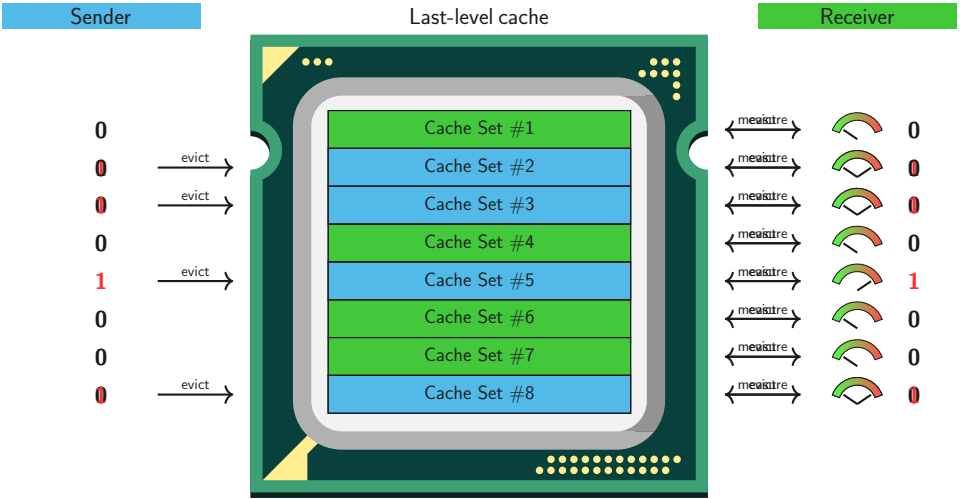

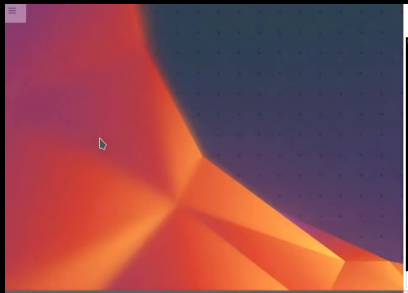

local — Konsole

File Edit View Bookmarks Settings Help

michael@michael-hp ~ %

local

sender (ec2) — Konsole

File Edit View Bookmarks Settings Help

ubuntu@ip-172-31-51-32 ~ %

sender (ec2)

receiver (ec2) — Konsole

File Edit View Bookmarks Settings Help

ubuntu@ip-172-31-58-32 ~ %

receiver (ec2)

sender (ec2) — Konsole <2>

File Edit View Bookmarks Settings Help

-= slurm 0.4.3 on ip-172-31-51-32 =-

```
                x
                x
               xx
              xxx
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
                xx
               xx
               x
               x
```

Active Interface: eth0 Interface Speed: unknown

Current RX Speed: 0.05 KB/s	Current TX Speed: 0.33 KB/s
Graph Top RX Speed: 0.98 KB/s	Graph Top TX Speed: 2.64 KB/s
Overall Top RX Speed: 0.98 KB/s	Overall Top TX Speed: 2.64 KB/s
Received Packets: 84	Transmitted Packets: 82
MBytes Received: 0.005 MB	MBytes Transmitted: 0.019 MB
Errors on Receiving: 0	Errors on Transmission: 0

sender (ec2)

receiver (ec2) — Konsole <2>

File Edit View Bookmarks Settings Help

-= slurm 0.4.3 on ip-172-31-58-32 =-

```
                x
                x
                x
                x
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
                xx
                x
                x
                x
```

Active Interface: eth0 Interface Speed: unknown

Current RX Speed: 0.05 KB/s	Current TX Speed: 0.29 KB/s
Graph Top RX Speed: 0.36 KB/s	Graph Top TX Speed: 0.84 KB/s
Overall Top RX Speed: 0.36 KB/s	Overall Top TX Speed: 0.84 KB/s
Received Packets: 38	Transmitted Packets: 36
MBytes Received: 0.002 MB	MBytes Transmitted: 0.010 MB
Errors on Receiving: 0	Errors on Transmission: 0

receiver (ec2)

HELLO FROM THE OTHER SIDE (DEMO): VIDEO STREAMING OVER CACHE COVERT CHANNEL

Line 17, Column 1 INSERT Soft Tabs: 4 (8) UTF-8 Normal

Other Microarchitectural Elements

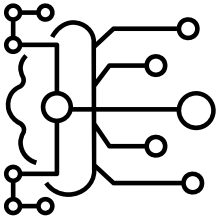

- Multiple other elements with timing differences
  - TLB
  - DRAM
  - Memory Bus
  - Execution Units
  - ...
- Many side-channel attacks exploiting them

Meta Data

- So far, only memory accesses
- Meta data, no actual data
- Sufficient to deduce data...
- ...if memory accesses are secret dependent

Building Block

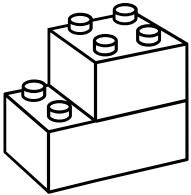

- Side channels can be part of an attack
- Also for conventional memory corruption attacks
- Side channels as building blocks
  - Required information (e.g., break ASLR)
  - Additional information (e.g., length of password)
  - Covertly transmit information
  - Transient-execution attacks

Meltdown

- Meltdown is a CPU vulnerabilities
- Discovered in 2017 by multiple independent teams
- Allows breaking the process isolation
- Side-channel attack is a core building block

Hardware Isolation

- Kernel is isolated from user space
- This isolation is a combination of hardware and software
- User applications cannot access anything from the kernel
- There is only a well-defined interface → syscalls

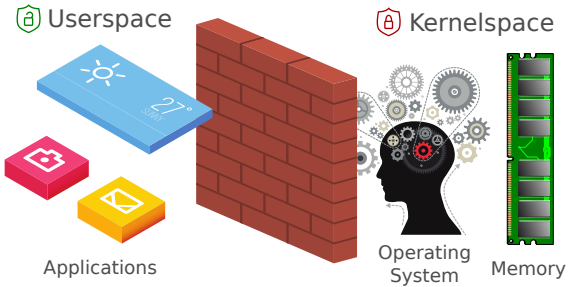

In-Order Execution

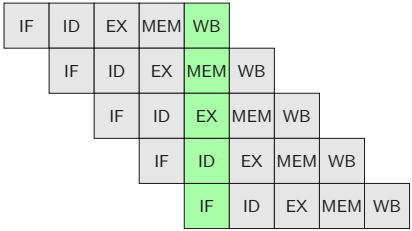

- Instructions are...
  - fetched (`IF`) from the L1 Instruction Cache
  - decoded (`ID`)
  - executed (`EX`) by execution units
- Memory access is performed (`MEM`)
- Architectural register file is updated (`WB`)

In-Order Execution

- Instructions are executed in-order
- Pipeline stalls when stages are not ready
- If data is not cached, we need to wait

Out-of-order Execution

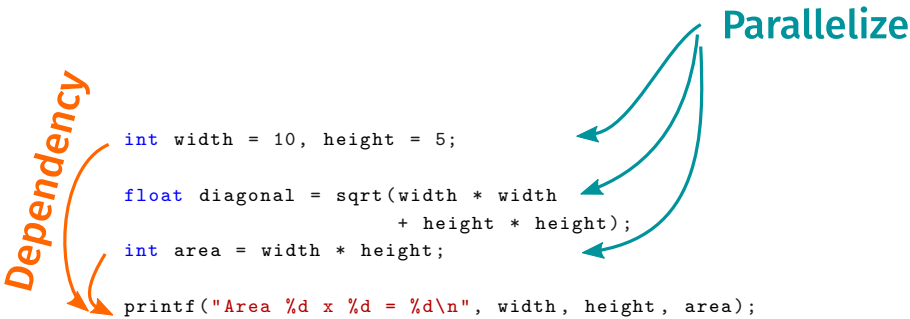

Out-of-Order Execution

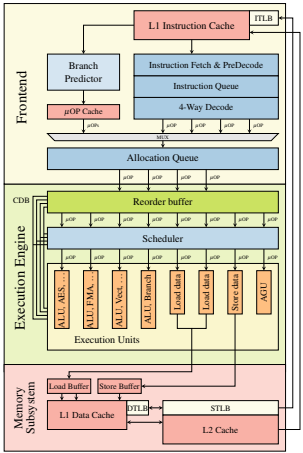

Instructions are

- fetched and decoded in the front-end
- dispatched to the backend
- processed by individual execution units

Out-of-Order Execution

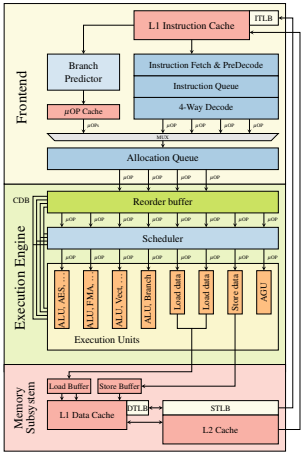

Instructions

- are executed out-of-order
- wait until their dependencies are ready
  - Later instructions might execute prior earlier instructions
- retire in-order
  - State becomes architecturally visible
- Exceptions are checked during retirement
  - Flush pipeline and recover state

The state does not become architecturally visible but . . .

Building the Code

- New code

```
*(volatile char*) 0;
array[84 * 4096] = 0;
```

- `volatile` because compiler was not happy

```
warning: statement with no effect [-Wunused-value]
         *(char*)0;
```

- Static code analyzer is still not happy

```
warning: Dereference of null pointer
         *(volatile char*)0;
```

- Flush+Reload over all pages of the array

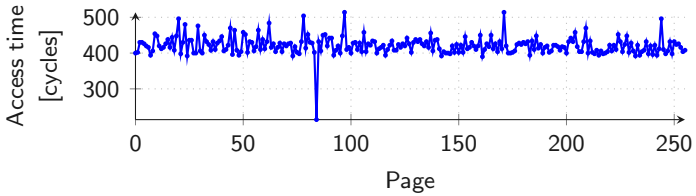

- "Unreachable" code line was actually executed
- Exception was only thrown afterwards

Building the Code

- Out-of-order instructions leave microarchitectural traces
  - We can see them for example in the cache
- Give such instructions a name: transient instructions
- We can indirectly observe the execution of transient instructions

Loading an address

Building the Code

- Add another layer of indirection to test

```
char data = *(char*) 0xffffffff81a000e0;
array[data * 4096] = 0;
```

- Then check whether any part of `array` is cached

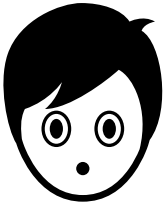

- Flush+Reload over all pages of the array

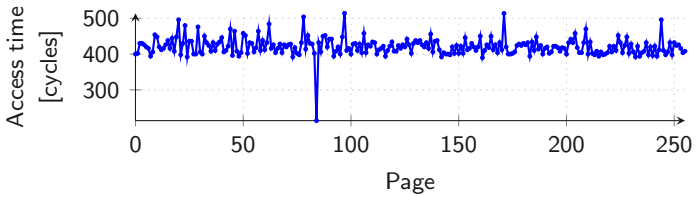

- Index of cache hit reveals data
- Permission check is in some cases not fast enough

Kernel Direct-Physical Map

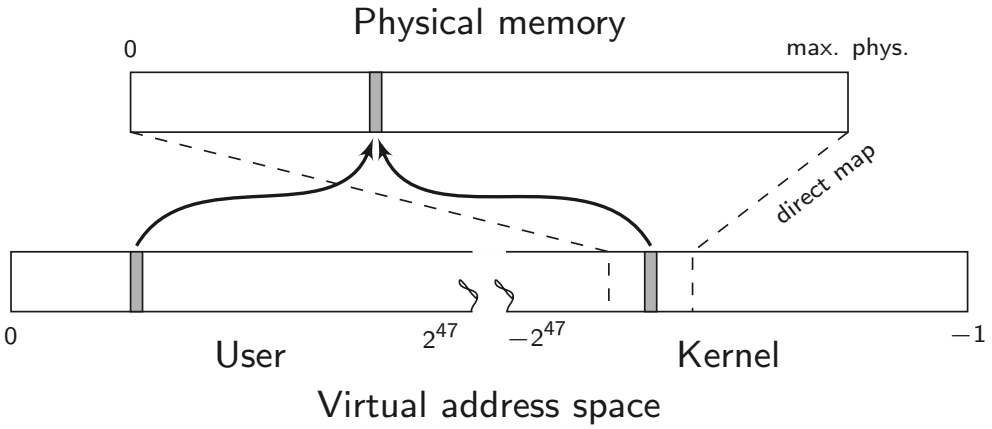

Meltdown

- Using out-of-order execution, we can read data at any address
- Index of cache hit reveals data
- Permission check is in some cases not fast enough
- Entire physical memory is typically accessible through kernel space


I SHIT YOU NOT
THERE WAS KERNEL MEMORY ALL OVER THE TERMINAL


There are no bugs,
just happy little accidents

Generalization

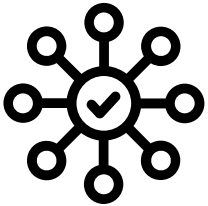

- Meltdown is a whole category of vulnerabilities
- Not only the user-accessible check
- Looking closer at the check...

Paging

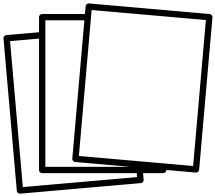

- CPU uses virtual address spaces to isolate processes
- Physical memory is organized in page frames
- Virtual memory pages are mapped to page frames using page tables

Address Translation on x86-64

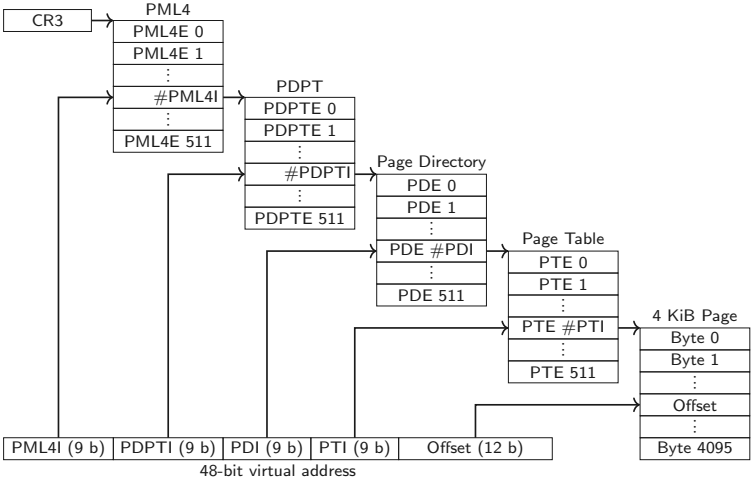

Page Table Entry

- User/Supervisor bit defines in which privilege level the page can be accessed

Page Table Entry

- Present bit is the next obvious bit

Foreshadow-NG

- An even worse bug $\rightarrow$ Foreshadow-NG/L1TF
- Exploitable from VMs
- Allows leaking data from the L1 cache
- Same mechanism as Meltdown
- Just a different bit in the PTE

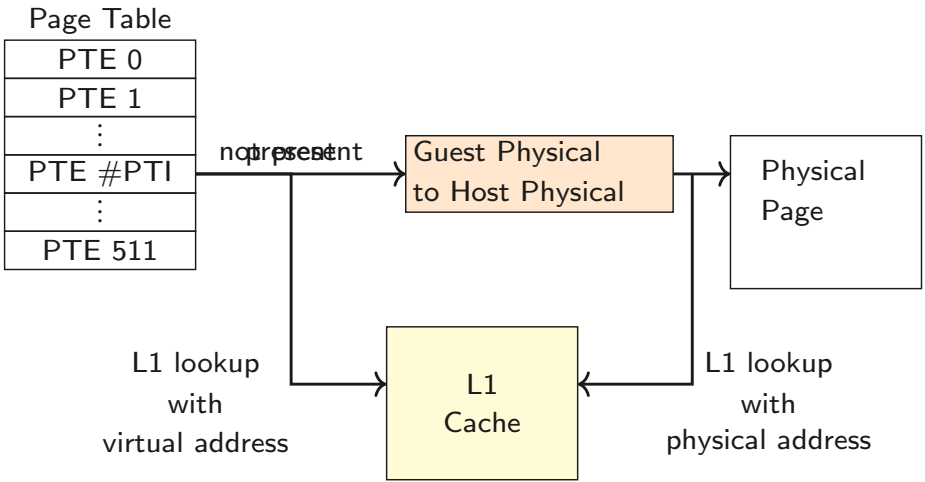
Page Table
PTE 0
PTE 1
PTE #PTI
PTE 511
not present
present
Guest Physical
to Host Physical
Physical
Page
L1 lookup
with
virtual address
L1
Cache
L1 lookup
with
physical address

Foreshadow-NG Fix

- KAISER/KPTI/KVA does not help
- Only software workarounds
  - → Flush L1 on VM entry
  - → Disable HyperThreading
- Workarounds might not be complete

Meltdown Variants

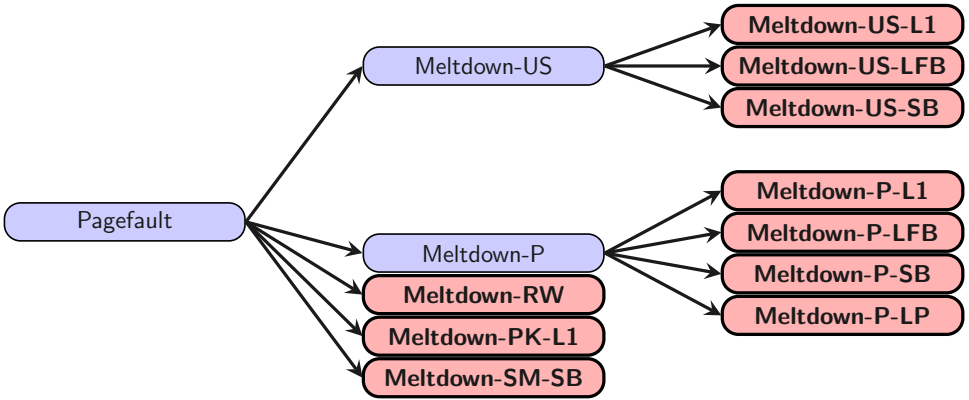

Meltdown Root Cause

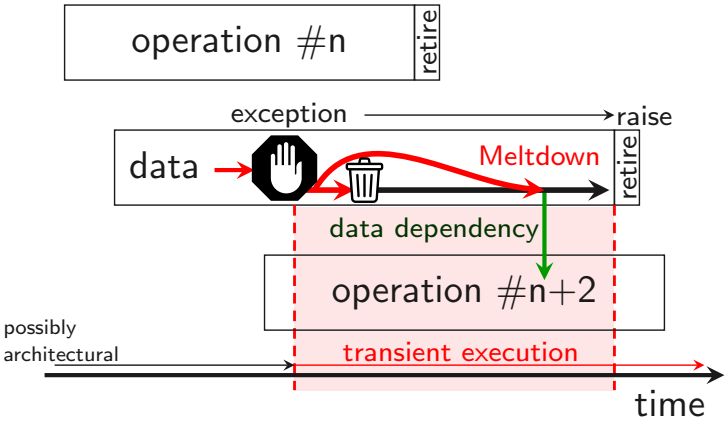


YOU GET A FAULT
AND YOU GET A FAULT.
EVERYONE GETS A FAULT
imgflip.com

Meltdown Tree

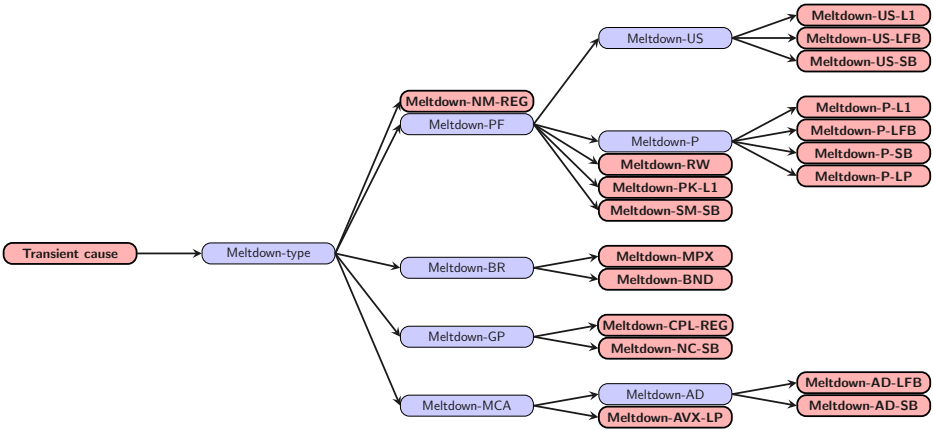

Latest Meltdown Variant: ZombieLoad

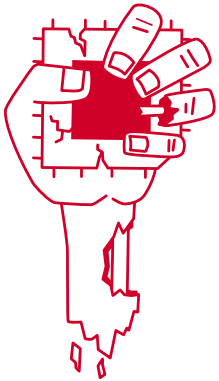

- Leaks from the fill buffer
- Crosses all privilege boundaries (Kernel, VM, SGX)
- Explored microcode assists as new type of faults
- Disadvantage: minimal control over leaked data


zombieload : zsh — Konsole <2>
File Edit View Bookmarks Settings Help
michael@hp /tmp/zombieload %
zombieload : zsh

Meltdown Outlook

- Meltdown is not a fully solved issue
- The tree is extensible
- More Meltdown-type issues to come
- Silicon fixes might not be complete

Transient-Execution Attacks

- Meltdown not the only transient-execution attacks
- Spectre is a second class of transient-execution attacks
- Instead of faults, exploit control (or data) flow predictions

Speculative Execution

- CPU tries to predict the future (branch predictor), ...
  - ... based on events learned in the past
- Speculative execution of instructions
- If the prediction was correct, ...
  - ... very fast
  - otherwise: Discard results

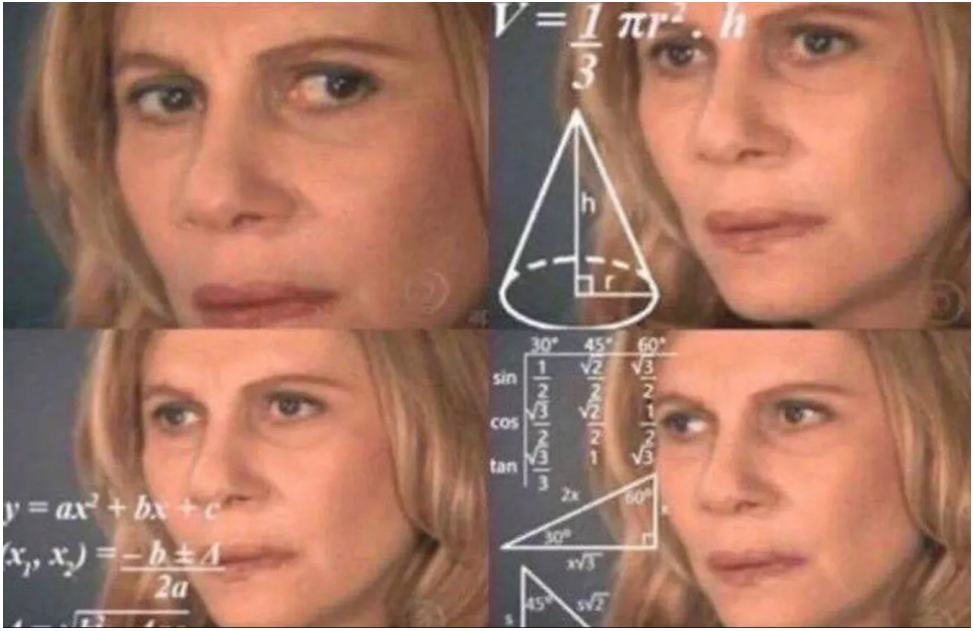
$V = \frac{1}{3}\pi r^2 \cdot h$
h
r
$y = ax^2 + bx + c$
$(x_1, x_2) = \frac{-b \pm \Delta}{2a}$
30° 45° 60°
sin $\frac{1}{2}$ $\frac{\sqrt{2}}{2}$ $\frac{\sqrt{3}}{2}$
cos $\frac{\sqrt{3}}{2}$ $\frac{\sqrt{2}}{2}$ $\frac{1}{2}$
tan $\frac{\sqrt{3}}{3}$ 1 $\sqrt{3}$
2x
60°
30°
x
$x\sqrt{3}$
45°
$x\sqrt{2}$
x

Spectre-PHT (aka Spectre Variant 1)

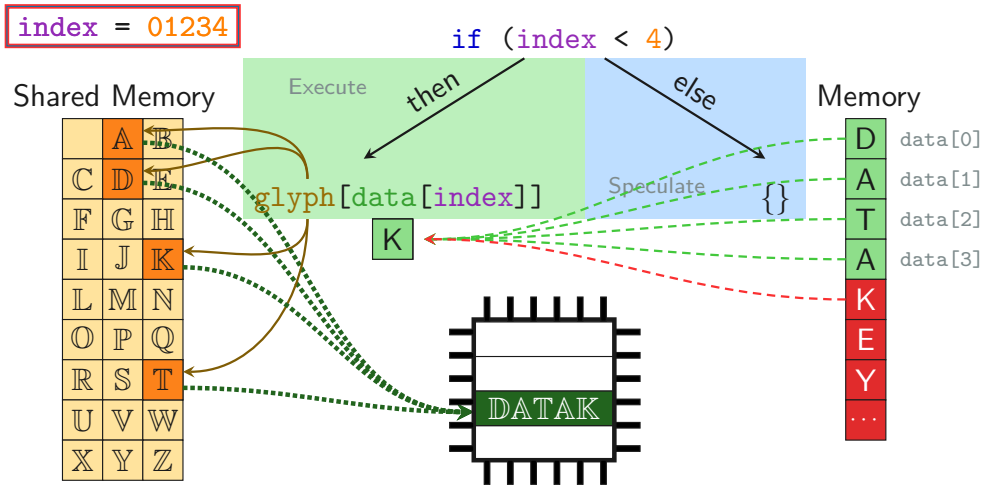

Spectre Root Cause

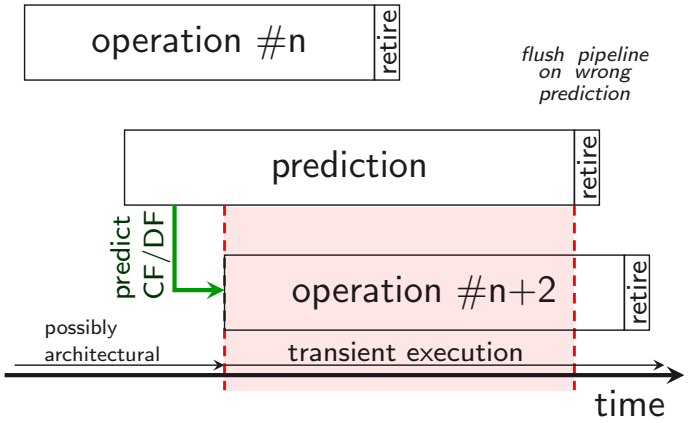

Spectre Root Cause

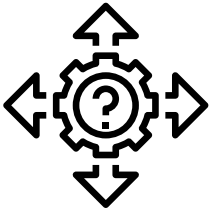

- Many predictors in modern CPUs
  - Branch taken/not taken (PHT)
  - Call/Jump destination (BTB)
  - Function return destination (RSB)
  - Load matches previous store (STL)
- Most are even shared among processes

Spectre Variants

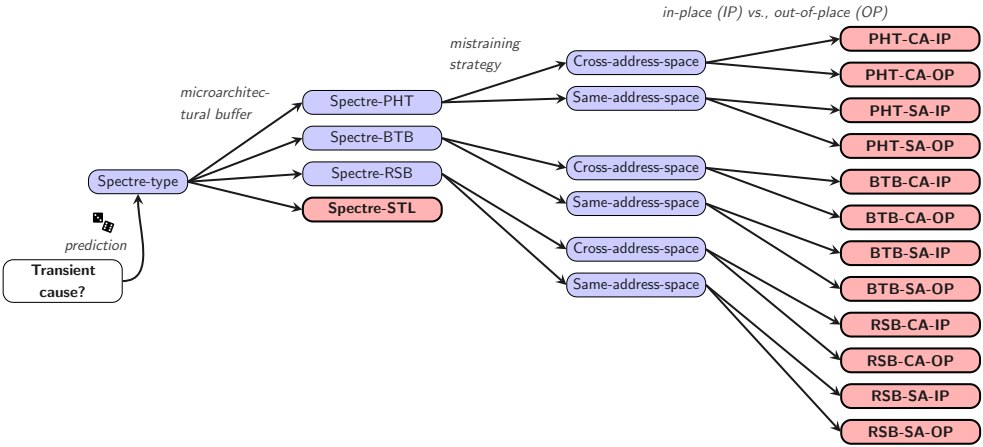

Spectre Fix

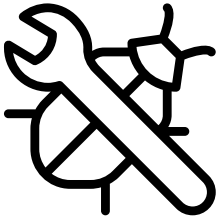

- Spectre is not a bug
- It is an useful optimization

→ Cannot simply fix it (as with Meltdown)

- Workarounds for critical code parts

Linux 4.19.4 & 4.14.83 Released With STIBP Code Dropped

Written by Michael Larabel in Linux Kernel on 24 November 2018 at 09:00 AM EST. 6 Comments

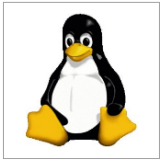

On Friday marked the release of the Linux 4.19.4 kernel as well as 4.14.83 and 4.9.139.

Greg Kroah-Hartman issued this latest round of stable point releases as basic maintenance updates. While these point releases don't tend to be too notable and generally go unmentioned on Phoronix, this round is worth pointing out since 4.19.4 and 4.14.83 are the releases that end up reverting the STIBP behavior that applied Single Thread Indirect Branch Predictors to all processes on supported systems. That is what was introduced in Linux 4.20 and then back-ported to the 4.19/4.14 LTS branches, which in turn hurt the performance a lot. So for now the code is removed.

As covered yesterday, there is improved STIBP code on the way for Linux 4.20 that by default just apply STIBP to SECCOMP threads and processes requesting it via prctl() but otherwise is off by default (that behavior can also be changed via kernel parameters).

Spectre Mitigations

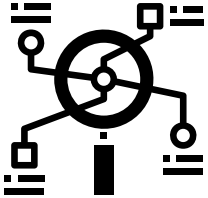

- Current mitigations are either incomplete or cost performance

→ More research required

- Both on attacks and defenses

→ Efficient defenses only possible when attacks are known

Leaking Data

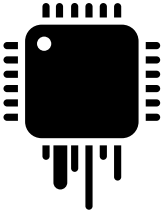

- Side channels so far
  - leak meta data
  - covertly transmit data
- As a building block
  - leak data
- What about modifying data?

DRAM organization

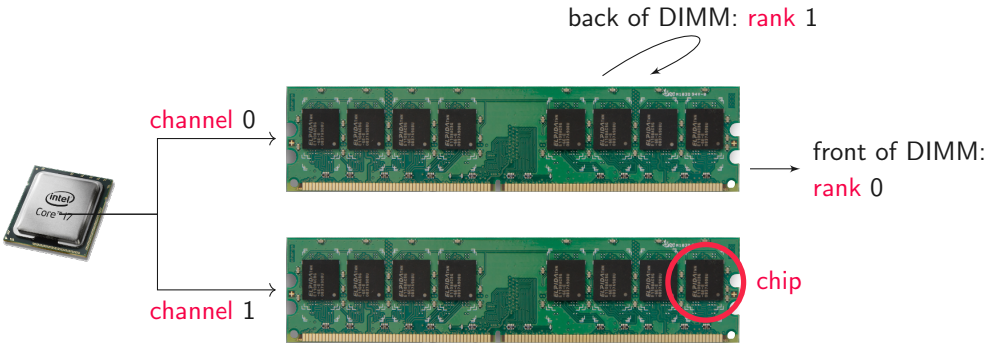

DRAM organization

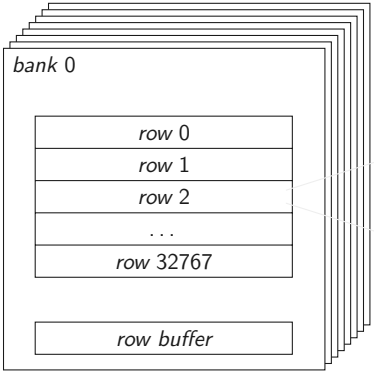

64k cells
1 capacitor,
1 transistor each

How reading from DRAM works

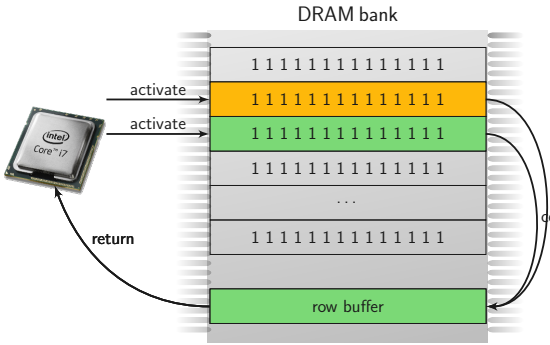

CPU wants to access row 2—again

⇒ row 2 already in row buffer

→ fast (row hit)

CPU wants to access row 2—again

⇒ **row 2 activated**

→ row 1 copied to row buffer

→ slow (row conflict)

row buffer = cache

Timing difference

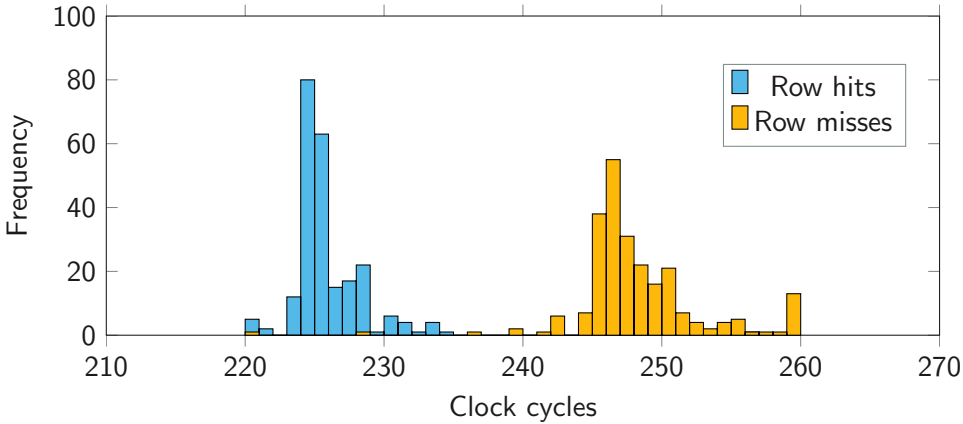

Rowhammer

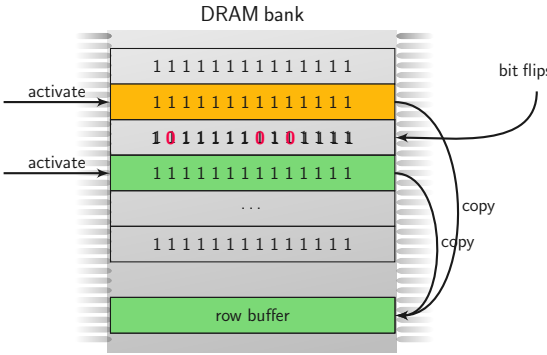

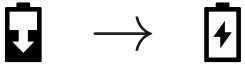

Cells leak faster upon proximate accesses $\rightarrow$ Rowhammer

How widespread is the issue?

- 85% affected (estimation 2014)
- 52% affected (estimation 2015)

- First believed to be safe
- We showed bit flips in 2016
- 67% affected (estimation 2016)

Modifying Bits

- Single bit flips allow
  - modifying instructions
  - breaking cryptography
  - changing permissions
  - crashing systems
  - ...
- In software, no permissions required

An Example

- Program containing conditional jump after password check:
  `je 80486c1 <check_password+0x44>`
- Machine code is
  `0x74 0x07` = `0b01110100 0b00000111`

- One bit changed: `0b01110101 0b00000111` = `0x75 0x07` =
  `jne 80486c1 <check_password+0x44>`
- Now only wrong passwords work → demonstrated on `sudo`

Future

- More attacks exploiting performance optimizations in hardware
  - New variants are disclosed frequently

Microarchitectural Data Sampling (MDS)

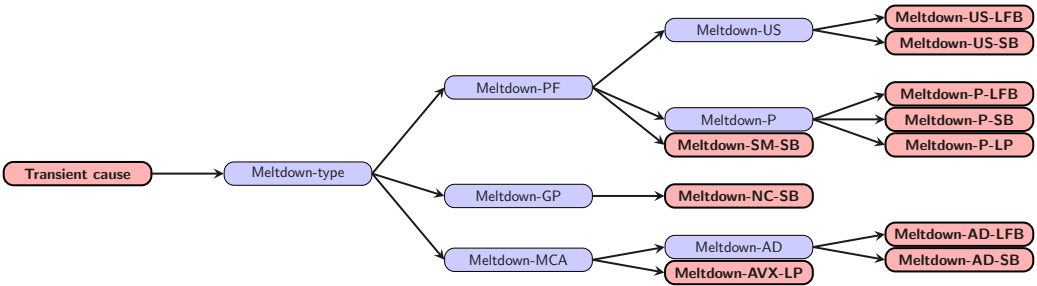

Transient-Execution Attacks

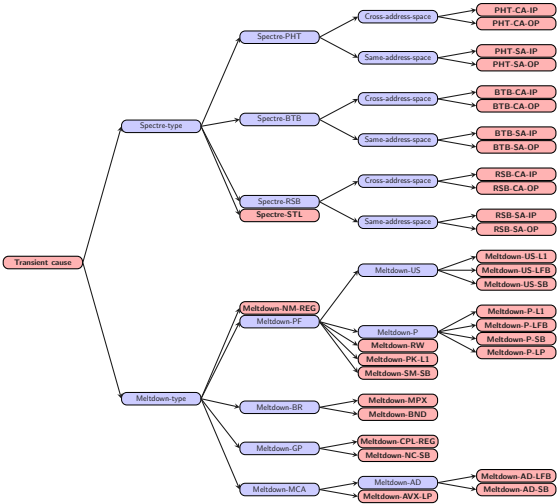

Transient-Execution Attacks

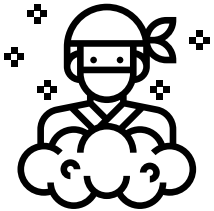

- Transient Execution Attacks are...
  - ...a novel class of attacks
  - ...extremely powerful
  - ...only at the beginning
- Many optimizations introduce side channels $\rightarrow$ now exploitable

A unique chance

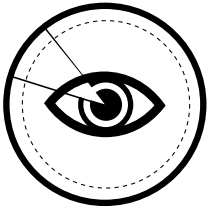

A unique chance to

- rethink processor design
- grow up, like other fields (car industry, construction industry)

Conclusion

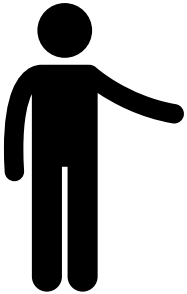

- Optimizations in hardware often lead to side channels
- Unknown and novel side channels are likely to exist
- Next to no permissions required for attacks
- Building countermeasures is extremely hard

More Information

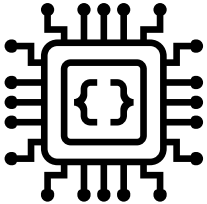

- Only an overview over some attacks
- Many more side-channel attacks
- Also some defenses, especially for crypto
- Master course: Embedded Security
- Talks from our group on YouTube: “InfoSec @ TU Graz”