

Information Security

System Security 3 - Physical Side-Channel and Fault Attacks

24.11.2023





“Human Side-Channel Analysis”



Attacks with Physical Access

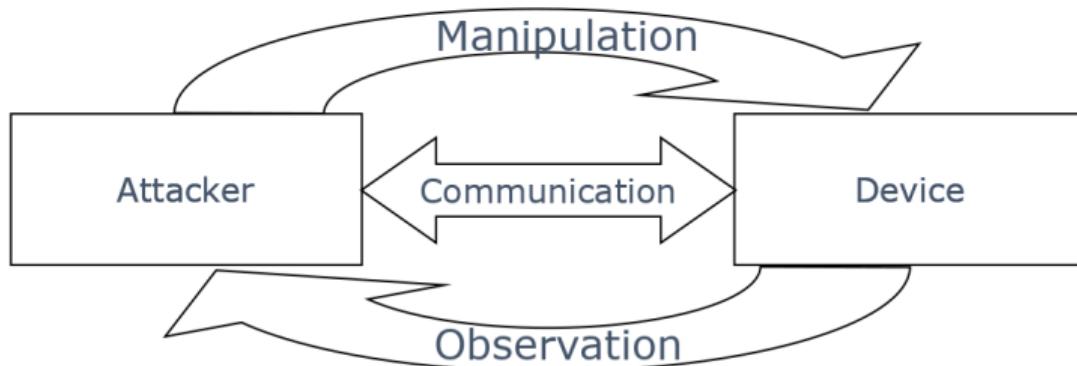


“If the attacker can execute code
... they have already won”

Applications Exposed to Physical Attacks



Physical Attack Principle



~~SECRET~~

(b) (3)-P.L. 86-36



Approved for Release by NSA on
09-27-2007, FOIA Case # 51633

TEMPEST: A Signal Problem

**The story of the discovery
of various compromising radiations
from communications and Comsec equipment.**

impractical. Hydraulic techniques—to replace the electrical—were tried and abandoned, and experiments were made with different types of batteries and motor generators, in attempts to lick the power-line problem. None was very successful.

During this period, the business of discovering new TEMPEST threats, or refining techniques and instrumentation for detecting, recording, and analyzing these signals, progressed more swiftly than the art of suppressing them. Perhaps the attack is more exciting than the defense—something more glamorous about finding a way to read one of these signals than going through the drudgery necessary to suppress that whacking great spike first seen in 1943. At any rate, when they turned over the next rock, they found the acoustic problem under it. Phenomenon No. 5.

Acoustics

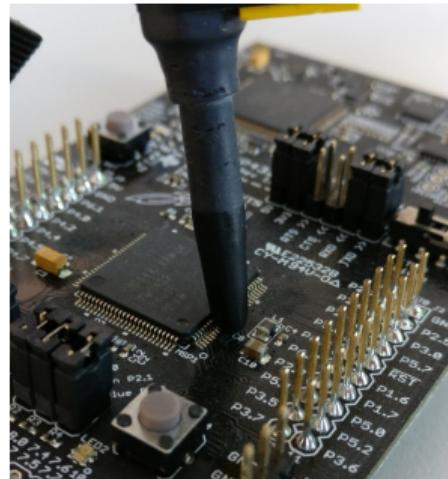
We found that most acoustic emanations are difficult to exploit if the microphonic device is outside of the room containing the source equipment; even a piece of paper inserted between, say, an offending keyboard and a pick-up

Physical Attacks: Categorization

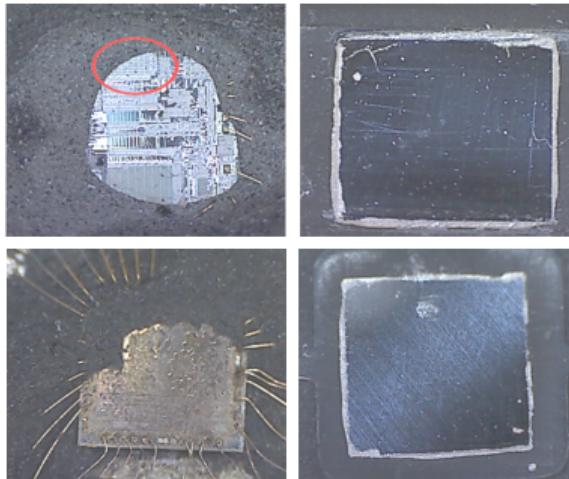


- Behavior of the attacker
 - Side-channel attack: passively observe physical properties
 - Fault attack: actively manipulate device to induce faults
- Degree of invasiveness
 - Non-invasive: Device is not altered physically
 - Semi-invasive: De-packaging, no electrical contact to internal signals
 - Invasive: No limits

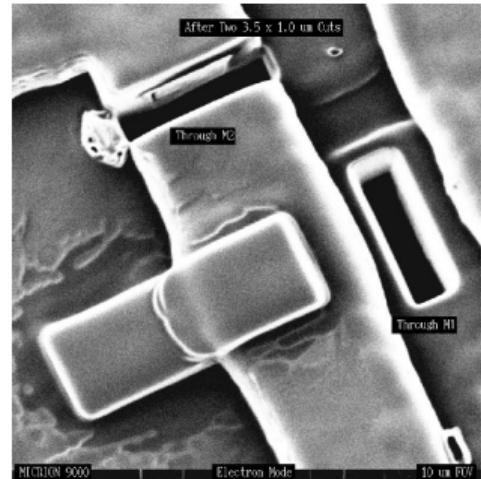
Degree of Invasiveness



Non-Invasive



Semi-Invasive



Invasive

Side-Channel Attacks



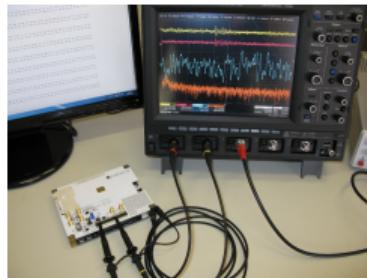
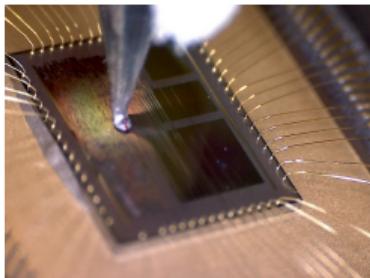
Basic Idea of Side-Channel Attacks



Any computation influences physical properties (meta-data)

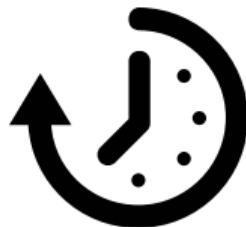
- Computations depend on secrets (data)
- We observe properties (meta-data) to infer secrets (data)

Side-Channel Attacks



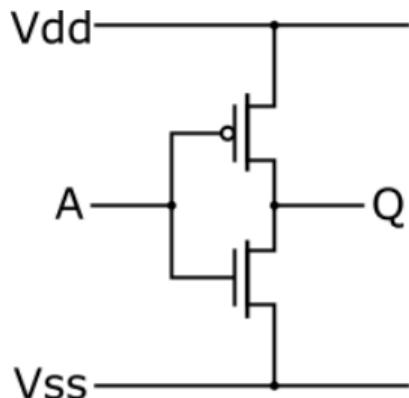
- Timing
- Power consumption
- EM emanations
- Sound
- ...

Timing Attacks on Cryptographic Implementations



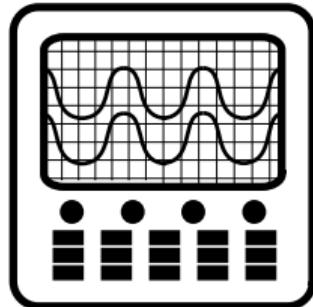
- Often overlooked / ignored
 - “outside of threat model”
 - implementation bugs
- Sometimes even on certified devices (e.g., Minerva and TPM-Fail)
→ Solution: Make everything constant-time?

CMOS Circuits

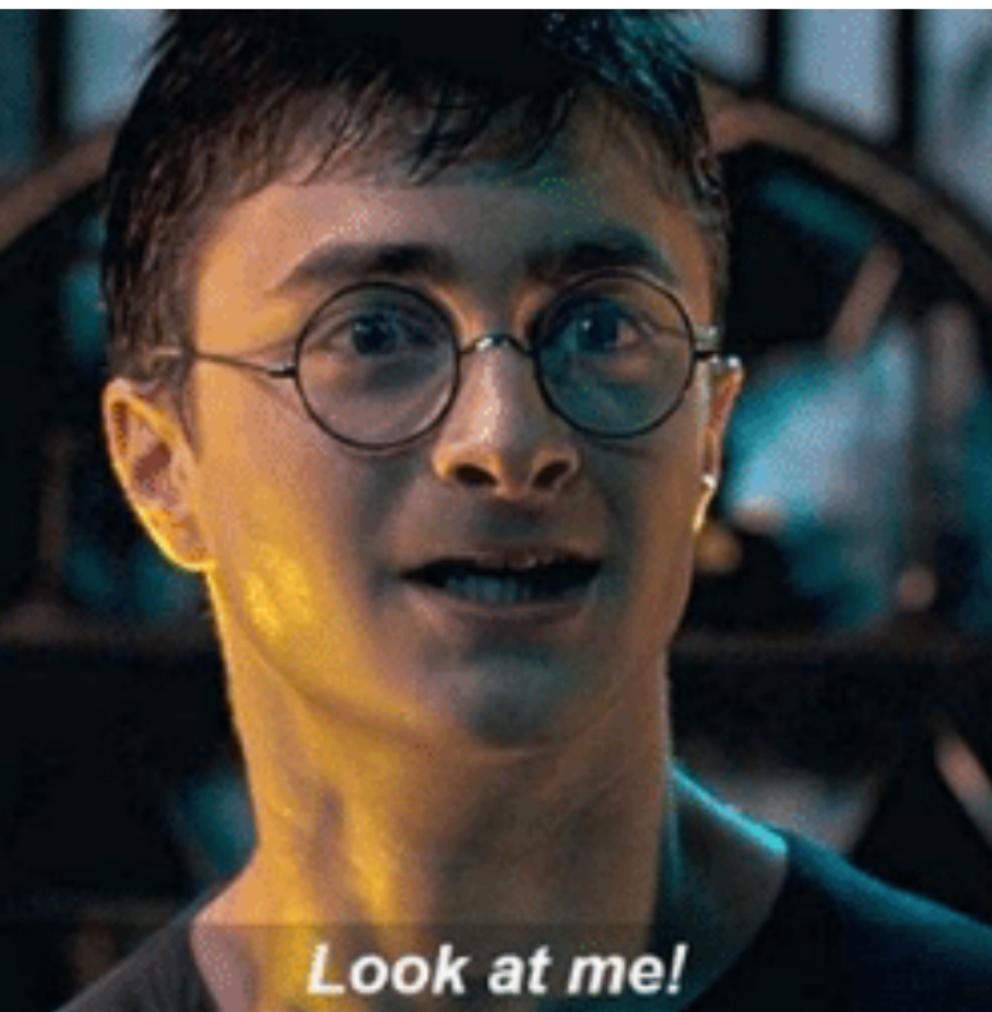


- Complementary Metal Oxide Semiconductor
- Today's digital circuits
- Nice properties:
 - high noise immunity
 - low power consumption
 - (Only switching draws power)
- Wait a second... switching draws power?

CMOS Circuits - Power Consumption



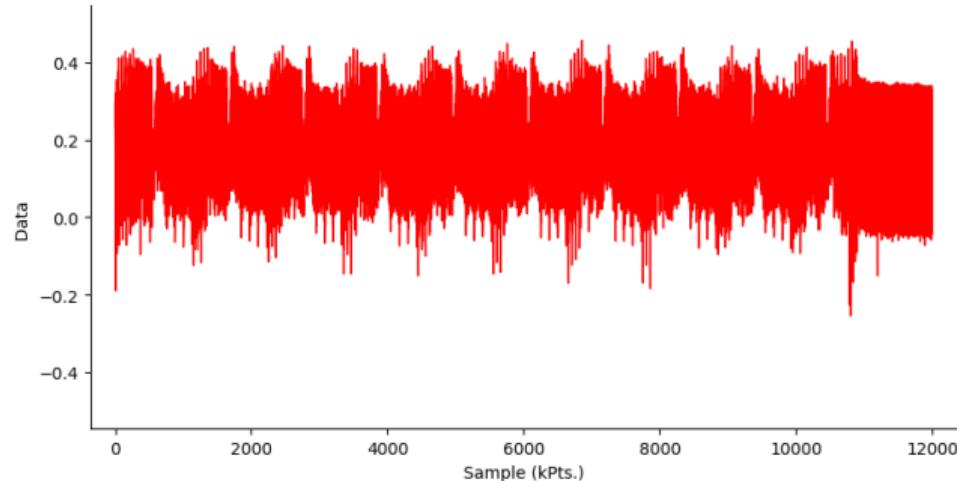
- Different instructions / data → different switching
- Idea: Measure power consumption during operation
 - sampling rate up to gigasamples (10^9 measurements per second)
 - a measured power-consumption curve is called a **power trace**
- First signal-processing step?



Look at me!



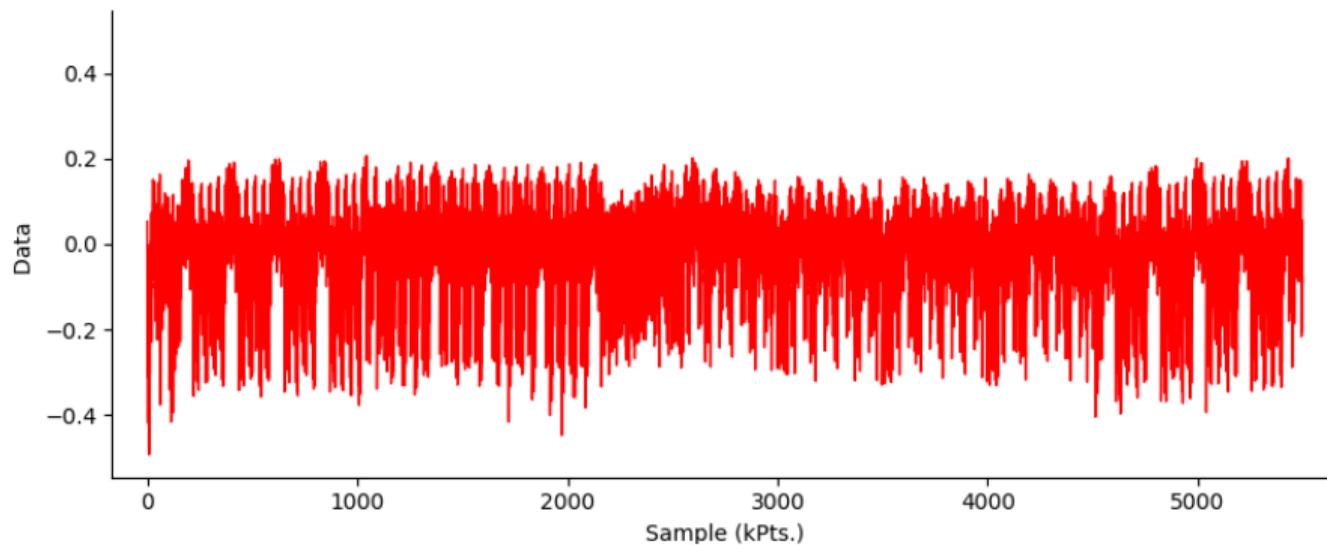
Power Consumption - An Example Trace



What do we see here?

10 rounds of AES

Power Consumption - Zoom in on Single Round



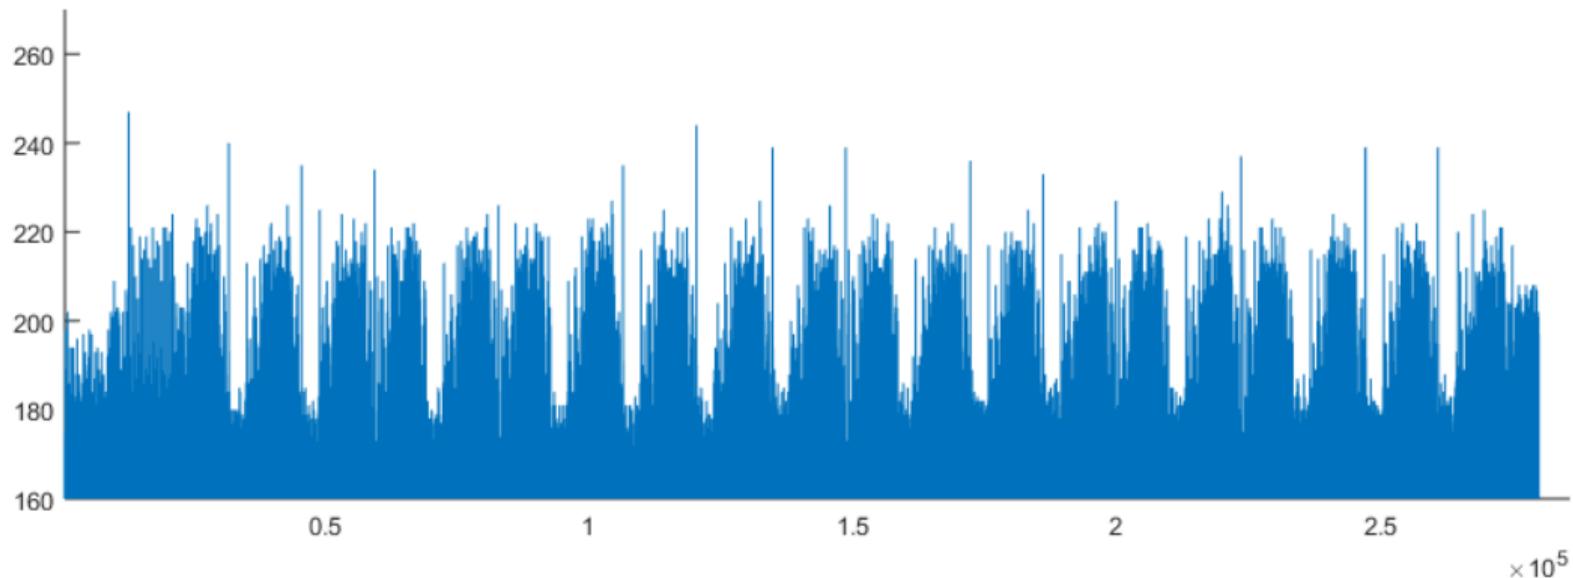
What do we see?



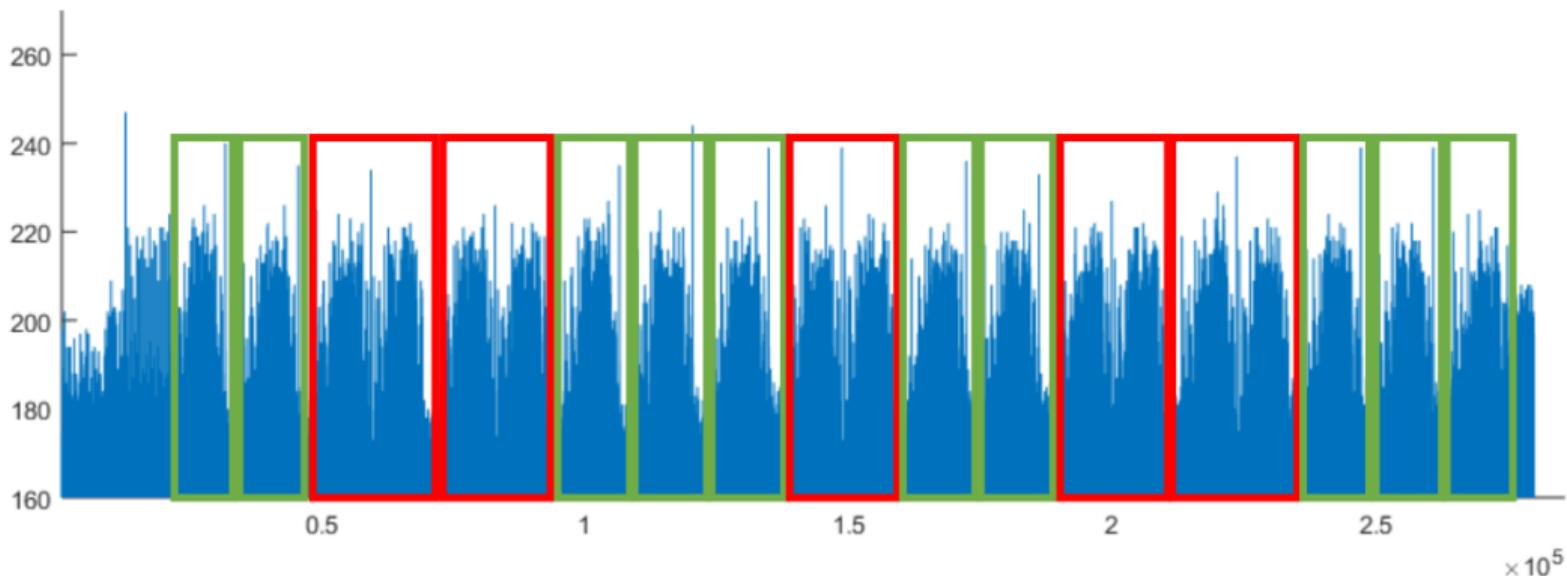
- Operations / Instructions
- Repeated patterns and variations of patterns
 - Loops, repeated operations, taken branches
 - Learn control flow / instruction sequence
- Jumps in power consumption profile
 - Memory accesses (especially EEPROM or flash programming)
 - Access to peripherals (e.g., co-processors, I/O)

→ Can we exploit that?

Our First Power-Analysis Attack



Our First Power-Analysis Attack



We see a power trace of an RSA exponentiation ($m = c^d \pmod n$)
How to get the key from that?

Excursion: Efficient Implementation of Modular Exponentiation

 RSA decryption: $m = c^d \bmod n$, where n has ≥ 2048 bits

 Efficient implementation?

- ⌚ Compute $(c^d) \bmod n$?
- ➔ c and d are also ≥ 2048 bits
- ⚠ c^d has more than 2^{2048} bits!

 Some reminders for modular arithmetic:

- $a \cdot b \bmod n = (a \bmod n) \cdot (b \bmod n) \bmod n$
- $c^{a+b} \bmod n = (c^a \bmod n) \cdot (c^b \bmod n) \bmod n$
- $c^{a \cdot b} \bmod n = (c^a \bmod n)^b \bmod n$

Excursion: Efficient Implementation of Modular Exponentiation

- Look at exponent d in binary: $d_i = i$ th bit of d , where $d_0 = \text{LSB}$
- Recursive decomposition of exponentiation:
 - We can write $d = 2 \cdot \lfloor d/2 \rfloor + (d \bmod 2) = 2 \cdot (d \gg 1) + d_0$
In the exponent, we get: $c^d = (c^{\lfloor d/2 \rfloor})^2 \cdot c^{d_0}$
 - But $c^{\lfloor d/2 \rfloor}$ is still way too large, so repeat:
 $c^{\lfloor d/2 \rfloor} = (c^{\lfloor d/4 \rfloor})^2 \cdot c^{\lfloor d/2 \rfloor \bmod 2} = (c^{\lfloor d/4 \rfloor})^2 \cdot c^{d_1}$
 - ... until $\lfloor d/2^x \rfloor = 1 \rightarrow c^{\lfloor d/2^x \rfloor} = c$
- Iterative version: Start at $\lfloor d/2^x \rfloor = 1$ and work our way up

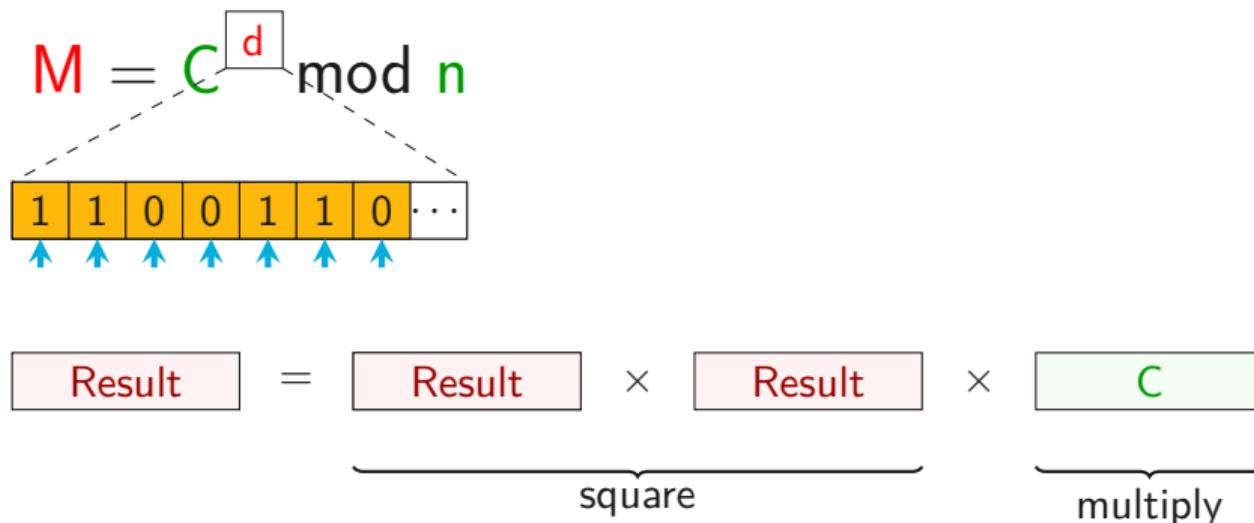
Excursion: Efficient Implementation of Modular Exponentiation

- Algorithm: Left-to-right *Square-and-Multiply* exponentiation

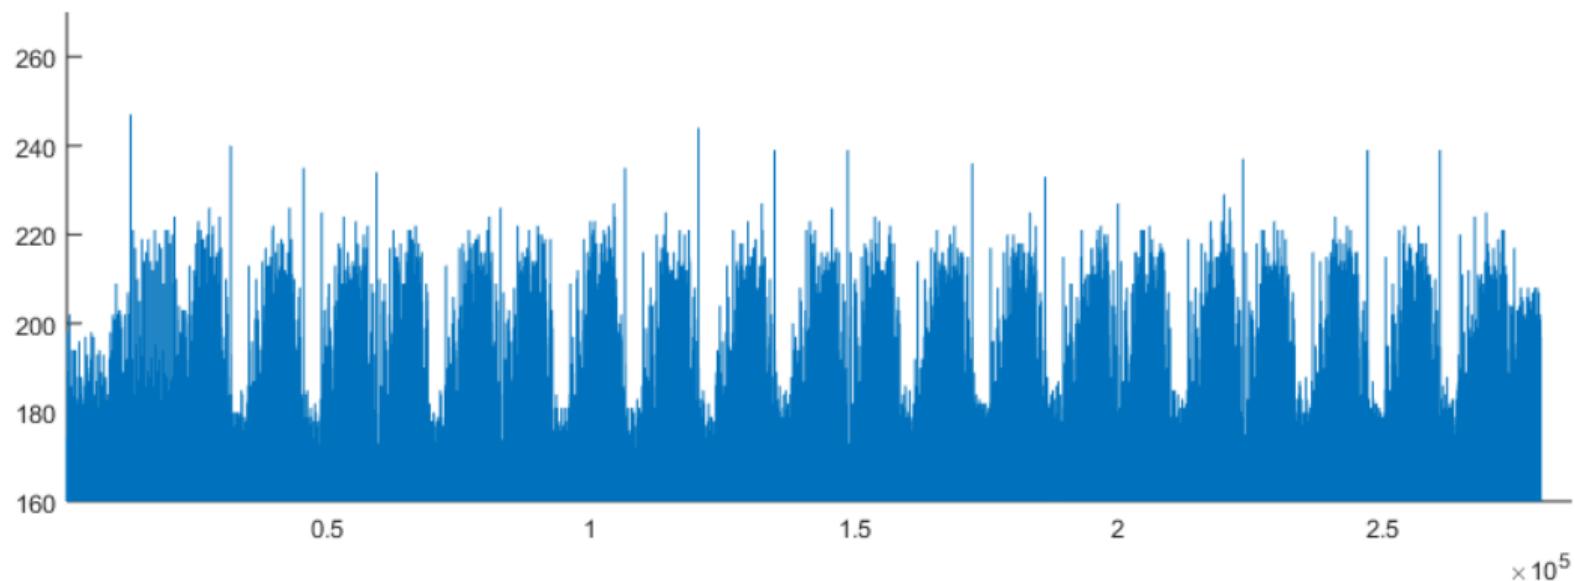
```
m ← 1           // init  
for i = 2047 downto 0 : // scan bits from MSB to LSB  
    m ← m2 mod n      // squaring:  $c^x = (c^{\lfloor x/2 \rfloor})^2 \cdot c^{x_0}$   
    if di = 1 then:      // if bit is set (else x0 = 0 → cx0 = 1, can skip mult.)  
        m ← m · c mod n  // then multiply:  $c^x = (c^{\lfloor x/2 \rfloor})^2 \cdot c^{x_0}$ 
```

- Example: $d = 26 = 11010_b \rightarrow c^{26} = (((((1^2 \cdot c)^2 \cdot c)^2)^2 \cdot c)^2)$

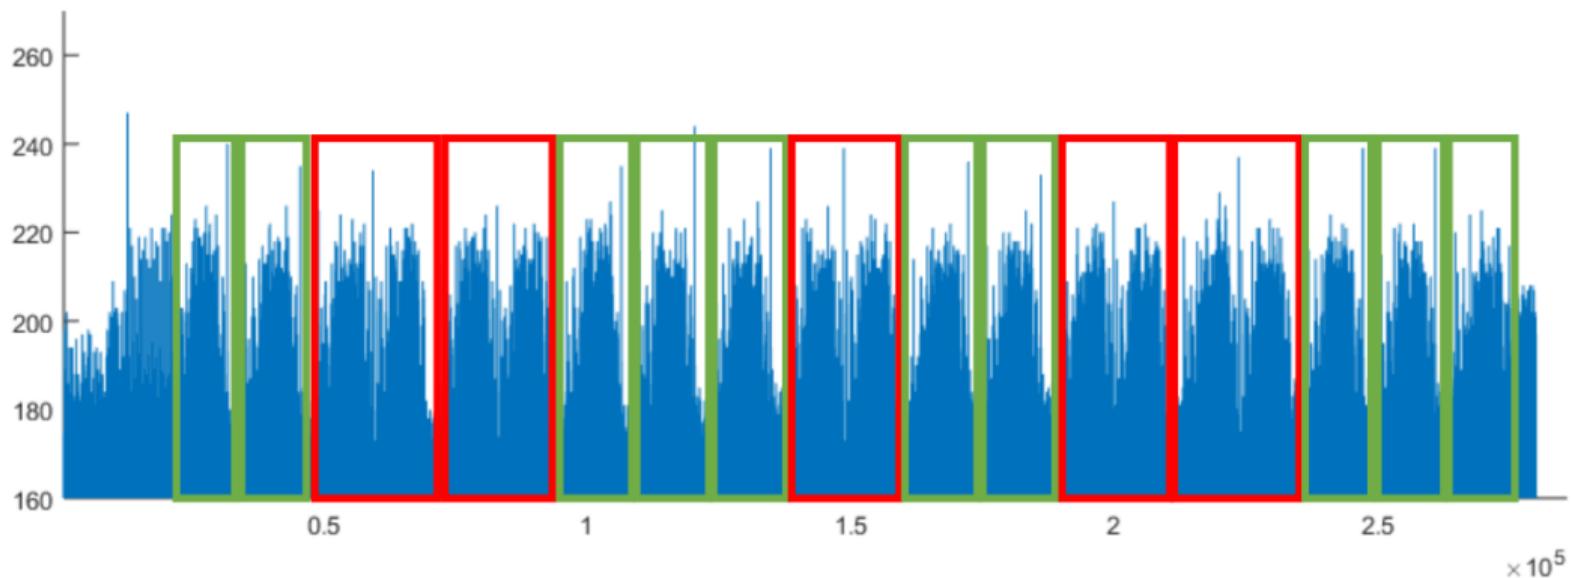
Excursion: Efficient Implementation of Modular Exponentiation



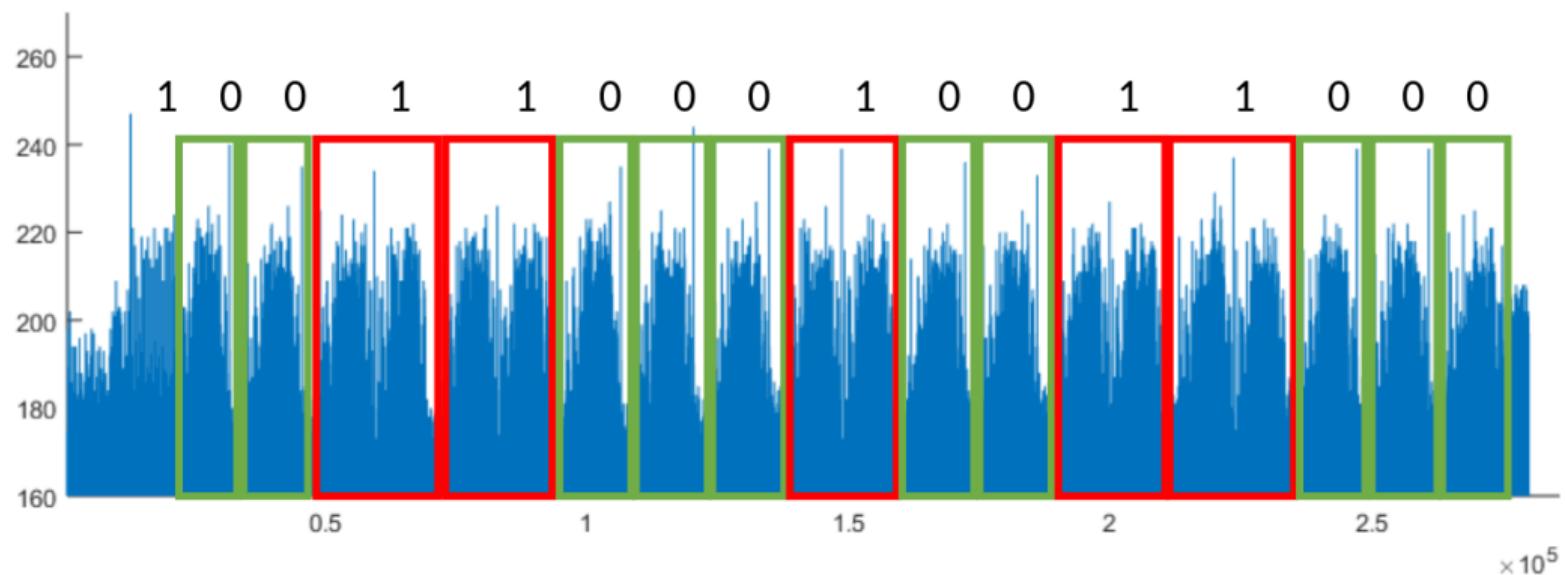
Our First Power-Analysis Attack - Key Recovery



Our First Power-Analysis Attack - Key Recovery



Our First Power-Analysis Attack - Key Recovery



Power Side Channel Countermeasures



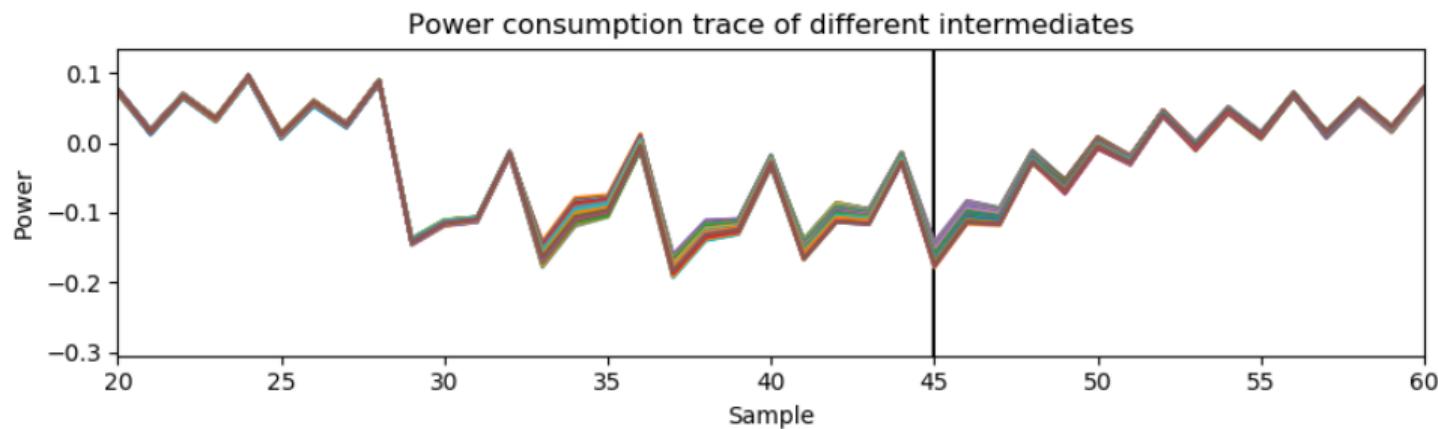
“Constant-time” means more than just *constant time*

- No branching on secret data: **constant runtime and control flow**
→ always same instruction sequence but different data
- More secure alternatives:
 - Constant-time exponentiation algorithms
 - Constant-time modular reduction
 - ...

And now: “Constant-time” → all problems solved?

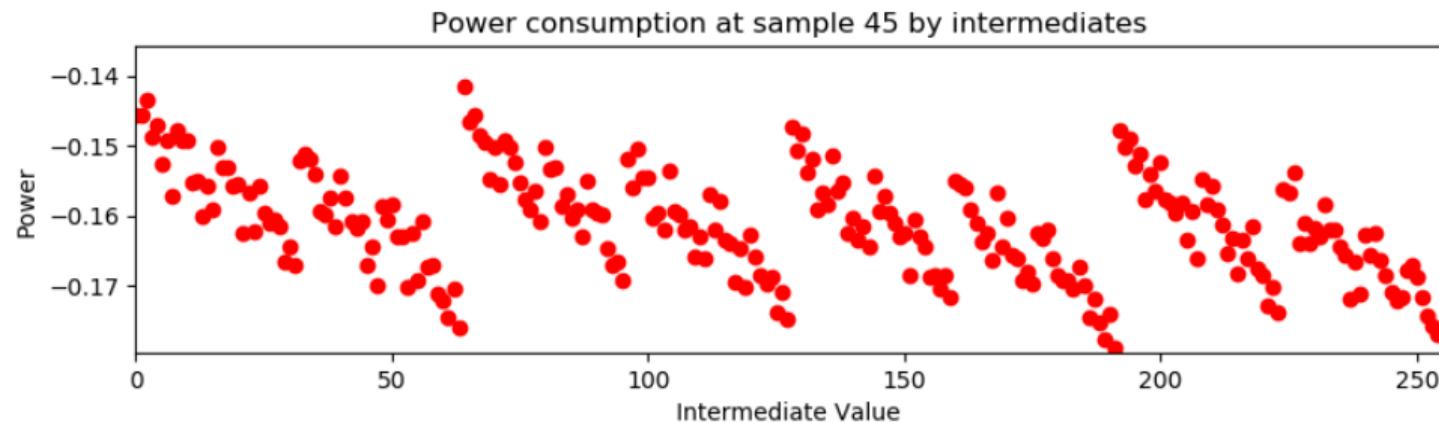
What about data differences?

CMOS Power Consumption: Data Dependency



Averaged power traces of a load instruction for values $\{0, 255\}$

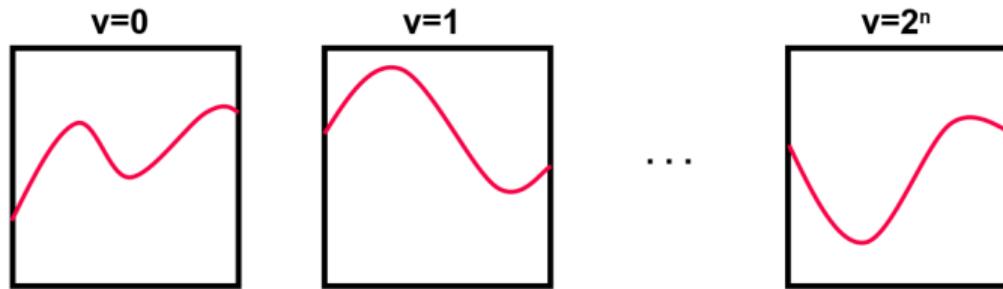
CMOS Power Consumption: A Closer Look



Different intermediate values → different power consumption

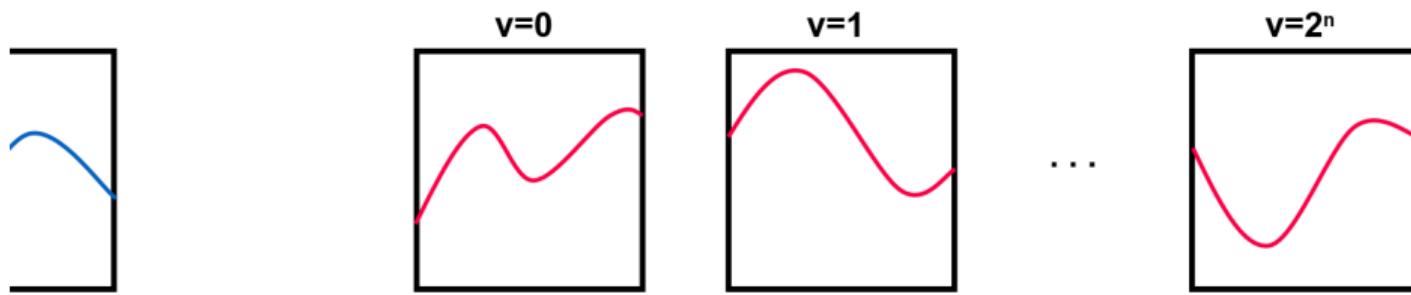
Record + match values = *Template Attack*

Template Attack - Phase 1 - Characterization / Profiling



- Profile power consumption for each possible value of intermediate v
- Record traces with all inputs known, group by v
- Profile == “Template”

Template Attack - Phase 2 - Attack / Exploitation



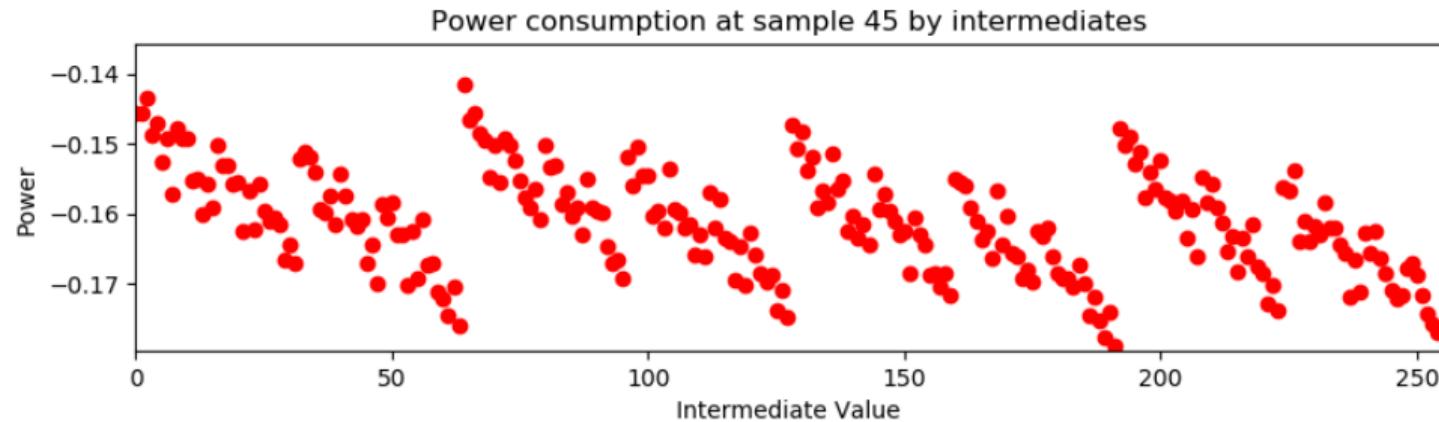
- Compare (match) measured traces to all templates
- Use v which fits best (compute probabilities)

Template Attacks: Conclusion



- Pro: Very powerful
 - Key recovery with single trace
 - Sometimes the only option (“we only have a single trace”)
- Contra: many prerequisites and detailed knowledge needed
 - When is secret processed?
 - What is the concrete algorithm?
 - Identical device / setup needed where you can control **all** inputs

Another Look at Intermediate Values



There is some kind of pattern...

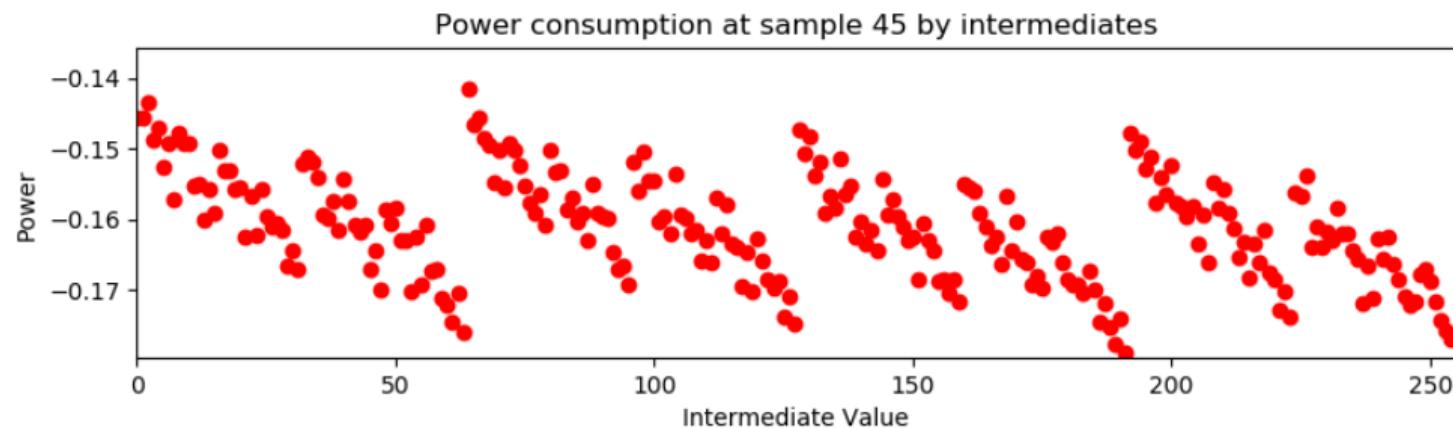
We can *model* the power consumption

Hamming Weight Power Model



- “Power consumption depends on switching”
- What’s stored before a value is stored? Assume 0
- Now the new value: each ‘1’-bit draws power → power is proportional to number of bits set
- number of bits set == Hamming weight

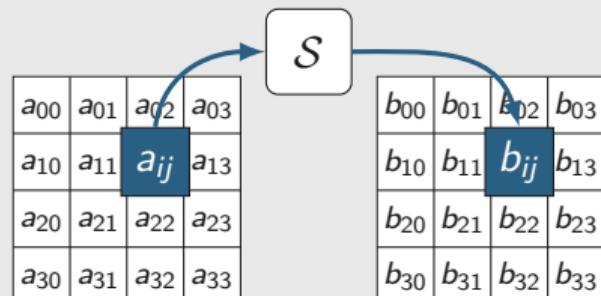
And Another Look at Intermediate Values



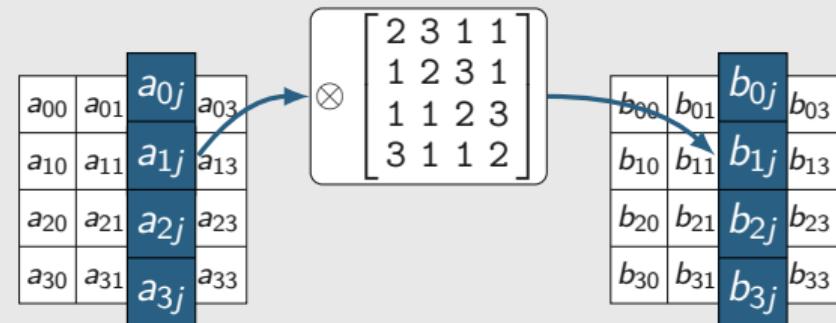
Many devices have similar power behavior → reuse power models
→ an attack without detailed knowledge of device an concrete implementation!

Reminder: AES-128 Block Cipher (10 Rounds)

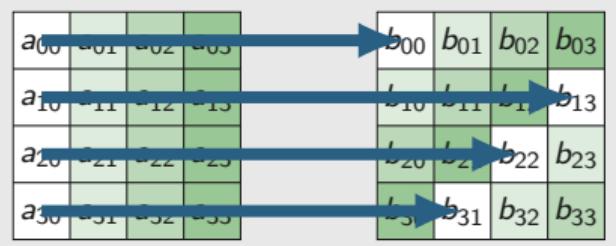
1. SubBytes (SB)



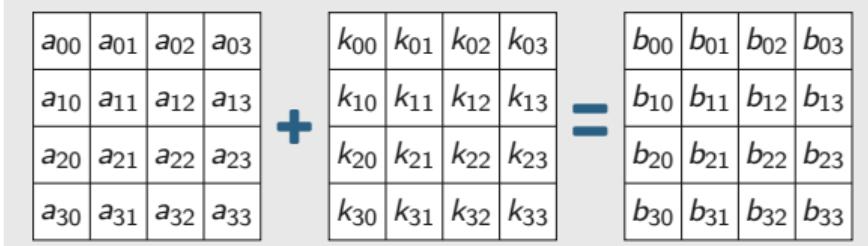
3. MixColumns (MC)



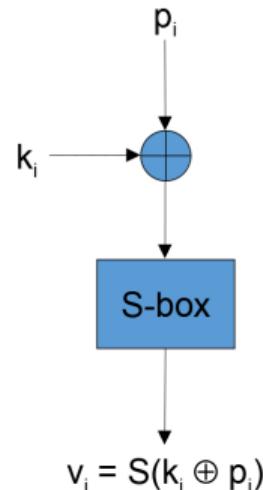
2. ShiftRows (SR)



4. AddRoundKey (AK)



Excursion: AES

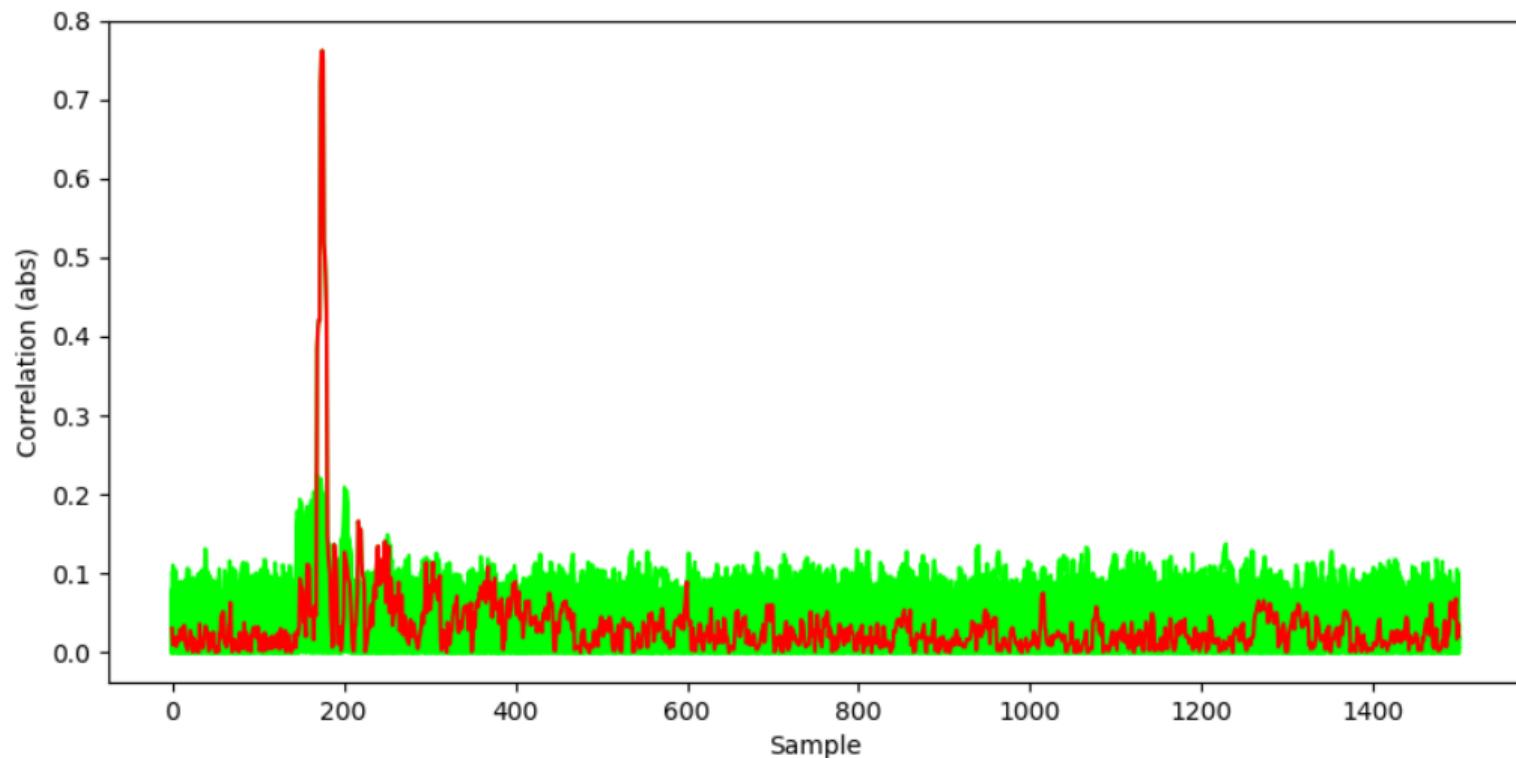


- First round: round key = key
- Other rounds: key schedule
 - key schedule is invertible

Differential Power Analysis (DPA)

Differential Power Analysis (DPA) in 5 Steps

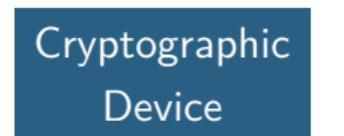
1. **Select intermediate** value that depends on a small number of key bits (subkey)
2. **Measure power** while querying device
3. **Enumerate all possible** subkey **values**
 - 2^8 key hypotheses
 - for each plaintext/ciphertext: predict intermediate for each key hypothesis
4. **Predict power** consumption of intermediate (power model, e.g., Hamming weight)
5. **Compare** prediction with measurement
 - pick key hypothesis that fits best
 - statistical hypothesis tests



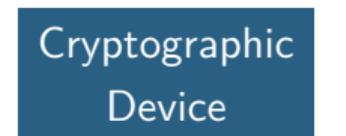
Countermeasures against Power Analysis



Unprotected



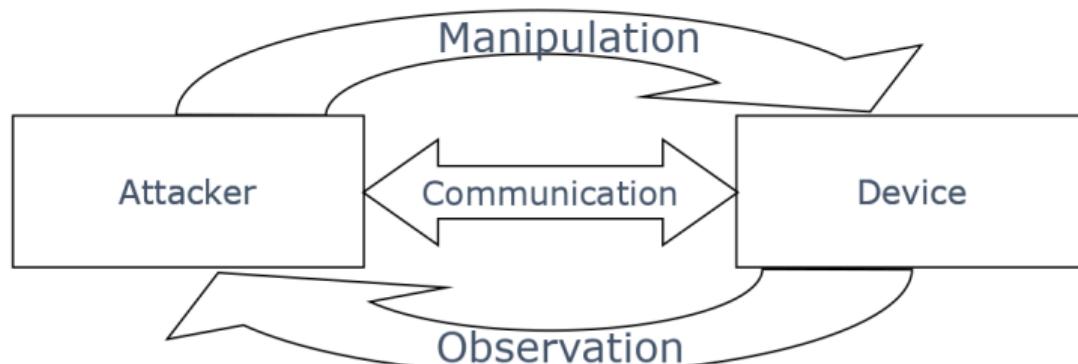
Hiding



Masking

Fault Attacks

Fault Attacks



Just listening is boring ...

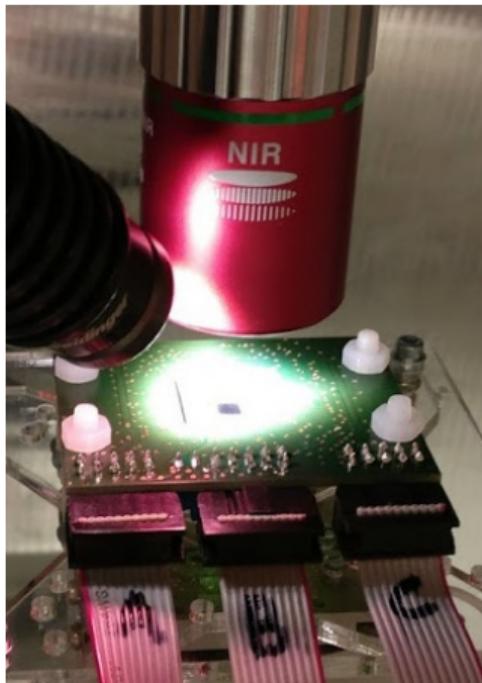
→ let's manipulate things more actively

Basic Idea of Fault Attacks



- Goal: manipulate device to compromise security
- Change behavior
 - Deactivate countermeasures / sensors
 - Skip PIN check
- Fault crypto algorithms
 - Compute faulty and correct ciphertexts
 - Use difference to reveal key

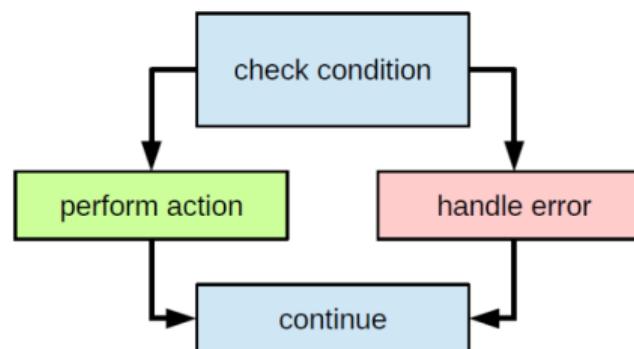
Fault Attack Techniques



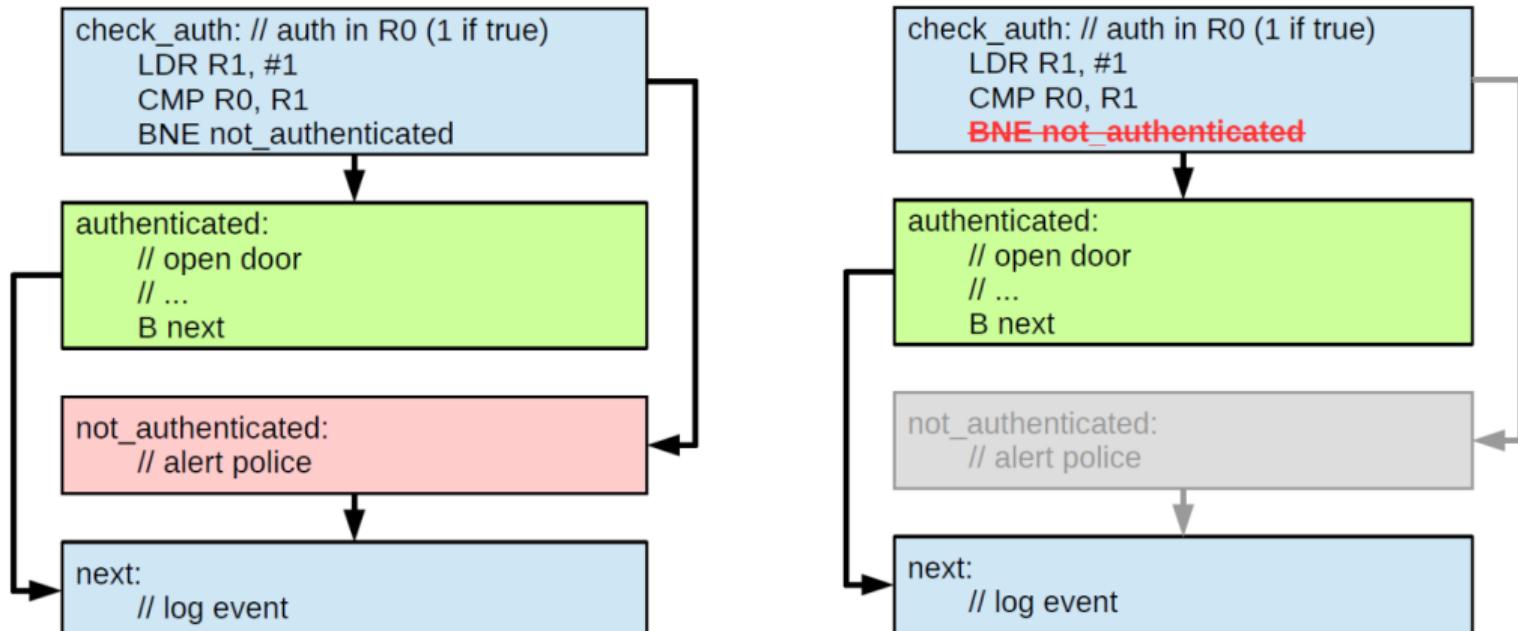
- Spike / glitch: clock, voltage, etc.
- Heat, Radiation, Laser
- Effects:
 - Instructions skipped
 - Data corrupted
 - ...

Example: PIN Check

```
unsigned pin = read_pin();
bool auth = tpm_check(pin);
if( auth ) {
    open_door();
} else {
    alert_police();
}
log_event();
```



Example: PIN Check - Skipping Attack



Real-World Example: Piracy



PayTV (early 2000s)

- vendors bricked pirated cards via firmware update
- insert endless loop in startup
- solution: glitch to escape loop (“unlooper device”)

Real-World Example: Piracy



Gaming devices

- Xbox360 reset hack
- voltage glitching on reset line
- execute untrusted modified firmware

Fault Attack on RSA

Excursion: Efficient RSA Signature with Chinese Remainder Theorem

✍ RSA signatures: $S = M^d \bmod n$, where $n = p \cdot q$

💡 Efficient implementation trick: Chinese Remainder Theorem (CRT)

↳ Compute signature result $\bmod p$ and $\bmod q$

$$S_p = S \pmod{p} = M^{d \bmod p-1} \pmod{p}$$

$$S_q = S \pmod{q} = M^{d \bmod q-1} \pmod{q}$$

⌚ ... and merge the results with CRT:

$$S = S_p \cdot (q^{-1} \bmod p) \cdot q + S_q \cdot (p^{-1} \bmod q) \cdot p \pmod{p \cdot q}$$

⌚ 2 exponentiations with half the bit-length and smaller exponents

Fault Attack on RSA Signatures with CRT (“Bellcore attack”)

- Compute signature twice and fault one computation \lightning

$$S = [S_p \cdot (\text{something } p)] \cdot q + [S_q \cdot (\text{some rest } q) \cdot p] \pmod{n}$$

$$S^\lightning = [S_p \cdot (\text{faulty mod } p)] \cdot q + [S_q \cdot (\text{some rest } q) \cdot p] \pmod{n}$$

$$S - S^\lightning = [\text{some garbage}] \cdot q + [0] \pmod{n}$$

- Get the secret q using

$$\gcd(S - S^\lightning, n) = \gcd([\text{some garbage}] \cdot q, p \cdot q) = q$$

```
bagger> dog Enclave/encl
```

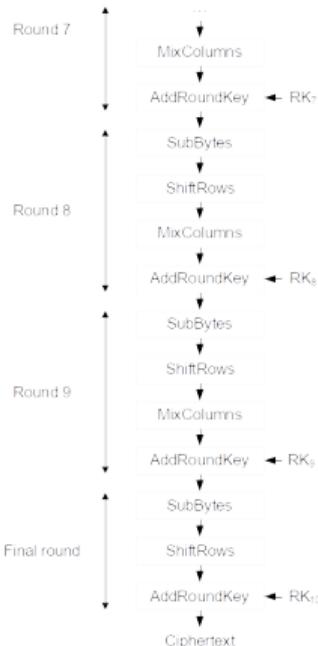
A close-up portrait of Simba, the young lion king from Disney's "The Lion King". He has his signature large, bushy orange-red mane and is looking slightly to the left with a neutral expression. The background is a bright blue sky with a few wispy white clouds.

MUFARSA

**My Undervolting
Fault Attack on RSA**

Fault Attack on AES

Differential Fault Attacks on AES



Faulting ciphertext?

- No.
- ciphertext difference does not depend on key

Differential Fault Attacks on AES



Faulting before AddRoundKey10?

- depends on faults
- not with bit flips (random or known)
- → fault propagates through \oplus :
 $c = v \oplus k$
 $c' = (v \oplus \Delta v) \oplus k = c \oplus \Delta v$
- → ciphertext difference still does not depend on key

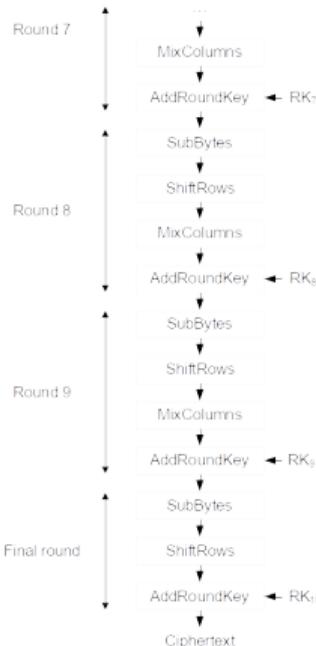
Differential Fault Attacks on AES



Faulting before ShiftRows10?

- same as before
- ShiftRows just rearranges bytes

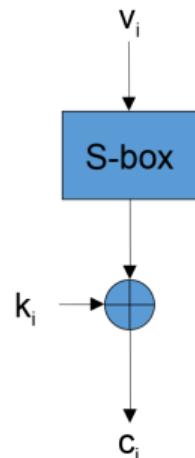
Differential Fault Attacks on AES



Faulting before SubBytes10?

- ...depends on faults
 - Able to flip 1 bit?
- Attacks possible

Single-Bit Fault before SubBytes

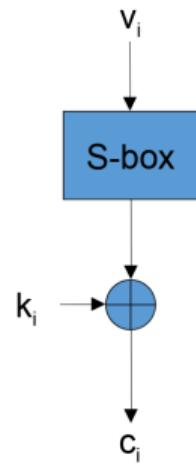


Receive correct and faulty ciphertext

Enumerate all 2^8 values for k_i

- compute back to v (for correct and fault, for all possible k_i)
- compute Δv for k_i
- check if Δv follows fault model (1 bit fault)
- indices can be different because of ShiftRows

Single-Bit Fault before SubBytes - Example



Correct Output = 1A

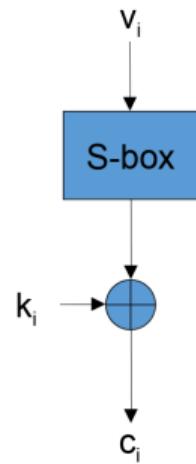
Faulty Output = 99

k: 0 1 2 3 4 5 6 7 8 ...

$C = 1a : S^{-1}(C \text{ xor } k) :$

$C' = 99 : S^{-1}(C' \text{ xor } k) :$

Single-Bit Fault before SubBytes - Example



Correct Output = 1A

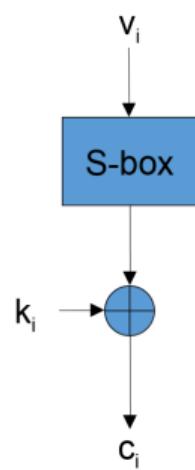
Faulty Output = 99

k: 0 1 2 3 4 5 6 7 8 ...

$C = 1a : S^{-1}(C \text{ xor } k) : 43\ 44\ 34\ 8e\ e9\ cb\ c4\ de\ 39\ ...$

$C' = 99 : S^{-1}(C' \text{ xor } k) : f9\ e2\ e8\ 37\ 75\ 1c\ 6e\ df\ ac\ ...$

Single-Bit Fault before SubBytes - Example



Correct Output = 1A

Faulty Output = 99

k: 0 1 2 3 4 5 6 7 8 ...

$C = 1a : S^{-1}(C \text{ xor } k) : 43\ 44\ 34\ 8e\ e9\ cb\ c4\ de\ 39\ ...$

$C' = 99 : S^{-1}(C' \text{ xor } k) : f9\ e2\ e8\ 37\ 75\ 1c\ 6e\ df\ ac\ ...$

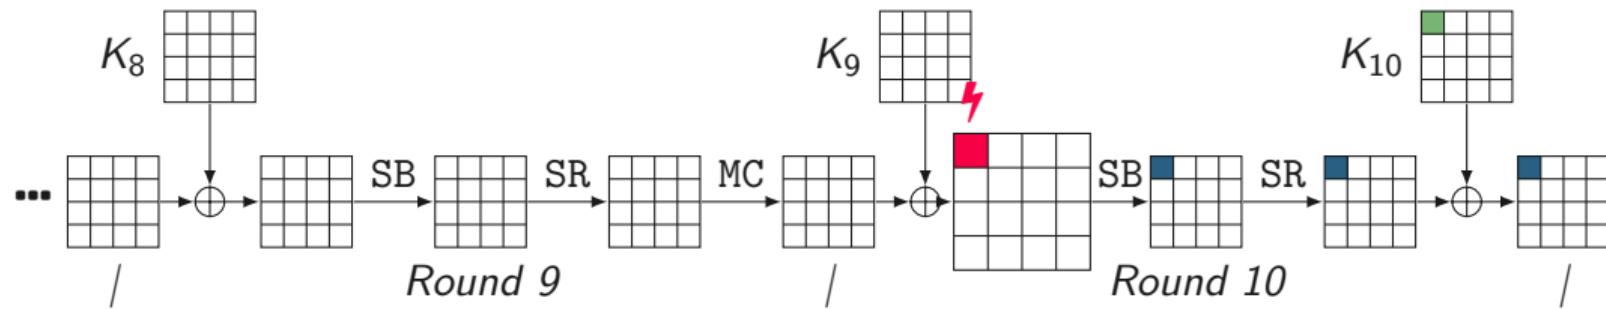
~~

Only few keys have this property → filter them

Use further C/C' pairs to get down to 1 key

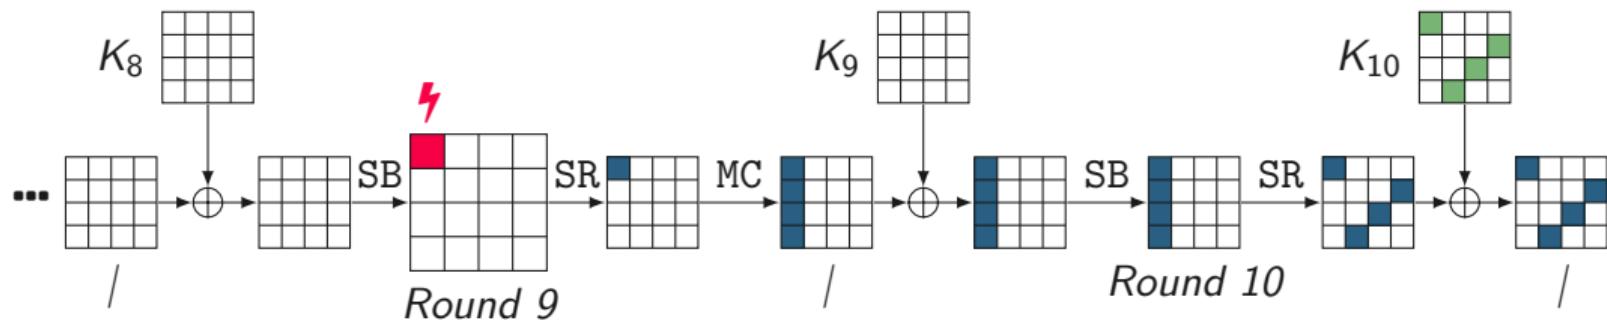
AES – Simple DFA (Summary)

- Assume the attacker can cause precise **1-bit flips** in Round 9 of AES, before S-box
- For each of 2^8 key guesses,
Test if the *partial decryption* produces the expected 1-bit flip.



AES – DFA on More Rounds

- Assume the attacker can cause imprecise 1-byte errors
- For each of 2^{32} key guesses,
Test if the *partial decryption* produces the expected 1-byte error.
(This can be optimized to require only 2 faulty encryptions to recover the full key)



AES-NI (New Instructions)

Instruction	Description
AESENC	Perform one round of an AES encryption flow
AESENCLAST	Perform the last round of an AES encryption flow
AESDEC	Perform one round of an AES decryption flow
AESDECLAST	Perform the last round of an AES decryption flow
AESKEYGENASSIST	Assist in AES round key generation
AESIMC	Assist in AES Inverse Mix Columns
PCLMULQDQ	Carryless multiply (CLMUL)

AES New Instructions - inside SGX

```
do
{
    i++;
    plaintext = <randomly generated>

    result1 = aes128_enc(plaintext);
    result2 = aes128_enc(plaintext);
} while (vec_equal_128(result1, result2) && i<iterations);
```

```
bagger> sudo ./aes-encrypt 100000 -262
```

```
|
```



Countermeasures for Fault Attacks

Analog Countermeasures

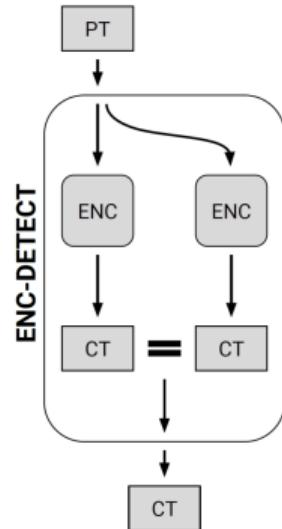


Detect anomalies^a

- Active fine wire meshes across IC → disruption is detected
- Power surge sensors
- Temperature sensors
- Light sensors

^aIBM 4767 Hardware Security Module battery-backed monitoring, meshes, light sensors, temperature sensors, etc. immediate deletion of keying material on tamper detection

Redundancy-based Countermeasure: Double Execution



Encrypt multiple times, compare result

- comparison at different granularities possible:
- encryption, single round, each operation, ...

But the attacker might be able to

- inject the same fault twice (difficult ...)
- or use more sophisticated methods (statistical attacks)

Take Aways



Attack → defense → next attack → next defense → ...

- different side channels, more sophisticated attacks
- a never-ending cat-and-mouse game

There is no “100% secure”, especially in the physical setting

- any device can be broken by a determined attacker

Our goal:

- Ensure that attack effort is much greater than the value of the secret
- or: Would you do an attack that costs millions to get a free tram ride?

Thanks to Peter Pessl for some of the slides!