

Information Security

System Security 4 - Sandboxing and Isolation

1 December 2023




- Bug-free code is hard to write
- **Impact** of exploits should be **minimized**
- Sometimes, untrusted code has to be executed
- **Restrict access** as much as possible

System Hardening

PoLP


»Every program and every privileged user of the system should operate using the least amount of privilege necessary to complete the job.«

- Jerome Saltzer, Communications of the ACM



- Important design decision
 - Only give **permissions** that are actually **needed**
 - Fewer permissions → fewer attack surfaces
- User account vs. admin account

Example: x86 Protection Rings




Least privileged



Most privileged




- Low-privilege user-space application can simply be executed
 - Drivers have high privileges (ring 2 to 0)
- Don't accept all drivers
- Only load drivers if they are signed by trusted vendor
- Root attacker cannot simply inject code into kernel




- UEFI supports **secure boot**
- UEFI ROMs, boot loader, kernel must be **signed**
- **Public key** in **firmware** to verify signatures
- Control-flow only handed over on successful verification

Example: Secure Boot Ubuntu




cert Microsoft UEFI CA certificate

cert Ubuntu CA certificate




- Bug in kernel allows accessing all user-space memory
 - Reduce the impact of kernel bugs
 - Explicitly enable/disable user-space access
- SMAP and SMEP



SMAP: Supervisor Mode Access Prevention


SMEP: Supervisor Mode Execution Prevention




- **SMEP** prevents **execution** of user-space code → never needed
- **SMAP** prevents **access** to user-space data → sometimes needed
- `stac` and `clac` instructions → enable/disable access
- **Every** user-space data **access** surrounded by `stac/clac`
- Supported in Linux, macOS, soon in Windows 10




- KAISER/KPTI is other way round
- **Unmap the kernel** in user space
- Kernel addresses are then **no longer present**
- Protection against **microarchitectural attacks** (e.g., Meltdown)



Kernel View



User View



context switch

Sandboxing



- A sandbox is a **restricted environment** for a program
- Resources of the process are strictly controlled:
 - own filesystem
 - no network connection
 - limited amount of memory
 - limited CPU time
 - ...
- Different approaches to sandboxing


- Multiple types of sandboxes
- Different advantages, disadvantages, and use cases



Language-level
Sandboxing



Rule-based
Execution




Container




Virtualization


Language-level Sandboxing




- Do not run native code
 - **Restrict** untrusted code on the **language** level
 - Languages without dangerous functionality (I/O, syscalls, ...)
- JavaScript, WebAssembly
- Access resources → ask user for permission



- Used in web browsers → website provides untrusted code
- Code cannot...
 - ...interact with the OS (syscalls)
 - ...communicate with other applications
 - ...access arbitrary memory (no pointers)
 - ...use unlimited memory
 - ...crash (memory safety)
- **No malicious** activity possible (in theory)




- Security guaranteed by the interpreter/runtime environment
- Interpreter does not provide dangerous functions
- Languages are **memory safe**
- A lot easier than sandboxing native code



- eBPF allows running **sandboxed user code** in kernel
- Originally to filter network packets
- Certain properties verified first:
 - Program **must terminate**
 - **No loops**/recursions (halting problem)
 - Jumps back only if they don't form loops
 - Call only to allowed functions
- Only loaded if **analyzed and verified**



- Runtime environments and interpreters are **complex**
 - Chrome JavaScript engine: \approx 1.9 million lines of code (2019)
 - Complexity introduces bugs \rightarrow sandbox escape
- \rightarrow Additionally sandbox the interpreter



- Chrome uses additional **site isolation**
 - Every **tab** is a **process**
- Exploited tab cannot access other tabs or browser

Rule-based Execution



- Rules what an application is allowed to do
- Usually multiple **rules** for an application
- Rules can be whitelists or blacklists
- Multiple rules are combined to a **policy**/profile



- Applications can use `seccomp-bpf` to restrict syscalls
 - First define which syscalls are required
 - Then block all other syscalls
 - Attacker is restricted to syscalls the application uses
- In many cases no `exec`



- seccomp-bpf is used by many (commercial) sandboxes
 - Docker
 - Firejail
 - Mbox
 - LXD
 - minijail
- It is even possible to **block** certain **syscall parameters** (e.g., no read except from standard input)

Seccomp Example

```
#include <stdio.h>
#include <seccomp.h>
#include <sys/prctl.h>
int main() {
    printf("step 1: init\n");
    prctl(PR_SET_NO_NEW_PRIVS, 1);
    prctl(PR_SET_DUMPABLE, 0);    // ptrace not allowed
    scmp_filter_ctx ctx;
    ctx = seccomp_init(SCMP_ACT_KILL); // blacklist everything
    // whitelist
    seccomp_rule_add(ctx, SCMP_ACT_ALLOW, SCMP_SYS(exit), 0);
    seccomp_rule_add(ctx, SCMP_ACT_ALLOW, SCMP_SYS(exit_group), 0);
    seccomp_rule_add(ctx, SCMP_ACT_ALLOW, SCMP_SYS(read), 0);
    seccomp_rule_add(ctx, SCMP_ACT_ALLOW, SCMP_SYS(write), 1,
                     SCMP_A0(SCMP_CMP_EQ, 1));
    seccomp_load(ctx);
    fprintf(stdout, "step 2: only 'write' to stdout\n");
    fprintf(stderr, "step 3: should be blocked\n");
}
```

- Compile with seccomp support

```
% gcc seccomp.c -lseccomp -o seccomp
```

- Run protected application


```
% ./seccomp  
step 1: init  
step 2: only 'write' to stdout  
[1]      23414 invalid system call  ./seccomp
```



- General approach: **mandatory access control** system (MAC)
- Applies to many resources, not only syscalls
- Rules have a
 - *Subject*: Process or thread
 - *Operation*: Access, write, execute, ...
 - *Object*: File, TCP port, shared memory, syscall, ...
- OS enforces policy (i.e., set of rules)



- Policies are created/installed by **administrator**
- Users cannot override policies
- Policies can be enforced (→ **kill application** on violation)...
- ...or just logged for later analysis



- In Windows as Mandatory Integrity Levels
- Implemented in Linux as **Linux Security Modules** (LSM)
- Different modules in the kernel
 - SELinux
 - AppArmor
 - Smack
 - TOMOYO Linux

Example Application

- Show man section 3 (C Library Functions)
- For example, man page of fopen

```
% c fopen
```

```
#include <stdio.h>
#include <unistd.h>
#include <string.h>

int main(int argc, char* argv[]) {
    if(argc > 1) {
        char* args[] = {"man", "3", argv[1], NULL};
        execvp(args[0], args);
    }
}
```

Example AppArmor Policy

```
#include <tunables/global>
```

```
/usr/bin/c {  
    #include <abstractions/base>  
    #include <abstractions/bash>  
    #include <abstractions/consoles>  
    #include <abstractions/evince>  
  
    /bin/dash ix,  
    /bin/less mrix,  
    /etc/groff/man.local r,  
    /etc/manpath.config r,  
    /usr/bin/c mr,  
    /lib/x86_64-linux-gnu/ld-*.so mr,  
    /usr/bin/groff mrix,  
    /usr/bin/groty mrix,
```

```
    /usr/bin/less mrix,  
    /usr/bin/locale mrix,  
    /usr/bin/man mrix,  
    /usr/bin/nroff mrix,  
    /usr/bin/nroff r,  
    /usr/bin/preconv mrix,  
    /usr/bin/tbl mrix,  
    /usr/bin/troff mrix,  
    /var/cache/man/oldlocal/index.db rk,  
    owner /home/*/.lessht r,
```

```
}
```



- Without policy: user can spawn shell from `man`
- Enforce policy

```
sudo aa-enforce /usr/bin/c
```


- Application still works, but “!/bin/bash” in `man` results in

```
sh: 1: /bin/zsh: Permission denied
```




- SELinux, AppArmor, and seccomp are **widely used**
- Not easy to create good policies...
- ...but **secure and efficient** for good policies
- Policies for popular applications can be found online

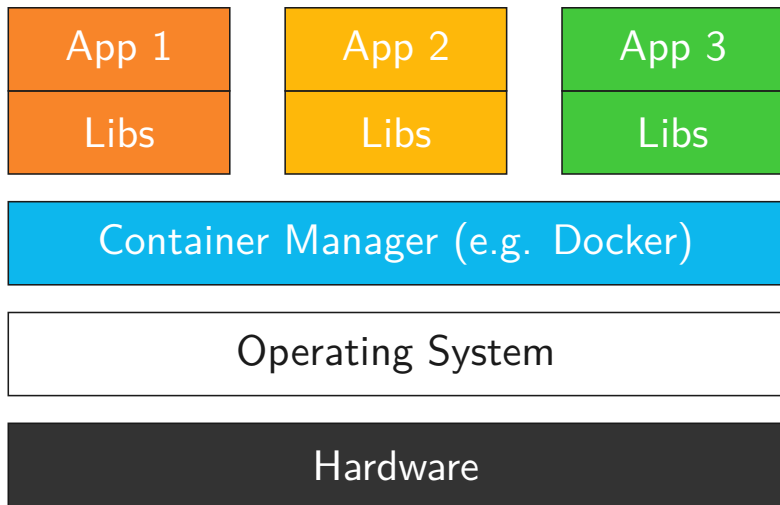
Container




- Containers are **operating-system-level virtualization**
- Allows multiple isolated user-space instances
- Every container is assigned resources (e.g., part of memory, folders, ...)
- Application in container can **only see assigned resources**




- One or multiple applications per container
- Own libraries but share the operating system
- File-system layer with copy on write
- “It works on my computer” → ship the environment






- Different **container manager**: Docker, OpenVZ, LXC, chroot, ...
- Require operating-system support
- Kernel is responsible for
 - Resource virtualization
 - Application isolation

→ **namespaces** and **cgroups** are the basis on Linux



- Namespaces **isolate system resources** between processes
- Default: all processes in same namespace
- Process can be started in new namespace
- **Limits** what the process (and it's children) can **see**
- Cannot interfere with other namespaces



- **Resources** which can be isolated using namespaces

Process ID Process sees only own and children processes

Mount Own mounts for process

Network Own network stack with virtual ethernet ports

IPC Interprocess communication isolation

UTS Own hostname and domain name

User ID Own set of users which can map to host users


Cgroup Hides the control group

Namespace Example

```
% top  
Tasks: 339 total, 1 running, 252 sleeping, 0 stopped  
KiB Mem : 24423136 total, 13416376 free, 5017528 used
```

Create process-id namespace and start shell

```
% sudo unshare --fork --pid --mount-proc /bin/bash  
$> top  
Tasks: 2 total, 1 running, 1 sleeping, 0 stopped  
KiB Mem : 24423136 total, 13043112 free, 5416304 used
```



- Control groups (cgroups) handle management and **accounting of resources**
- 12 different controllers (e.g., CPU, memory, I/O, ...)
- Every controller can have multiple cgroups
- A process (and its children) are in **one cgroup per controller**
- Controller and cgroups are in `/sys/fs/cgroup/`


Control Group Example

```
#include <stdio.h>
#include <unistd.h>
#include <sys/stat.h>
int main() {
    mkdir("/sys/fs/cgroup/memory/ml", 0600); // create cgroup
    FILE* f =
        fopen("/sys/fs/cgroup/memory/ml/memory.limit_in_bytes", "w");
    fprintf(f, "%d", 64*1024*1024); // 64MB memory limit
    fclose(f);
    f = fopen("/sys/fs/cgroup/memory/ml/cgroup.procs", "w");
    fprintf(f, "%d", getpid()); // restrict own pid
    fclose(f);

    setgid(1000); setuid(1000); // drop privileges
    execv("/bin/bash", NULL);
}
```


Control Group Example

```
% sudo swapoff -a
% sudo ./memlimit
$> whoami
mschwarz
$> stress -m 1 --vm-bytes 60000000
stress: info: [5432] dispatching hogs: 0 cpu, 0 io, 1 vm, 0 hdd
$> stress -m 1 --vm-bytes 70000000
stress: info: [5434] dispatching hogs: 0 cpu, 0 io, 1 vm, 0 hdd
stress: FAIL: [5434] (415) <-- worker 5435 got signal 9
stress: WARN: [5434] (417) now reaping child worker processes
stress: FAIL: [5434] (451) failed run completed in 0s
$> dmesg
Memory cgroup out of memory: Kill process 5435 (stress) score 908 or
      sacrifice child
Killed process 5435 (stress) total-vm:76600kB, anon-rss:59184kB, file-rss
      :196kB, shmem-rss:0kB
```



- Control groups limit **physical** resources
- Namespaces isolate **system** resources (including cgroups)
- Combine both → restrict resources for process(es)
- Basis of nearly all containers on Linux (e.g. Docker)

Installing Docker

```
curl -fsSL get.docker.com -o get-docker.sh  
sudo sh get-docker.sh
```

Using Docker:

```
docker run --rm -it ubuntu bash
```


Starts a shell inside a container



- Only use containers from **trusted sources**
- Limit the number of shared resources
- Keep **host** system up to date and **patched**
- Add **seccomp** (limit functionality) → additional security
- Unauthorized users should not interact with container manager



- No additional OS → small overhead
- Fast start-up time
- Many containers can run on one host
- No complicated configuration of policies

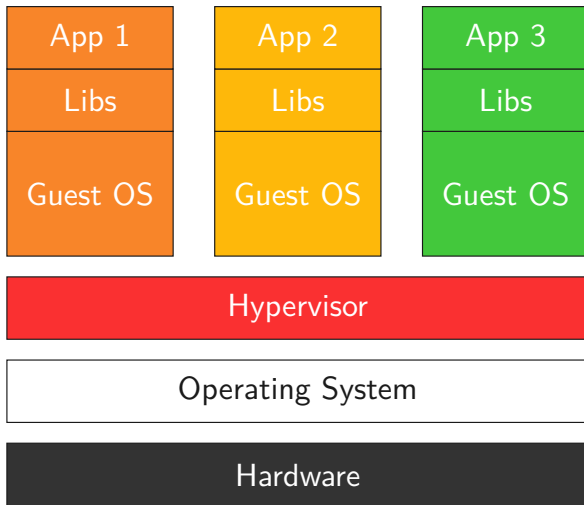



- **Shared kernel** of all containers and host
 - All containers must use same OS
- Kernel bugs are exploitable from containers
- **Exploiting the kernel** allows breaking out...
 - ...and taking over the whole host

Virtualization



- Do **not share kernel** anymore
- Emulate entire system → Virtual machine
- Process runs inside **own operating system**
- No access to host





- Different **types of hypervisors**

Bare metal Run directly on the hardware (e.g., Xen)

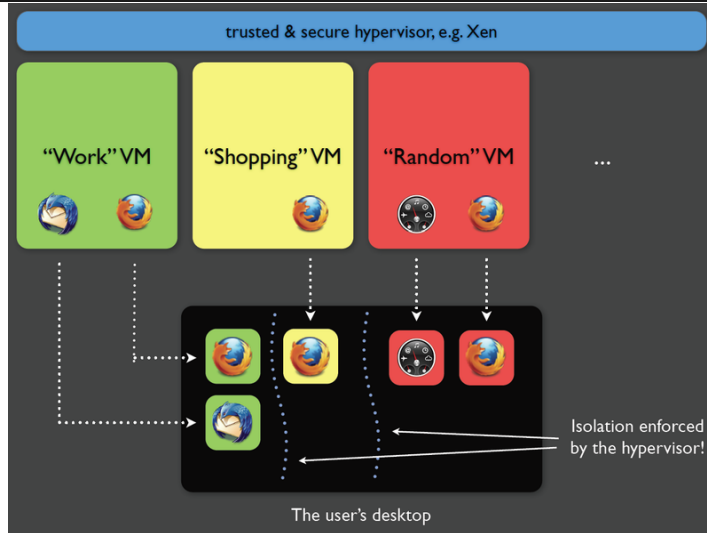
Hosted Run on top/as part of the host OS
(e.g., VirtualBox, KVM)

- Hypervisors **emulates** the machine **hardware**, e.g., graphic card
- OS and application are unaware of running inside VM




- Qubes OS is a security-focused OS
- Provides **security through isolation**
- Multiple security domains, isolated by hypervisor
- All **applications** run inside (different) **VMs**
- Malicious software is limited to one domain

Qubes OS Security Domains





- Virtualization provides **best isolation**
- Considered secure
- Applications not limited in functionality




- Large **resource overhead** compared to containers
- Requires a guest OS for every isolated application
- **Runtime overhead** (e.g., paging, traps to hypervisor)
- Still not 100 % secure




- VM escape: breaking out of a VM → interact with hypervisor
- Access to host and all other machines
- VM escape usually using memory safety violations
 - Mostly: bugs in drivers of emulated devices
 - Extremely powerful, but complicated to mount



- Multiple ways of sandboxing applications
- Higher security → often more overhead
- Sandboxing mechanisms can be combined
- Generic defense → damage control
- Sometimes only solution (e.g. legacy software)



- Interaction of sandboxing with side-channel attacks?
- Run on the same hardware → shared resources
- Often just require memory accesses and timer
- Available in most sandboxes
- No real protection against side-channel attacks



Microarchitectural attacks shown from

- JavaScript (Spectre, Prime+Probe, Rowhammer, ...)
- eBPF (Spectre)
- Docker (Prime+Probe, ...)
- VMs (Spectre, Prime+Probe, Rowhammer, ZombieLoad, ...)

Some **only** work from VMs → Foreshadow-NG


Isolation



- Sandboxes assume **trusted system** and **untrusted application**
- Protects the system from harm
- Sometimes, we want to **protect** the **application** from the system
- Assumption: untrusted system, trusted application
- **Isolation** of application



- Applications for isolation:
 - Working with **sensitive data** (e.g., passwords, money)
 - Distrusting the cloud provider
 - Intellectual property (e.g., algorithms)
 - Rights management (**DRM**)
- Ensures security even against active attacks



- Requires some form of **hardware support**
- Well-known isolation: user space - kernel space
- Protects OS against malicious applications
- Enforced by the hardware (→ page table)
- Similar concepts to protect application from OS




- Trusted computing base (TCB) is everything required to guarantee security
- Has to be trusted
- No security without a TCB
- Exploiting TCB → undermine entire security
- TCB should be as small as possible




- CPU and firmware usually in the TCB
 - Kernel and system programs usually in TCB
- Protected by the hardware (→ protection rings)
- For sandboxes: sandbox in TCB
 - What if we don't want to trust so many elements?


Trusted-Execution Environments



- Secure area of a CPU
- **Integrity and confidentiality** guarantees for code and data
- Hardware still shared with other applications
- (Nearly) no performance impacts



- Assumptions in TEEs:
 - Attacker **controls** the **OS**
 - Only the **CPU** is **trusted** (\rightarrow TCB)
- TEE **memory** is **encrypted** and inaccessible to OS
- TEE has access to OS




- Implementations for various CPUs
 - Intel: Software Guard Extension (SGX) and Management Engine (ME)
 - ARM and AMD: TrustZone
- Widely used in mobile phones




- Netflix uses Widevine DRM
- DRM in TrustZone
- Video is directly drawn on screen
- No app (not even root) can access video data

Execution Flow (SGX)






- Hardware-assisted protection of **sensitive data**
- Small **overhead**
- Could be abused for **malicious** software
- Bad code in TEEs is still **exploitable**
- No protection against **side-channel attacks**



- Interaction of TEE with side-channel attacks?
- Run on the same hardware → shared resources
- Stronger attacker: malicious operating system
- No real protection against side-channel attacks



- **Microarchitectural attacks** shown on SGX via
 - Branch predictors
 - Caches
 - Interrupt latency
 - Page tables
 - Exceptions (cf. Foreshadow)
 - Transient-execution attacks (cf. ZombieLoad)
- Considered **out-of-scope**



- Enclaves are black boxes
- Protected from all applications and OS
- What if they contain malicious code?
- Can we hide zero days?




- **Side-channel** attacks from SGX (Prime+Probe) → steal secrets from system
- **Fault** attacks from SGX (Rowhammer) → manipulate system/denial of service
- **Return-oriented programming** from SGX → break out of enclave

Intel's Statement

[...] Intel is aware of this research which is based upon assumptions that are outside the threat model for Intel SGX. The value of Intel SGX is to execute code in a protected enclave; however, Intel SGX does not guarantee that the code executed in the enclave is from a trusted source [...]

Hardware Isolation




- TEE is not fully isolated (→ shared hardware)
- Hardware security modules (HSM) are physically isolated
- Dedicated hardware, nothing shared
- Can be an external device or a plug-in card




Internal HSM




Photo by Wileyfh / CC BY



- Protection of **high-value** cryptographic keys
- Untrusted environment
- Attacker controls OS and has **physical access**
- Attacker tries to actively **attack** HSM



- Contains a **crypto processor** for
 - Secure key generation and management
 - Digital signatures
 - Data encryption/decryption
- Physical and logical **protection** of data
- Sometimes secure **timestamp** and strong **random** number generator



- **PKI** environments (e.g., certification authorities)
 - Store and handle asymmetric keys
- Card payment systems (**banks**)
 - Manage smart cards
 - Authorize transactions
- **Cryptocurrency** wallets
- Handy-Signatur



- Isolation allows **protecting applications** in hostile environments
- Less shared resources → better isolation
- Isolation is sometimes similar to sandboxing...
- ...but mostly an orthogonal problem
- Can be **combined** for isolation in both directions
- Choose the methods which best fit your **threat model**

