

# Principal Component Analysis

---

Machine Learning 1 — Lecture 6

25<sup>th</sup> April 2023

Robert Peharz

Institute of Theoretical Computer Science

Graz University of Technology

- So far, we considered **linear models** for **supervised learning**
- That is, we predicted some **target**  $y$  from an **input vector**  $\mathbf{x}$ :
  - **linear regression** (for regression)
  - **logistic regression** (for classification)
  - both can be made non-linear, by replacing  $\mathbf{x}$  with a non-linear transformation  $\phi(\mathbf{x}): \mathbb{R}^D \mapsto \mathbb{R}^K$
- Today, we will use a linear model for an **unsupervised learning** technique: **principal component analysis (PCA)**
- Thus, we will have a data set

$$\mathcal{D} = \left\{ \mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \right\}$$

of input vectors, but **no target**

- Goal of unsupervised learning: find “interesting structure” in the data

PCA is a technique for **dimensionality reduction**.

Generally, dimensionality reduction means that we replace our data

$$\mathcal{D} = \left\{ \mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \right\}$$

containing  $D$ -dimensional vectors with a new dataset

$$\mathcal{D}' = \left\{ \mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots, \mathbf{z}^{(N)} \right\}$$

of  $K$ -dimensional vectors,  $K \ll D$ , where  $\mathcal{D}'$  captures “certain characteristics” of  $\mathcal{D}$ .

Replace  $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$  with  $\mathcal{D}' = \{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(N)}\}$ .


**Potential target characteristics (qualitative):**

- if  $\mathbf{x}^{(i)}$  and  $\mathbf{x}^{(j)}$  are close, then also  $\mathbf{z}^{(i)}$  and  $\mathbf{z}^{(j)}$  should be close (**locality preserving**)
- $\mathbf{z}^{(i)}$  should allow for a good reconstruction of  $\mathbf{x}^{(i)}$ , i.e. there is some **decoder function**  $g(\mathbf{z}^{(i)}) \approx \mathbf{x}^{(i)}$
- $\mathcal{D}'$  captures as much information or variance as possible of  $\mathcal{D}$

### **Applications of dimensionality reduction:**

- lossy compression
- feature extraction (as input for other ML models)
- data visualization


**Iris Setosa**



petal

sepal

**Iris Versicolor**



petal

sepal

**Iris Virginica**



petal

sepal

Recall the **Iris dataset**, containing 4 features:

"sepalwidth", "sepalwidth", "petallength", "petalwidth", "class"

5.1, 3.5, 1.4, 0.2, Iris-setosa

4.9, 3.0, 1.4, 0.2, Iris-setosa

...

7.0, 3.2, 4.7, 1.4, Iris-versicolor



6.4, 3.2, 4.5, 1.5, Iris-versicolor

...


6.3, 3.3, 6.0, 2.5, Iris-virginica

5.8, 2.7, 5.1, 1.9, Iris-virginica

How to simultaneously plot 4 input features?



Computing the **first two principal components** transforms the data into a 2-dimensional space, which can be nicely visualized:





Principal Component Analysis: Formulation

Computing the PCA


Example: MNIST Digits

Linear Discriminant Analysis


Vectors  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_K$  are called **linearly independent**, when none of them can be expressed as a linear combination of the other vectors. That is, for each  $1 \leq i \leq K$  and any coefficients  $z_k$

$$\mathbf{x}_i \neq \sum_{k=1, k \neq i}^K z_k \mathbf{x}_k$$

Linearly dependent vectors




Linearly independent vectors



Let  $B = (\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_K)$  be a  $D \times K$ -matrix with orthonormal vectors as columns, where  $K < D$ . The **linear subspace** spanned by  $B$  is defined as all possible linear combinations of  $B$ 's columns:

$$\mathbf{V} = \left\{ \mathbf{v} = B\mathbf{z} \mid \mathbf{z} = (z_1, \dots, z_K)^T \in \mathbb{R}^K \right\}$$



Let  $B = (\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_K)$  be a matrix with  $K$  orthonormal vectors as columns, spanning a  $K$ -dimensional subspace  $\mathbf{V}$  of  $\mathbb{R}^D$ .

We can **project** an arbitrary vector  $\mathbf{x} \in \mathbb{R}^D$  onto  $\mathbf{V}$  by

- computing the **projection coefficients**  $B^T \mathbf{x} =: \mathbf{z} \in \mathbb{R}^K$
- computing the **projection/reconstruction**  $B\mathbf{z} =: \hat{\mathbf{x}} \in \mathbf{V}$
- thus,  $\hat{\mathbf{x}} = \underbrace{BB^T}_{\text{projection matrix}} \mathbf{x}$
- $\hat{\mathbf{x}}$  is the **closest point** in  $\mathbf{V}$  (in Euclidean distance) to  $\mathbf{x}$
- the **residual**  $\mathbf{r} = \mathbf{x} - \hat{\mathbf{x}}$  is always orthogonal to  $\hat{\mathbf{x}}$

Idea of PCA: Learn an “interesting” subspace  $\mathbf{V}$ .

Equivalently, learn a  $D \times K$ -matrix  $B$  of orthonormal vectors capturing  $K$  “interesting directions.”

# **Principal Component Analysis: Formulation**


---

# Data

PCA is an **unsupervised** technique, i.e. we have a training dataset of feature vectors

$$\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

We can interpret the data as a “point cloud” in  $\mathbb{R}^D$ :



## Step 1: Centering the Data

For PCA we assume that the average of the data is  $\bar{\mathbf{x}} = \mathbf{0}$ :


$$\bar{\mathbf{x}} := \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} = \mathbf{0}$$

If not, we simply remove the average from each training example:

$$\mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$$

This effectively moves the center of the point cloud to the origin.







Thus, we will w.l.o.g. assume that  $\bar{\mathbf{x}} = \mathbf{0}$ .

# Goal of PCA

- Find an “interesting subspace”  $\mathbf{V}$  of the data
- This amounts to finding a matrix  $B = (\mathbf{b}_1, \dots, \mathbf{b}_K)$  of  $K$  orthonormal basis vectors
- The basis vectors  $\mathbf{b}_k$  define  $K$  “interesting directions”
- $K$  is a **hyper-parameter** selected by the user
- For now, let's consider the special case  $K = 1$ , i.e. we want to find the “most interesting direction” in the data
- How to define “interestingness”?



Which of the three directions  $b$ ,  $b'$ ,  $b''$  is the most “interesting” one? Why?



$b'$  is the direction where

1. the **variance of the projected data is maximal**
2. the **sum of squared projection errors is minimal**

Let  $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$  be a centered training set of  $D$ -dimensional vectors.

**Among all unit vectors**, let  $\mathbf{b}_1$  be the unit vector where

- the **variance of the projected data**  
 $\text{var} \left( \left\{ z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)} \right\} \right)$  is **maximal**, where

$$z_1^{(i)} = \mathbf{b}_1^T \mathbf{x}^{(i)}$$

or, **equivalently**, where

- the sum of **squared projection errors**

$$\sum_{i=1}^N \|\mathbf{x}^{(i)} - \hat{\mathbf{x}}^{(i)}\|_2^2 = \sum_{i=1}^N \|\mathbf{x}^{(i)} - \mathbf{b}_1 z_1^{(i)}\|_2^2$$

is **minimal**.

The vector  $\mathbf{b}_1$  is called the **first principal direction** of dataset  $\mathcal{D}$ .

The values  $\{z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)}\}$  are called the **first principal component (scores)** of  $\mathcal{D}$ .

- the scores  $z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)}$  are a 1-dimensional representation of the original samples  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}$
- the vector  $\hat{\mathbf{x}}^{(i)} = \mathbf{b}_1 z_1^{(i)} = \mathbf{b}_1 \mathbf{b}_1^T \mathbf{x}^{(i)}$  is the projection of  $\mathbf{x}^{(i)}$  onto  $\mathbf{b}_1$  (or, the best reconstruction of  $\mathbf{x}^{(i)}$  given  $z_1^{(i)}$ )
- $\mathbf{b}_1$  is the **direction in which the data varies the most**, or the **direction of highest energy** – often, although not always, this direction carries the “most useful information” in the data

- the equivalence between **maximal variance** and **minimal projection/reconstruction error** (not proven here) is an example of **duality**, which occurs often in optimization
- by subtracting the projections  $\hat{\mathbf{x}}^{(i)}$  from the original samples  $\mathbf{x}^{(i)}$  and finding the principal component in the **residual vectors**  $\mathbf{r}^{(i)} = \mathbf{x}^{(i)} - \hat{\mathbf{x}}^{(i)}$ , one finds the **second principal component**, and so forth
- we will see in the following, however, that it is actually quite simple to find the first  $K$  principal components in one step

## Computing the PCA

---



- Assume that the dataset  $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$  is already centered, i.e.  $\bar{\mathbf{x}} = \mathbf{0}$
- We want to find the **first principal direction**  $\mathbf{b}_1$
- To simplify notation, we denote it by  $\mathbf{b}$ , dropping the subscript

## Problem Setup

- The PCA scores are given as

$$z^{(i)} = \mathbf{b}^T \mathbf{x}^{(i)}$$

- We wish to maximize the **empirical variance**

$$\text{var}(z^{(i)}) = \frac{1}{N} \sum_{i=1}^N (z^{(i)})^2 = \frac{1}{N} \sum_{i=1}^N (\mathbf{b}^T \mathbf{x}^{(i)}) (\mathbf{x}^{(i)T} \mathbf{b})$$

- We require that  $\mathbf{b}$  is a **unit vector**, i.e.

$$\mathbf{b}^T \mathbf{b} = 1, \quad \sum_{i=1}^D b_i^2 = 1$$

(note that if  $\mathbf{b}$  was not constrained to be a unit vector,  $\text{var}(z^{(i)})$  would grow to infinity with the norm of  $\mathbf{b}$ )

## Problem Setup

We can transform the problem as follows:

$$\begin{aligned}\max_{\mathbf{b}^T \mathbf{b} = 1} \text{var}(z^{(i)}) &= \frac{1}{N} \sum_{i=1}^N \mathbf{b}^T \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \mathbf{b} \\ &= \mathbf{b}^T \underbrace{\left( \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)}_{\text{=: covariance matrix } C} \mathbf{b}\end{aligned}$$


# Covariance Matrix

The (empirical) **covariance matrix**  $C$  is a  $D \times D$  matrix containing all possible empirical (co-)variances in the data:

$$C = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}.$$

(**Recall** that  $\mathbf{X}$  is the design matrix, whose rows are the  $N$  training examples).

- the diagonal element  $C_{dd}$  is the **empirical variance** of the  $d^{\text{th}}$  dimension:  $\frac{1}{N} \sum_i x_d^{(i)} x_d^{(i)}$
- the off-diagonal element  $C_{de}$  is the **empirical co-variance** between  $d^{\text{th}}$  and  $e^{\text{th}}$  dimension:  $\frac{1}{N} \sum_i x_d^{(i)} x_e^{(i)}$
- $C$  is always a symmetric matrix



## PCA: Optimization Problem

With the covariance matrix, the PCA problem can compactly be written as

$$\max_{\mathbf{b}^T \mathbf{b} = 1} \mathbf{b}^T \mathbf{C} \mathbf{b}$$

What is the optimal  $\mathbf{b}$  here? Evidently,  $\mathbf{b}$  must be a vector somehow connected with  $\mathbf{C}$ .

- Key and solution to the PCA problem: **eigen decomposition** of  $C$
- For any matrix  $C$ , when it holds for some vector  $\mathbf{v}$  and a scalar  $\lambda$  that

$$C\mathbf{v} = \lambda\mathbf{v}$$

then we call  $\mathbf{v}$  an **eigen vector** of  $C$  and  $\lambda$  its corresponding **eigen value**

- In this case, also  $\alpha\mathbf{v}$  is an eigen vector for any  $\alpha \neq 0$ ; thus, we consider only normalized eigen vectors

For any **symmetric**  $D \times D$ -matrix  $C$ , there are  $D$  **orthonormal** eigen vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_D$  and  $D$  sorted eigen values  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_D$ , such that  $C$  can be written as

$$C = VEV^T$$

where  $V = (\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_D)$  is an **orthonormal matrix** and  $E$  is diagonal:

$$E = \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_D \end{pmatrix}$$

For covariance matrices it further holds that all  $\lambda_i \geq 0$ ; they are so-called **positive definite matrices**.



# PCA Solution

Using the eigen decomposition we can write the PCA problem as

$$\max_{\mathbf{b}^T \mathbf{b} = 1} \mathbf{b}^T \mathbf{C} \mathbf{b} = \underbrace{\mathbf{b}^T \mathbf{V}}_{\mathbf{a}^T} \mathbf{E} \underbrace{\mathbf{V}^T \mathbf{b}}_{\mathbf{a}} = \sum_{i=1}^D \lambda_i a_i^2$$

- Since  $\mathbf{b}$  is a unit vector and  $\mathbf{V}$  is orthonormal, also  $\mathbf{a}$  is a unit vector, i.e.  $\sum_{i=1}^D a_i^2 = 1$
- Evidently,  $a_i^2 = (\mathbf{v}_i^T \mathbf{b})^2 \geq 0$
- Thus, the values  $a_i^2$  are **non-negative** and **sum to one**
- The variance is

$$\lambda_1 a_1^2 + \lambda_2 a_2^2 + \cdots + \lambda_D a_D^2$$

- When is this maximal?

- The variance becomes maximal when  $a_1^2 = 1$ :

$$\lambda_1 1 + \lambda_2 0 + \cdots + \lambda_D 0$$

- Since  $a_i = (\mathbf{v}_i^T \mathbf{b})$ , this exactly happens if  $\mathbf{b} = \mathbf{v}_1$ !
- **Thus, the first eigen vector  $\mathbf{v}_1$  of the covariance matrix is the first principal direction  $\mathbf{b}_1$**
- The corresponding eigen value  $\lambda_1$  is the variance which is “captured” in this direction (**explained variance**)


- More generally, the first  $K$  eigen vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_K$  are the **first  $K$  principal directions**
- The sum of their corresponding eigen values is the variance explained by the first  $K$  principal components:

$$\text{"explained variance"} = \sum_{i=1}^K \lambda_i$$


## **Example: MNIST Digits**

---

- collection of 60,000 images of handwritten digits
- images are  $28 \times 28$  pixels, with 256 gray scales
- white = 0, black = 255




By following a fixed pixel order (e.g. row-wise) we can convert each image into a vector. By “uprolling” the vector, we get back an image.




The first step is to center the data, i.e. subtract the average vector (image) from each sample:

$$\bar{\mathbf{x}} = \frac{1}{60,000} \sum_{i=1}^{60,000} \mathbf{x}^{(i)}, \quad \forall i: \mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$$

$\bar{\mathbf{x}} =$



$$C = \frac{1}{60,000} \sum_{i=1}^{60,000} \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$





$$V, E = \text{eigen}(C)$$

The first  $K$  columns of  $V$  contain the first  $K$  **principal directions**.  
We store them in a  $784 \times K$  matrix  $B$ .

**Note:** Some implementations return the eigenvalues in arbitrary order.  
In this case, sort them in descending order and sort the columns of  $V$  correspondingly.




black pixels correspond to large positive values, white pixels correspond to small negative values


The matrix  $B$  projects the data into a  $K$ -dimensional subspace:

$$\mathbf{Z} = \mathbf{X}\mathbf{B}$$


$K = 2$




$K = 3$



Reconstructions of four samples (first column) with 1, 10, 50, 100, 200 and 500 PCA components:



Reconstruction error (measured in mean squared error) and explained variance ( $\sum_i \lambda_i$ ) over  $K$ :



→ can be used to select  $K$

**Given:** dataset  $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$ , parameter  $K$

1. Center data

- $\bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$
- for all  $i$ :  $\mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$

2. Compute covariance matrix

$$C = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$

3. Compute eigen decomposition:  $V, E \leftarrow \mathbf{eigen}(C)$

4. Let  $B$  be the matrix consisting of the first  $K$  columns of  $V$

5. **return**  $B$

- The **principal component scores** (low-dimensional features) are then given as  $\mathbf{Z} = \mathbf{X}B$  ( $N \times K$ -matrix)
- Projections/reconstructions are given as  $\hat{\mathbf{X}} = \mathbf{Z}B^T$
- Add  $\bar{\mathbf{x}}$  again to the reconstructions

# Linear Discriminant Analysis

---


# Supervised Dimensionality Reduction

- PCA is **unsupervised**, i.e. there are no targets
- If one additionally has target values, one might use them to guide dimensionality reduction
- A classical method is Fisher's **linear discriminant analysis (LDA)**, using class information



# Separating Classes

- Let  $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$  be a binary classification dataset, i.e. each  $y^{(i)} \in \{-1, 1\}$
- We want to learn a linear function  $z = f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$  projecting  $D$ -dimensional vectors onto the real line
- This function should allow us to separate the two classes well
- How to select  $\mathbf{w}$ ?



**Note:**  $\|\mathbf{w}\|$  is an arbitrary scaling, hence we can assume  $\|\mathbf{w}\| = 1$ .

# Separating the Means?

**First idea:** Maximize the distance between projected means

- Let  $\mathbf{m}^-$  and  $\mathbf{m}^+$  be the means of the negative and positive classes, respectively:

$$\mathbf{m}^- = \frac{1}{N^-} \sum_{i: y^{(i)}=-1} \mathbf{x}^{(i)} \quad \mathbf{m}^+ = \frac{1}{N^+} \sum_{i: y^{(i)}=+1} \mathbf{x}^{(i)}$$


Here  $N^-$  and  $N^+$  are the number of positive and negative samples, respectively.

- Let  $\mathbf{m}^-$  and  $\mathbf{m}^+$  be the projected means, i.e.

$$\mathbf{m}^- = \mathbf{w}^T \mathbf{m}^- \quad \mathbf{m}^+ = \mathbf{w}^T \mathbf{m}^+$$

- It can be shown, that  $\mathbf{w} \propto \mathbf{m}^- - \mathbf{m}^+$ , i.e. the vector pointing from one mean to the other, maximizes the distance between the projected means  $\mathbf{m}^-$  and  $\mathbf{m}^+$

# Separating the Means?



- The idea works to a certain extent, but apparently we are not doing a perfect job
- Considerable overlap between the classes in projected space
- The problem is that we ignored the **covariance structure** of the point clouds

- Let  $z^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)}$  be the  $i^{\text{th}}$  projected sample
- Let again  $m^-$  and  $m^+$  be the projected means
- Further, let  $v^-$  and  $v^+$  be the **within-class variances** of the projected data:

$$v^- = \frac{1}{N^-} \sum_{i: y^{(i)} = -1} \left( z^{(i)} - m^- \right)^2 \quad v^+ = \frac{1}{N^+} \sum_{i: y^{(i)} = +1} \left( z^{(i)} - m^+ \right)^2$$

- The **Fisher criterion** is

$$J(\mathbf{w}) = \frac{(m^- - m^+)^2}{v^- + v^+}$$

- Increasing  $J(\mathbf{w})$ 
  - increases the quadratic distance between  $m^-$  and  $m^+$
  - and/or decreases variance (spread) within each class


- Fisher criterion  $J(\mathbf{w})$  has an analytic maximum, denoted as **linear discriminant analysis (LDA)**
- Let  $C^-$  and  $C^+$  be the covariance matrices of the negative and positive class in the original space

$$C^- = \frac{1}{N^-} \sum_{i: y^{(i)} = -1} (\mathbf{x}^{(i)} - \mathbf{m}^-)(\mathbf{x}^{(i)} - \mathbf{m}^-)^T$$
$$C^+ = \frac{1}{N^+} \sum_{i: y^{(i)} = +1} (\mathbf{x}^{(i)} - \mathbf{m}^+)(\mathbf{x}^{(i)} - \mathbf{m}^+)^T$$

- Then the maximum of  $J(\mathbf{w})$  is

$$\mathbf{w}_{LDA} \propto (C^- + C^+)^{-1}(\mathbf{m}^- - \mathbf{m}^+)$$

In contrast to the previous solution, LDA perfectly separates the two classes in this example:



- Similar to PCA, LDA uses covariance structure to determine a linear subspace
- Main difference is that PCA is unsupervised, while LDA is supervised
- LDA can be generalized to many classes as well, yielding a  $(|\mathcal{C}| - 1)$ -dimensional subspace