


Physical Addresses Only




- Programmer needs to manage physical memory space
 - Inconvenient & hard
 - Harder when you have multiple processes
- Difficult to support code and data relocation
 - Addresses are directly specified in the program
- Difficult to support multiple processes
 - Protection and isolation between multiple processes
 - Sharing of physical memory space
- Difficult to support data/code sharing across processes

Abstraction Virtual [[Memory]]

- **Programmer** sees **virtual memory**
 - Can assume the memory is “infinite”
- Reality: **Physical memory** size is much smaller than what the programmer assumes
- **The system** (system software + hardware, cooperatively) maps **virtual memory addresses** to **physical memory**
 - The system automatically manages the physical memory space **transparently to the programmer**
- + Programmer does not need to know the physical size of memory nor manage it → A small physical memory can appear as a huge one to the programmer → Life is easier for the programmer
- -- More complex system software and architecture
- Programmer does not deal with physical addresses
- Each process has its own mapping from virtual → physical addresses
- Enables
 - Code and data to be located anywhere in physical memory (**relocation**)
 - Isolation/separation of code and data of different processes in physical memory (**protection and isolation**)
 - Code and data sharing between multiple processes (**sharing**)

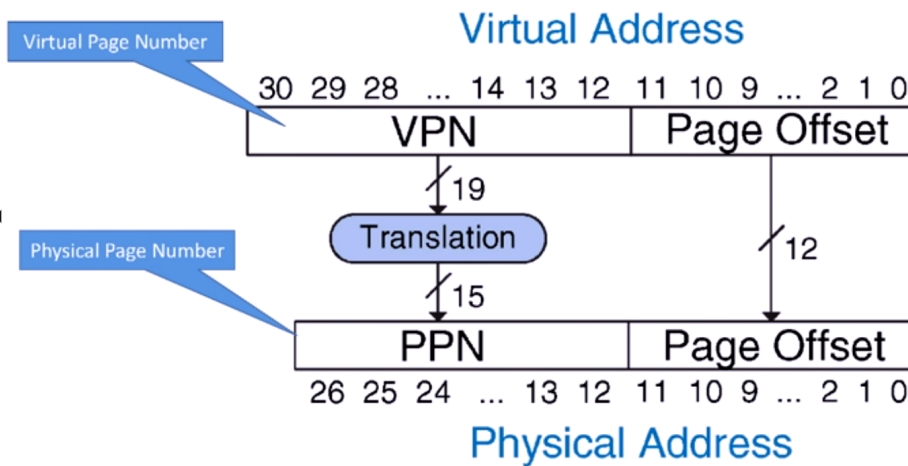
Virtual to Physical Address Translation



- **Address Translation:** The hardware converts virtual addresses into physical addresses via an OS-managed lookup table (page table)
- similar to [[Caching]]

- Virtual address space divided into **pages**
- Physical address space divided into **frames**
- A virtual page is mapped to
 - A physical frame, if the page is in physical memory
 - A location in disk, otherwise
- If an accessed virtual page is not in memory, but on disk
 - Virtual memory system brings the page into a physical frame and adjusts the mapping → this is called **demand paging**

Cache	Virtual Memory
Block	Page
Block Size	Page Size
Block Offset	Page Offset
Miss	Page Fault
Tag	Virtual Page Number



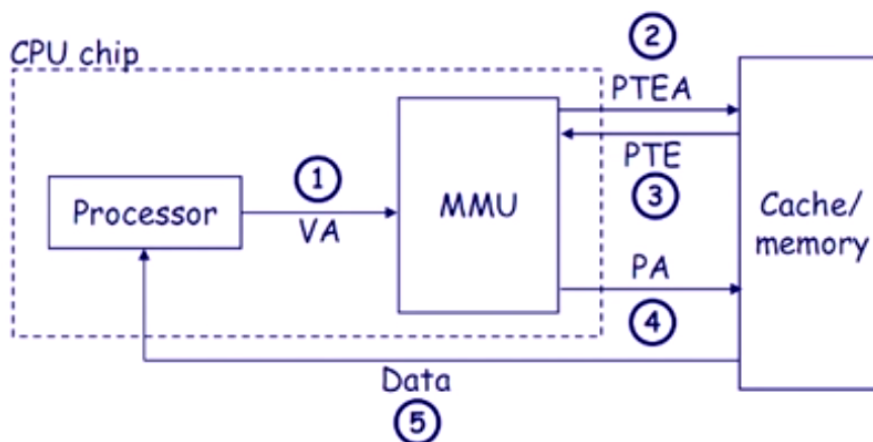
- Each **page table entry** has:
 - **Valid bit**: whether the virtual page is located in physical memory (if not, it must be fetched from the hard disk)
 - **Physical page number**: where the virtual page is located in physical memory
 - (Replacement policy, dirty bits)

Challenges

- Challenge 1: **Page table is large**
 - at least part of it needs to be located in physical memory
 - solution: multi-level (hierarchical) page tables
- Challenge 2: **Each instruction fetch or load/store requires at least two memory accesses:**
 1. one for address translation (page table read)
 2. one to access data with the physical address (after translation)
- Translation Lookaside Buffer TLB
 - **Idea: Cache the page table entries (PTEs) in a hardware structure in the processor to speed up address translation**
 - Small cache of most recently used translations (PTEs)
 - Reduces number of memory accesses required for *most* instruction fetches and loads/stores to only one

Virtual Memory Support

- Virtual memory **requires both HW+SW support**
 - Page Table is in memory
 - Can be cached in special hardware structures called Translation Lookaside Buffers (TLBs)



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to L1 cache
- 5) L1 cache sends data word to processor