

-

Terminology

- lock
 - logical synchronization of TXs
 - blocks access to db objects
- latch
 - physical synchronization of access
 - blocks access to shared data structures
- pessimistic concurrency control
 - assumes error will happen
 - thus locks schemes
 - lock-based database scheduler
 - full serialization of TXs
- optimistic concurrency control
 - assumes error will not happen
 - no locks but validation phase afterwards
 - * check of conflicts
 - timestamp-based database schedulers
- mixed concurrency control
 - combines PCC and OCC
 - might return synchronization errors (deadlocks)


Serializability Theory

Operations of Transaction T_j

- Read and write operations of A by T_j : $r_j(A)$ $w_j(A)$
- Abort of transaction T_j : a_j (unsuccessful termination of T_j)
- Commit of transaction T_j : c_j (successful termination of T_j)

Schedule S

- Operations of a transaction T_j are executed in order
- Multiple transactions may be executed concurrently
- Schedule describes the total ordering of operations



Equivalence of Schedules S_1 and S_2

- Read-write, write-read, and write-write dependencies on data object A executed in same order:

$$r_i(A) <_{S_1} w_j(A) \Leftrightarrow r_i(A) <_{S_2} w_j(A)$$

$$w_i(A) <_{S_1} r_j(A) \Leftrightarrow w_i(A) <_{S_2} r_j(A)$$

$$w_i(A) <_{S_1} w_j(A) \Leftrightarrow w_i(A) <_{S_2} w_j(A)$$

Example Serializable Schedules

- Input TXs
 T_1 : BOT $r_1(A)$ $w_1(A)$ $r_1(B)$ $w_1(B)$ c_1
 T_2 : BOT $r_2(C)$ $w_2(C)$ $r_2(A)$ $w_2(A)$ c_2
- Serial execution
 $r_1(A)$ $w_1(A)$ $r_1(B)$ $w_1(B)$ c_1 $r_2(C)$ $w_2(C)$ $r_2(A)$ $w_2(A)$ c_2
- Equivalent schedules
 $r_1(A)$ $r_2(C)$ $w_1(A)$ $w_2(C)$ $r_1(B)$ $r_2(A)$ $w_1(B)$ $w_2(A)$ c_1 c_2
 $r_1(A)$ $w_1(A)$ $r_2(C)$ $w_2(C)$ $r_1(B)$ $w_1(B)$ $r_2(A)$ $w_2(A)$ c_1 c_2
- Wrong schedule
 $r_1(A)$ $r_2(C)$ $w_2(C)$ $r_2(A)$ $w_1(A)$ $r_1(B)$ $w_1(B)$ $w_2(A)$ c_1 c_2

Serializability Graph (conflict graph)

- Operation dependencies (read-write, write-read, write-write) aggregated
- Nodes: transactions; edges: transaction dependencies
- Transactions are serializable (via topological sort) if the graph is acyclic

- Given two transactions T_1 and T_2 , which pairs of the following three schedules are equivalent? Explain for each pair (S_1 - S_2 , S_1 - S_3 , S_2 - S_3) why they are equivalent or non-equivalent. [5 points]


- $T_1 = \{r_1(a), r_1(c), w_1(a), w_1(c)\}$
- $T_2 = \{r_2(b), w_2(b), r_2(c), w_2(c)\}$

Schedules

- $S_1 = \{r_1(a), r_1(c), w_1(a), w_1(c), r_2(b), w_2(b), r_2(c), w_2(c)\} = \{T_1, T_2\}$
 → $S_1 \equiv S_2$ (equivalent, because $r_2(b), w_2(b)$ independent of T_1)
- $S_2 = \{r_1(a), r_2(b), r_1(c), w_1(a), w_2(b), w_1(c), r_2(c), w_2(c)\}$
 → $S_1 \not\equiv S_3$ (transitive)
- $S_3 = \{r_1(a), r_2(b), r_1(c), w_1(a), w_2(b), r_2(c), w_1(c), w_2(c)\}$
 → $S_2 \not\equiv S_3$ (non-equivalent, because $w_1(c), r_2(c)$ of c in different order)

Locking Schemes

- exclusive/write x-lock
 - only current lock may write
- shared/read s-lock
 - current lock and other locks may read
- multi-granularity-locking
 - abuses hierarchy of db objects
 - intentional x/s-lock



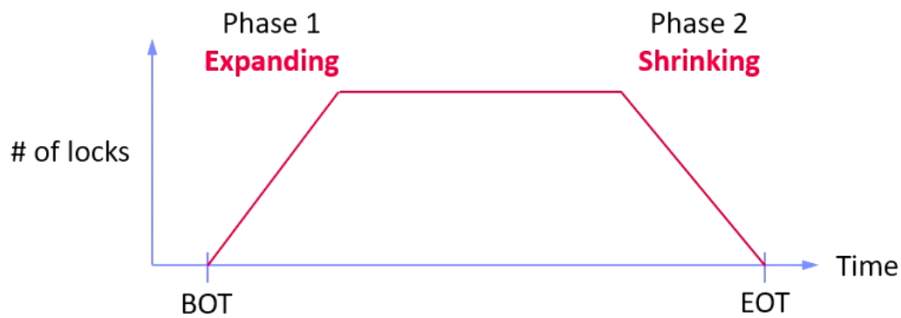
- lock compatibility

	None	S	X	IS	IX
S	Yes	Yes	No	Yes	No
X	Yes	No	No	No	No
IS	Yes	Yes	No	Yes	Yes
IX	Yes	No	No	Yes	Yes

Two-Phase Locking

- concurrency protocol that guarantees serializable
- pessimistic concurrency control
- expanding phase
 - acquires locks needed by TX


- shrinking phase
 - release locks acquired by TX



- potential problems fixed

- **Strict 2PL (S2PL) and Strong Strict 2PL (SS2PL)**

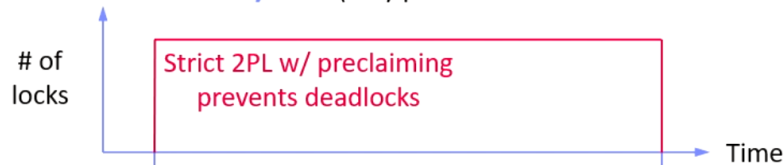
- **Problem:** Transaction rollback can cause (**Dirty Read**)
 - Release all X-locks (S2PL) or X/S-locks (SSPL) **at end of transaction (EOT)**




- deadlock

- **Strict 2PL w/ pre-claiming (aka conservative 2PL)**

- Problem: incremental expanding can cause deadlocks for interleaved TXs
 - **Pre-claim all necessary locks** (only possible if entire TX known + **latches**)



*



*

▪ #1 Deadlock Prevention

- Guarantee that deadlocks can't happen
- E.g., via **pre-claiming** (but overhead and not always possible)

DEADLOCK, as this will never happen



▪ #2 Deadlock Avoidance

- Attempts to avoid deadlocks before acquiring locks via timestamps per TX
- **Wound-wait** (T1 locks something held by T2 → if $T1 < T2$, restart T2)
- **Wait-die** (T1 locks something held by T2 → if $T1 > T2$, abort T1 but keep TS)

▪ #3 Deadlock Detection

- Maintain a wait-for graph of blocked TX (similar to serializability graph)
- Detection of cycles in graph (on timeout) → abort one or many TXs

*

Timestamp Ordering

- optimistic concurrency control
- low overhead scheme if conflicts are rare
- TXs get timestamp at BOT
- each data object has read and write timestamp
- timestamp comparison to validate access/abort
- no locks but latches

- **Read Protocol $T_j(A)$**
 - If $TS(T_j) \geq writeTS(A)$: **allow read**, set $readTS(A) = \max(TS(T_j), readTS(A))$
 - If $TS(T_j) < writeTS(A)$: **abort T_j** (older than last modifying TX)
- **Write Protocol $T_j(A)$**
 - If $TS(T_j) \geq readTS(A)$ AND $TS(T_j) \geq writeTS(A)$: **allow write**, set $writeTS(A) = TS(T_j)$
 - If $TS(T_j) < readTS(A)$: **abort T_j** (older than last reading TX)
 - If $TS(T_j) < writeTS(A)$: **abort T_j** (older than last modifying TX)
- **BEWARE:** Timestamp Ordering requires additional handling of dirty reads, and concurrent transactions in general (e.g., via abort or versions)

Optimistic Concurrency Control

- **Read Phase**
 - Initial reads from DB, **repeated reads and writes into TX-local buffer**
 - Maintain **ReadSet(T_j)** and **WriteSet(T_j)** per transaction T_j
 - TX seen as read-only transaction on database
- **Validation Phase**
 - Check read/write and write/write conflicts, **abort on conflicts**
 - BOCC (Backward-oriented concurrency control) – check all older TXs T_i that finished (EOT) while T_j was running ($EOT(T_i) \geq BOT(T_j)$)
 - **Serializable:** if $EOT(T_i) < BOT(T_j)$ or $WSet(T_i) \cap RSet(T_j) \neq \emptyset$
 - **Snapshot isolation:** $EOT(T_i) < BOT(T_j)$ or $WSet(T_i) \cap WSet(T_j) \neq \emptyset$
 - FOCC (Forward-oriented concurrency control) – check running TXs
- **Write Phase**
 - Successful TXs: propagate TX-local buffer into the database and log
 - Unsuccessful TXs: discard the TX-local buffer