

Failure Types

- transaction failures
 - violated integrity constraints
 - R1-Recovery: partial undo of uncommitted tx
- system failures
 - HW/OS system crash, power outage, ...
 - kill all current TX
 - does not lose persistent data
 - R2-Recovery: partial redo of committed TX
 - R3-Recovery: global undo of uncommitted TX
- media failure
 - hard disk errors (non-restorable)
 - lose persistent data ==> need back up
 - R4-Recovery: global redo of all committed TX

Database Transaction Log

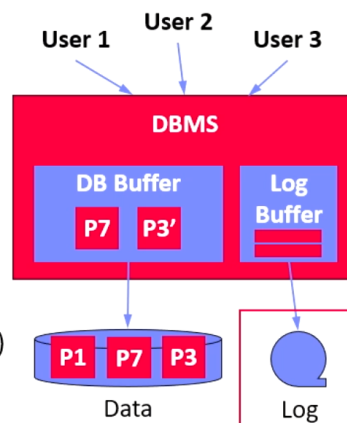
Database Architecture

- **Page-oriented storage** on disk and in memory (DB buffer)
- Dedicated **eviction algorithms**
- Modified in-memory pages marked as dirty, flushed by cleaner thread
- **Log**: append-only TX changes
- Data/log often placed on different devices and periodically archived (backup + truncate)

Write-Ahead Logging (WAL)

- The log records representing changes to some (dirty) data page must be on **stable storage before the data page** (UNDO - atomicity)
- **Force-log on commit** or full buffer (REDO - durability)
- **Recovery**: forward (REDO) and backward (UNDO) processing
- Log sequence number (LSN)

[C. Mohan, Donald J. Haderle, Bruce G. Lindsay, Hamid Pirahesh, Peter M. Schwarz: ARIES: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using



Logging Types

#1 Logical (Operation) Logging

- REDO: **log operation (not data)** to construct after state
- UNDO: **inverse operations** (e.g., increment/decrement), not stored
- Non-determinism** cannot be handled, more flexibility on locking

#2 Physical (Value) Logging

- REDO: **log REDO (after) image** of record or page
- UNDO: **log UNDO (before) image** of record or page
- Larger space overhead** (despite page diff) for set-oriented updates

```
UPDATE Emp
SET Salary=Salary+100
WHERE Dep='R&D';
```

Restart Recovery (ARIES)

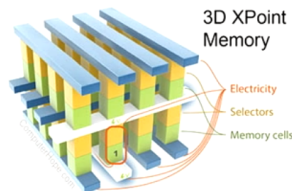
- Conceptually: take database checkpoint and replay log since checkpoint
- Operation and value locking**; stores log seq. number (LSN, PageID, PrevLSN)
- Phase 1 Analysis**: determine winner and loser transactions
- Phase 2 Redo**: replay all TXs in order **[repeating history]** → **state at crash**
- Phase 3 Undo**: replay uncommitted TXs (losers) in reverse order

Recovery on Storage Class Memory

Excursus: Recovery on Storage Class Memory

Background: Storage Class Memory (SCM)

- Byte-addressable, persistent memory** with higher capacity, but latency close to DRAM
- Examples**: Resistive RAM, Magnetic RAM, Phase-Change Memory (e.g., **Intel 3D XPoint**)



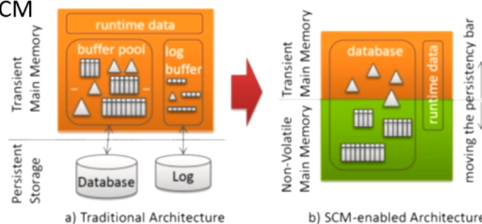
[Credit: <https://computerhope.com>]

SOFORT: DB Recovery on SCM

- Simulated DBMS prototype on SCM
- Instant recovery by trading TX throughput vs recovery time **(% of data structures on SCM)**



[Ismail Oukid, Wolfgang Lehner, Thomas Kissinger, **Thomas Willhalm**, Peter Bumbulis: Instant Recovery for Main Memory Databases. **CIDR 2015**]



Write-Behind Logging (for hybrid SCM)

- Update persistent data (SCM) on commit, log change metadata + timestamps → **1.3x**

[Joy Arulraj, Matthew Perron, Andrew Pavlo: Write-Behind Logging. **PVLDB 2016**]

