


## Table Scan vs Index Scan

- for highly selective predicates
  - index scan asymptotically way better than table scan
- index scan higher per tuple overhead
  - break even at ~5% output ratio
  - index scan better if filter ratio below ~5%
- multi-column predicates
  - fetch/RID-list intersection

### Table Scan




### Index Scan




## Use Cases for Indexes


### Lookups / Range Scans



### Unique Constraints



### Index Nested Loop Joins



### Aggregates (count, min/max)



## Create Index

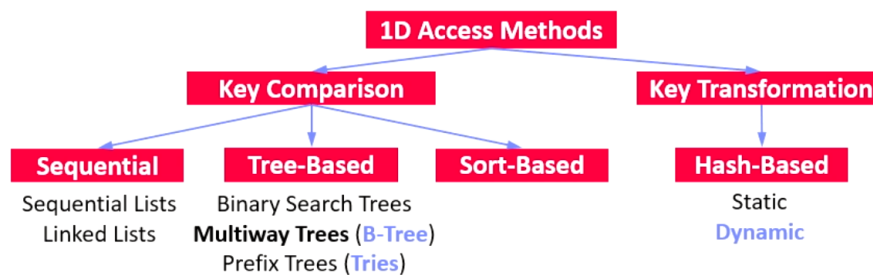
- create secondary (nonclustered) index on set of attributes
  - clustered: tuples sorted by index
  - nonclustered: sorted attribute with tuple references

```
CREATE INDEX ixStudLname
ON Students USING btree
(Lname ASC NULLS FIRST);
```

```
DROP INDEX ixStudLname;
```

- 
- allows specifying uniqueness, order, indexing method
- postgresSQL methods
  - [[Binary Search and BTree]]
  - hash
  - gist
  - gin

### Classification of Index Structures



### ND Access Methods

- Linearization of ND key space + 1D indexing
- Multi-dimensional trees and hashing (e.g.,
- Spatial index structures (e.g., R tree)