

Join Ordering Problem


- given a join query graph, find the optimal join ordering
- usually NP-hard
 - polynomial algorithms exist for special cases
- search space sizes

	Chain (no CP)			Star (no CP)		Clique / CP (cross product)		
	left-deep	zig-zag	bushy	left-deep	zig-zag/bushy	left-deep	zig-zag	bushy
n	2^{n-1}	2^{2n-3}	$2^{n-1}C(n-1)$	$2(n-1)!$	$2^{n-1}(n-1)!$	$n!$	$2^{n-2}n!$	$n! C(n-1)$
5	16	128	224	48	384	120	960	1,680
10	512	~131K	~2.4M	~726K	~186M	~3.6M	~929M	~17.6G

C(n) ... Catalan Numbers


Join Order Search Strategies

- tradeoff between optimal plan and compile time



-
- naive full enumeration
 - infeasible for large queries
- exact dynamic programming
 - guarantees optimal plan
 - often too expensive (beyond 20 relations)
 - bottom-up/top-down approach
- greedy/heuristic algorithms
 - tests out some
 - ignore worst

- further optimization on best
 - * e.g. random mutations (as with genetics)
- approximate algorithms



Greedy Join Ordering

- does not always return optimal join ordering
- algorithm:
 - calculate cost of each two table join combination
 - calculate costs of previous join with next table
 - repeat until every table is used
- example

▪ Part \bowtie Lineorder \bowtie Supplier \bowtie σ (Customer) \bowtie σ (Date), **left-deep plans**

#	Plan	Costs	#	Plan	Costs
1	Lineorder \bowtie Part	30M	3	((Lineorder \bowtie σ (Date)) \bowtie σ (Customer)) \bowtie Part	120M
	Lineorder \bowtie Supplier	20M		((Lineorder \bowtie σ (Date)) \bowtie σ (Customer)) \bowtie Supplier	105M
	Lineorder \bowtie σ (Customer)	90K	4	((((Lineorder \bowtie σ (Date)) \bowtie σ (Customer)) \bowtie Supplier) \bowtie Part	135M
	Lineorder \bowtie σ (Date)	40K			
	Part \bowtie Customer	N/A			
			
2	((Lineorder \bowtie σ (Date)) \bowtie Part)	150K			
	((Lineorder \bowtie σ (Date)) \bowtie Supplier)	100K			
	((Lineorder \bowtie σ (Date)) \bowtie σ (Customer))	75K			

Note: Simple $O(n^2)$ algorithm for left-deep trees;
 $O(n^3)$ algorithms for bushy trees existing (e.g., GOO)

Dynamic Programming Join Ordering

- exact enumeration via dynamic programming
 - tries to find optimal substructures first
 - overlapping subproblems allow for memoization

* reuse already retrieved data

- e.g. DPSize
 - bottom-up
 - split into independent subproblems
 - solve subproblems
 - combine solutions

▪ **Example**

		Q1+Q1		Q1+Q2, Q2+Q1		Q1+Q3, Q2+Q2, Q3+Q1	
Q1	Plan	Q2	Plan	Q3	Plan	Q4	Plan
{C}	Tbl, IX	{C,L}	$L \bowtie C, C \bowtie L$	{C,D,L}	$(L \bowtie C) \bowtie D, D \bowtie (L \bowtie C), (L \bowtie D) \bowtie C, C \bowtie (L \bowtie D)$	{C,D,L,P}	$((L \bowtie C) \bowtie D) \bowtie P, P \bowtie ((L \bowtie C) \bowtie D)$
{D}	Tbl, IX	{D,L}	$L \bowtie D, D \bowtie L$	{C,L,P}	$(L \bowtie C) \bowtie P, P \bowtie (L \bowtie C), (P \bowtie L) \bowtie C, C \bowtie (P \bowtie L)$	{C,D,L,S}	...
{L}	...	{L,P}	$L \bowtie P, P \bowtie L$	{C,L,S}	...	{C,L,P,S}	...
{P}	...	{L,S}	$L \bowtie S, S \bowtie L$	{D,L,P}	...	{D,L,P,S}	...
{S}	...	{C,D}	N/A	{D,L,S}	...	Q1+Q4, Q2+Q3, Q3+Q2, Q4+Q1	
		{L,P,S}	...	Q5	Plan
						{C,D,L,P,S}	...

[[Plan Optimization]]