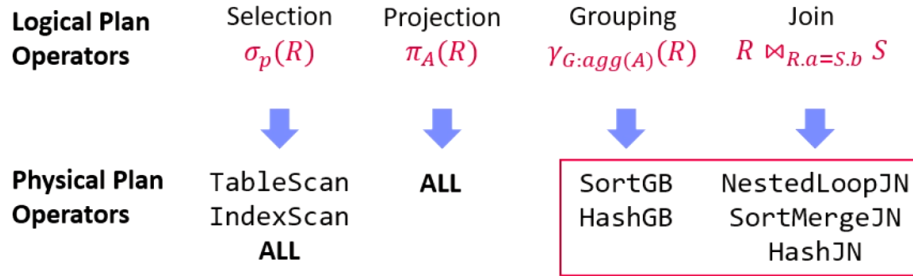


## Overview Plan Operators

- different operators for different data and query characteristics



- Lecture 07
This Lecture

## Nested Loop Join


- most general join operator
- no ordering/indexing
- slow

- Algorithm** (pseudo code)
 

```

for each s in S
  for each r in R
    if( r.RID  $\theta$  s.SID )
      emit concat(r, s)
      
```

How to implement **next()**?



- Complexity**
  - Complexity: Time:  **$O(N * M)$** , Space:  **$O(1)$**
  - Pick smaller table as inner if it fits entirely in memory (buffer pool)

## Block/Index Nested Loop Join

### Block Nested Loop Join

- Avoid I/O by blocked data access
- Read blocks of  $b_R$  and  $b_S$  R and S pages
- Complexity unchanged but potentially much fewer scans

```


for each block  $b_R$  in R
  for each block  $b_S$  in S
    for each r in  $b_R$ 
      for each s in  $b_S$ 
        if( r.RID  $\theta$  s.SID )
          emit concat(r, s)
      
```

### Index Nested Loop Join

- Use index to locate qualifying tuples ( $=$ ,  $>$ ,  $<$ ,  $\leq$ ,  $\geq$ )
- Complexity (for equivalence predicates):  
Time:  **$O(N * \log M)$** , Space:  **$O(1)$**

```

for each r in R
  for each s in  $S.IX(\theta, r.RID)$ 
    emit concat(r, s)
      
```




## Sort Merge Join

### Overview

- **Sort Phase:** sort the input tables R and S (w/ external sort algorithm)
- **Merge Phase:** step-wise merge with lineage scan

### Algorithm (Merge, PK-FK)

```
Record next() {
  while( curR!=EOF && curS!=EOF ) {
    if( curR.RID < curS.SID )
      curR = R.next();
    else if( curR.RID > curS.SID )
      curS = S.next();
    else if( curR.RID == curS.SID ) {
      t = concat(curR, curS);
      curS = S.next(); //FK side
      return t;
    }
  }
  return EOF;
}
```



### Complexity

- Time (unsorted vs sorted):  $O(N \log N + M \log M)$  vs  $O(N + M)$
- Space (unsorted vs sorted):  $O(N + M)$  vs  $O(1)$

## Hash Join

### Overview

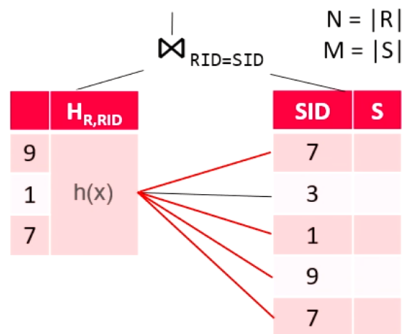
- **Build Phase:** read table S and build a hash table  $H_S$  over join key
- **Probe Phase:** read table R and probe  $H_S$  with the join key

### Algorithm (Build+Probe, PK-FK)

```
Record next() {
  // build phase (first call only)
  while( (r = R.next()) != EOF )
    Hr.put(r.RID, r);

  // probe phase
  while( (s = S.next()) != EOF )
    if( Hr.containsKey(s.SID) )
      return concat(Hr.get(s.SID), s);

  return EOF;
}
```



### Complexity

- Time:  $O(N + M)$ , Space:  $O(N)$
- Classic hashing: p in-memory partitions of  $H_r$  w/ p scans of R and S


## Group By Types

### Recap: Classification of Aggregates (04 Relational Algebra)

- Additive, semi-additive, additively-computable, others

### Sort Group-By

- Similar to sort-merge join (Sort, GroupAggregate)
- Sorted group output



### Hash Group-By

- Similar to hash join (HashAggregate)
- Higher temporary memory consumption
- Unsorted group output
- #1 w/ **tuple grouping**
- #2 w/ **direct aggregation** (e.g., count)
- Beware:** cache-unfriendly if many groups ( $\text{size}(H) > \text{L2/L3 cache}$ )

