


Memory safety

- spatial violation
 - buffer overflow/read
 - out-of-bounds reads
 - null pointer dereference
- temporal violation
 - use after free
 - double free
 - use of uninitialized memory
- stack frame layout

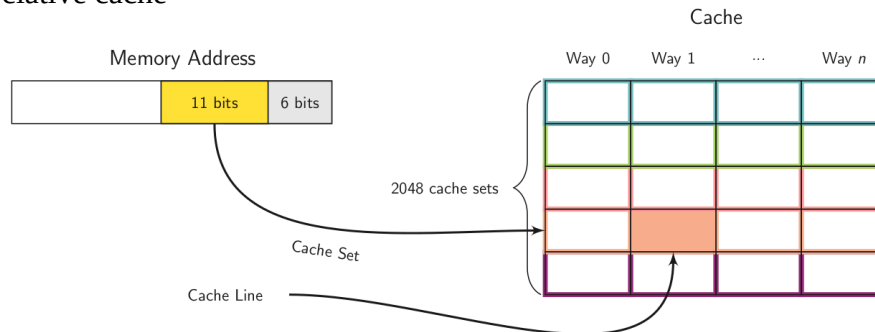


-
- prevent (some) buffer overflows with
 - stack canary
 - * random number stored below stack frame and buffer
 - * before returning check for overwrite
 - safe and unsafe stack
 - * store buffers on unsafe stack
 - ASLR
 - * randomizes location of memory
 - * requires large enough randomization range
 - * addresses must not be leaked
- control flow integrity CFI
 - Control-flow **graph** must be **correctly constructed**
 - **Function pointers** cannot be protected if destination set is large
 - Some functions (e.g., library functions) have **many call locations** and therefore return locations
-


Side-channel attacks

- passively observe physical properties
- information leakage through side-effects
 - power consumption

- execution time
 - * preventable through constant runtime and control flow
- CPU caches
 - * fast on cache hit, slow on miss
- set associative cache



- each memory address has a designated cache set to be cached in
- Flush+Reload
 - requires shared memory
 - determine which memory locations have been accessed by measuring the time to access it
 - * attacker flushes cache line
 - * victim might access memory locations of this cache line
 - * attacker reloads the cache line and measures the time
 - ◆ short => victim accessed the location
 - keystrokes can be retrieved because they cause code execution in shared library (e.g. libgdk)
- covert channel
 - two processes communicate over secret channel



transient execution attacks

- meltdown
 - read data at any address using out-of-order execution

```
char data = *(char*) 0xffffffff81a000e0;
array[data * 4096] = 0;
```

- *
 - combined with Flush+Reload
 - index of cache hit reveals data
 - preventable with KAISER/KPTI
 - * unmap kernel pages in user space
 - * kernel addresses are no longer present
- NG-Foreshadow
 - leak data from L1 cache of host into VM
- spectre
 - exploit control flow predictions
 - speculative execution runs

```
if (x < array1_size)
    y = array2[array1[x] * 256];
```

- *
 - * array1 is a valid buffer
 - * x is a way too big and causes an overread into the victim's memory
 - * array2 is uncached shared memory
- crashes because it accesses not allowed memory
 - * still accesses the page (speculative execution) and rollbacks the operation
 - * accessing page at array1[x] causes a cache hit

Fault attack

- actively manipulate device to induce faults
- Row hammer
 - accesses to nearby DRAM rows cause cell to leak energy which may cause bit flips
- Glitching/Skipping attacks
 - may corrupt data/skip instructions
- Undervolting
- countermeasures
 - detect anomalies
 - * active fine wire meshes across IC
 - * power surge/temperature/light sensors?
 - double execution
 - * unlikely to produce same fault twice

Sandboxing and Isolation


- Principle of Least Privilege

»Every program and every privileged user of the system should operate using the least amount of privilege necessary to complete the job.«

- e.g. x86 Protection Rings
- drivers have higher privileges than user-space applications
 - * only accept drivers signed by trusted vendor
 - * root attacker cannot inject code into kernel

- Secure Boot

- UEFI ROMs, boot loader, kernel must be signed
- public key in firmware to verify signatures
- execute after verification



- Supervisor Mode Access/Execution Prevention

- prevent access to user-space data
- prevent execution of user-space code
- opposite of KAISER/KPTI

- Sandboxing

- restricted environment to execute program in
- resources strictly controlled
 - * own filesystem
 - * no network connection
 - * limited memory
 - * limited CPU time
 - * ...
- Language-Level-Sandboxing
 - * restrict untrusted code on the language level
 - ◆ e.g. JavaScript
 - * no dangerous functionality (I/O, syscalls, ...)
 - ◆ ask user for permission if needed
 - * interpreter
 - * eBPF verifies certain properties before executing code
 - ◆ termination

- ◆ no loops/recursion
 - halting problem
 - ◆ jumps may not form loops
 - ◆ only allowed functions
- Rule-based Execution
 - * define what an application is allowed to do
 - ◆ white/blacklists
 - * e.g. seccomp
 - * good policies are hard to create but secure and efficient
- Container
 - * OS-level virtualization
 - * isolated user-space instances
 - * each container is assigned resources
 - ◆ memory
 - ◆ folder
 - ◆ ...
 - * only see assigned resources
 - * shared OS, separate libraries/dependencies
 - * Control Groups?
 - * Namespaces
 - ◆ isolate system resources between processes
 - cannot see other processes
 - own mount
 - own network stack
 - ...
- Virtualization
 - * no shared kernel since process runs in own OS
 - * emulate entire system => massive overhead
 - ◆ bare metal
 - run directly on hardware
 - ◆ hosted
 - on top of host OS
 - * VM escape
 - ◆ access to host and other VMs
- Isolation
 - isolate application from system
 - * trusted application on untrusted system
 - applications

- * sensitive data
- * distrust against cloud provider
- * intellectual property
- * rights management
- Trusted Computing Base TCB
 - * CPU and firmware usually
 - * kernel and system programs usually too
 - ◆ protected by protection rings
- Trusted-Execution Environment
 - * secure area within CPU
 - * guarantee integrity and confidentiality for code + data
 - * still shared hardware
 - * small overhead
 - * CPU is the only TCB
 - * memory encrypted and inaccessible to OS
 - * does not protect against side channel attacks
- Hardware Isolation with Hardware Security Modules HSM
 - * external, dedicated hardware
 - ◆ nothing shared
 - * protects high-value cryptographic keys
 - * crypto processor for
 - ◆ key generation/management
 - ◆ signatures
 - ◆ data en/decryption
 - ◆ strong RNG
 - ◆ secure timestamp